



Streaming

Yalun Lin Hsu & Yuchen Cheng

2020-10-12

MapReduce

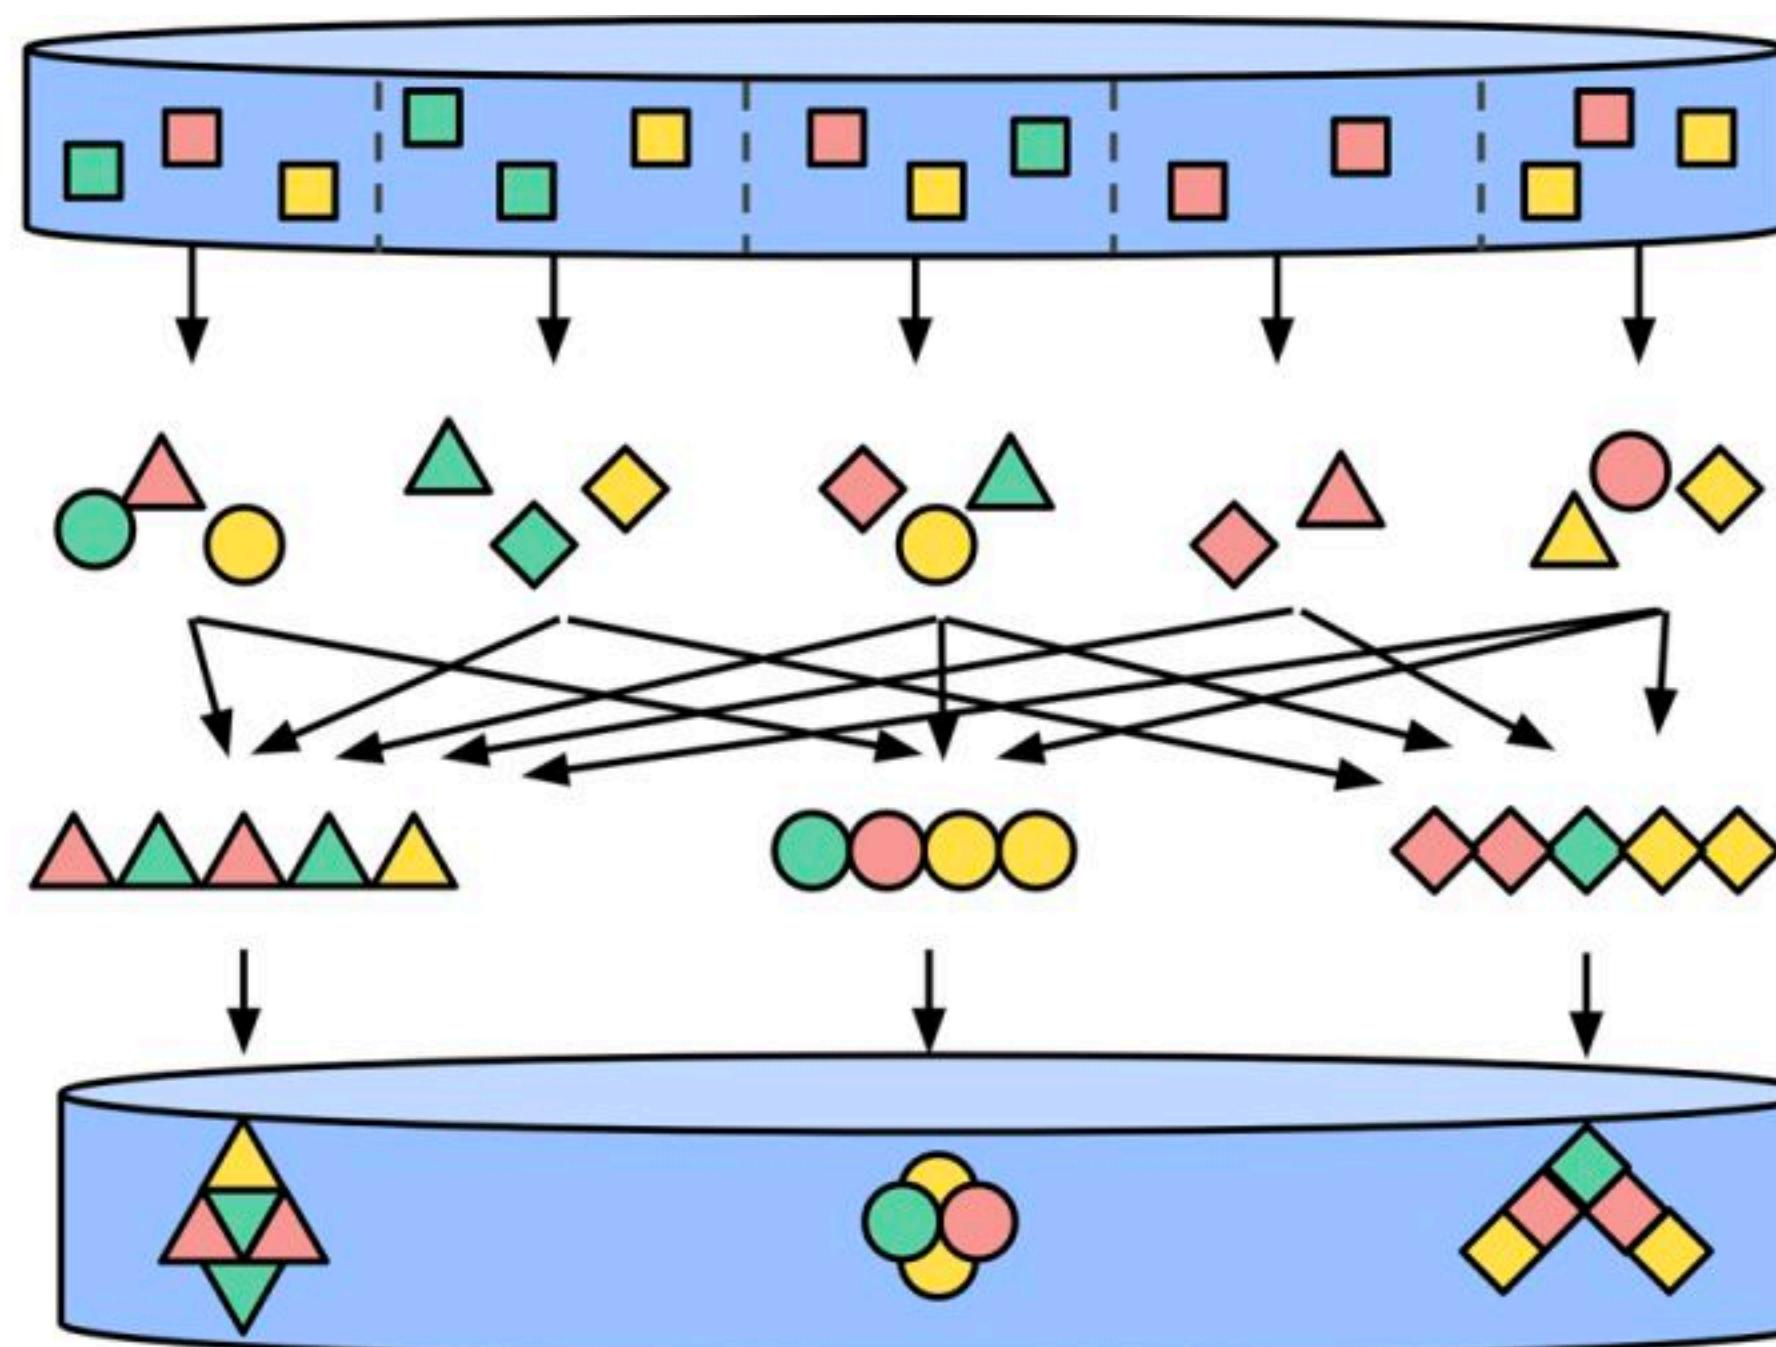
(Prepare)

Map

(Shuffle)

Reduce

(Produce)



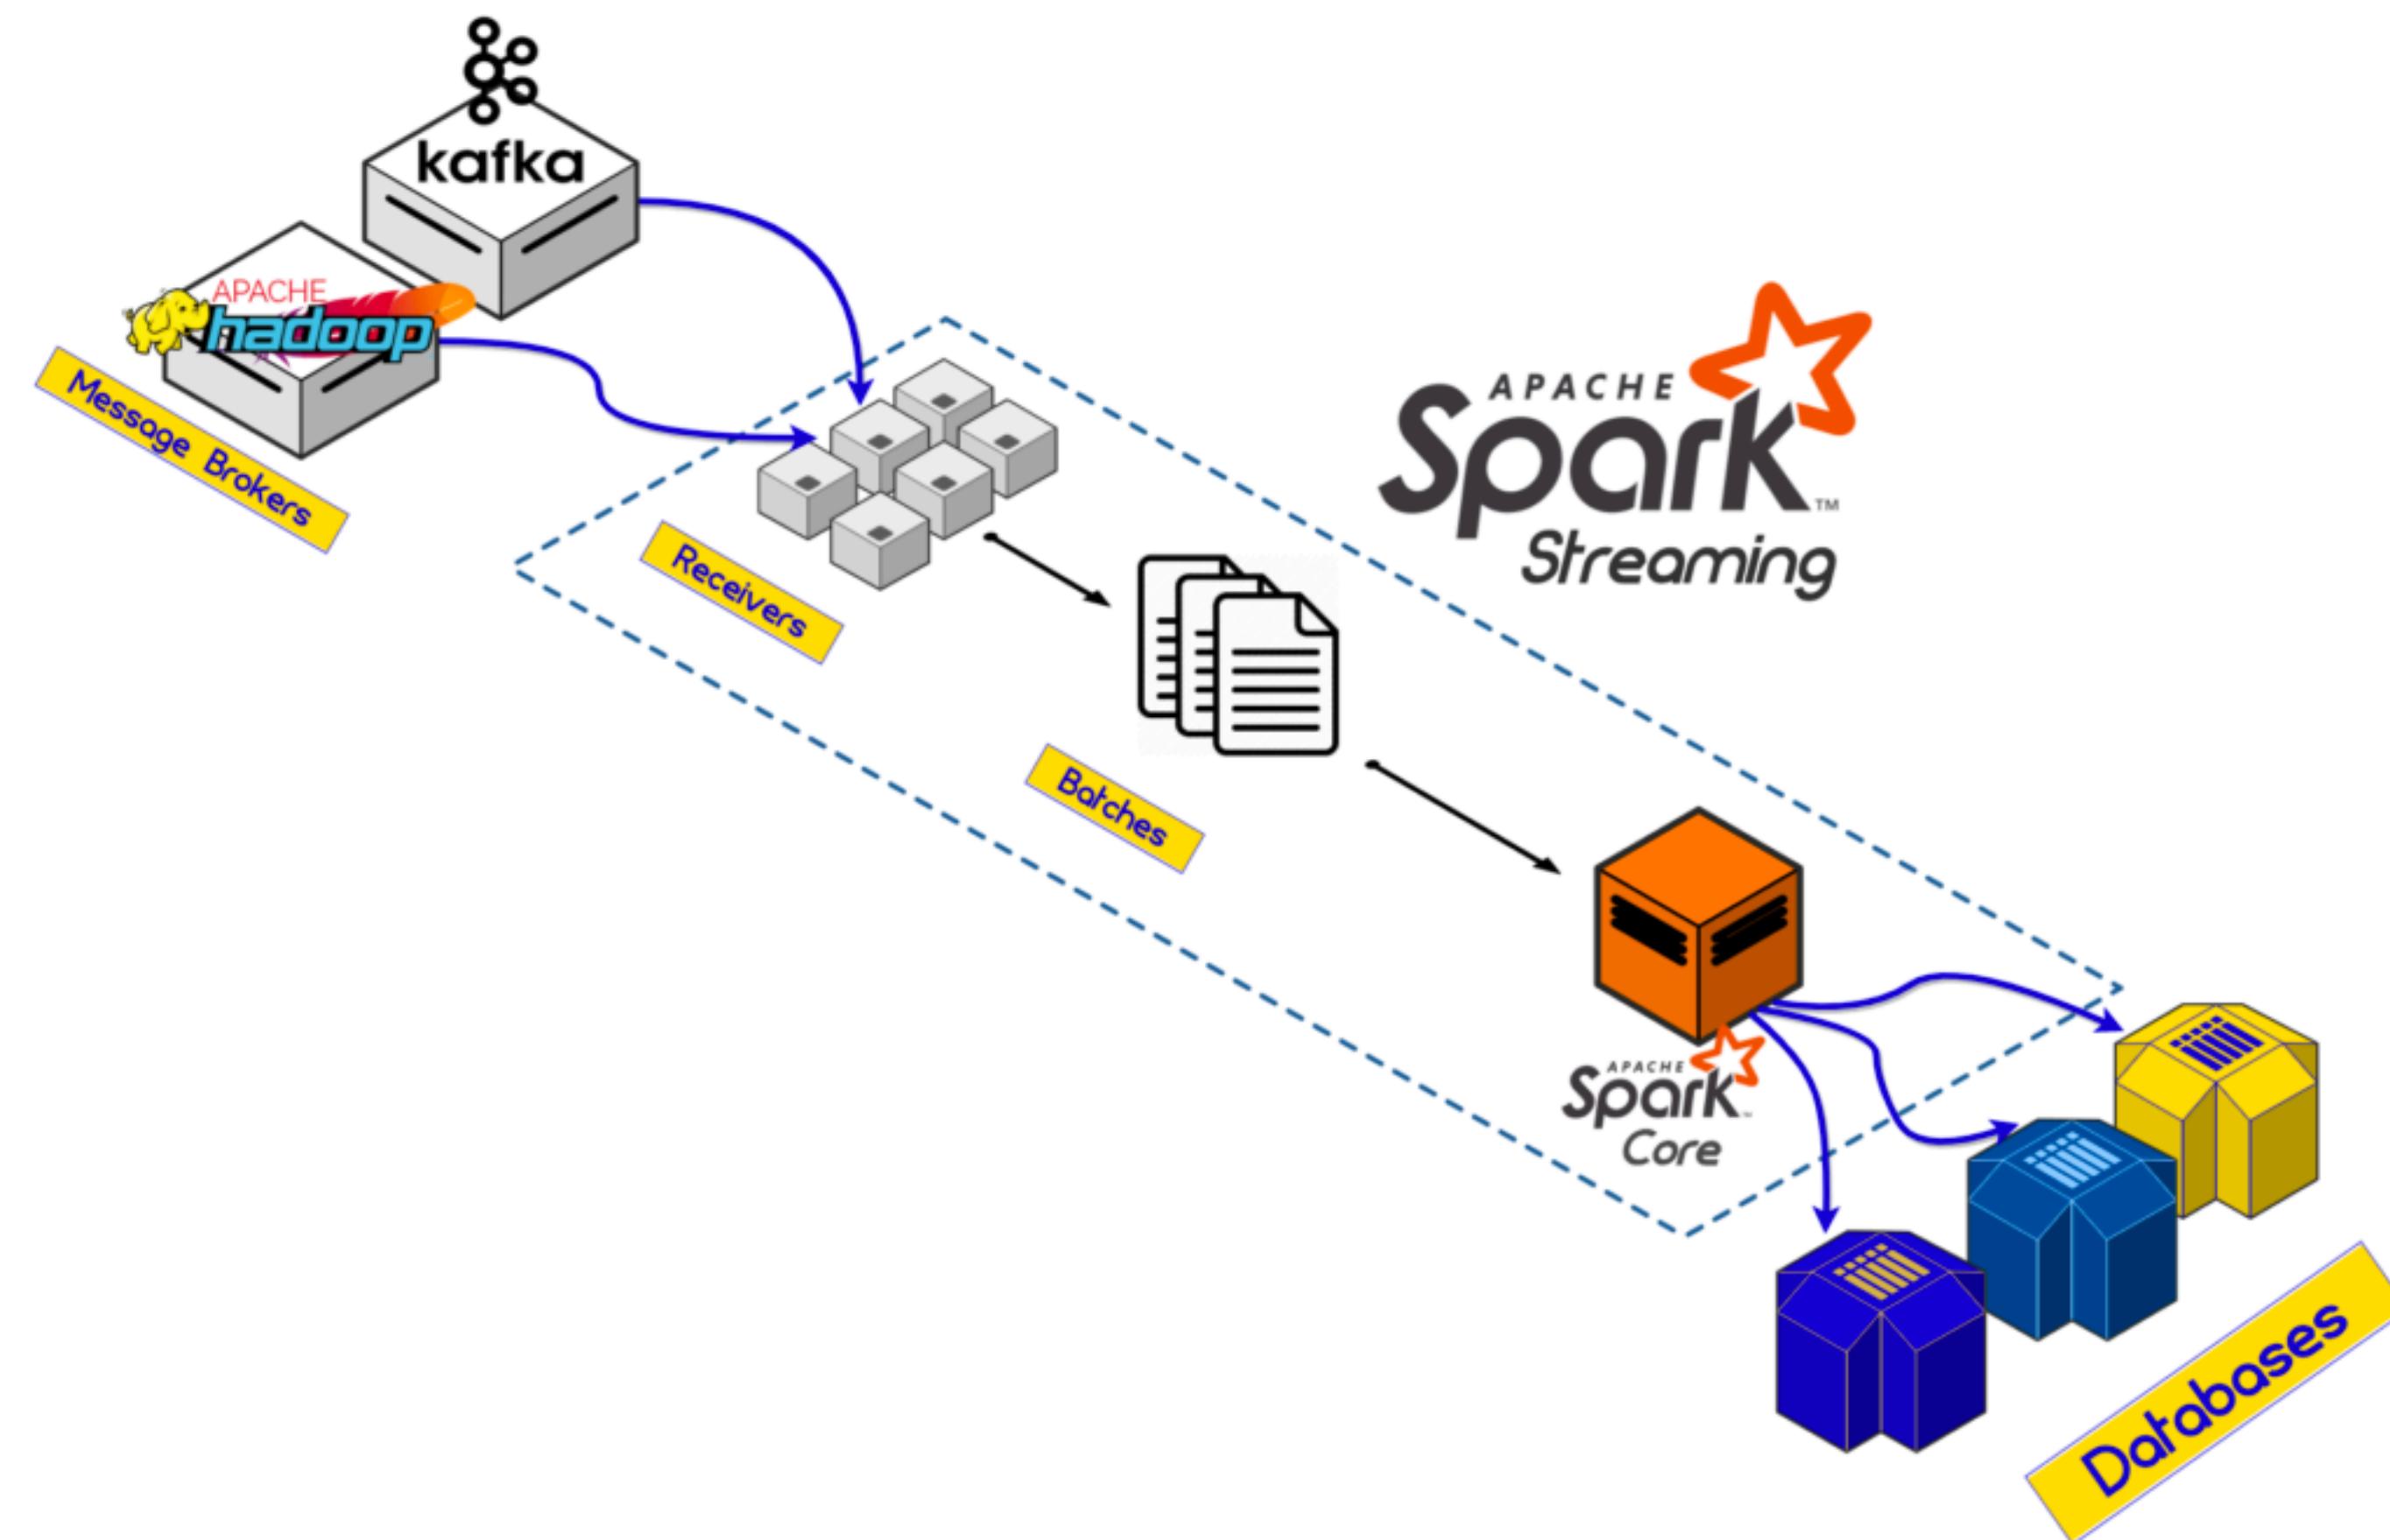
Hadoop & Spark



VS

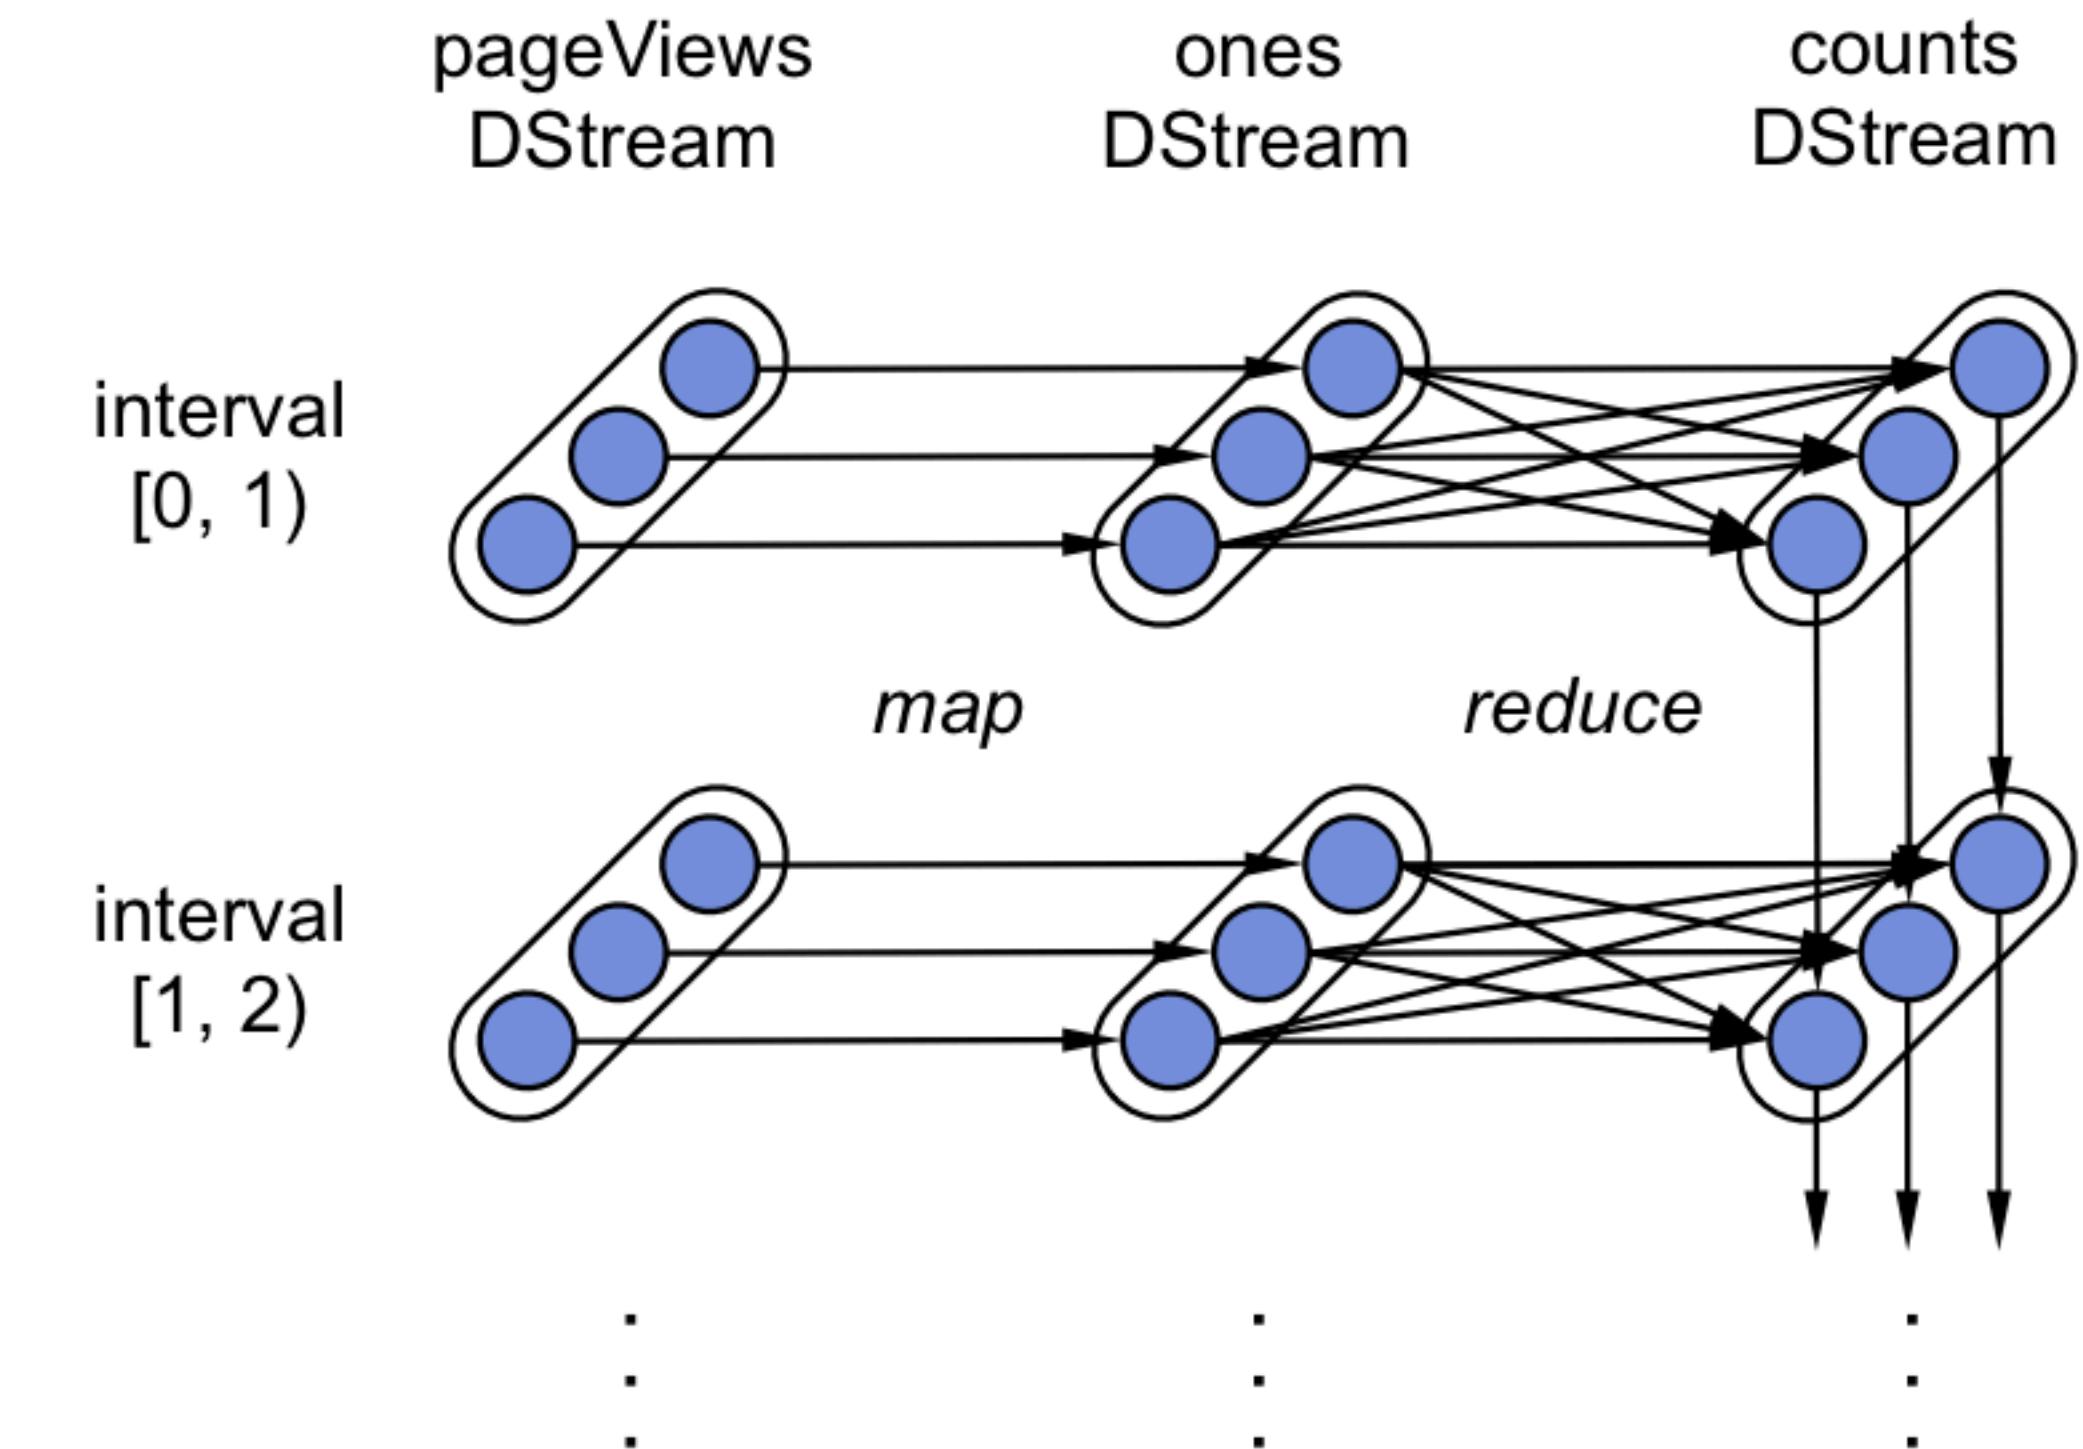


Spark Streaming

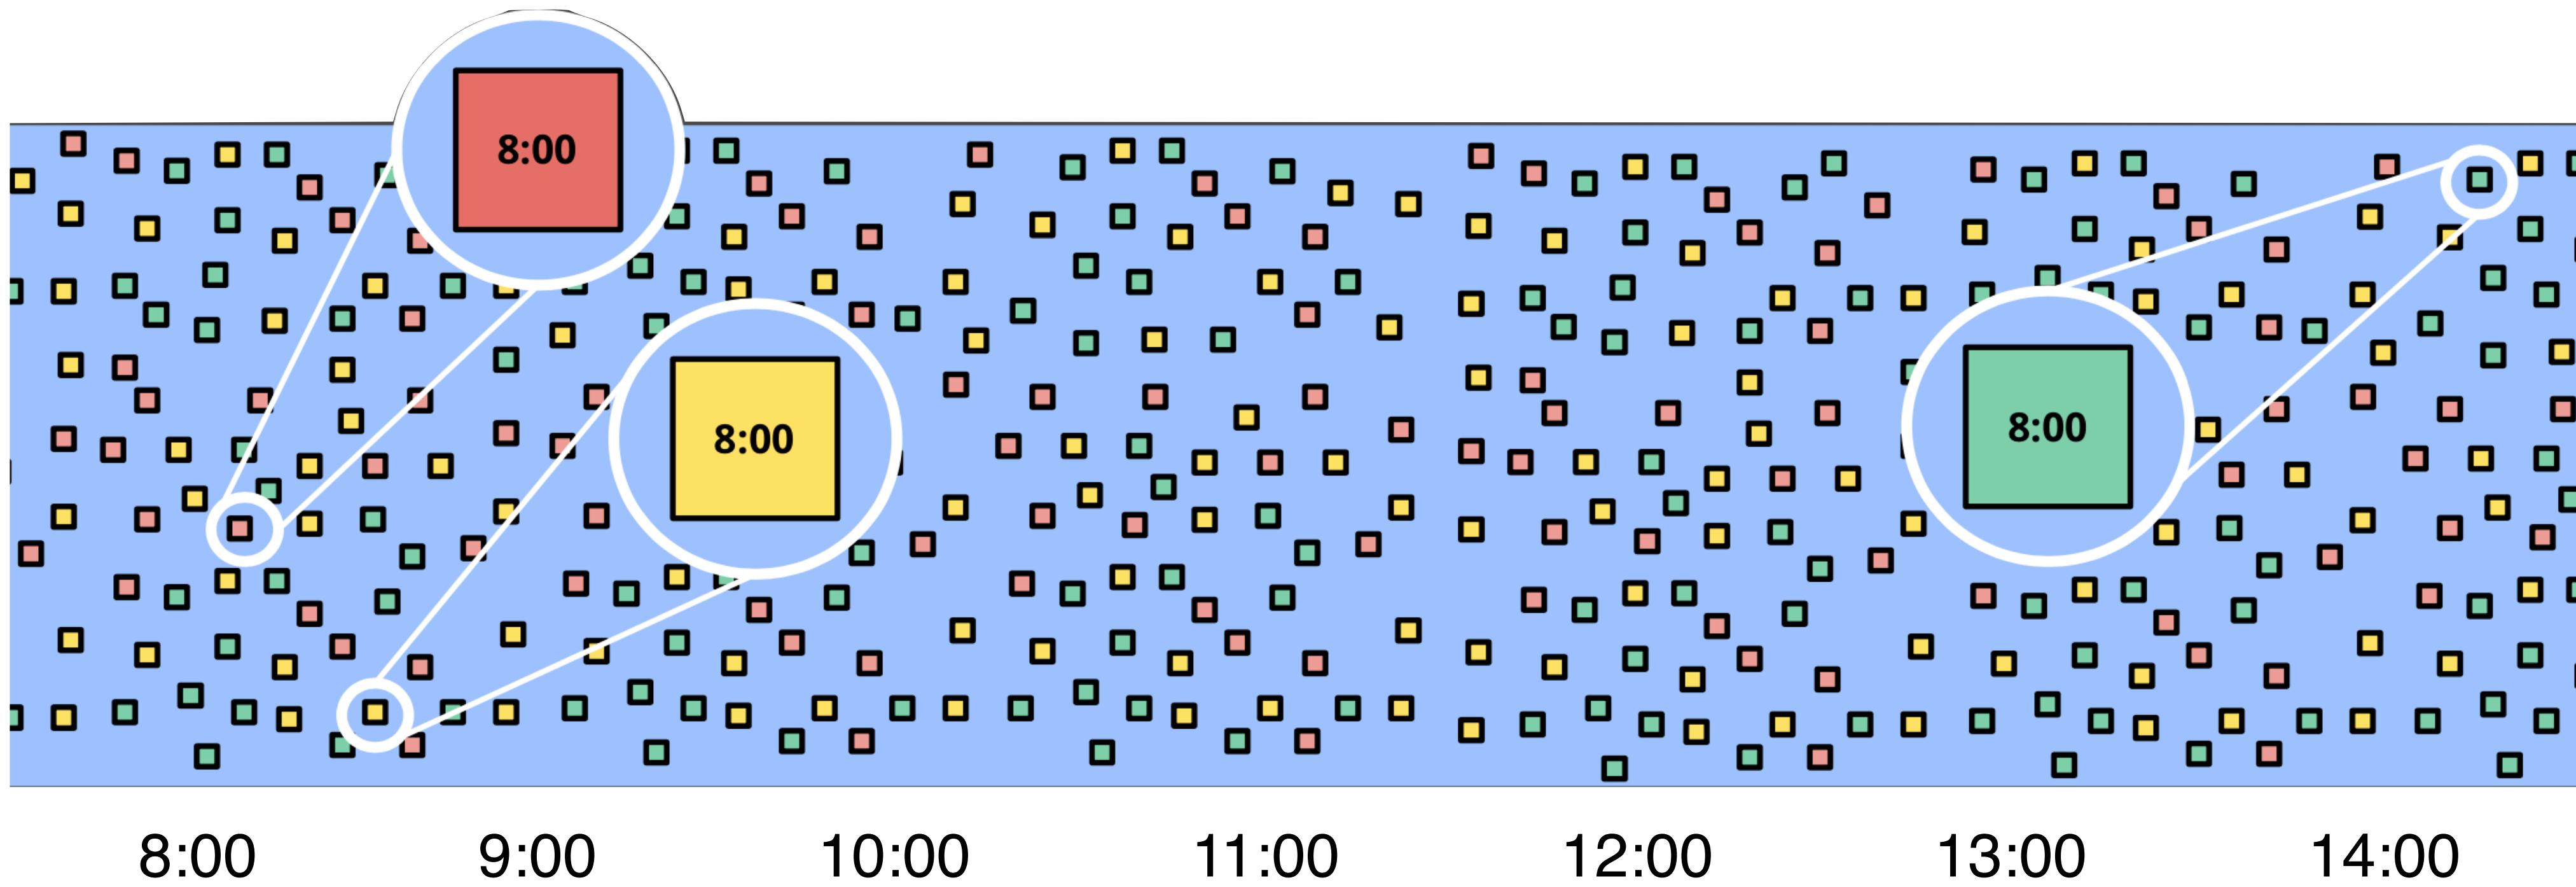


Discretized Streams

```
pageViews = readStream("http:// ... ", "1s")
ones = pageViews.map(event => (event.url, 1))
counts = ones.runningReduce((a, b) => a + b)
```



What is Streaming?



Data can be infinitely big with unknown delays.

Back to Discretized Streams

Spark Streaming 1.x is not **TRUE** Streaming

Beat Batch, Build Stream Dynasty

- **Tools for reasoning about time**

- The Dataflow Model: A Practical Approach to Balancing Correctness, Latency, and Cost in Massive-scale, Unbounded, Out-of-order Data Processing

- **Correctness**

- Lightweight Asynchronous Snapshots for Distributed Dataflows

Beat Batch, Build Stream Dynasty

- **Tools for reasoning about time**

- The Dataflow Model: A Practical Approach to Balancing Correctness, Latency, and Cost in Massive-scale, Unbounded, Out-of-order Data Processing

- **Correctness**

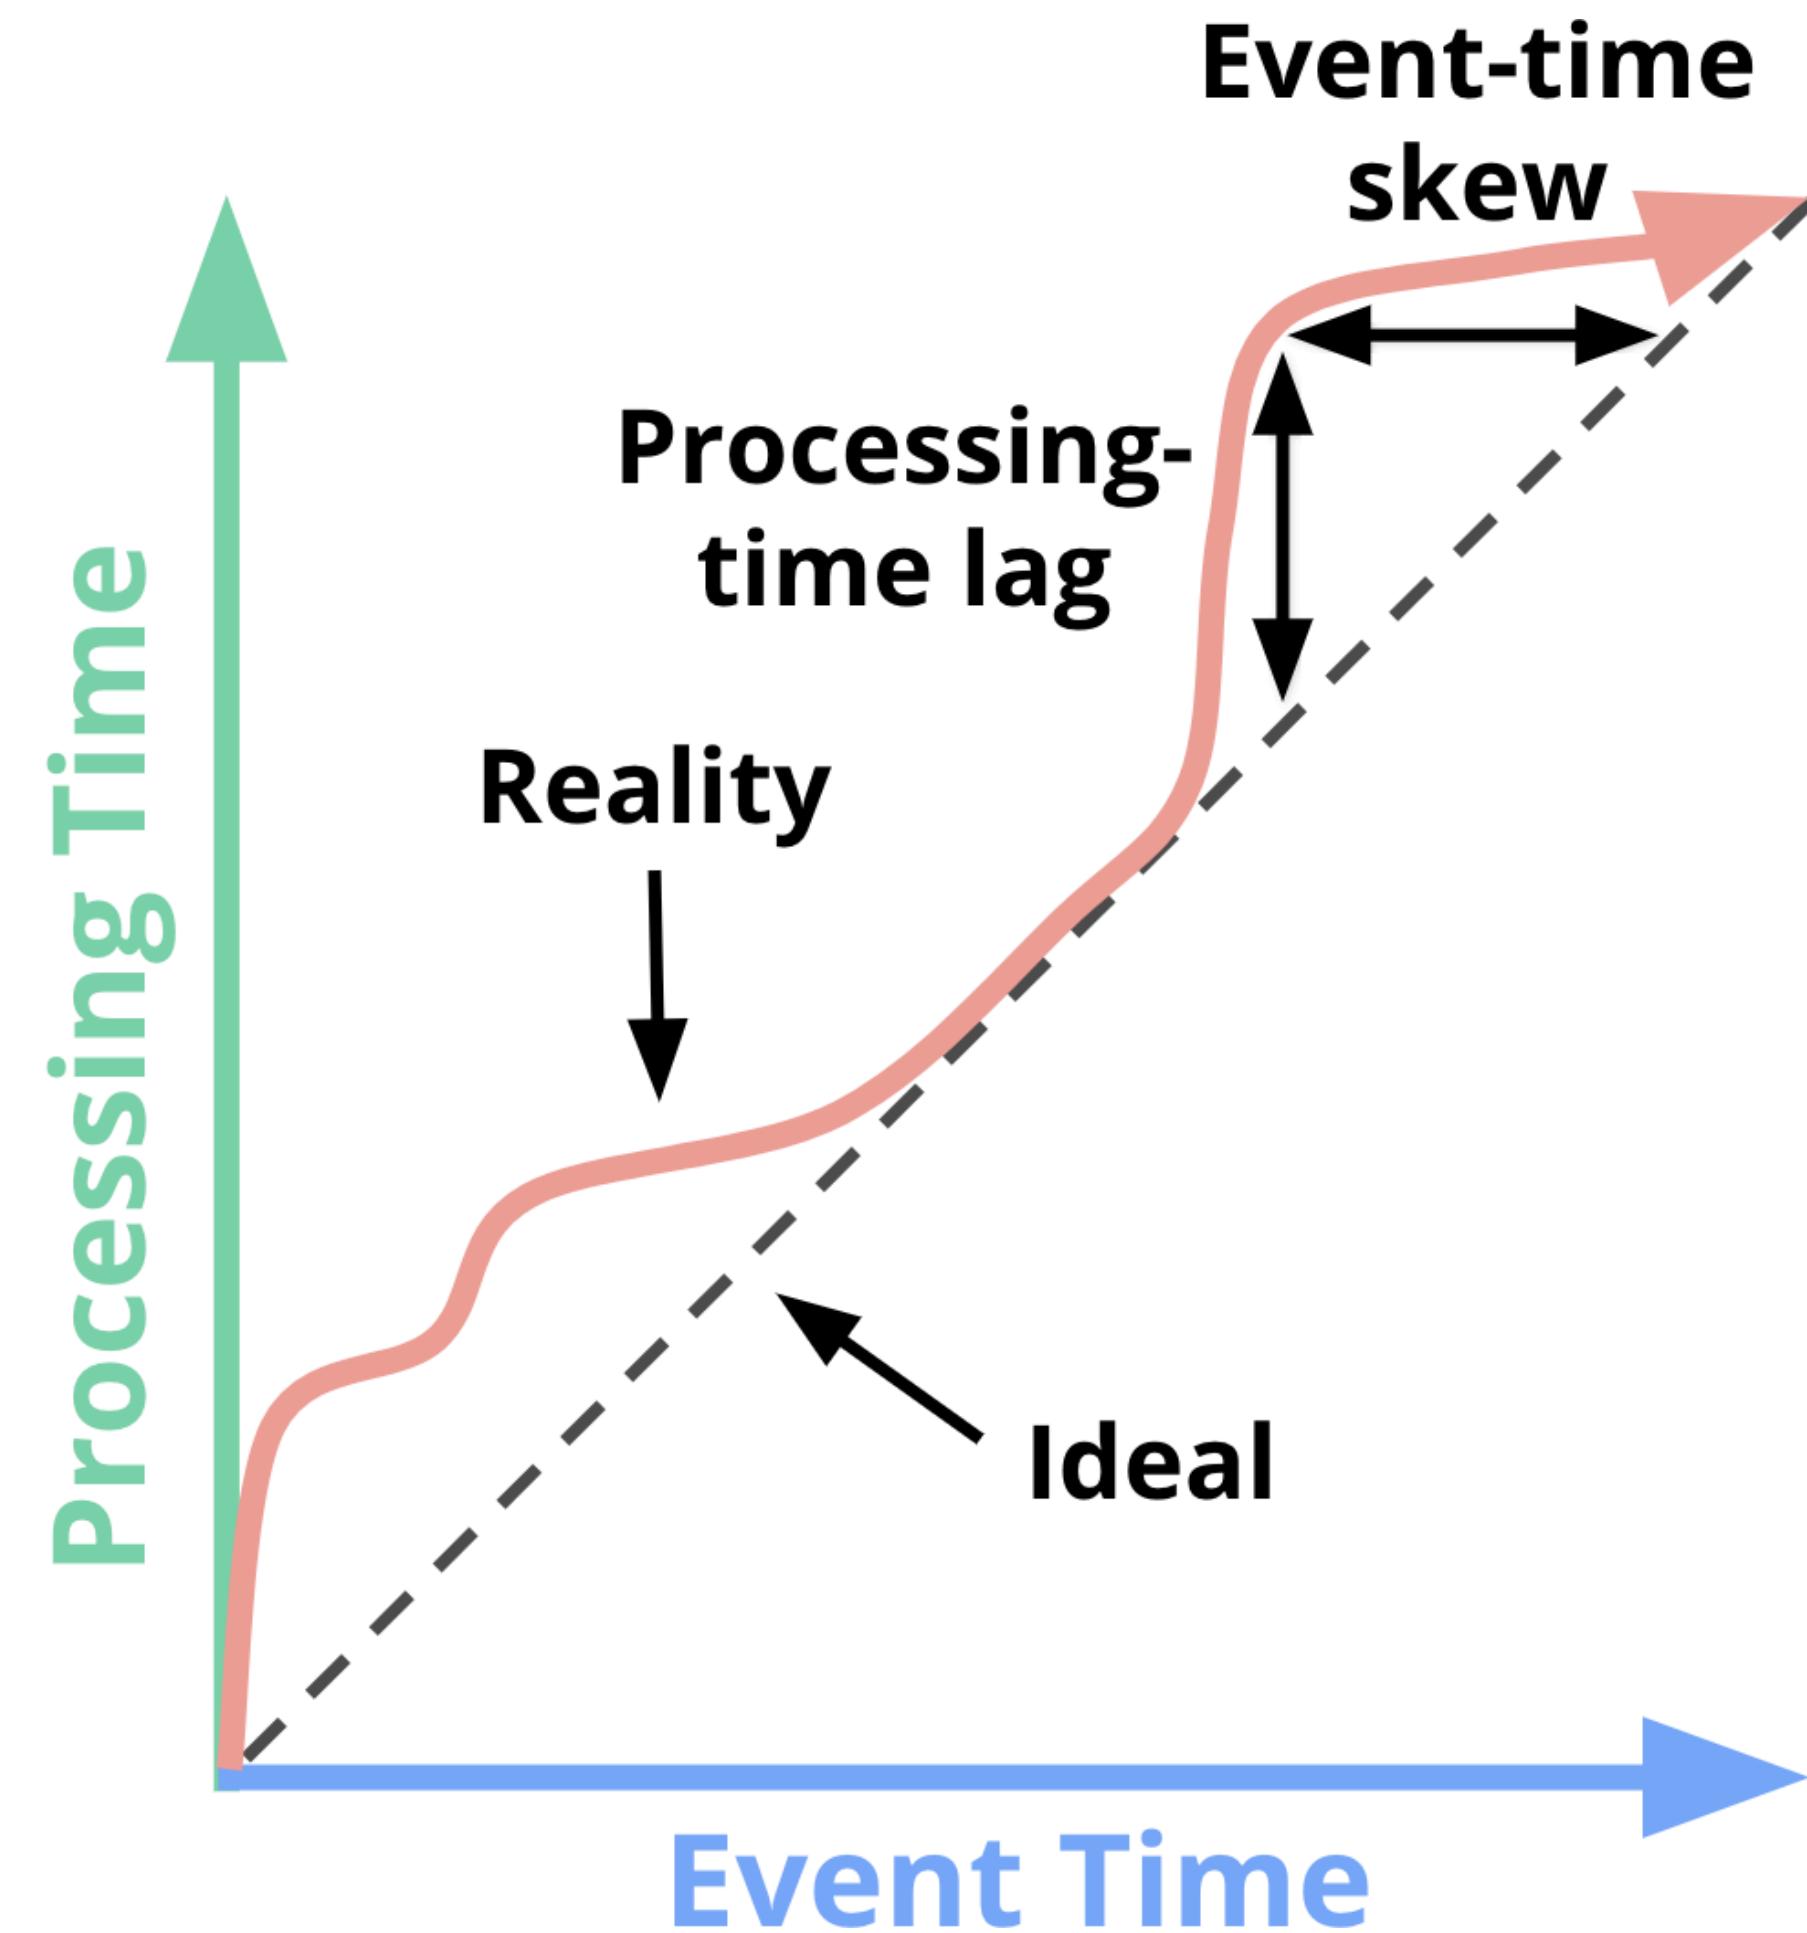
- Lightweight Asynchronous Snapshots for Distributed Dataflows

The Dataflow Model

A Practical Approach to Balancing Correctness, Latency, and Cost in
Massive-scale, Unbounded, Out-of-order Data Processing

<https://www.usenix.org/conference/atc18/presentation/oakes>

Time-domain Mapping



The Dataflow Model

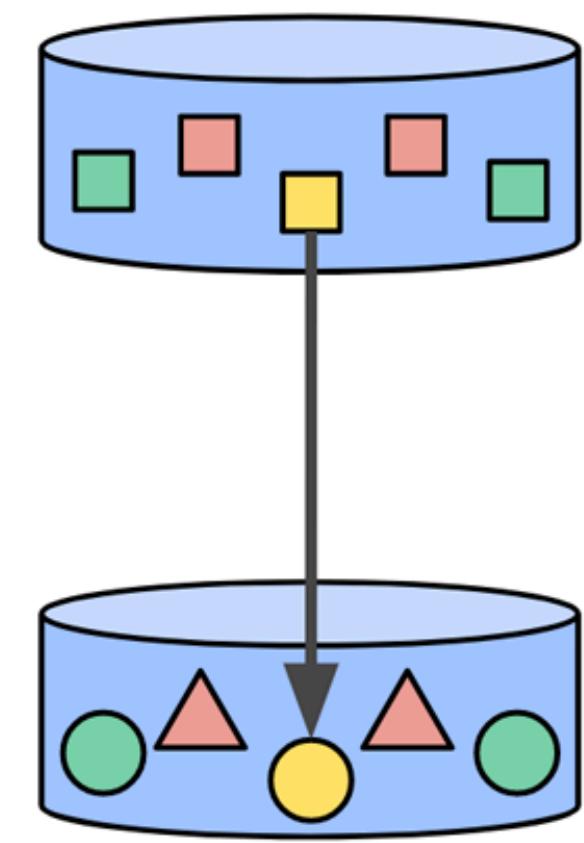
What are you computing?

Where in event time?

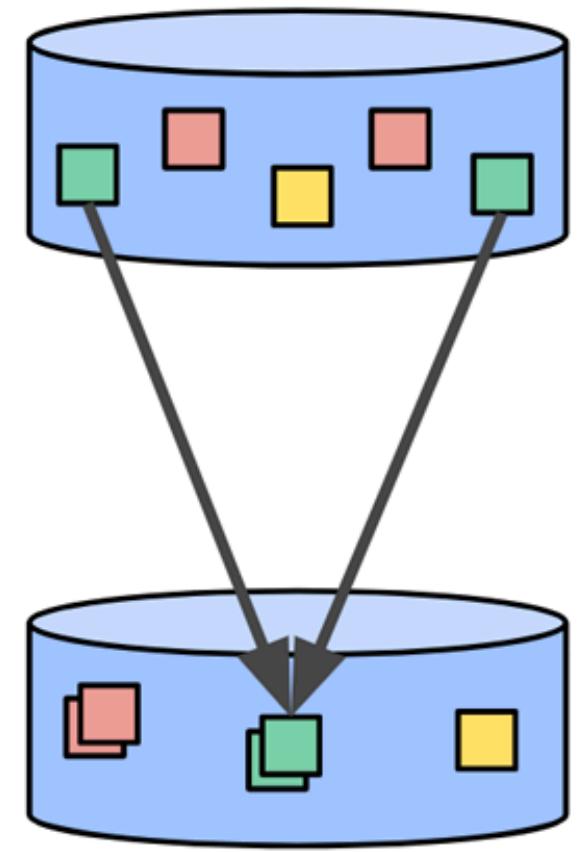
When in processing time?

How do refinements relate?

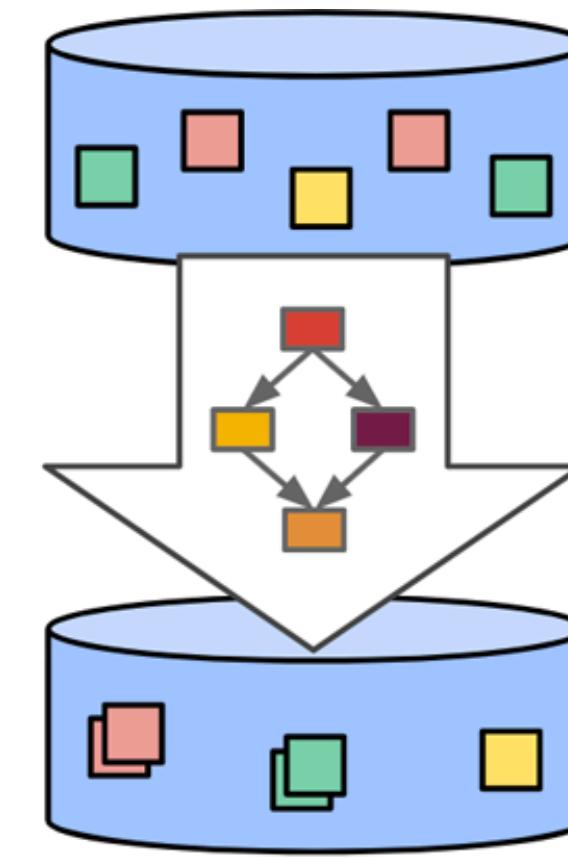
What are you computing?



Element-wise



Grouping

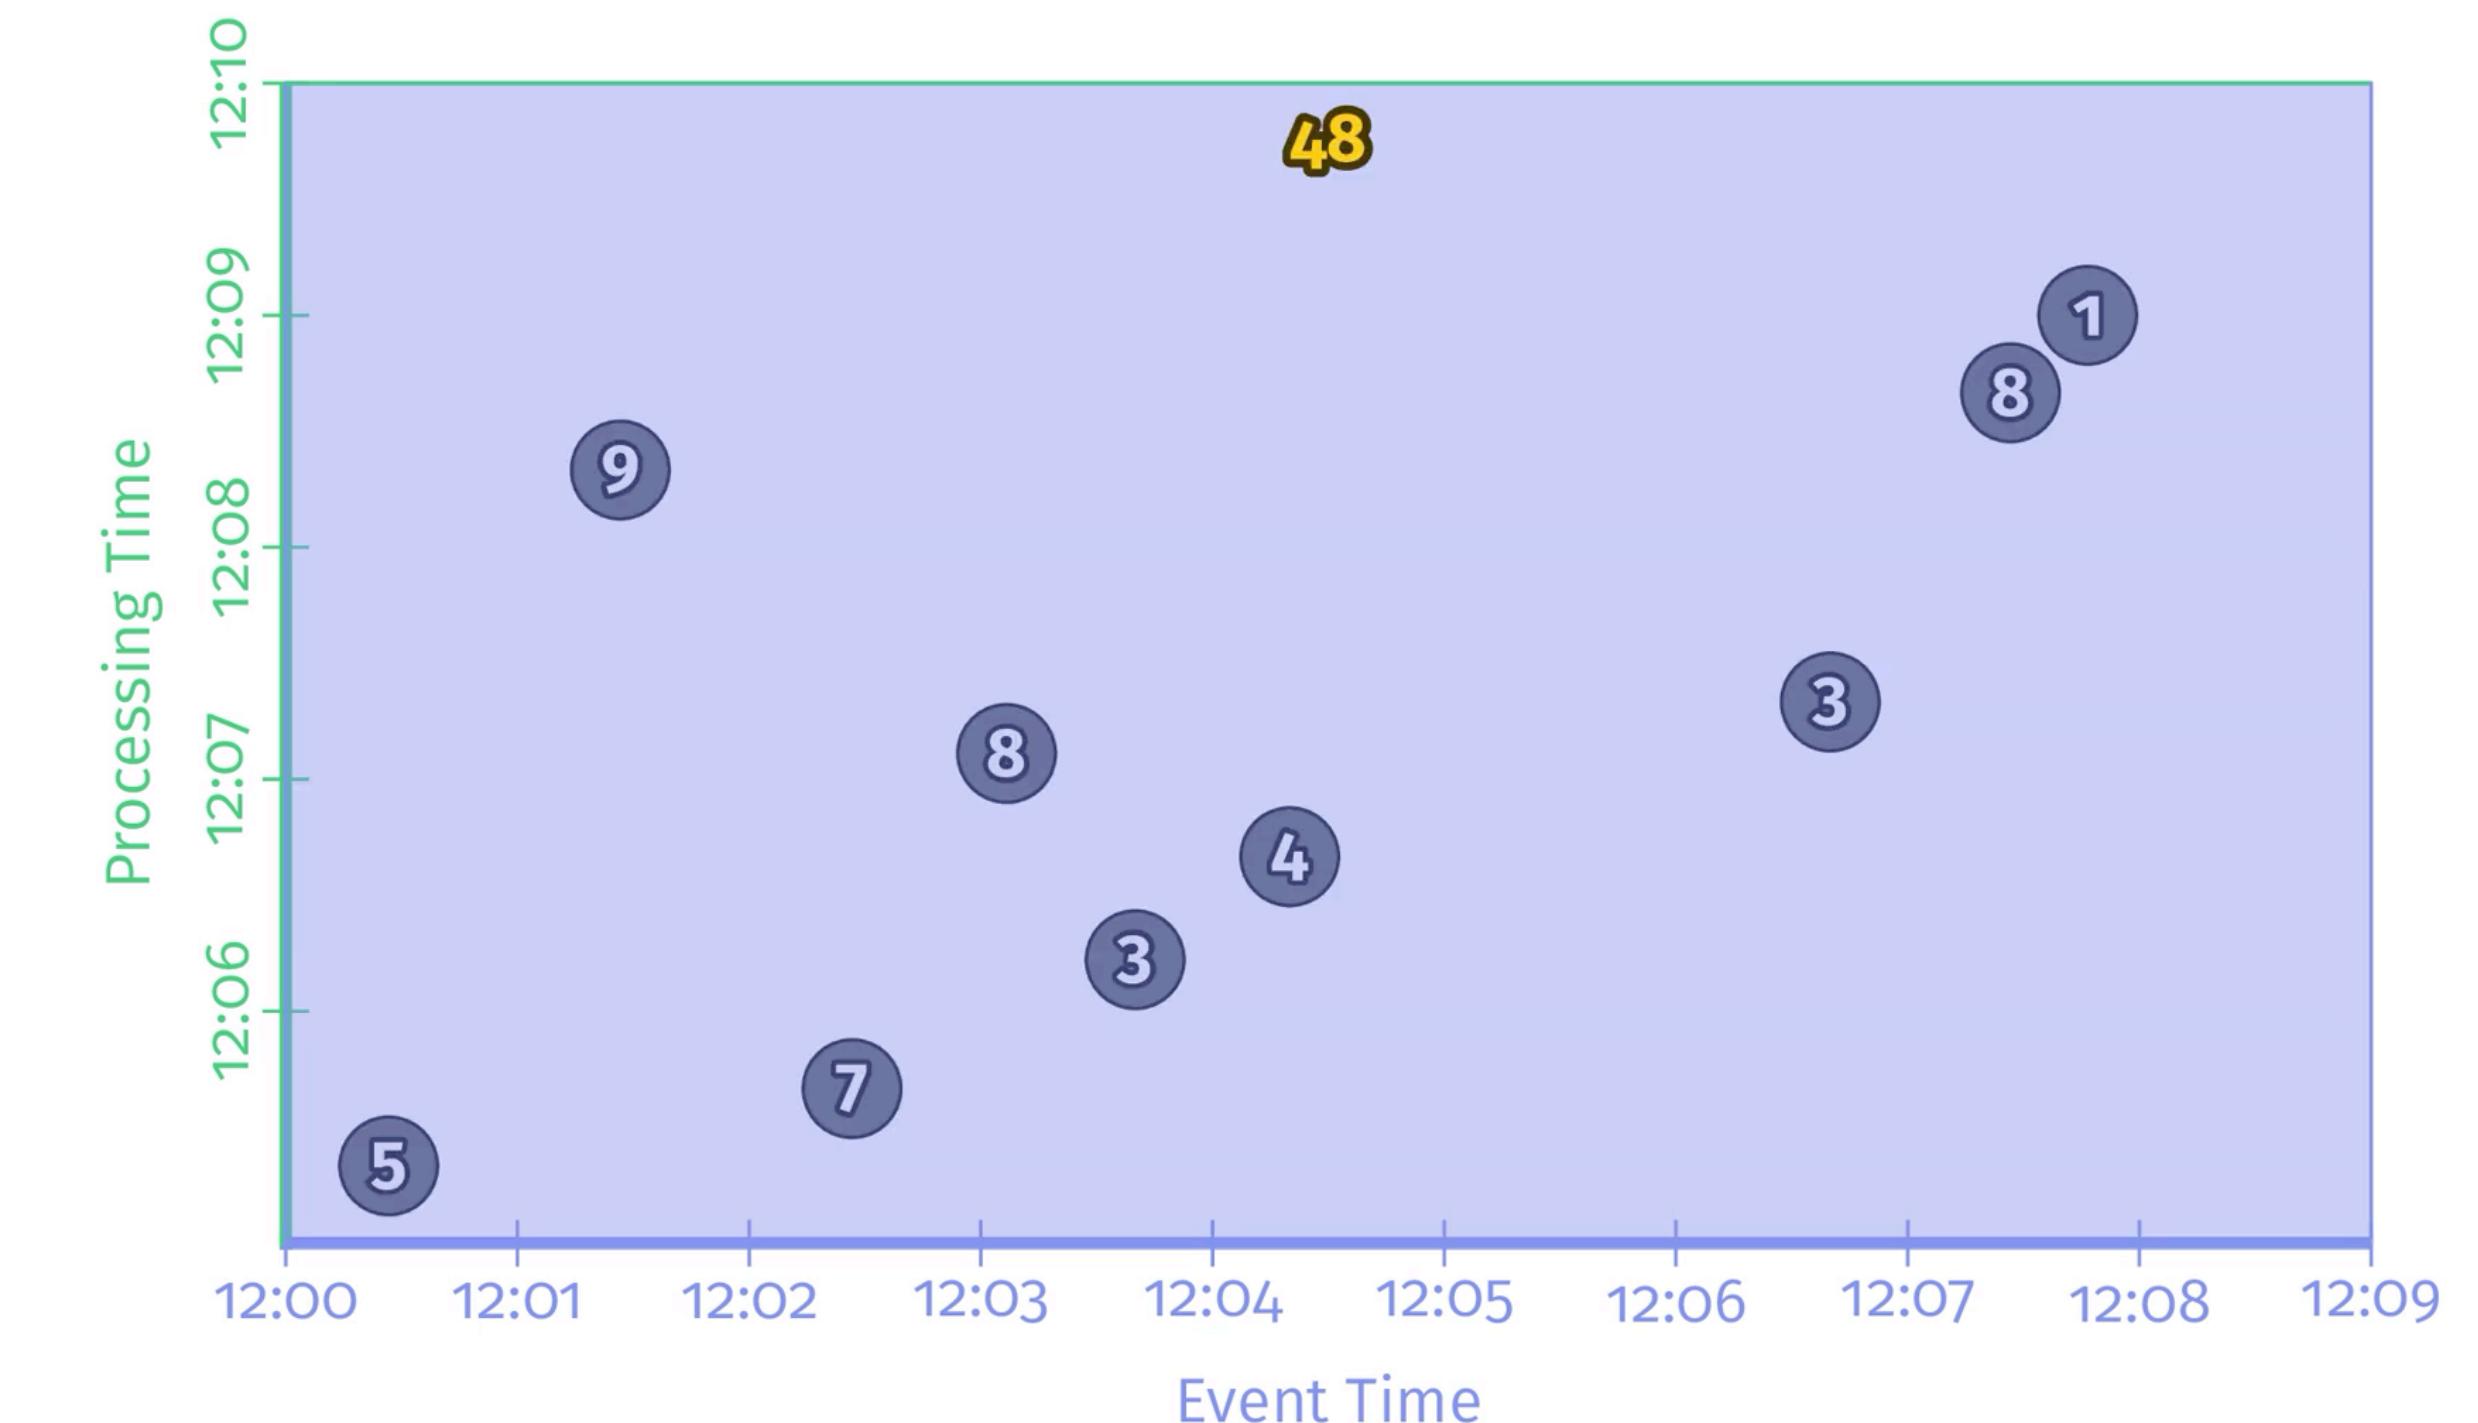


Composite

Types of Transformations

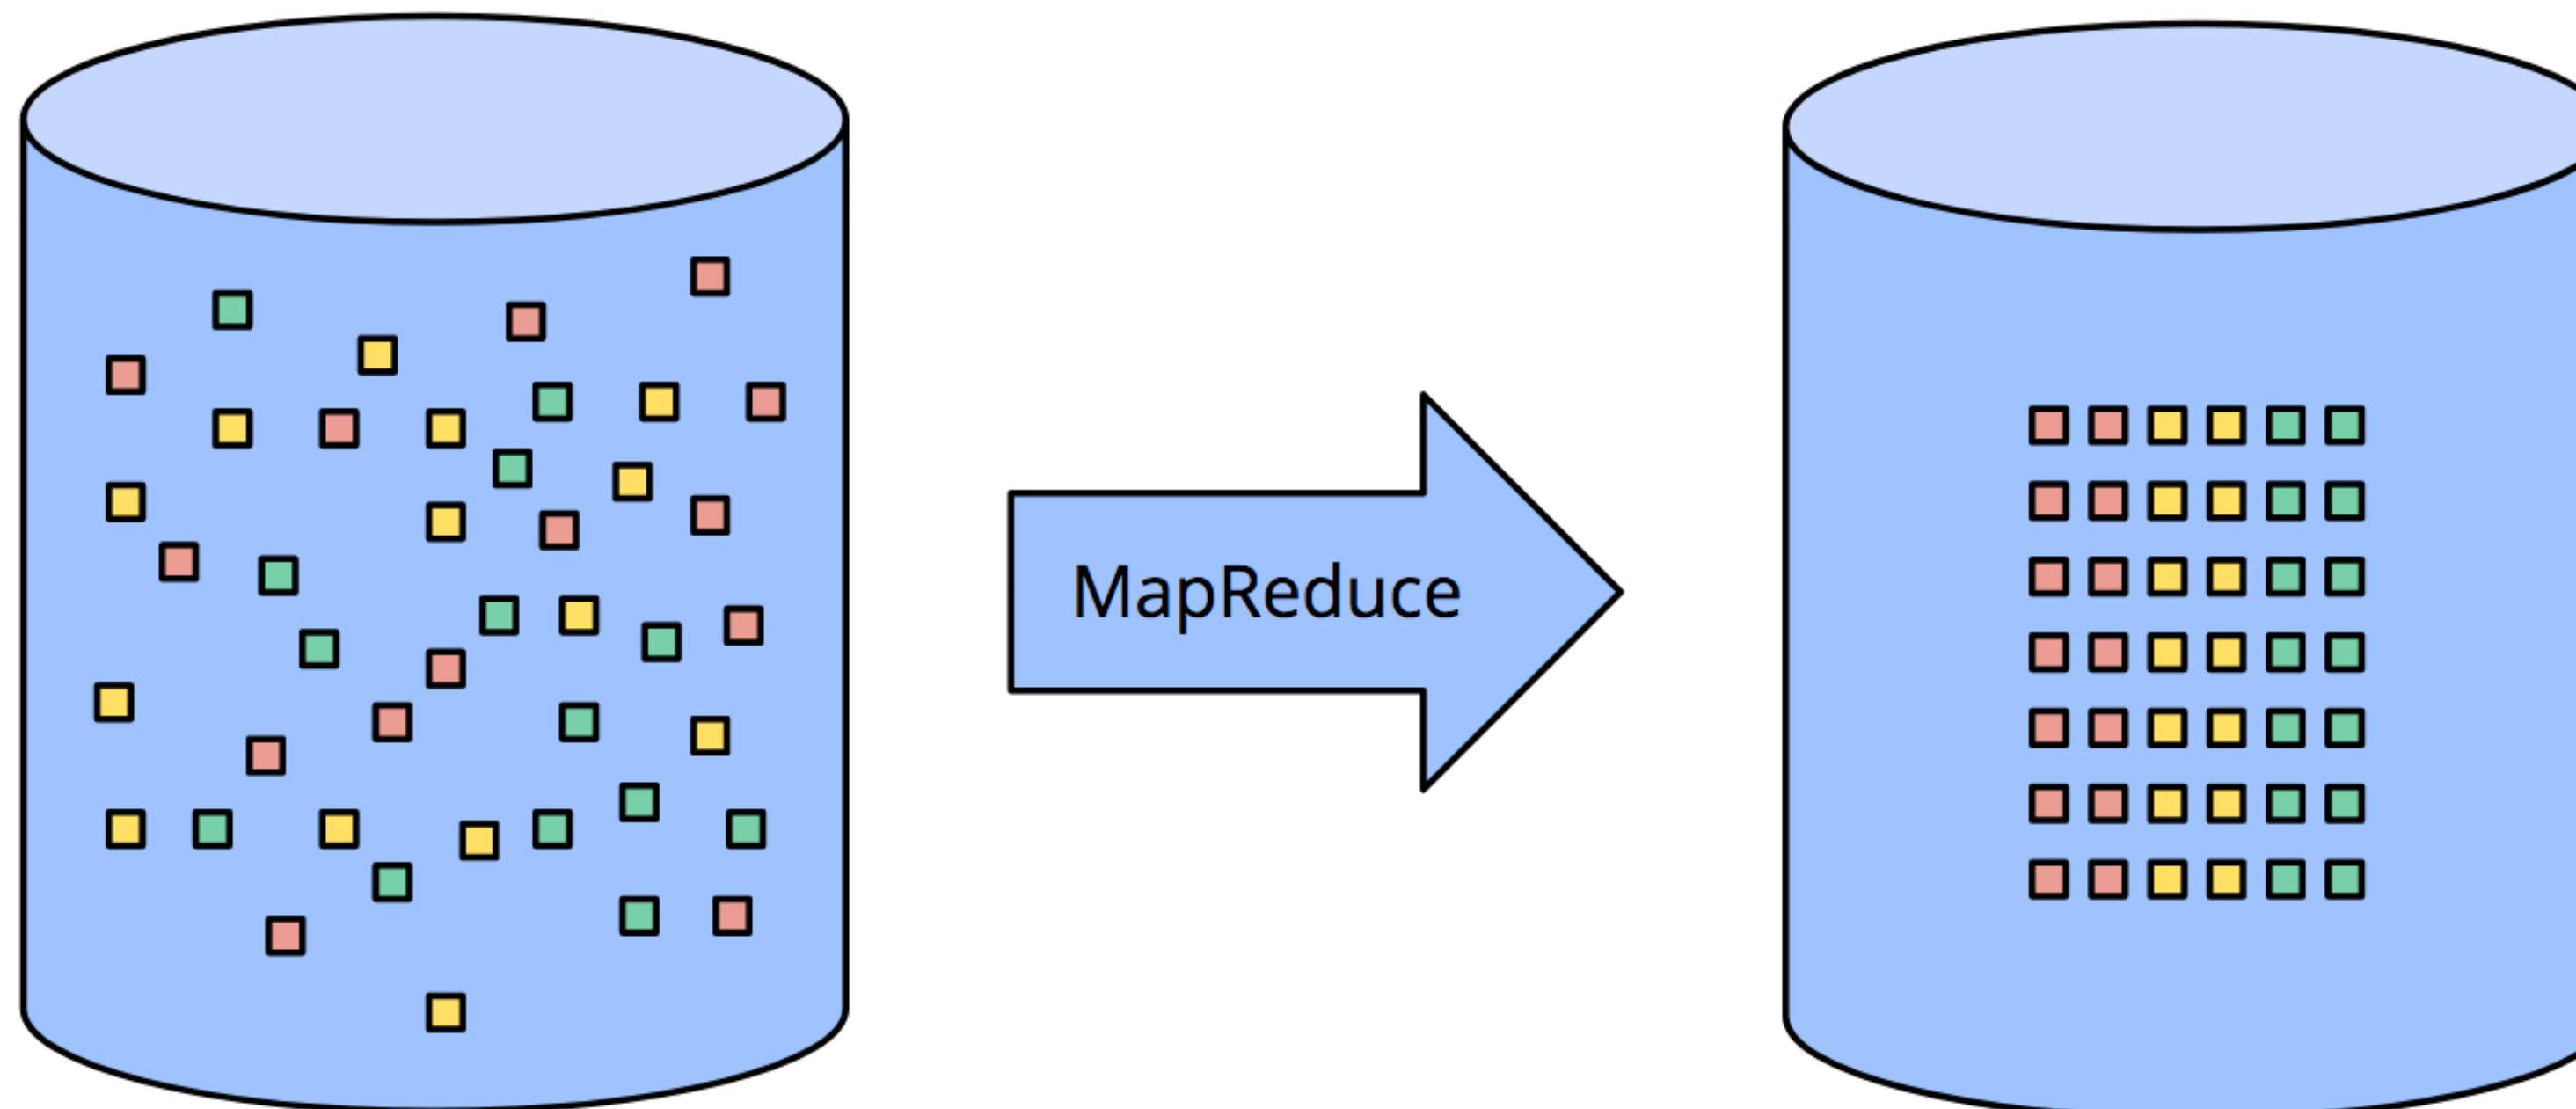
What are you computing?

```
PCollection<KV<String, Integer>> input = IO.read( ... );  
PCollection<KV<String, Integer>> output = input  
    .apply(Sum.integersPerKey());
```



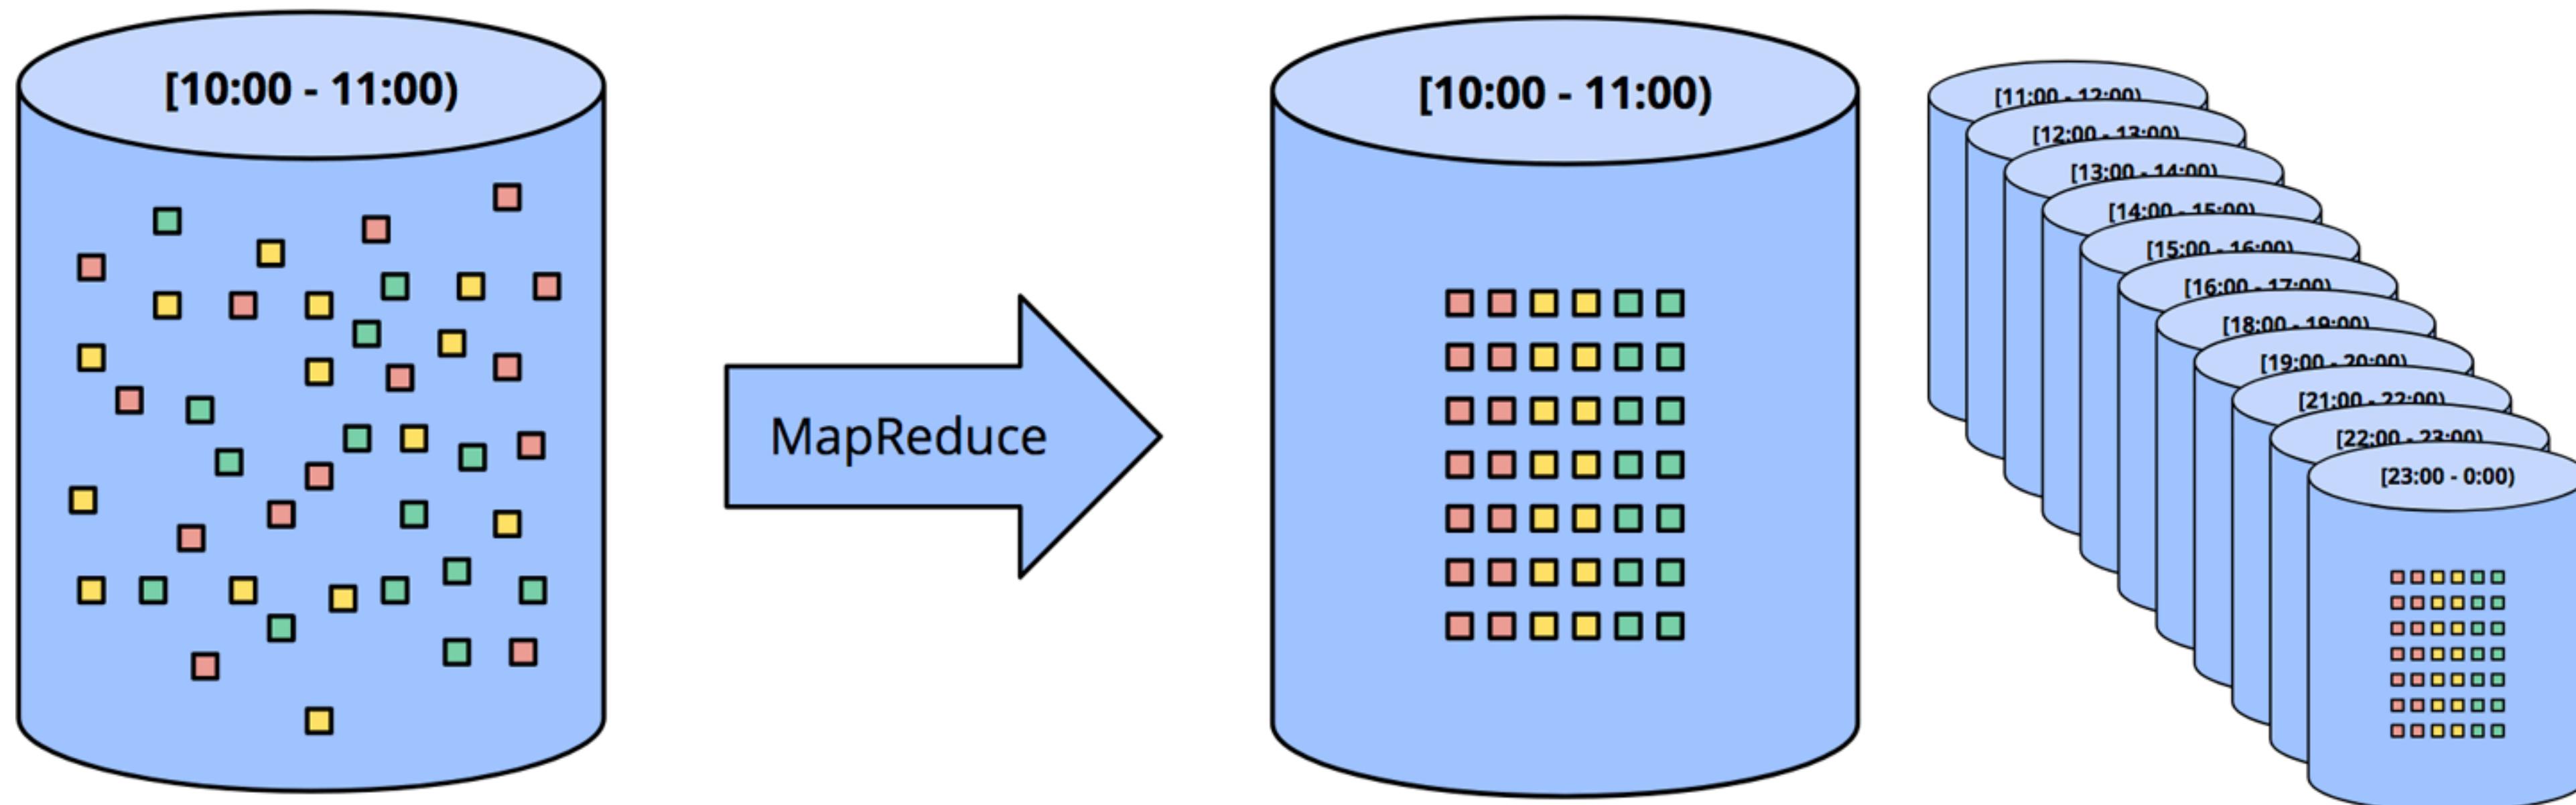
Where in event time?

Batch Example



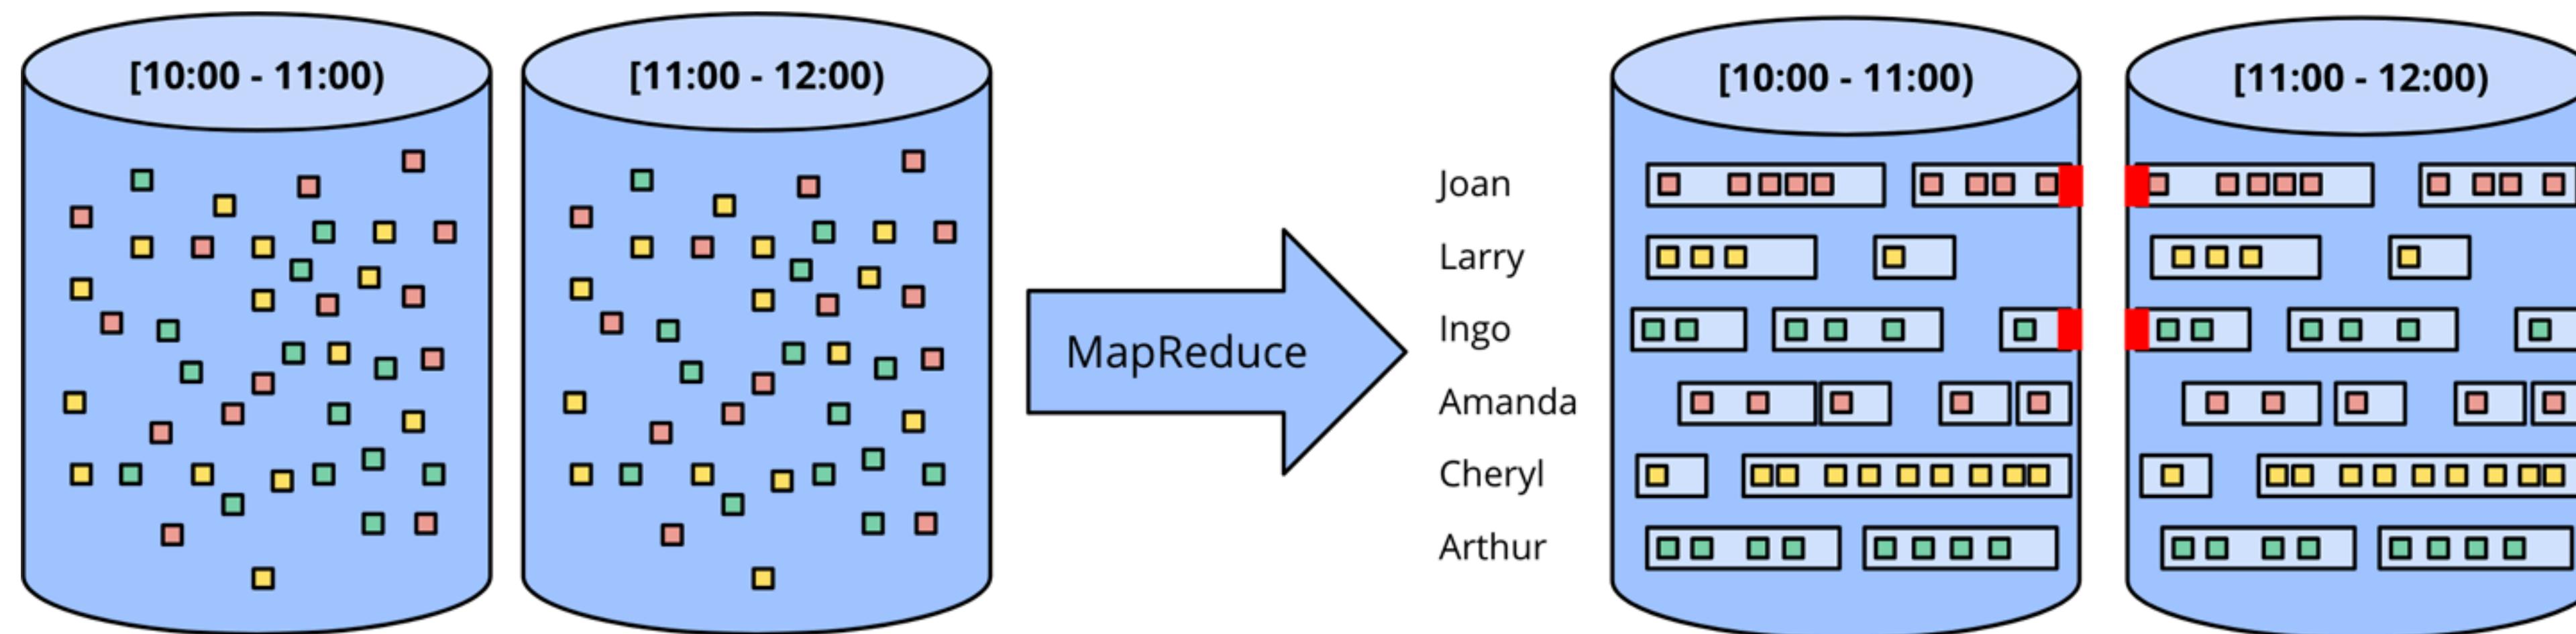
Where in event time?

Batch Example



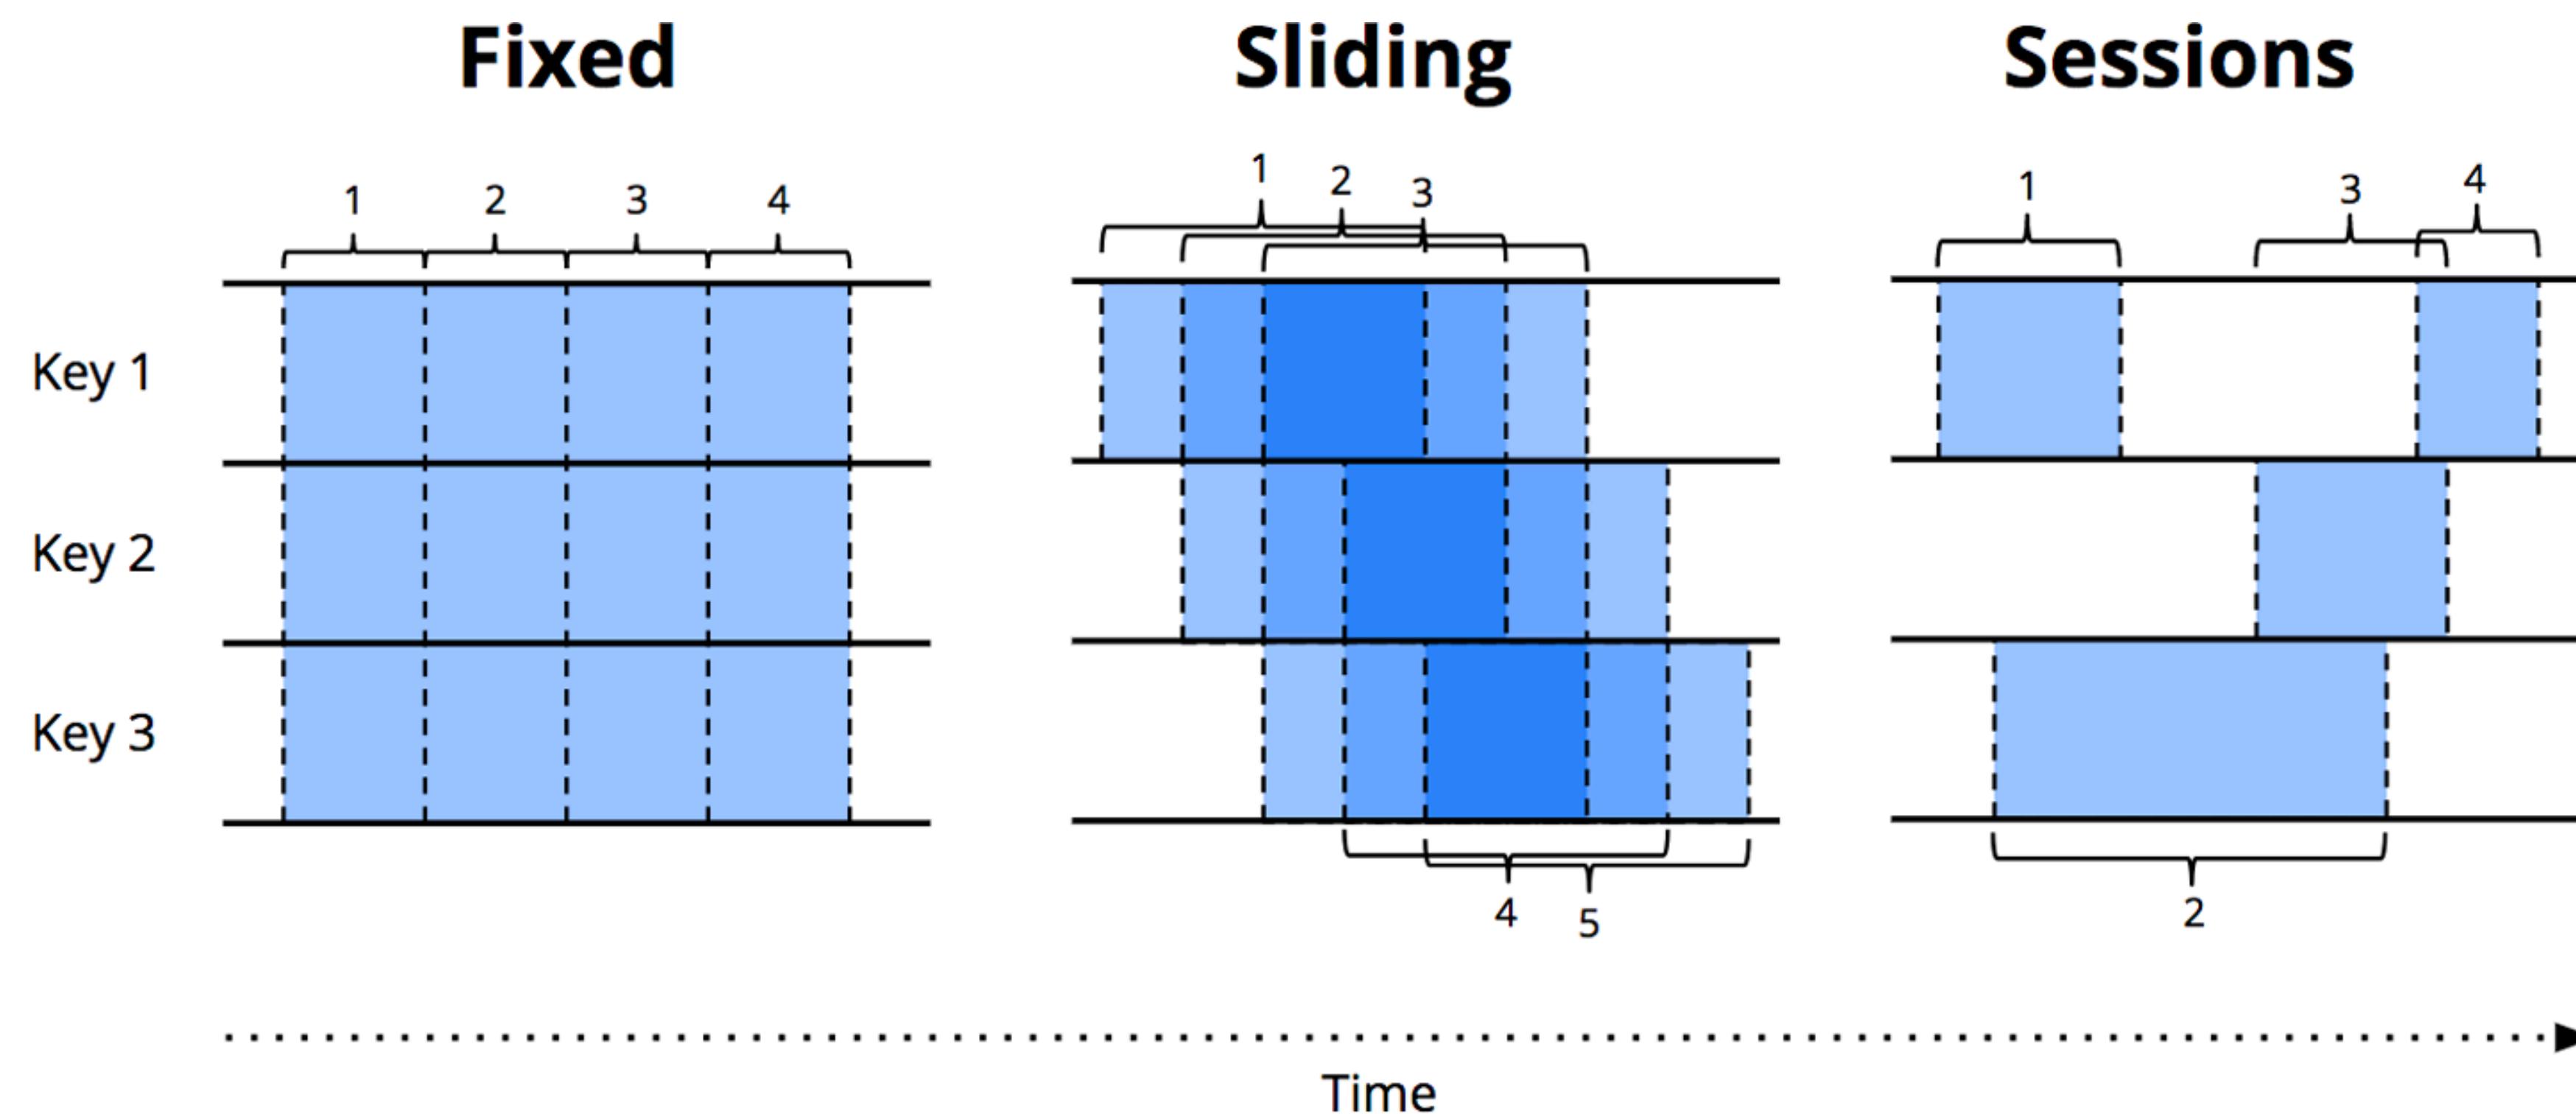
Where in event time?

Batch Example



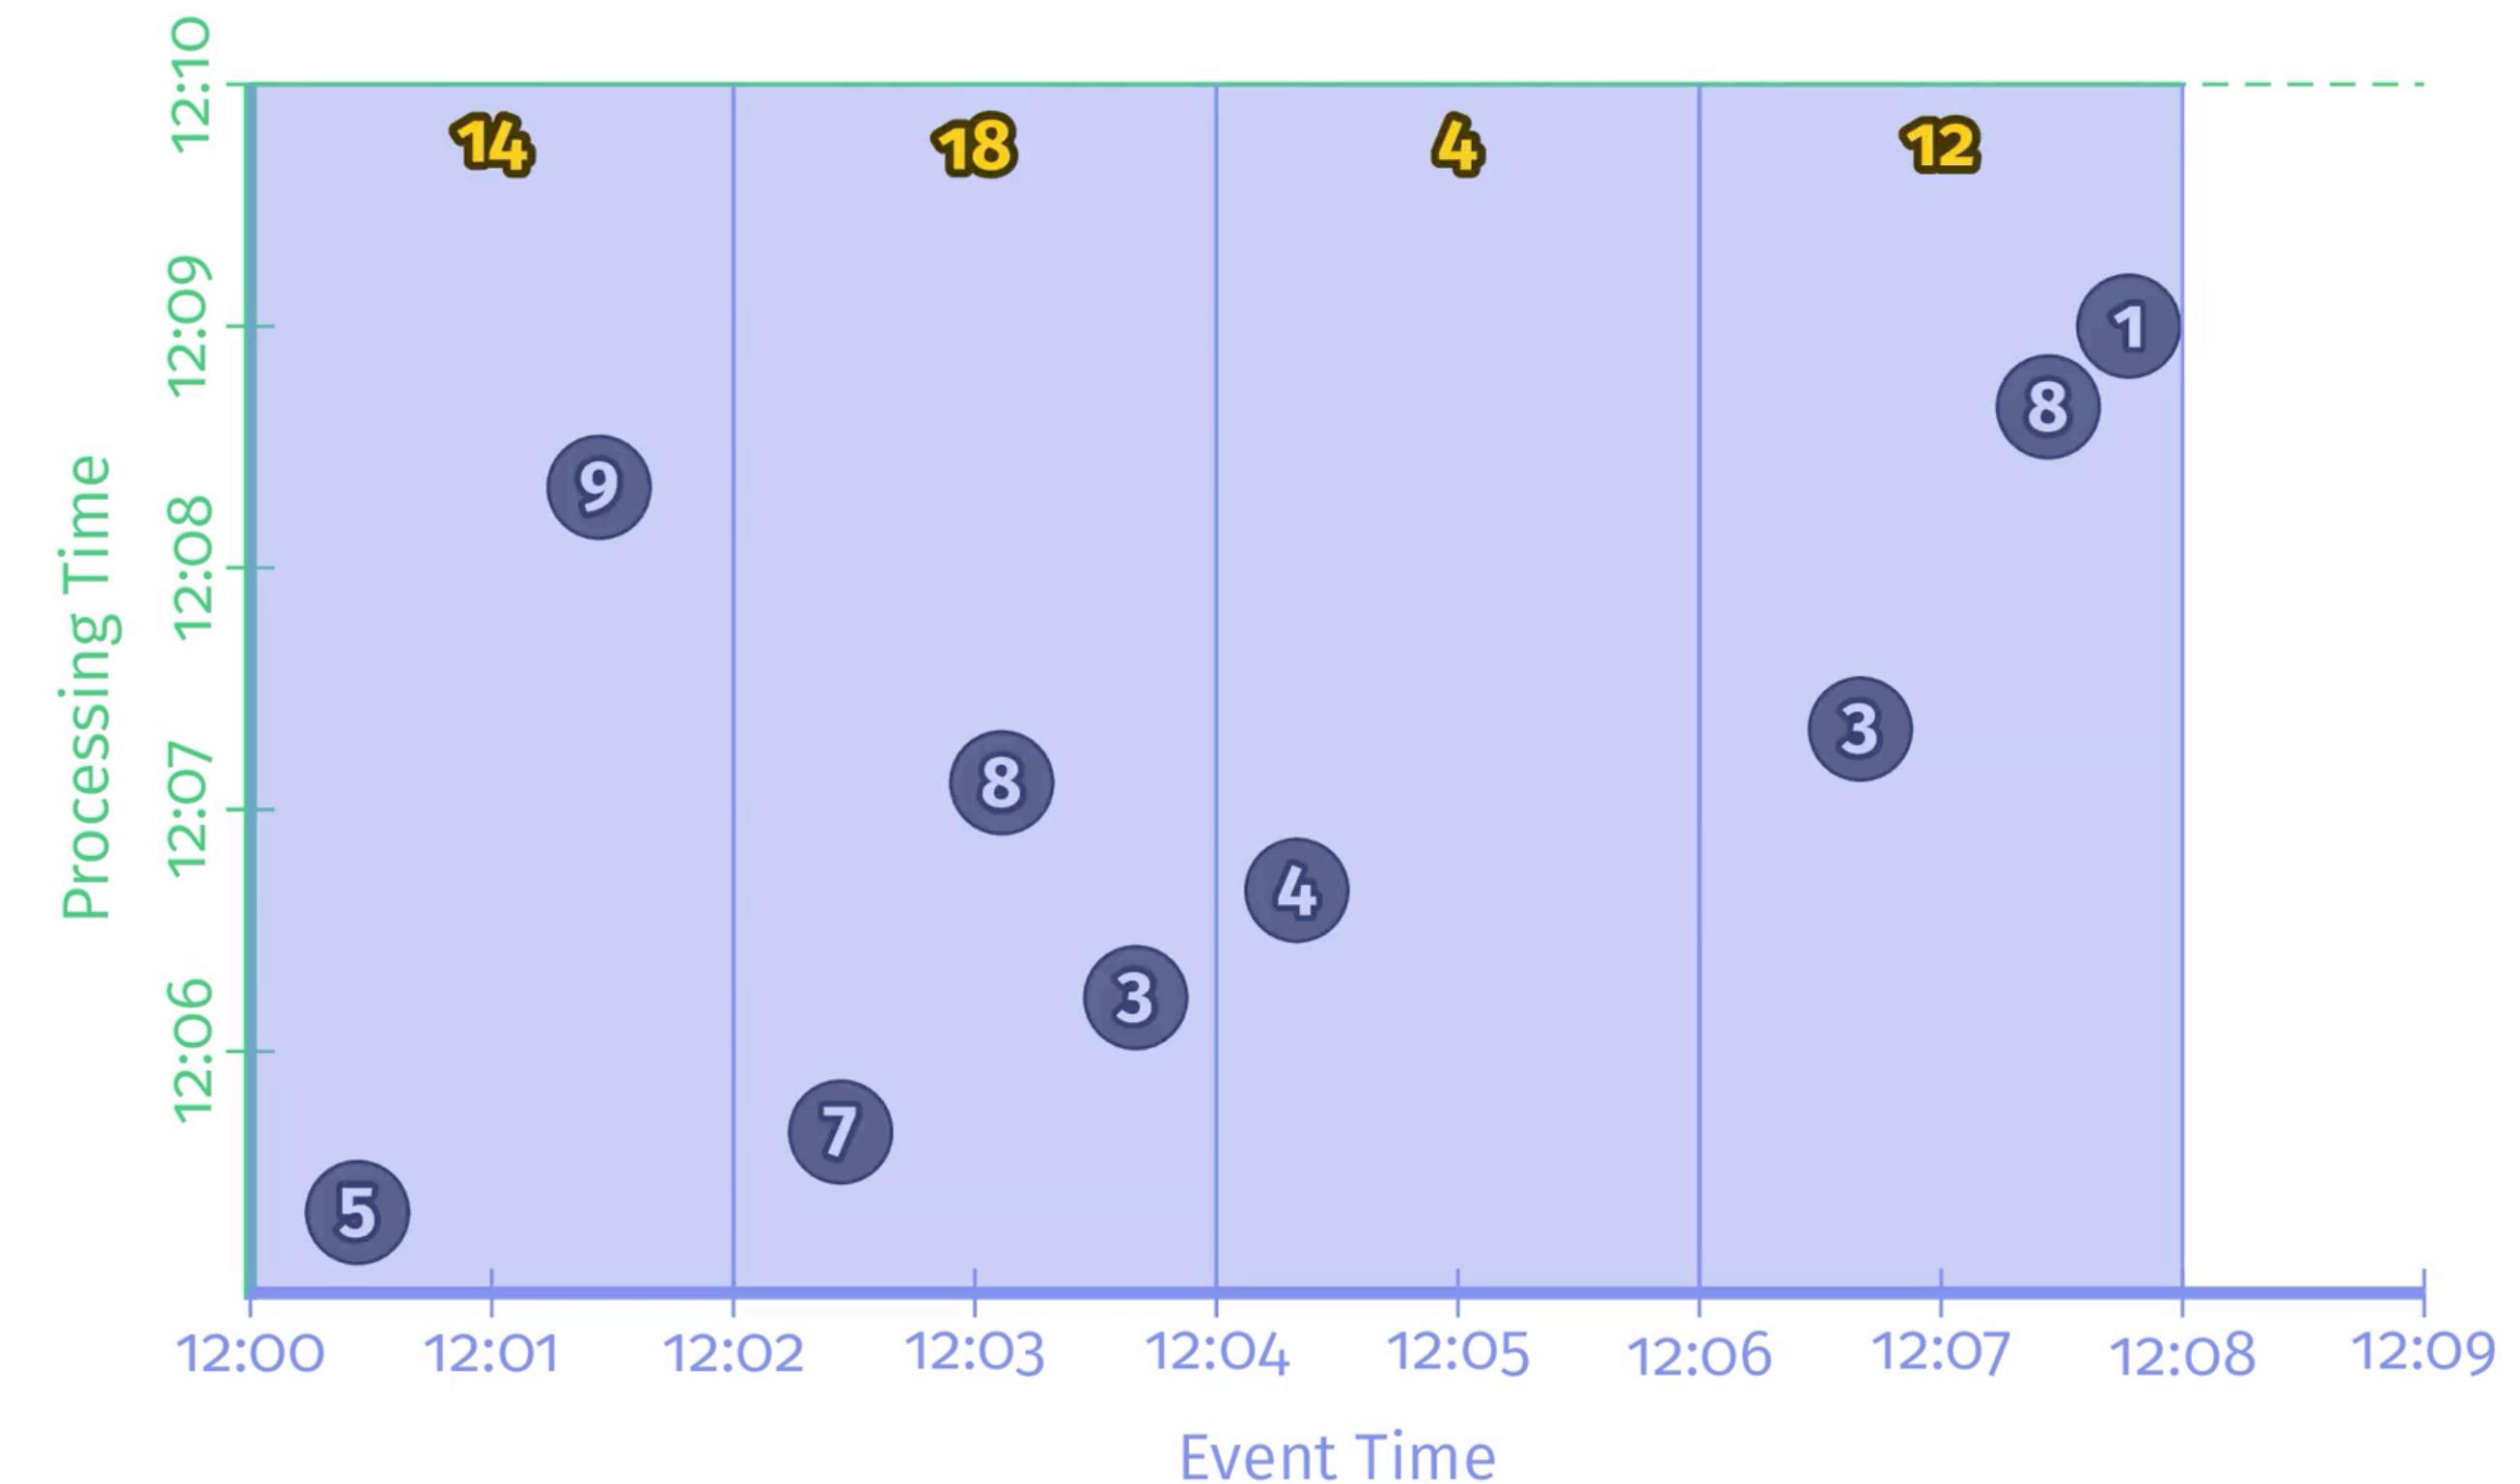
Where in event time?

Windowing



Where in event time?

```
PCollection<KV<Team, Integer>> scores = input  
    .apply(Window.into(FixedWindows.of(TWO_MINUTES)))  
    .apply(Sum.integersPerKey());
```

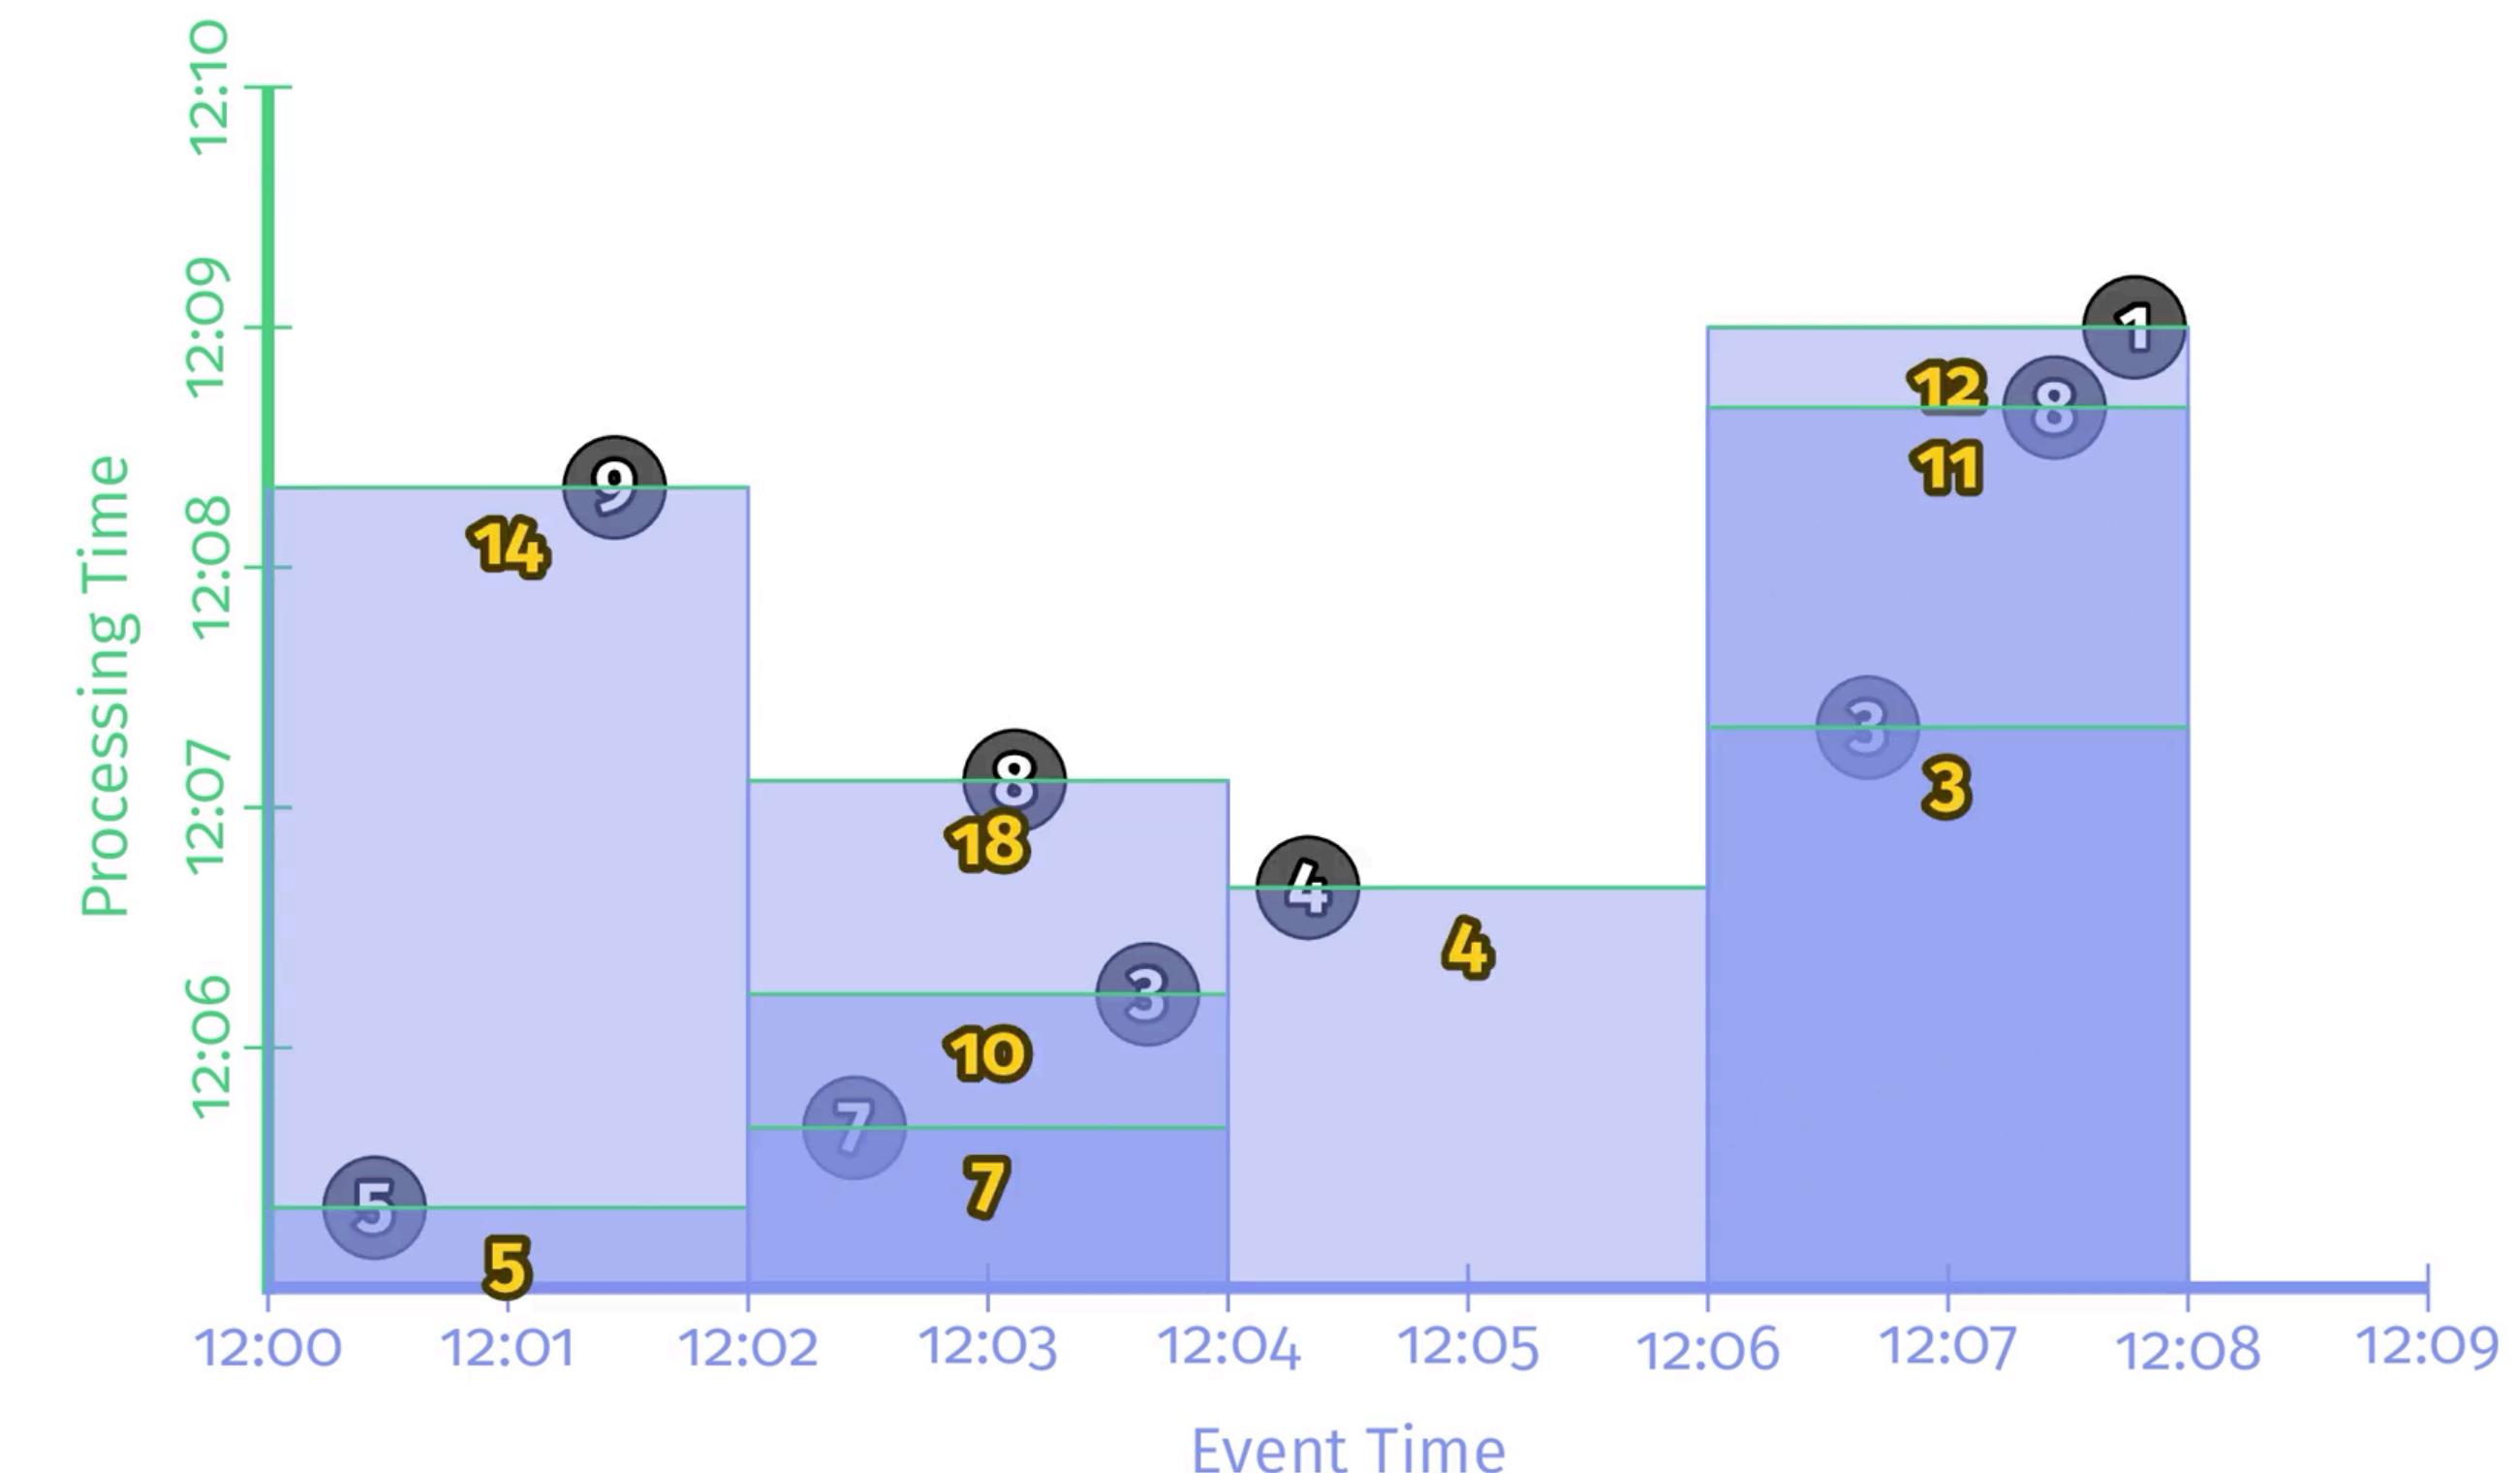


When in processing time?

Repeated Update Triggers

```
PCollection<KV<Team, Integer>> scores = input
| apply(Window.into(FixedWindows.of(TWO_MINUTES))
| | | | | .triggering(Repeatedly(AfterCount(1))));
```

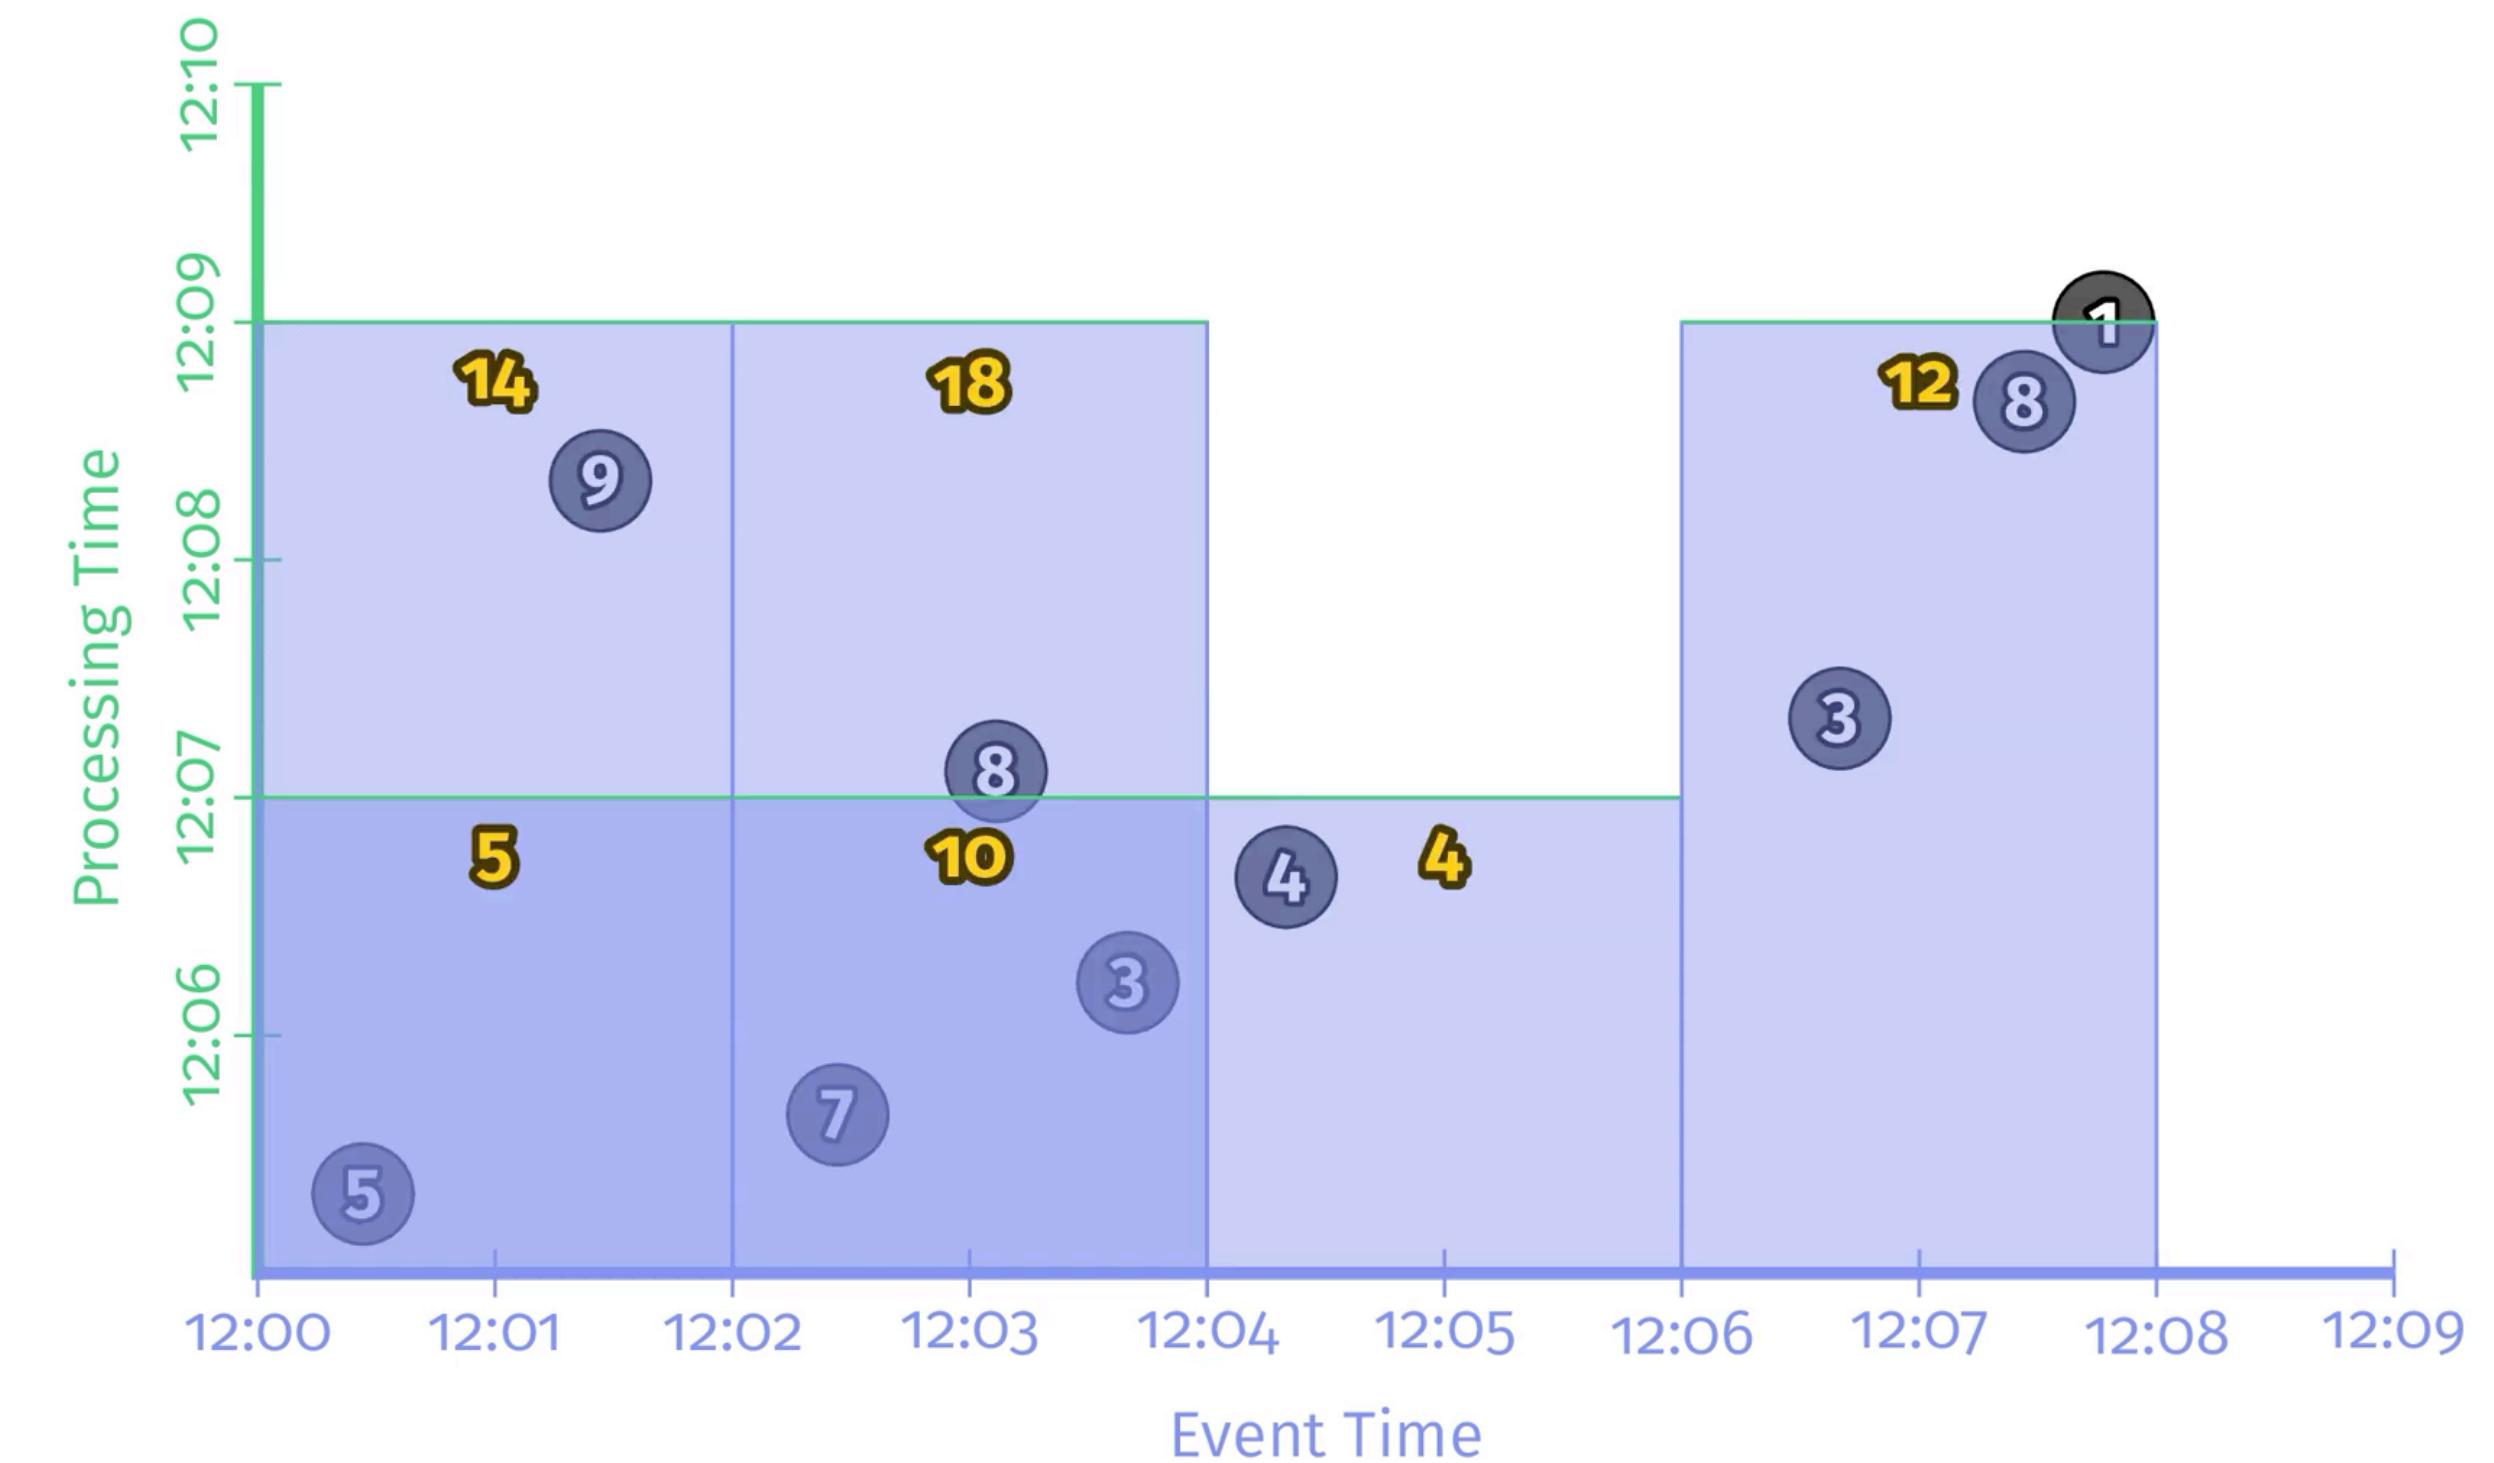
.apply(Sum.integersPerKey());



When in processing time?

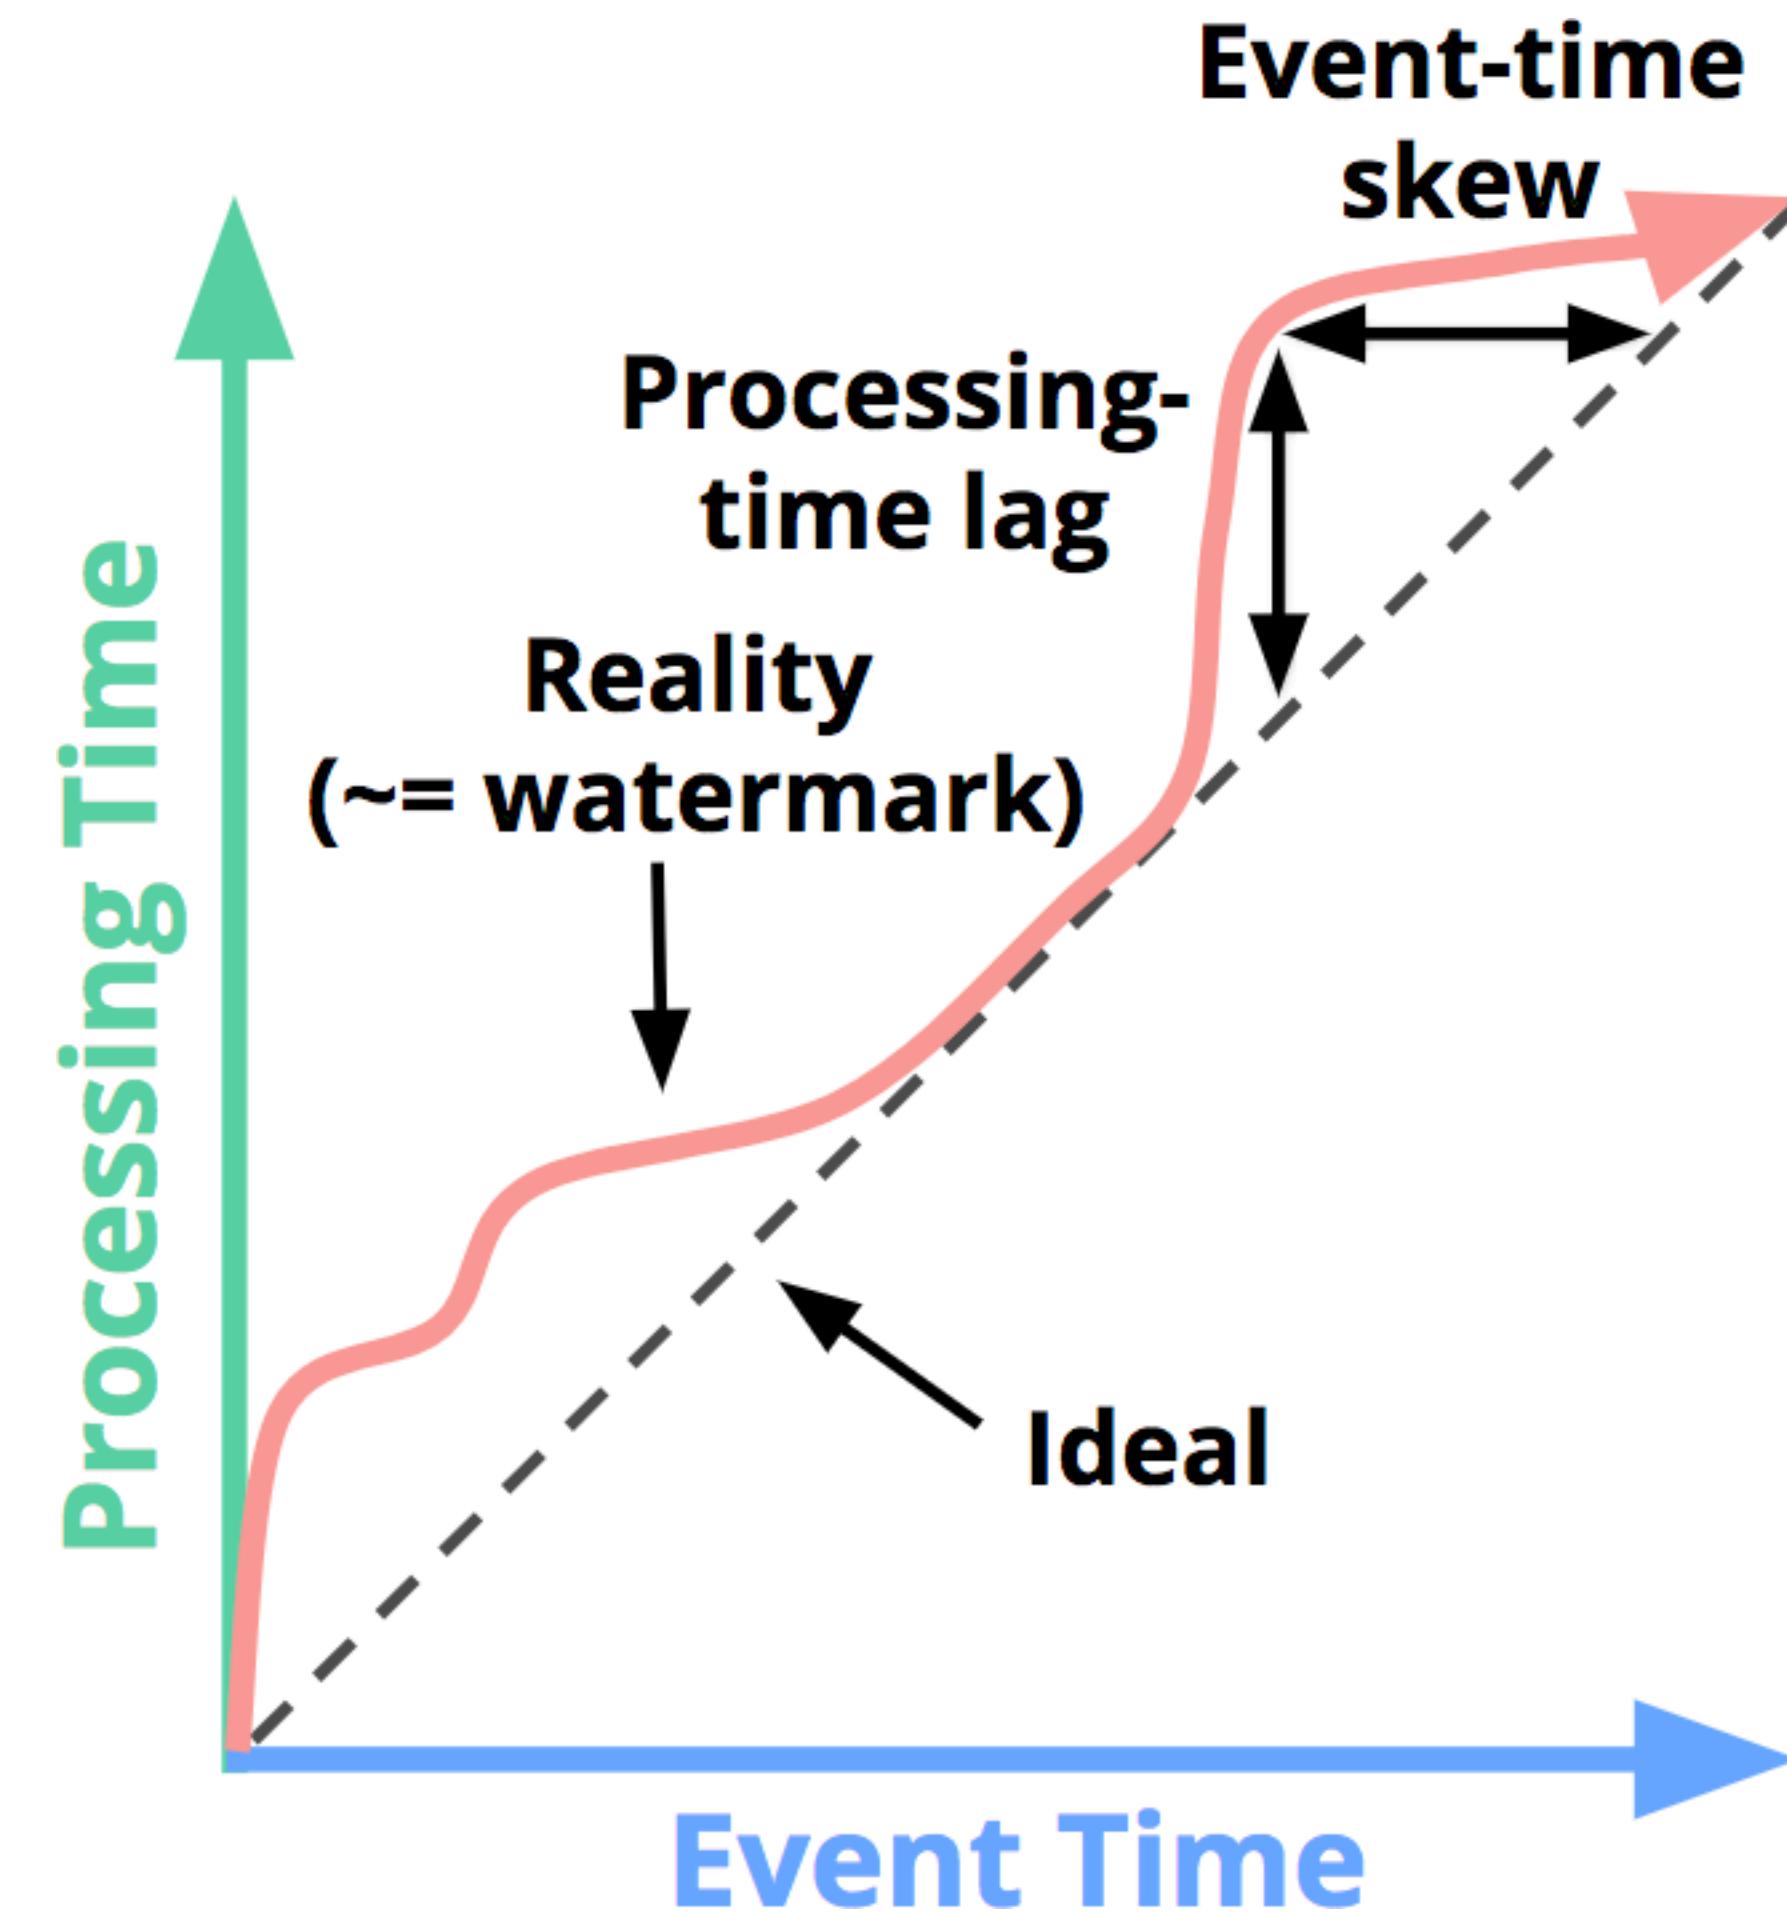
Repeated Update Triggers

```
PCollection<KV<Team, Integer>> scores = input
    .apply(Window.into(FixedWindows.of(TWO_MINUTES))
        .triggering(Repeatedly(AlignedDelay(TWO_MINUTES))))
    .apply(Sum.integersPerKey());
```



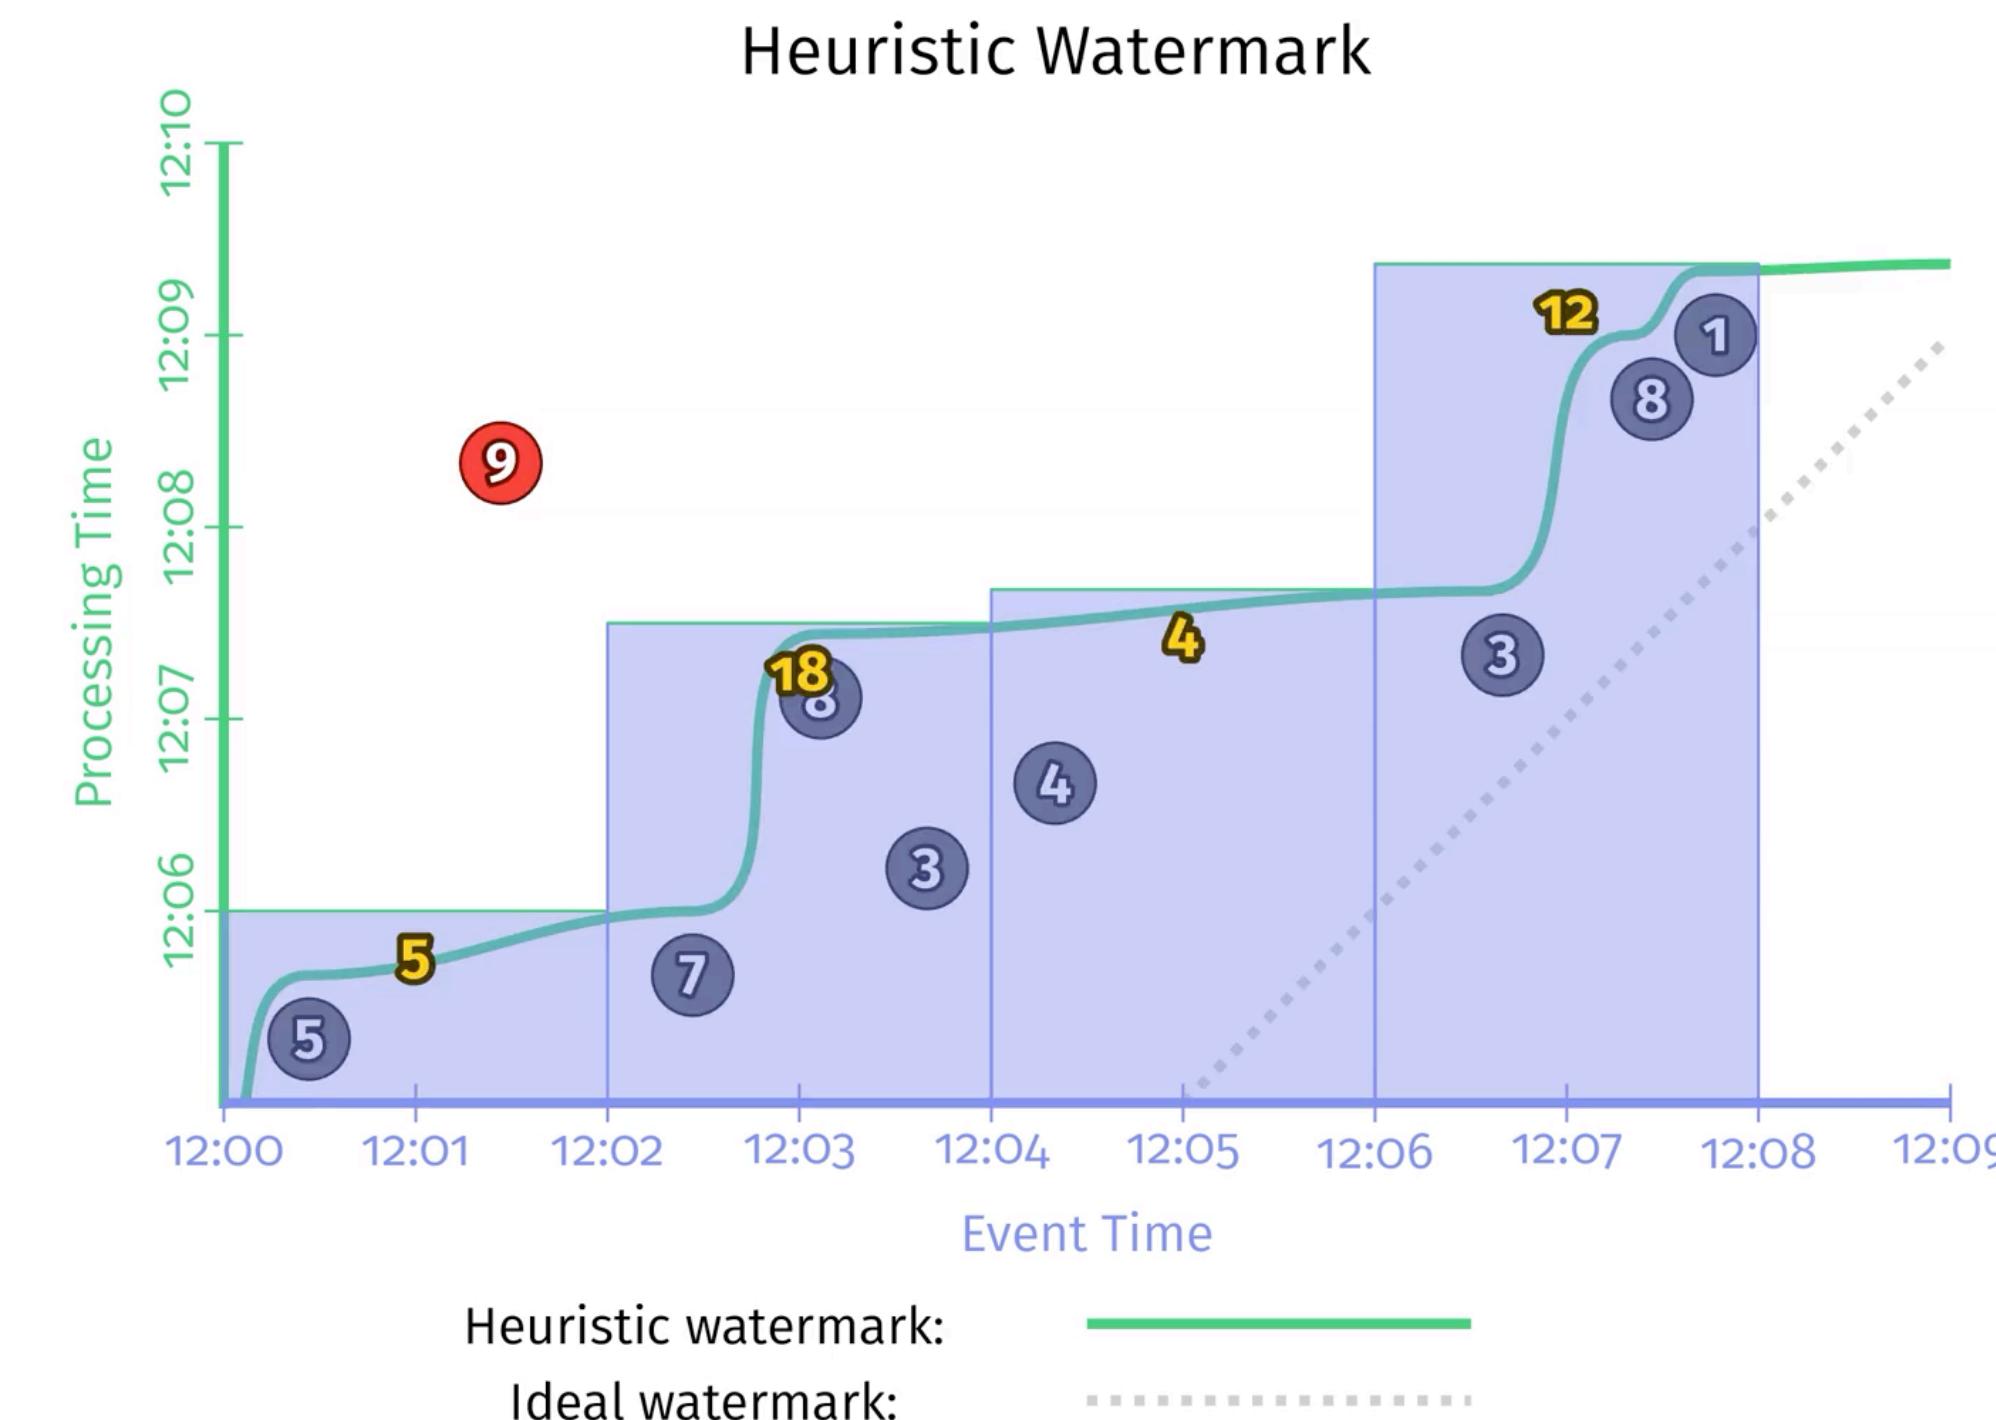
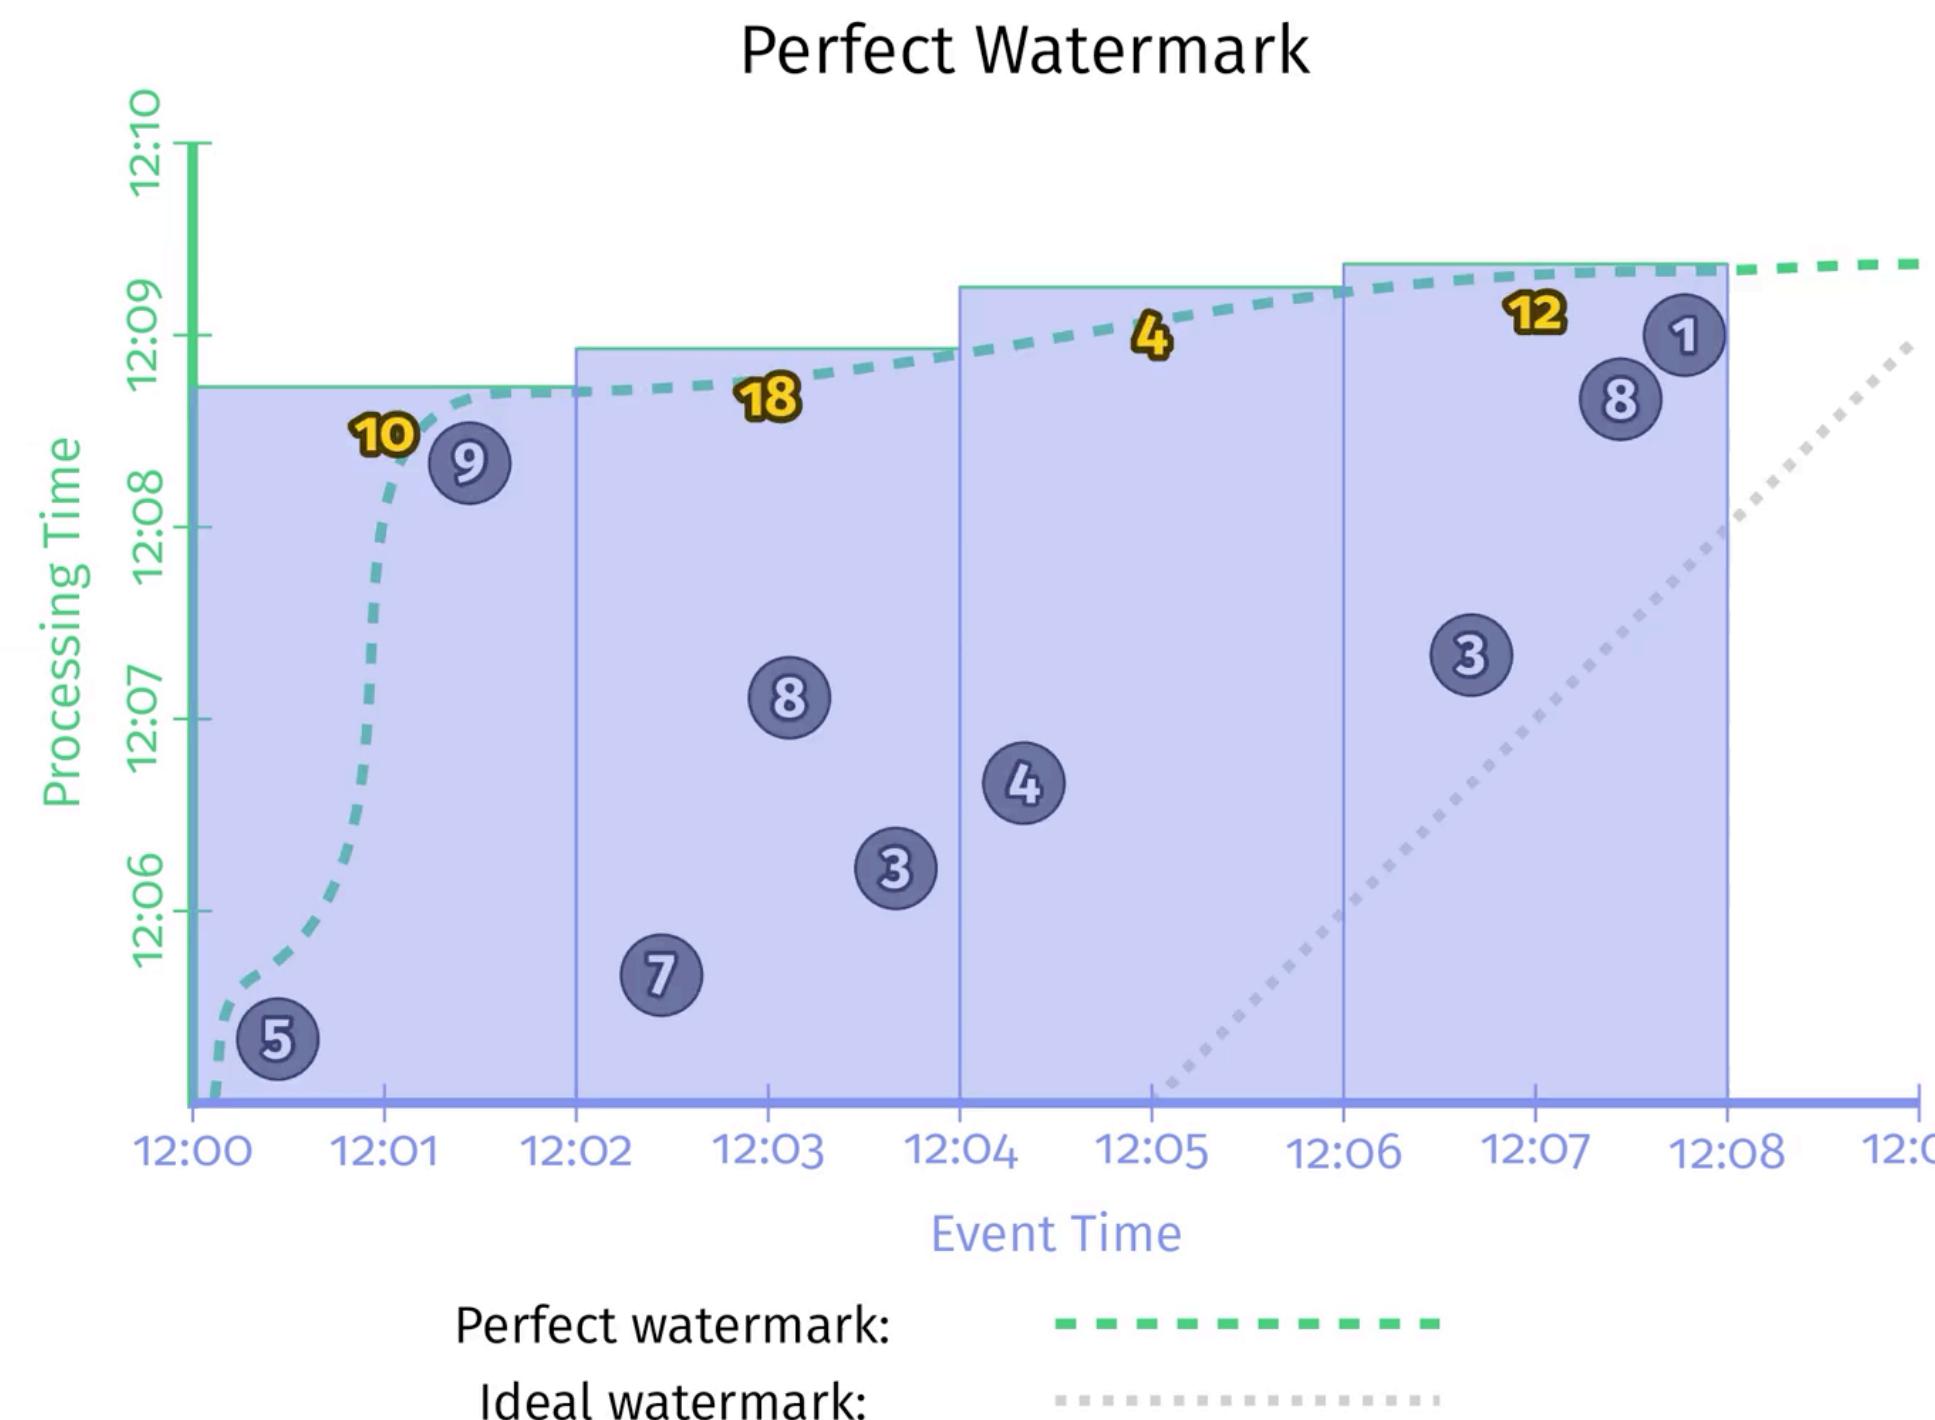
When in processing time?

Watermarks – Completeness Triggers



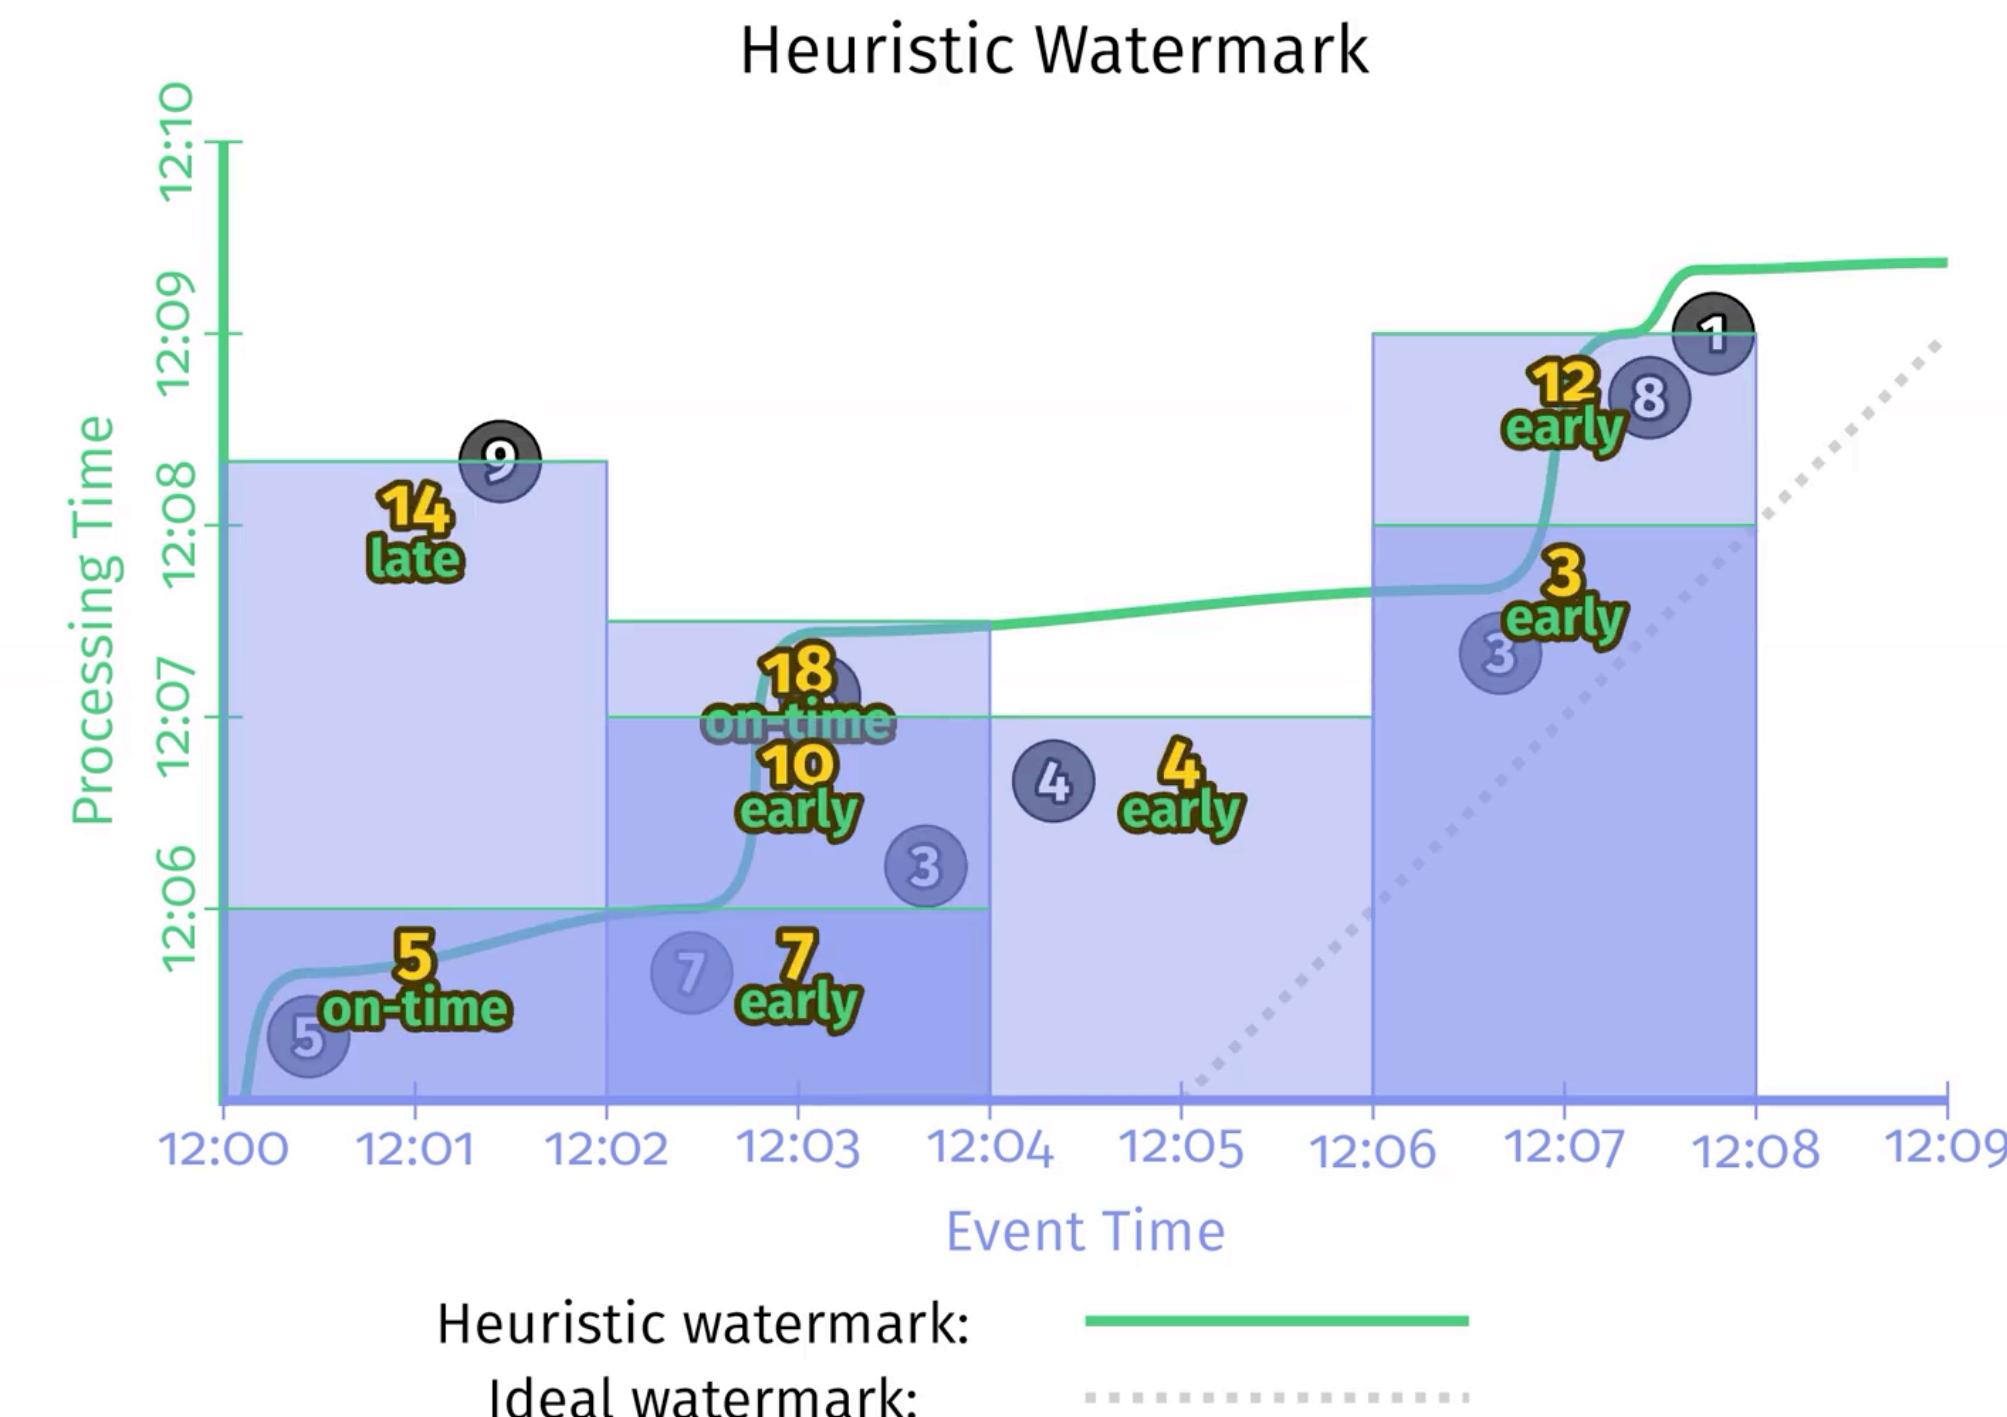
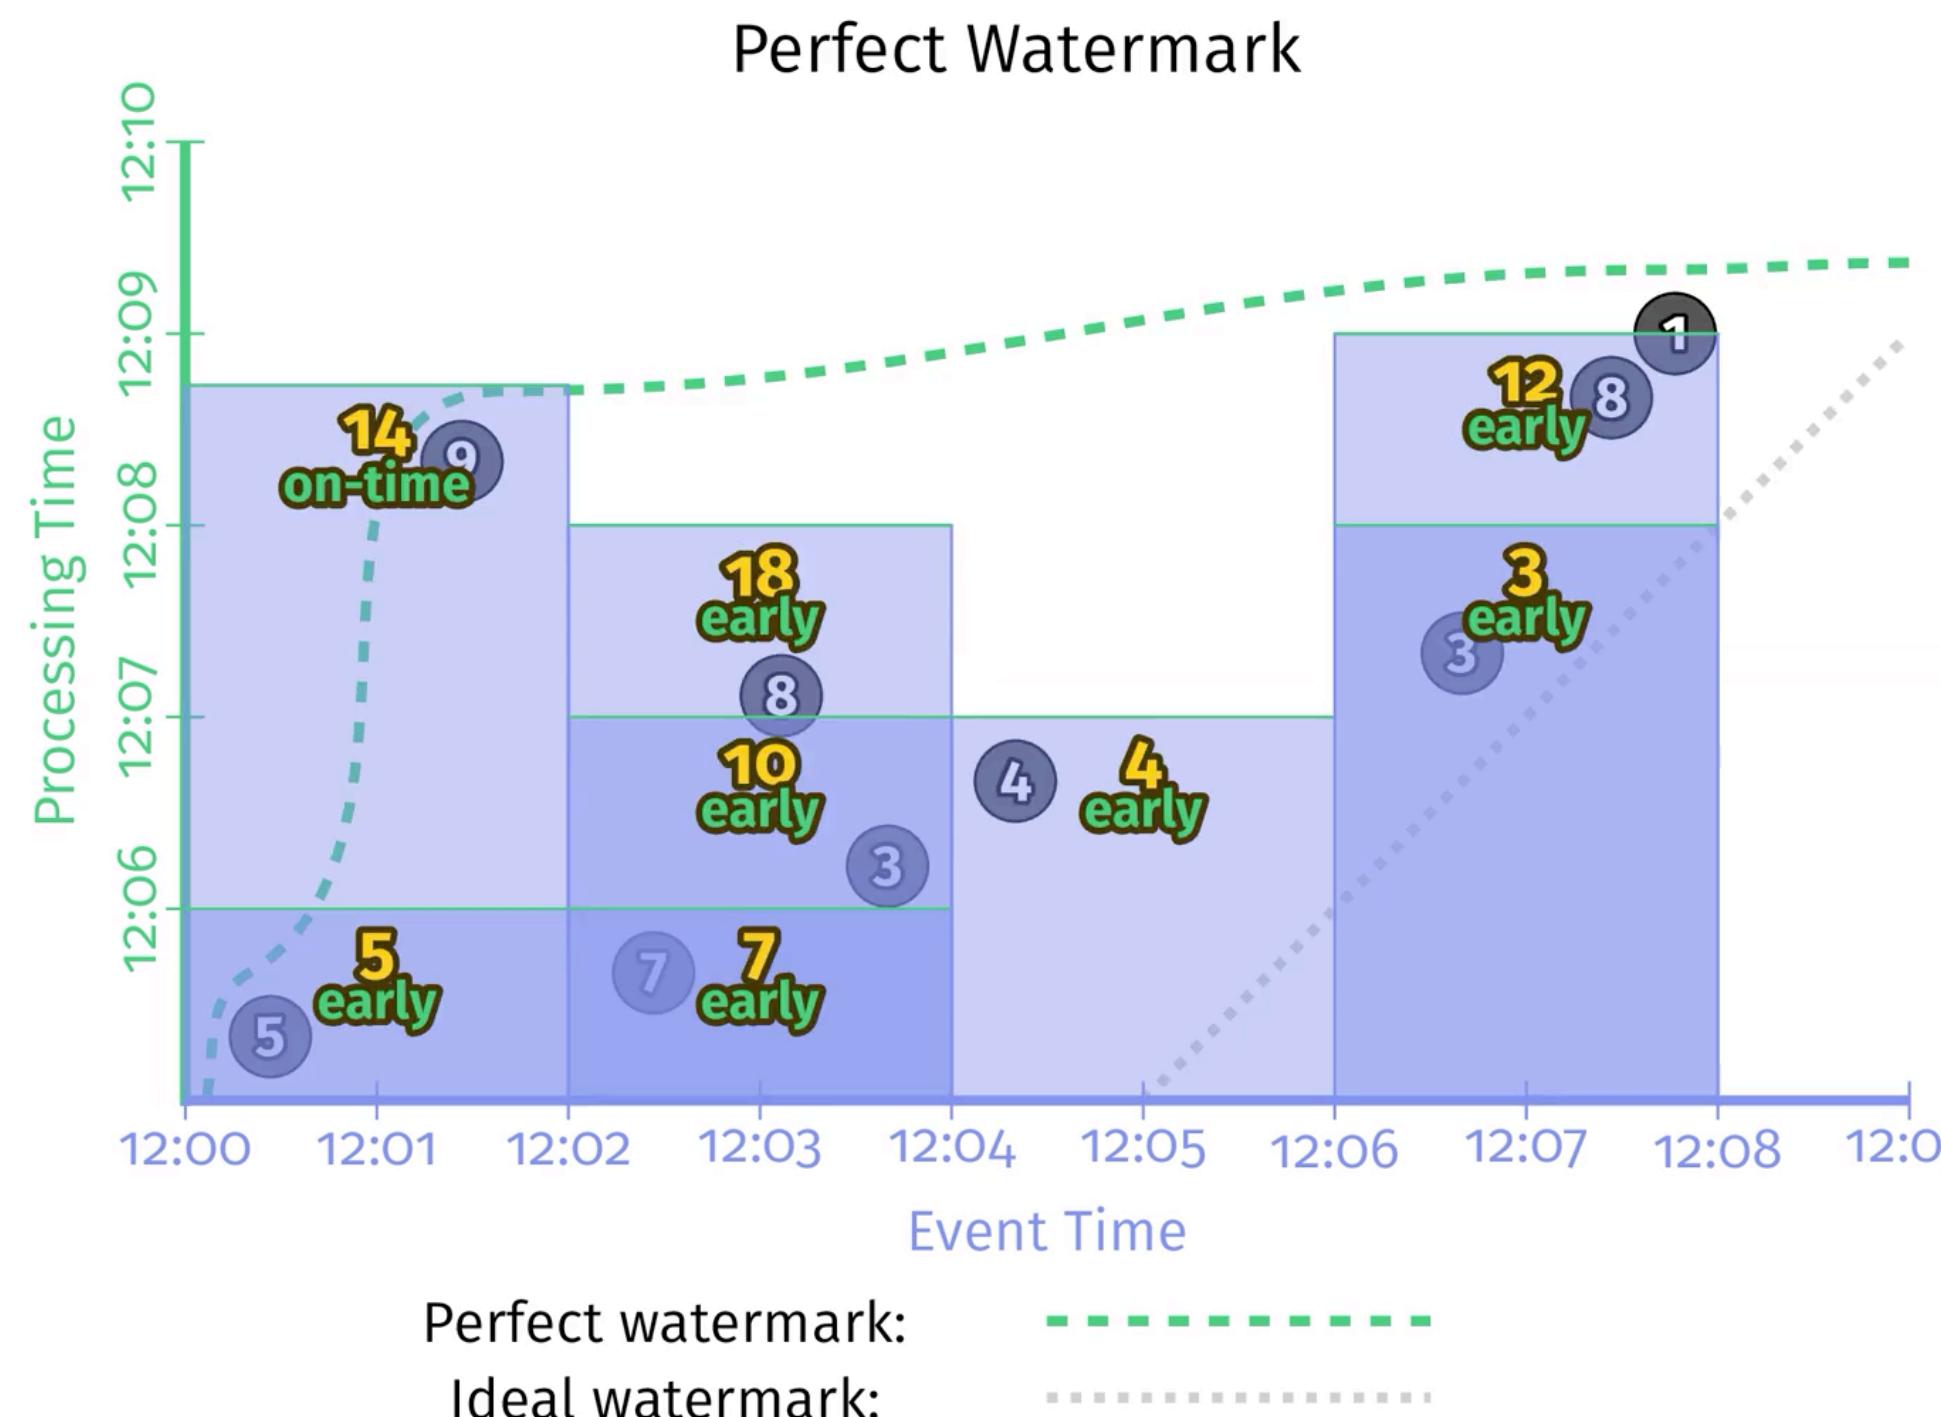
When in processing time?

Watermarks – Completeness Triggers



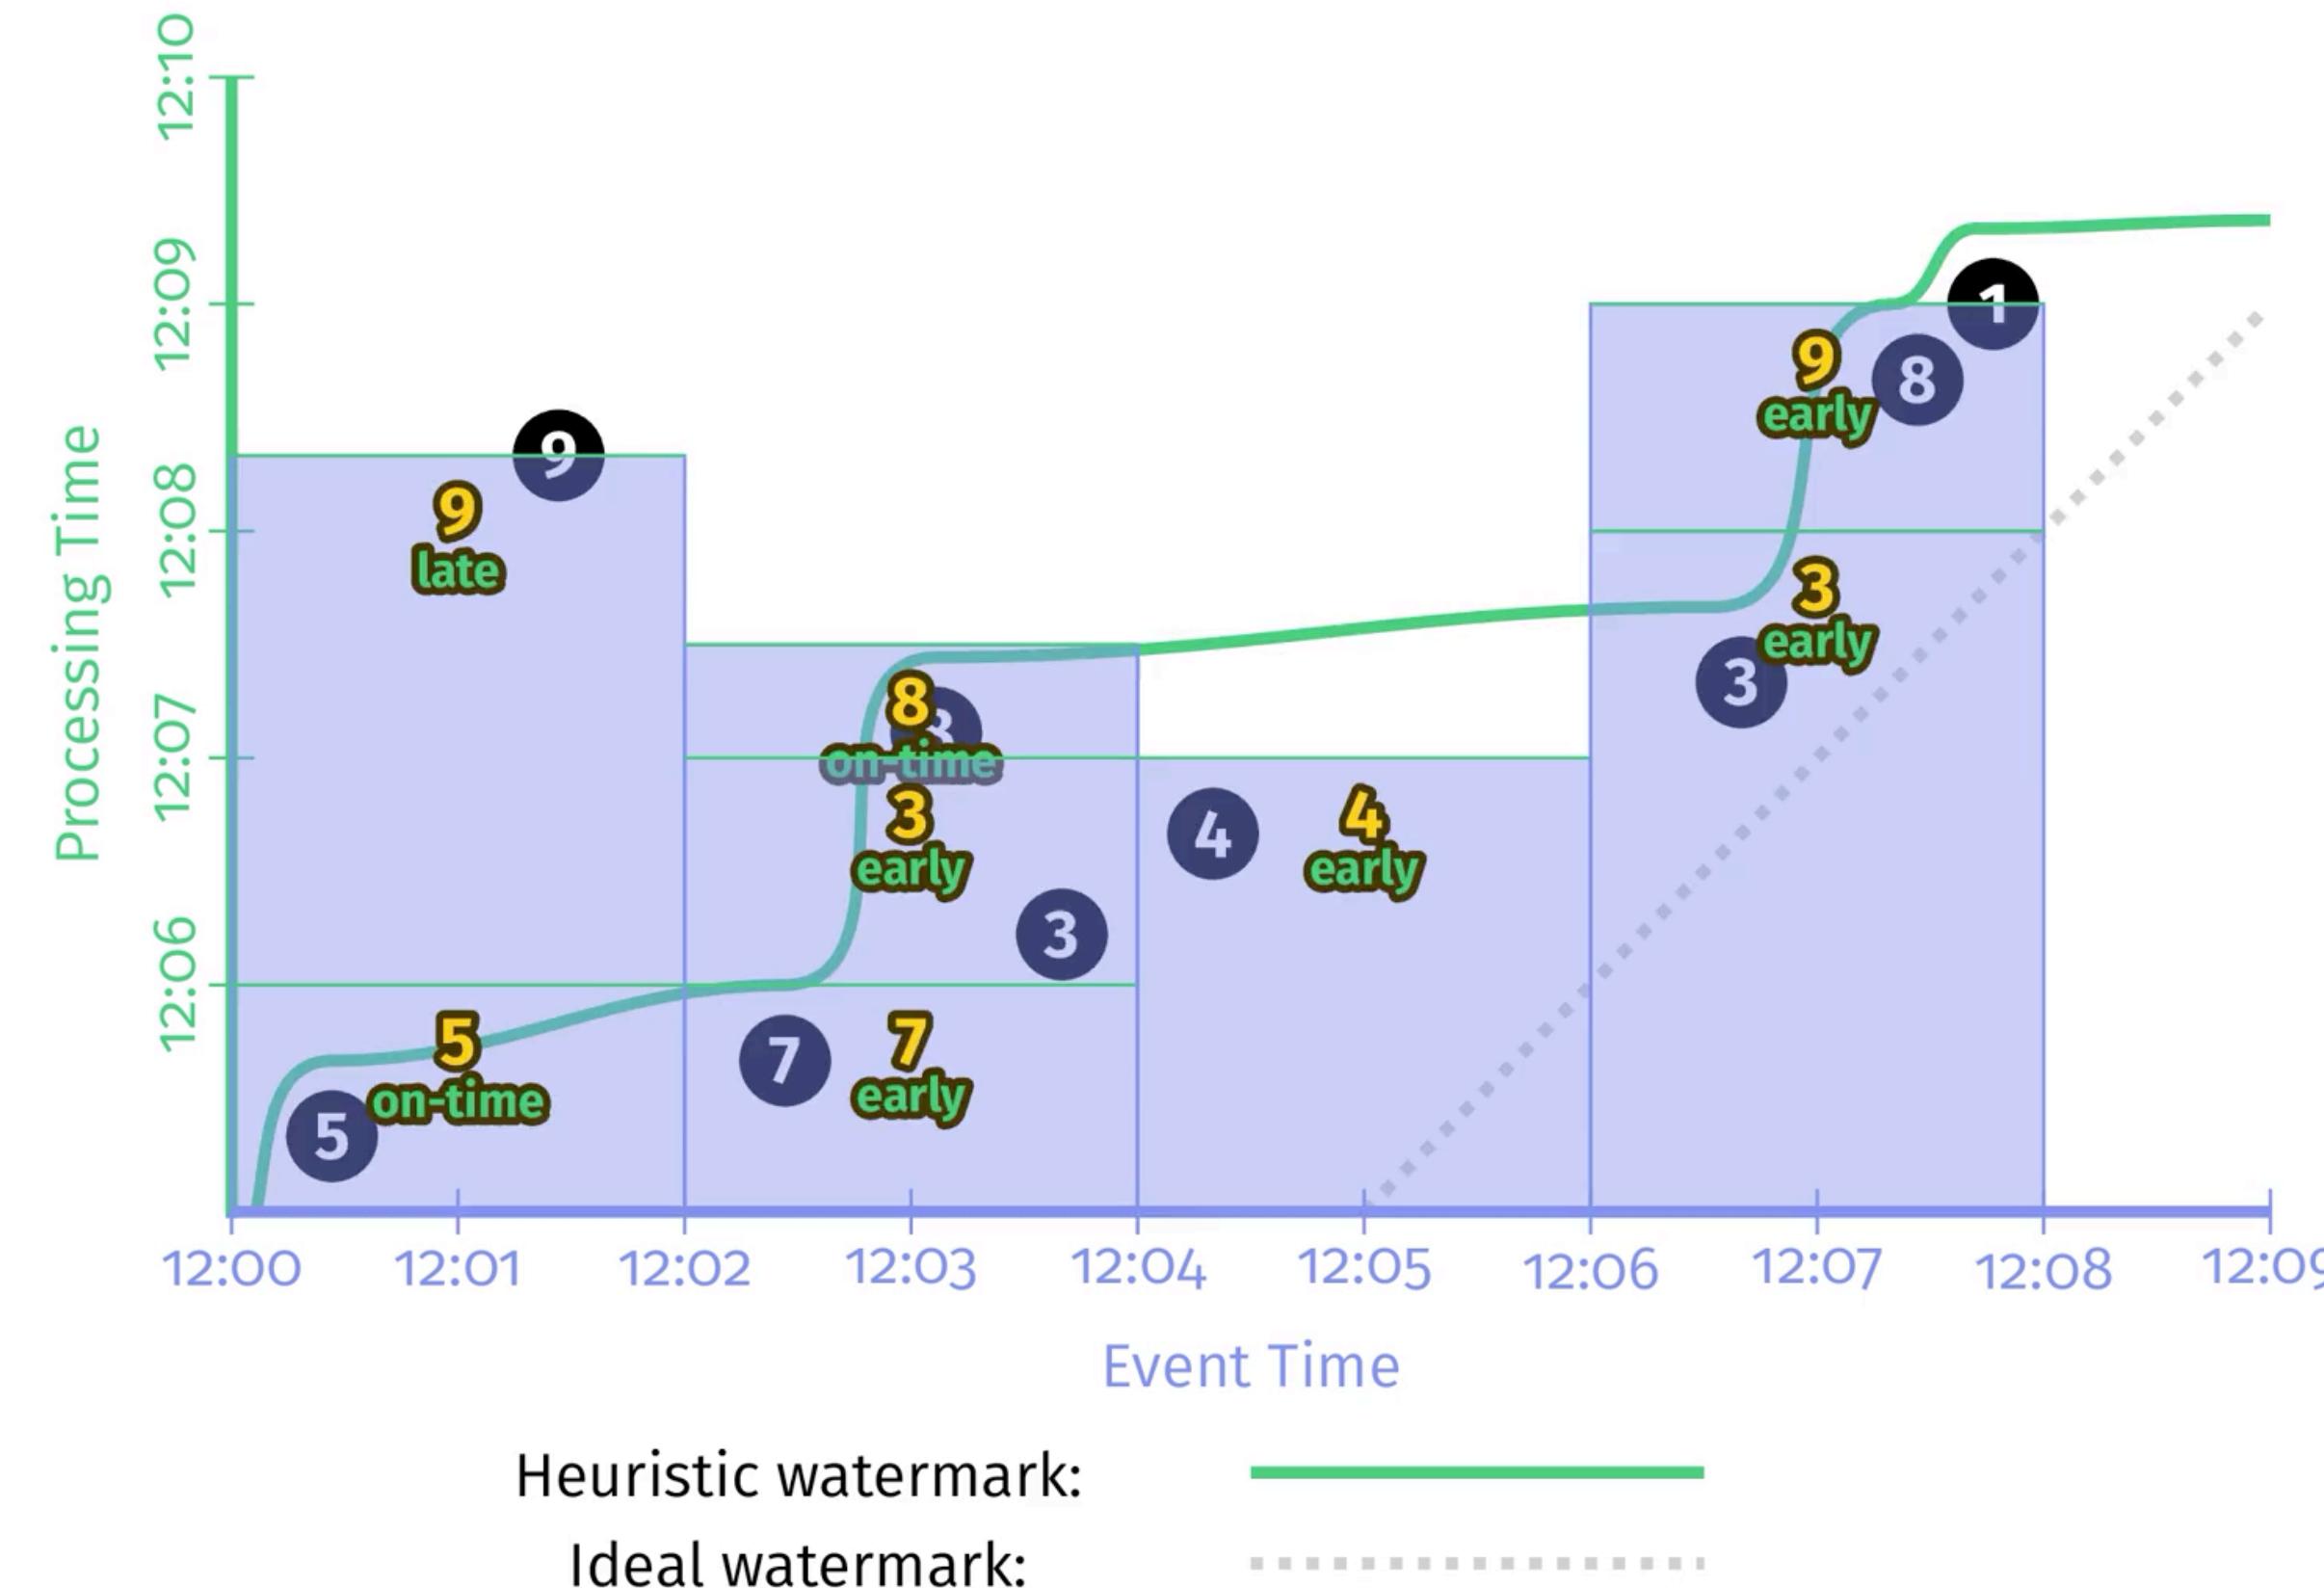
When in processing time?

Early/On-time/Late Trigger



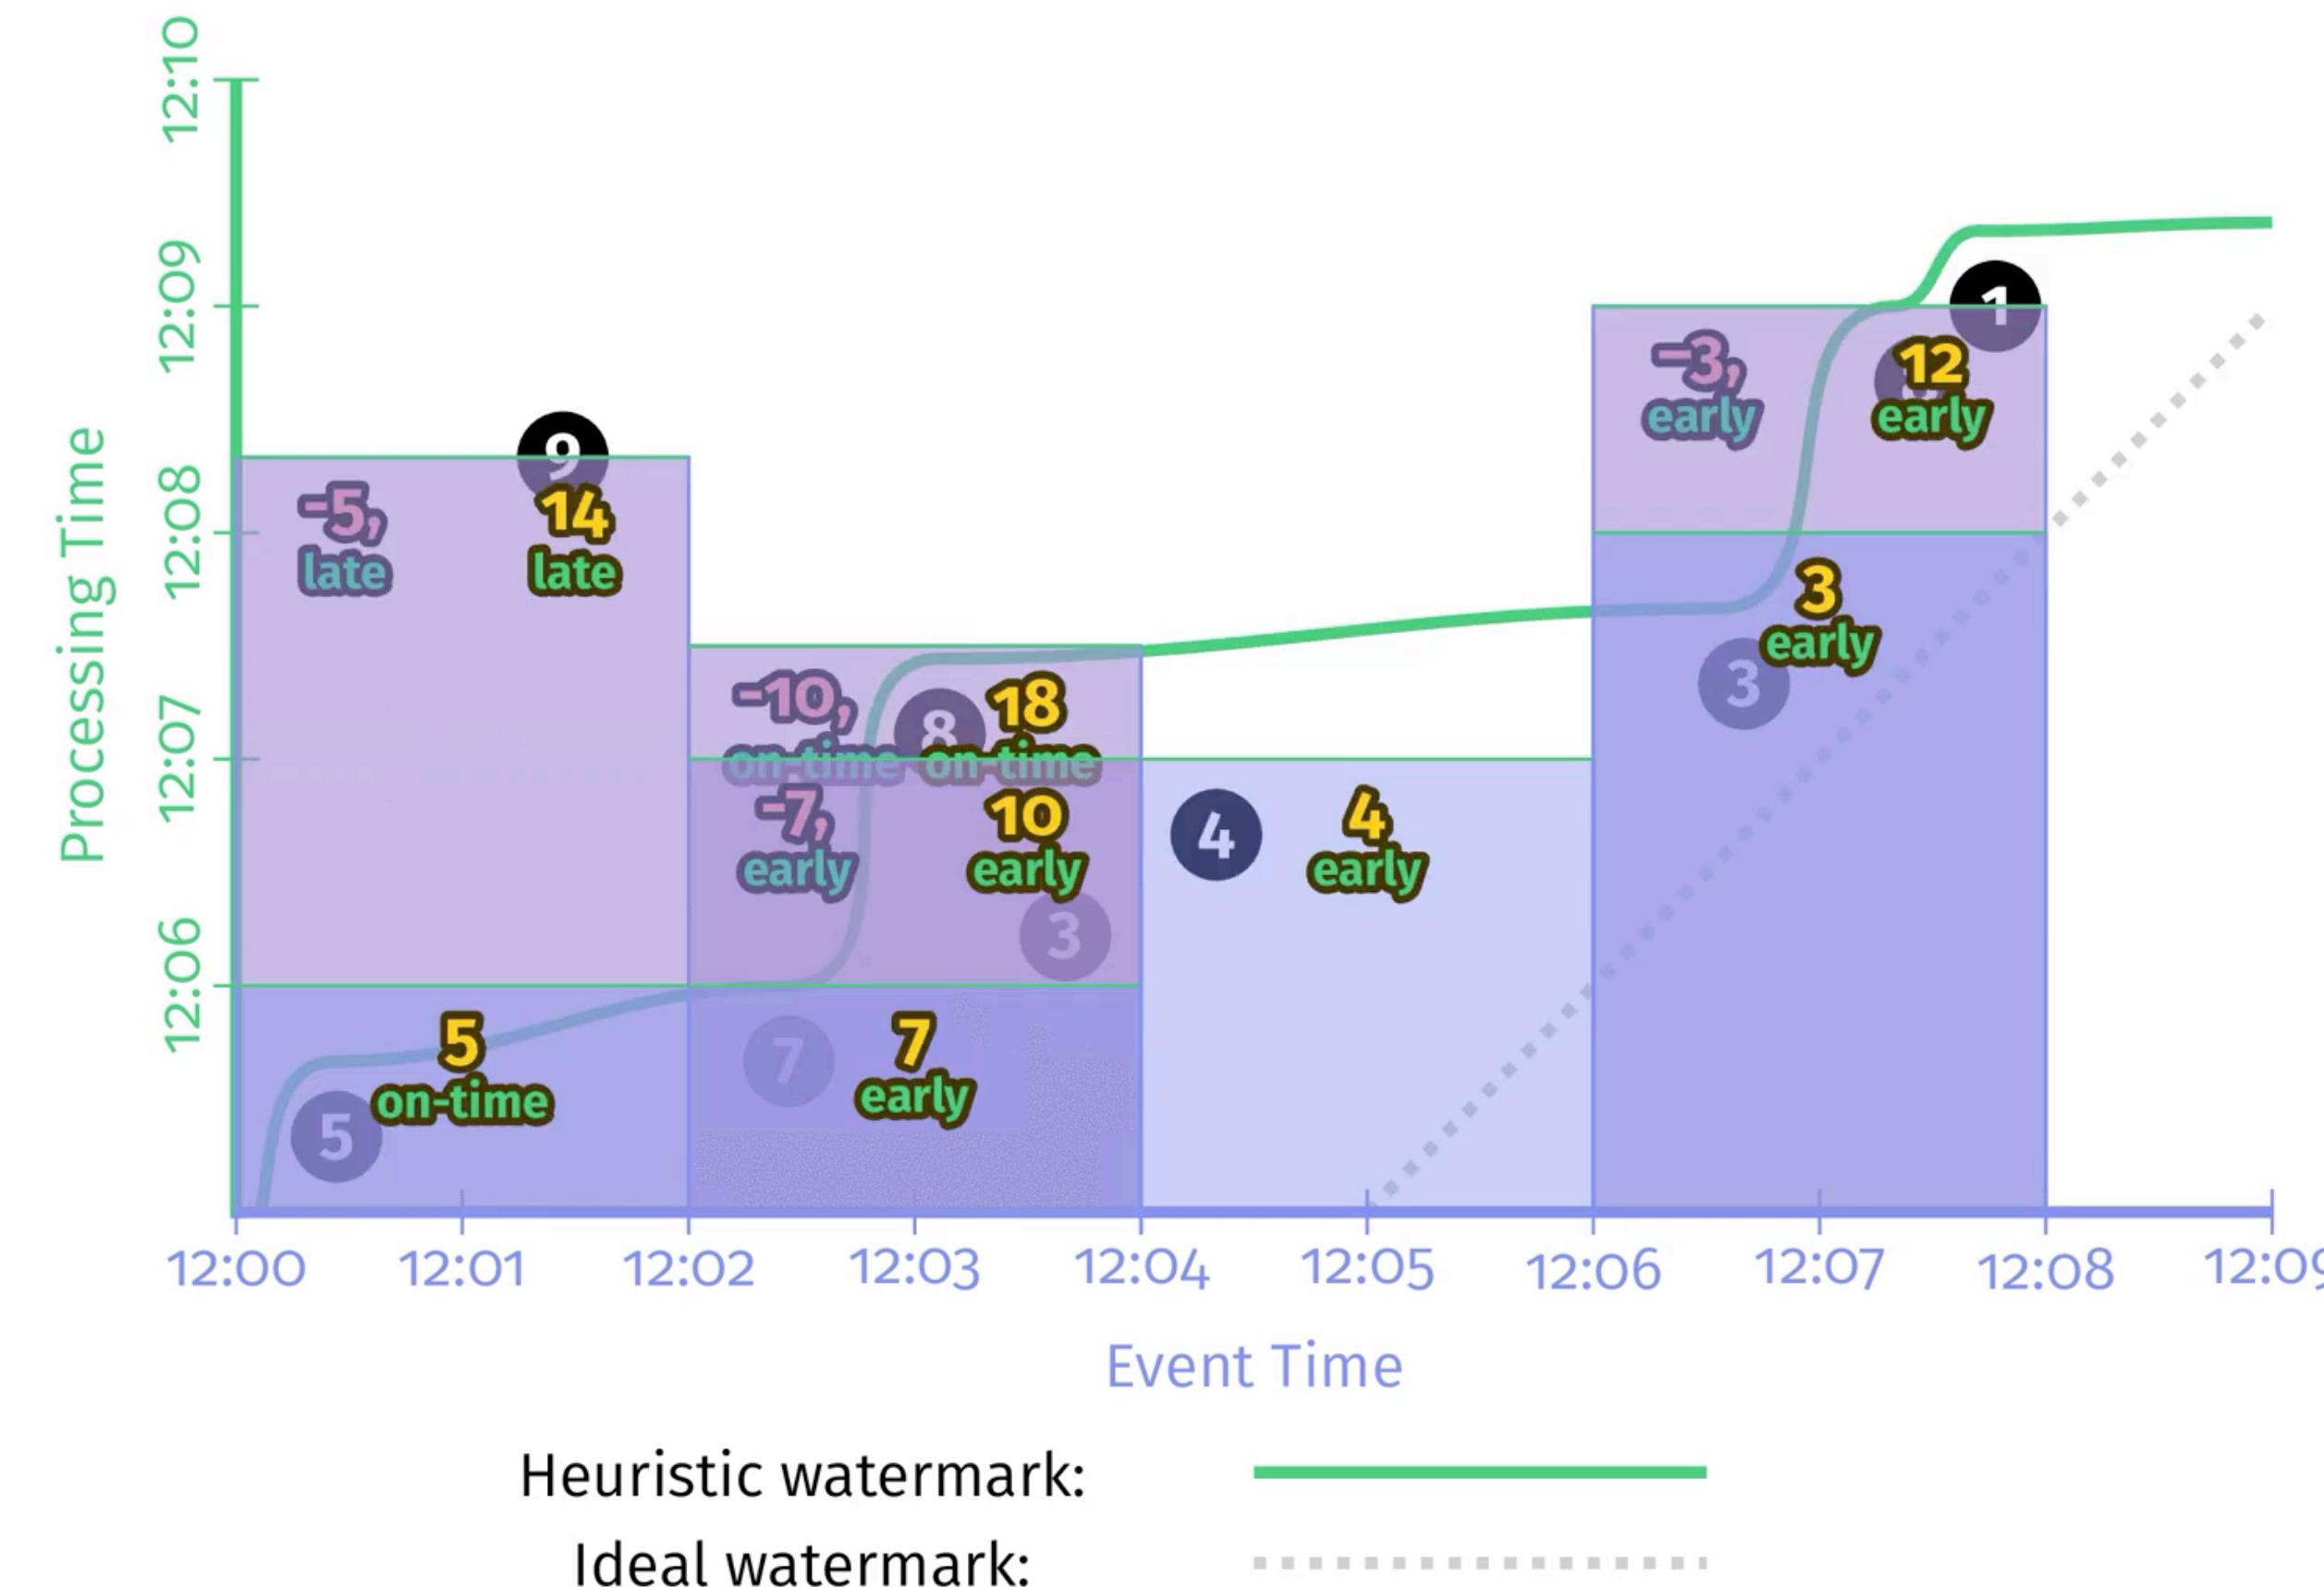
How do refinements relate?

Discarding



How do refinements relate?

Accumulating & Retracting



The Dataflow Model Summary

What are you computing?

- Pipeline Code

Where in event time?

- Windowing

When in processing time?

- Triggers & Watermark

How do refinements relate?

- Discarding, Accumulating and Accumulating&Retracting

Beat Batch, Build Stream Dynasty

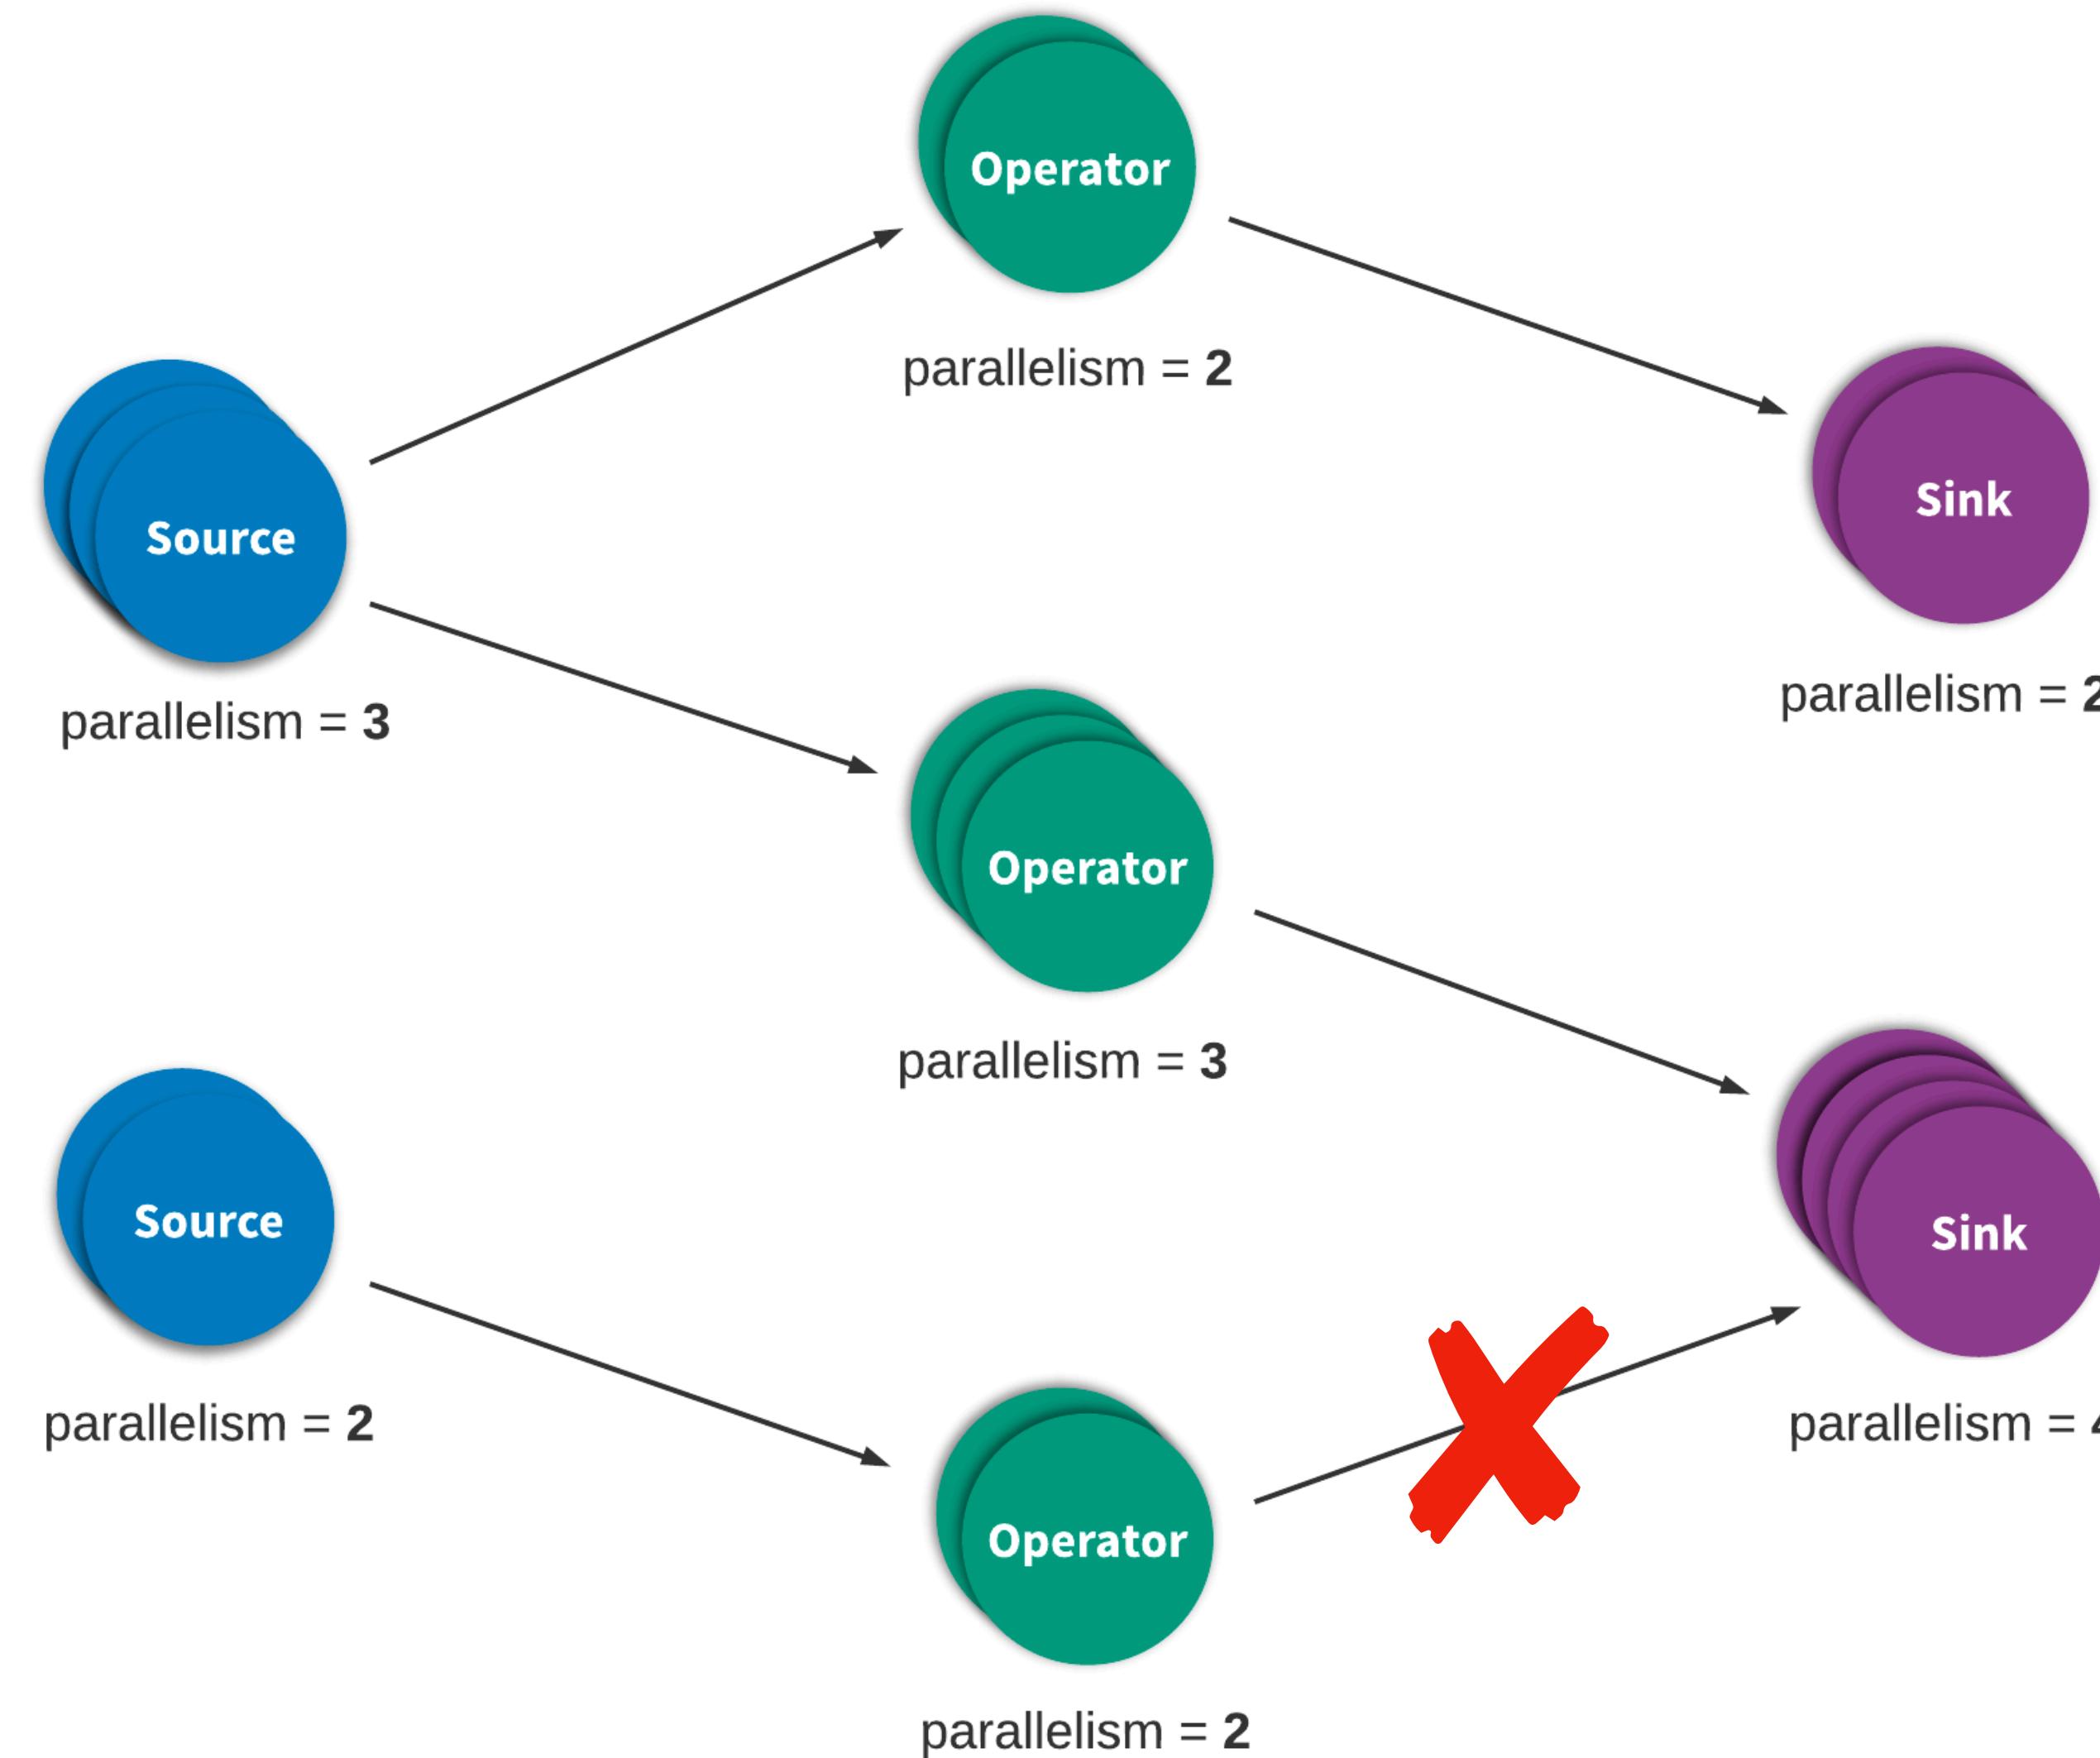
- **Tools for reasoning about time**

- The Dataflow Model: A Practical Approach to Balancing Correctness, Latency, and Cost in Massive-scale, Unbounded, Out-of-order Data Processing

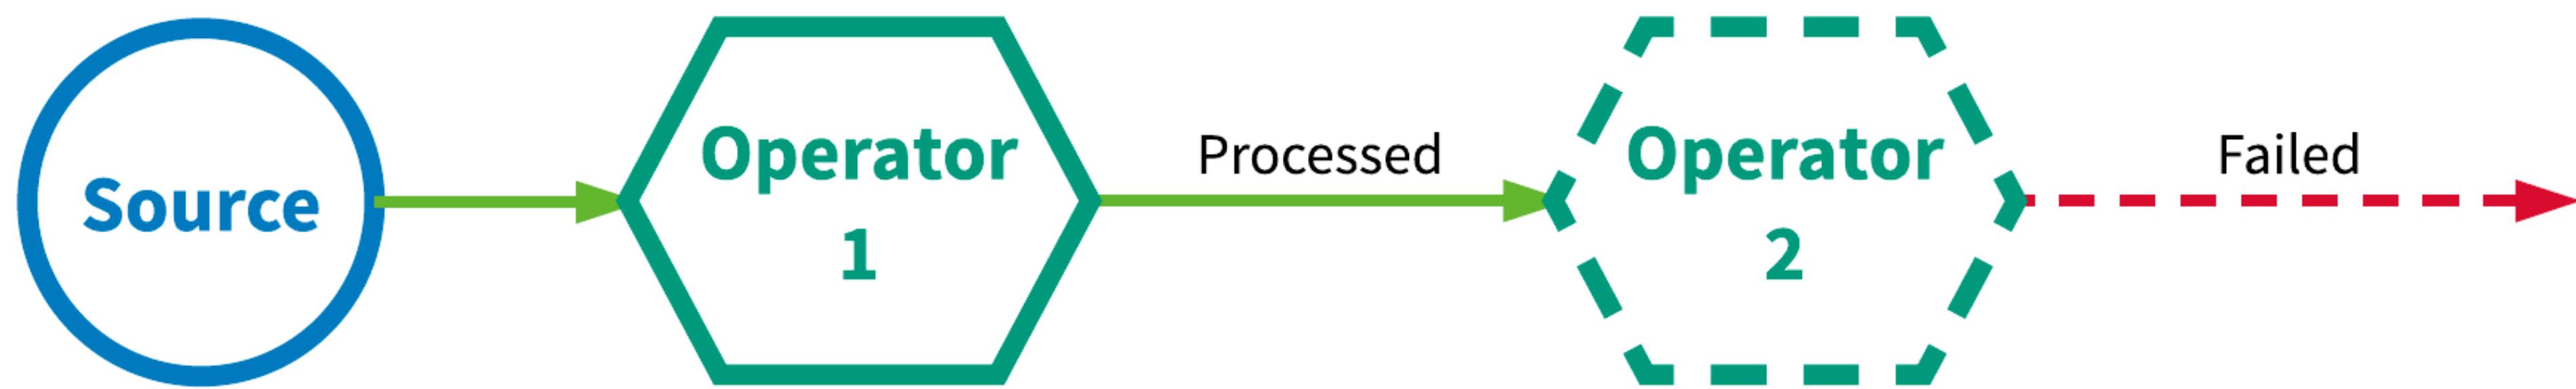
- **Correctness**

- Lightweight Asynchronous Snapshots for Distributed Dataflows

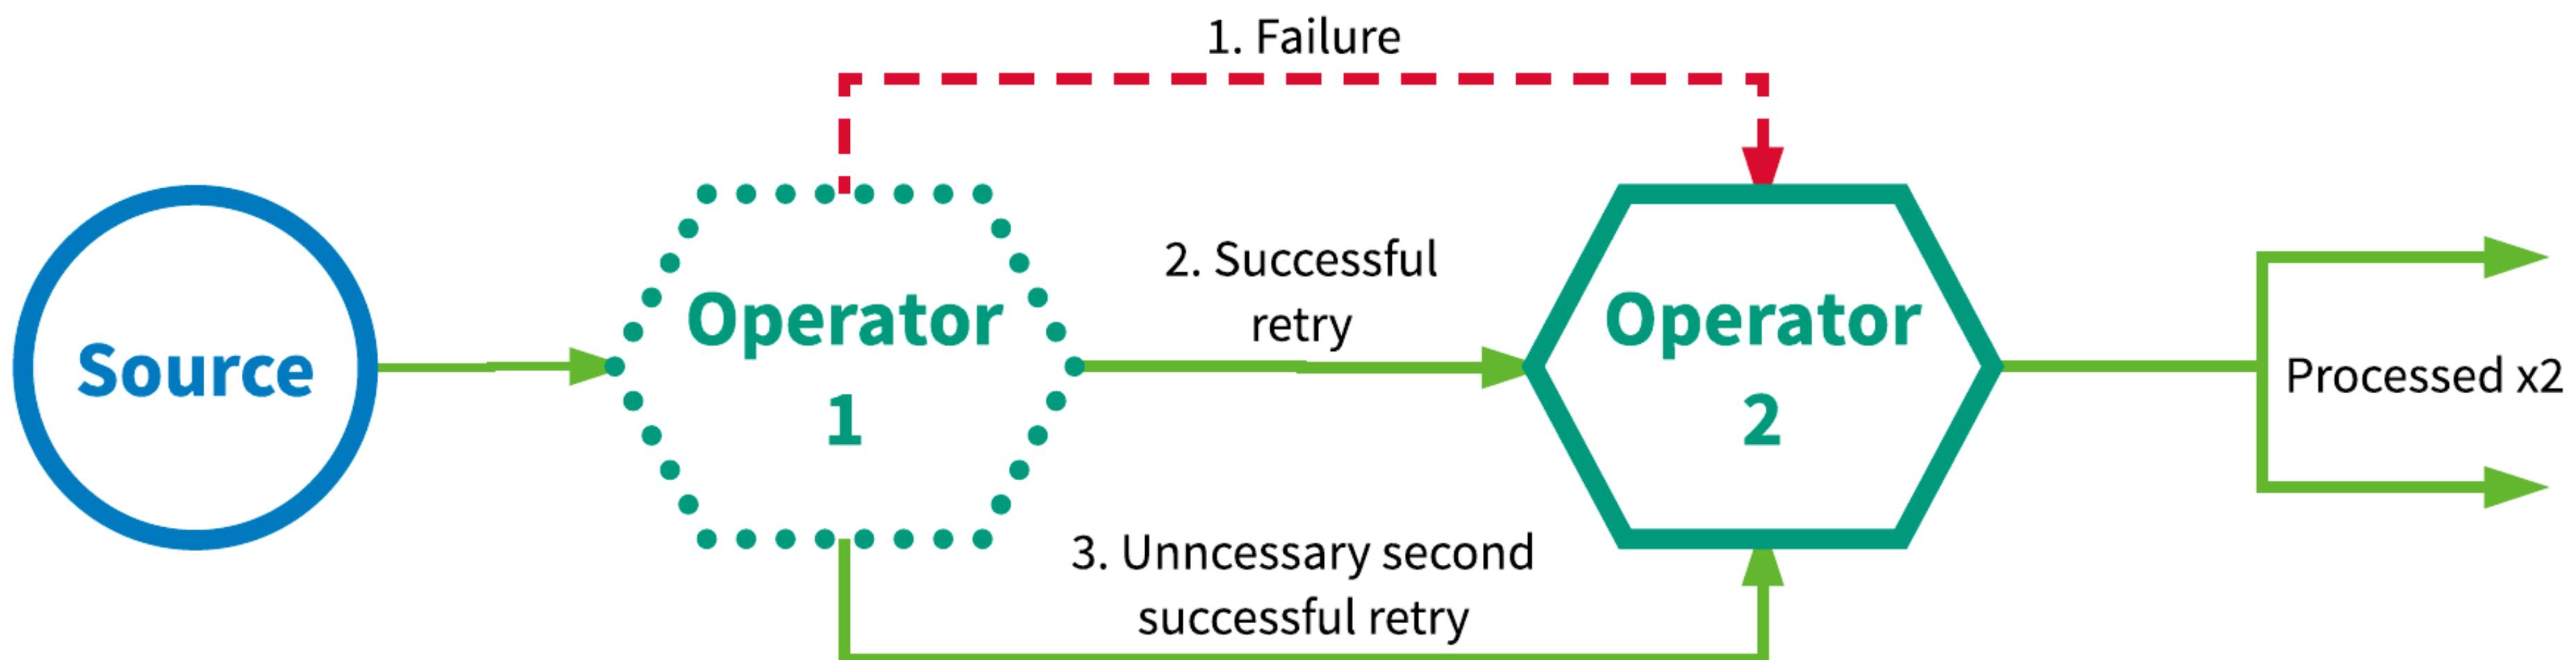
Correctness



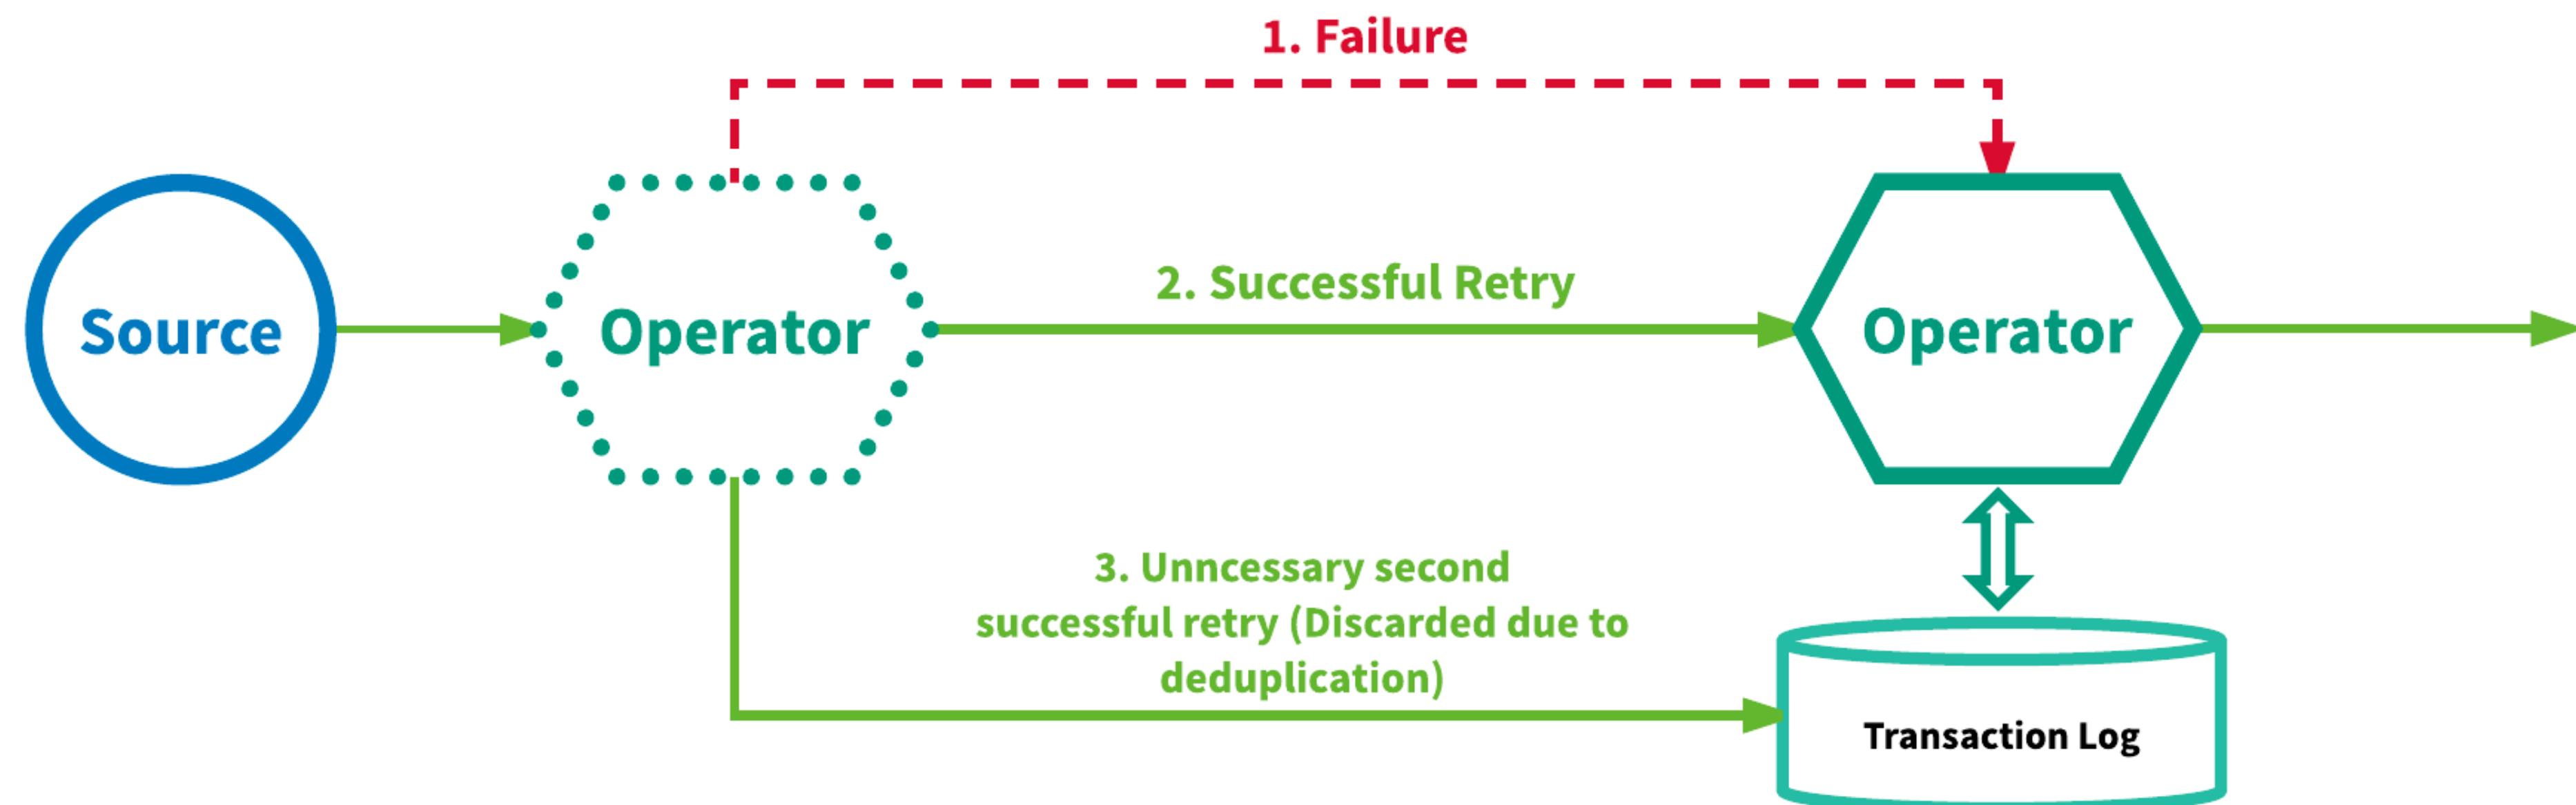
At-most-once



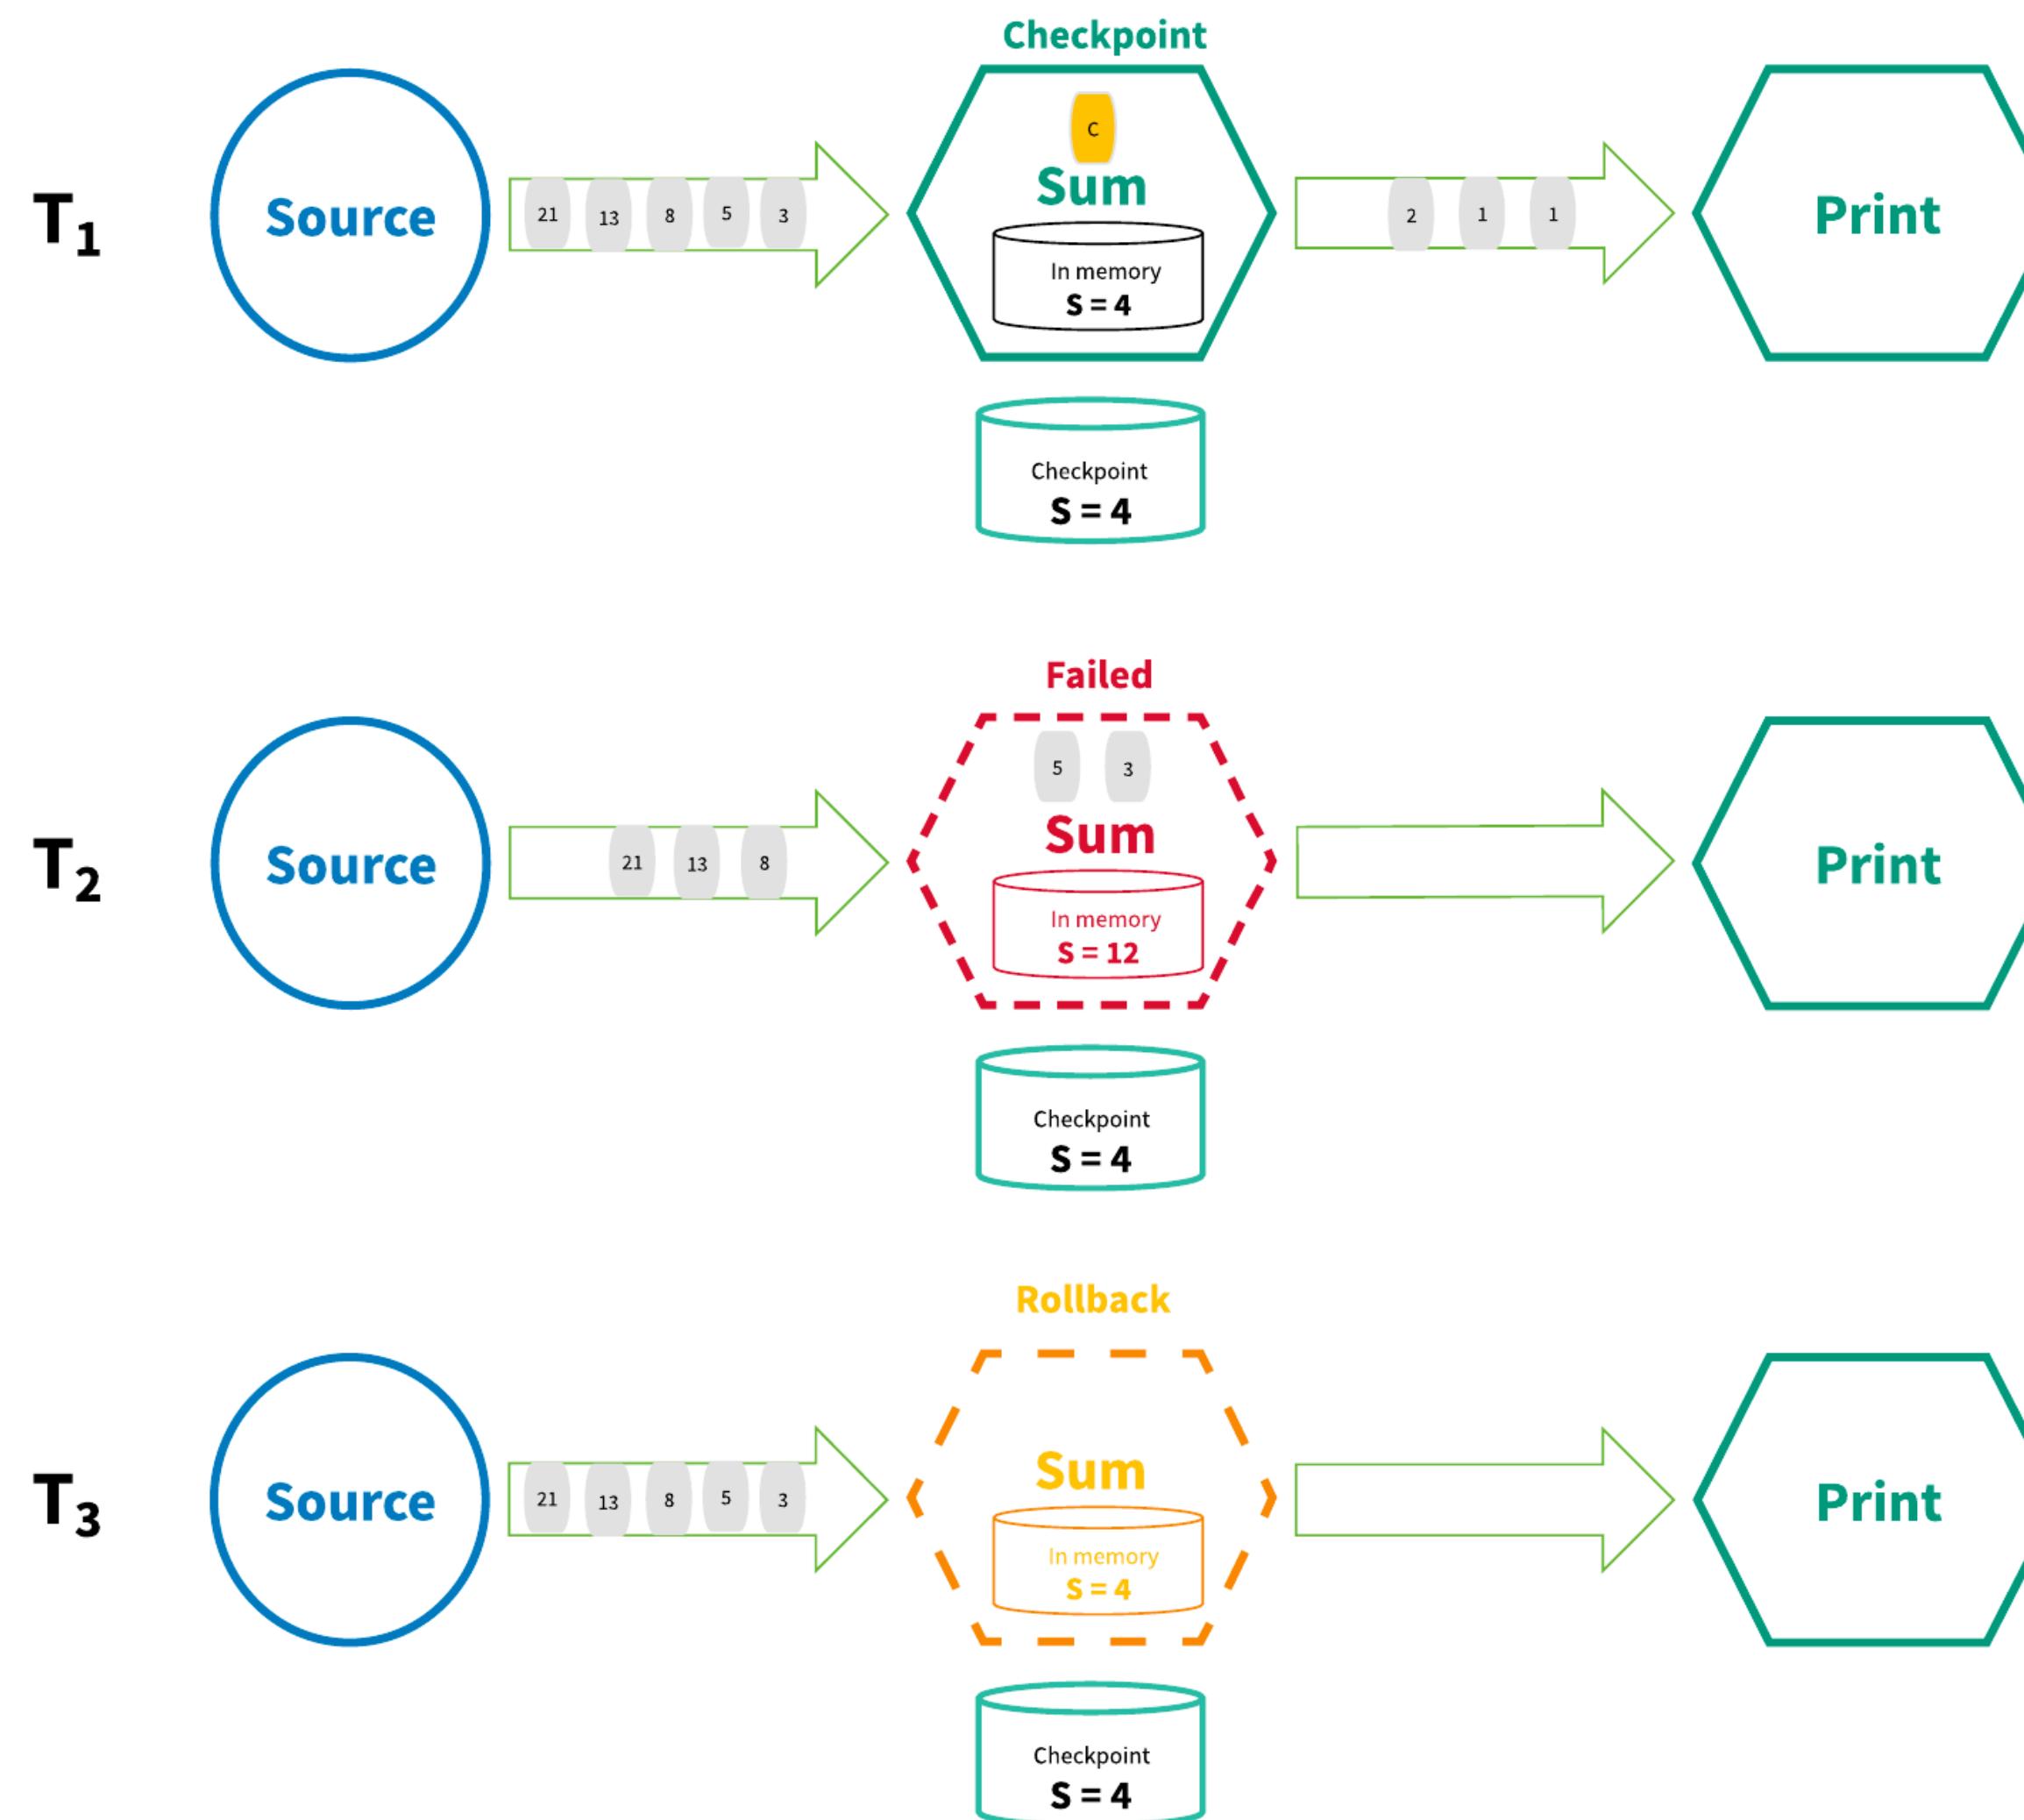
At-least-once



Exactly-once



Exactly-once



ABS

Lightweight Asynchronous Snapshots for Distributed Dataflows

<https://arxiv.org/abs/1506.08603>

Flink System Architecture

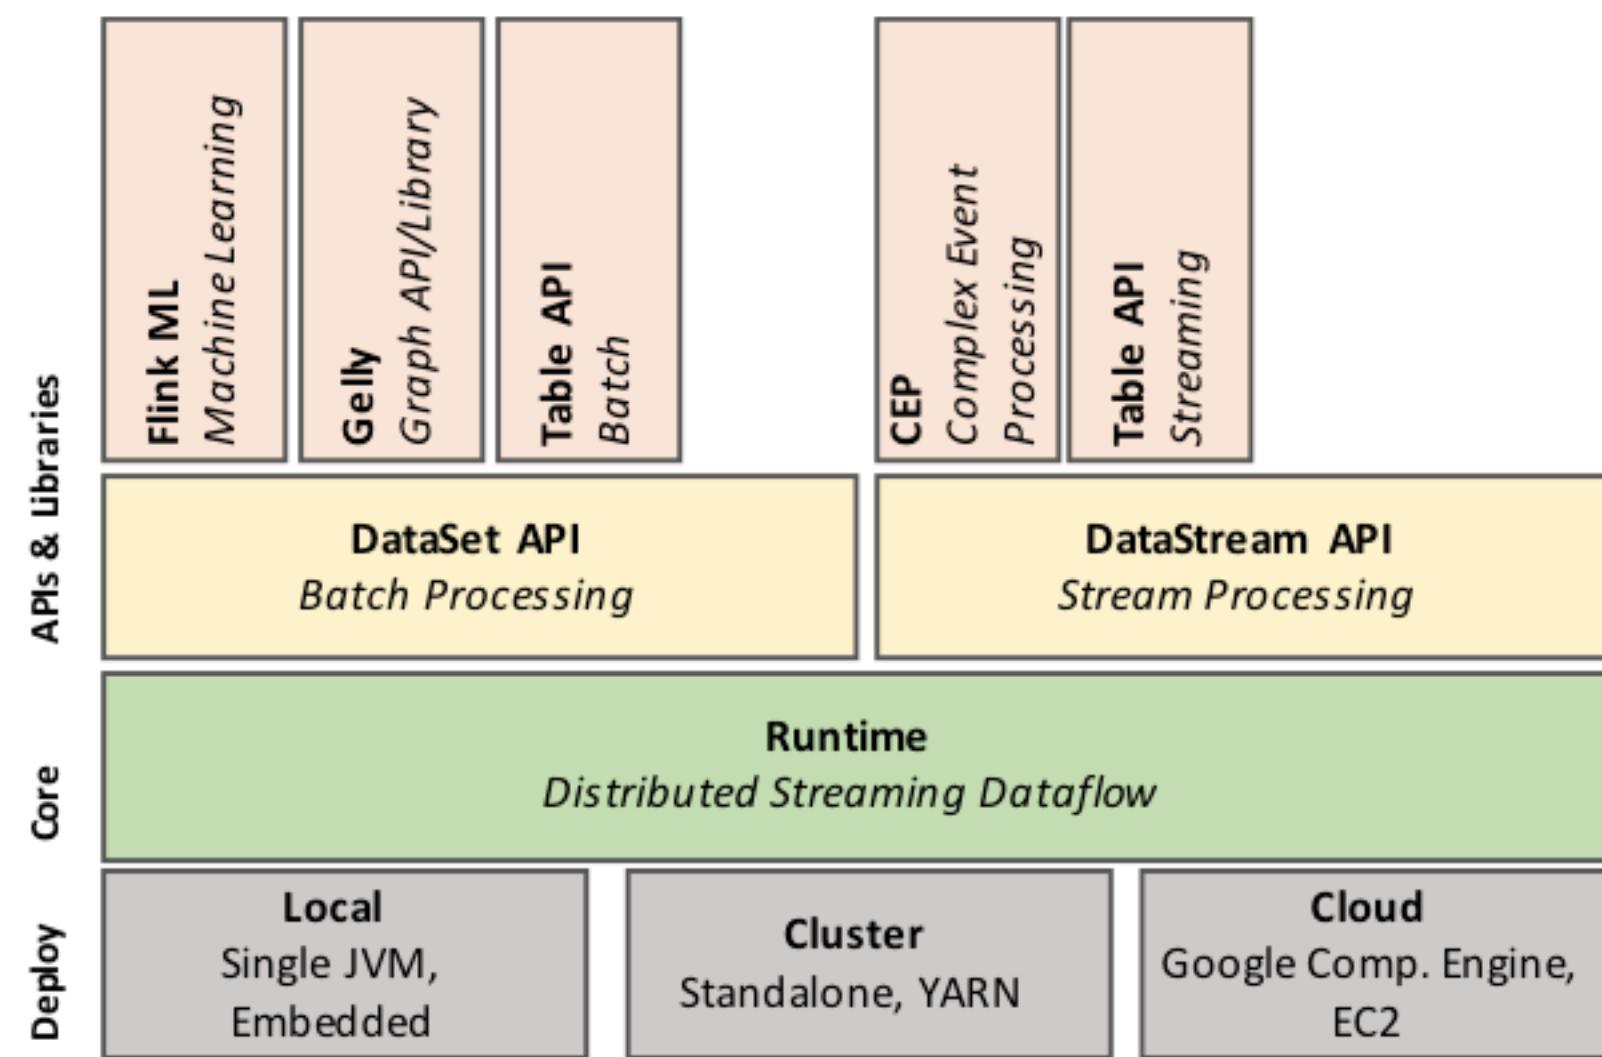


Figure 1: The Flink software stack.

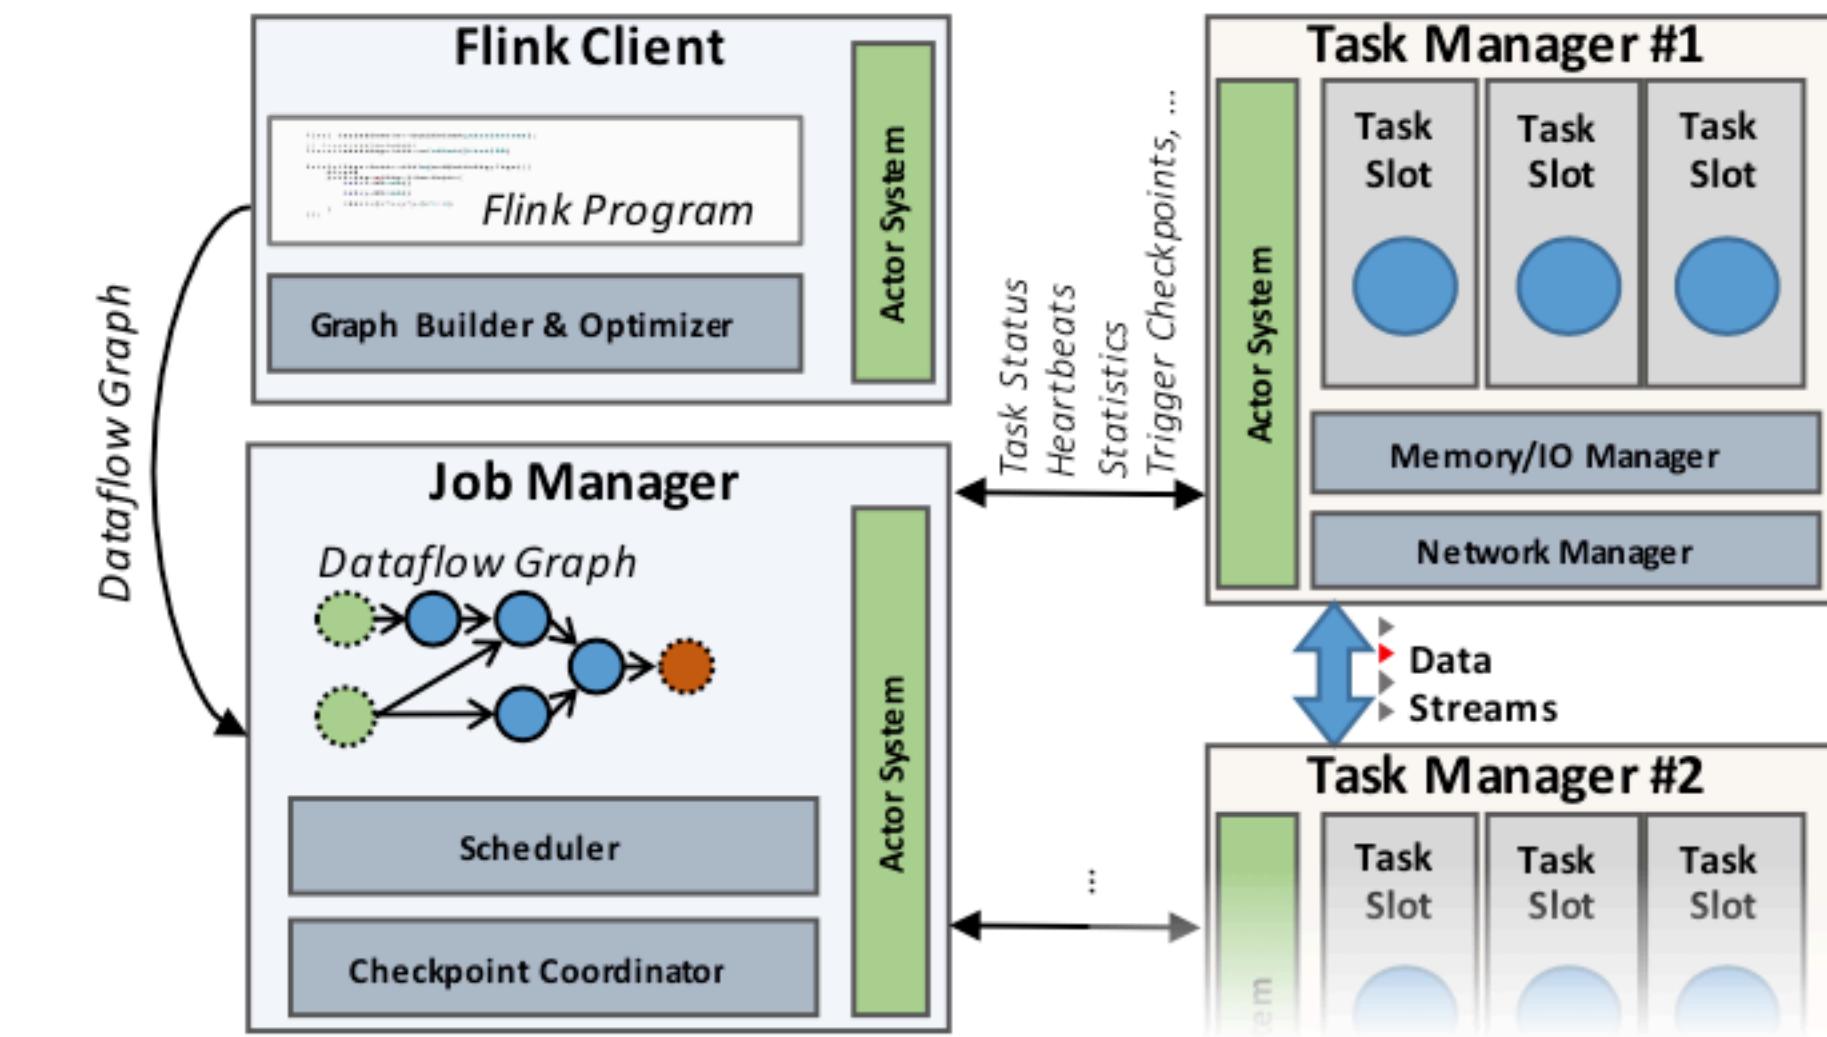
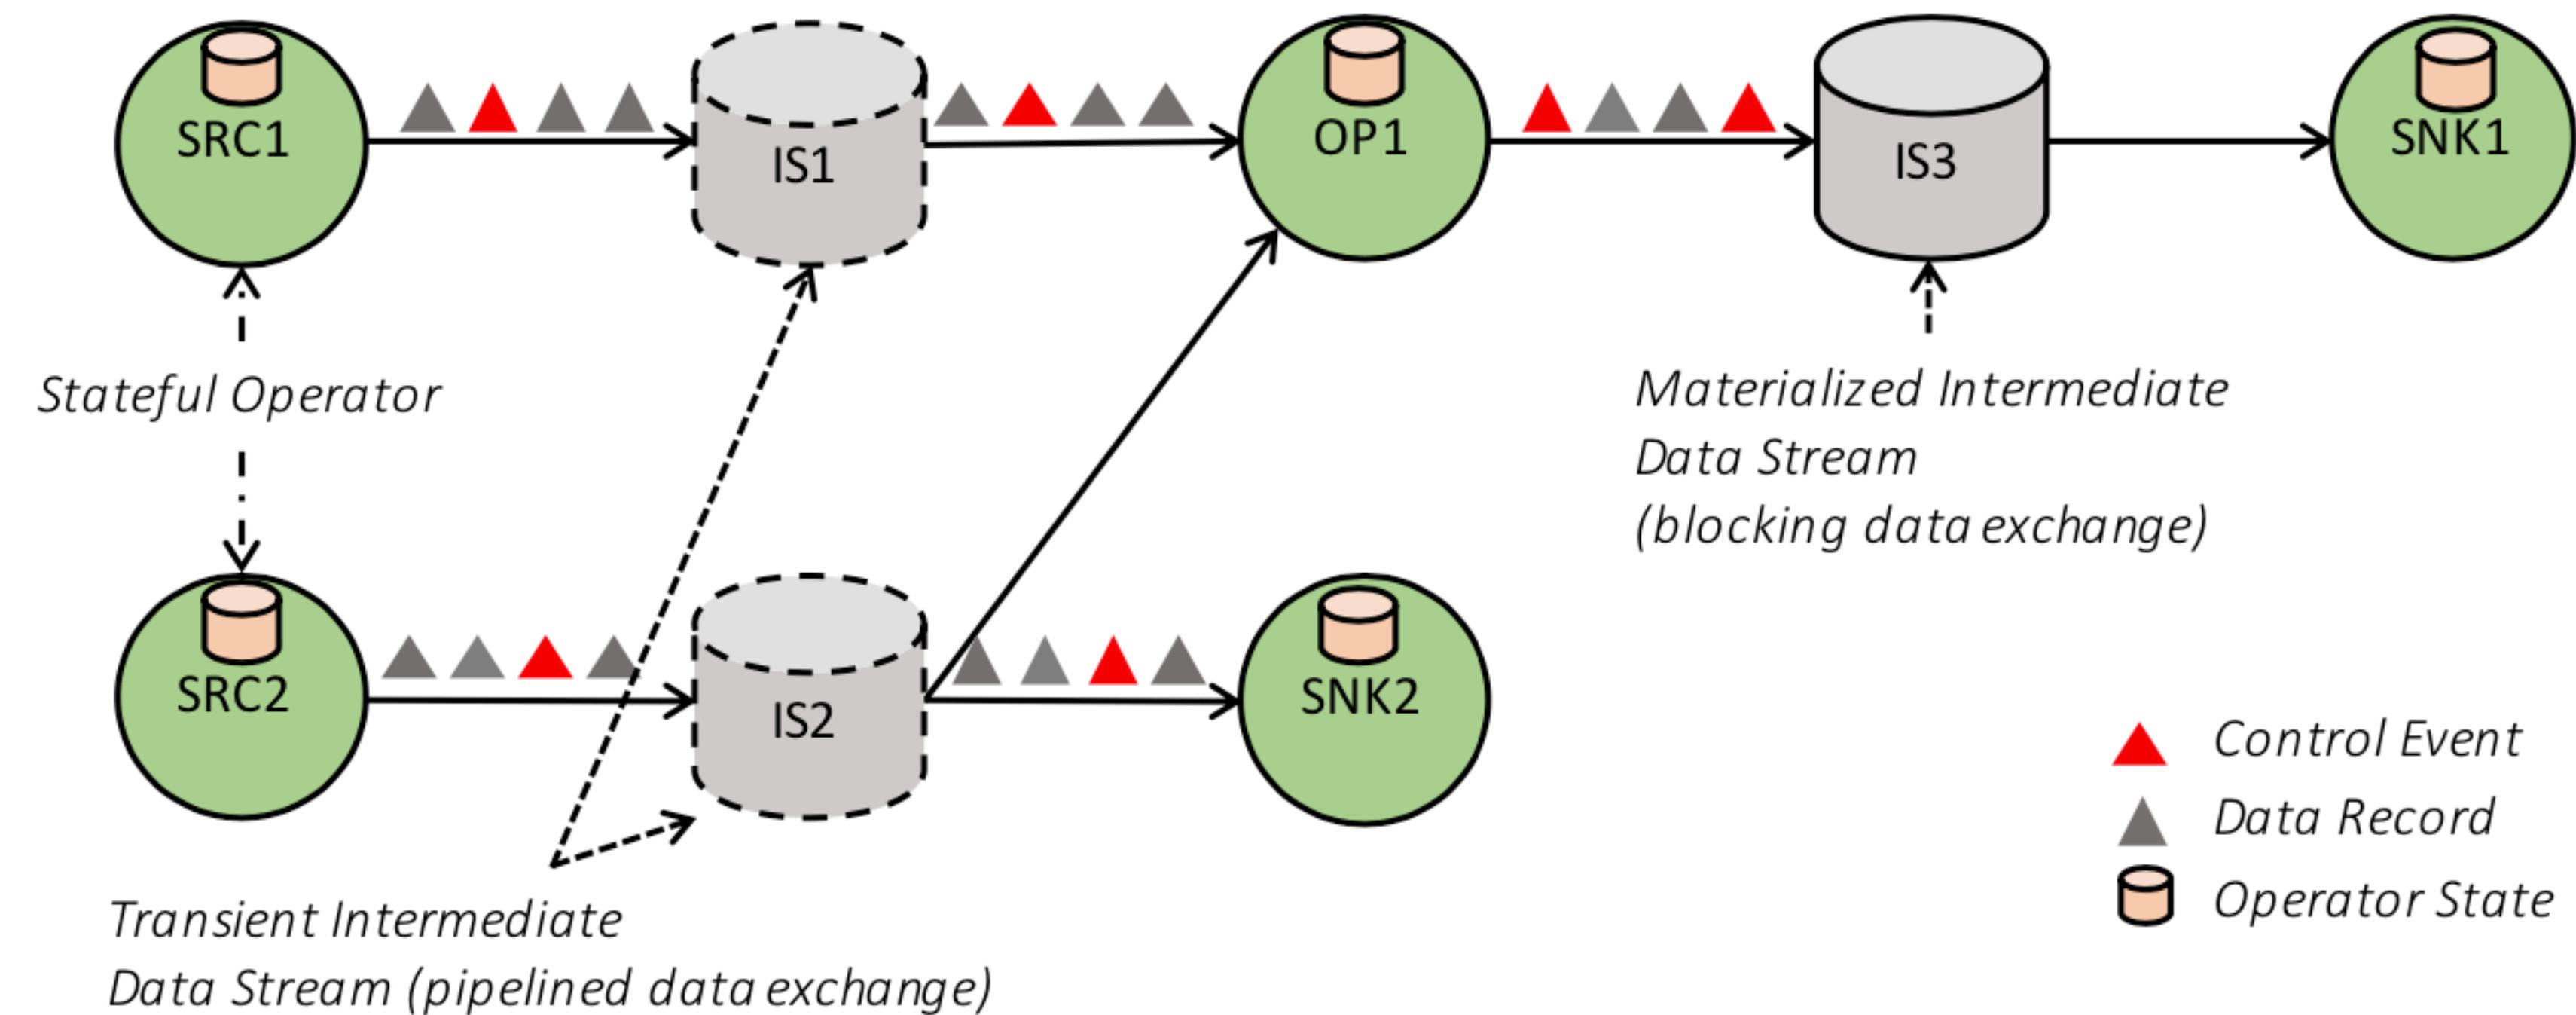


Figure 2: The Flink process model.

Apache Flink: Stream and Batch Processing in a Single Engine

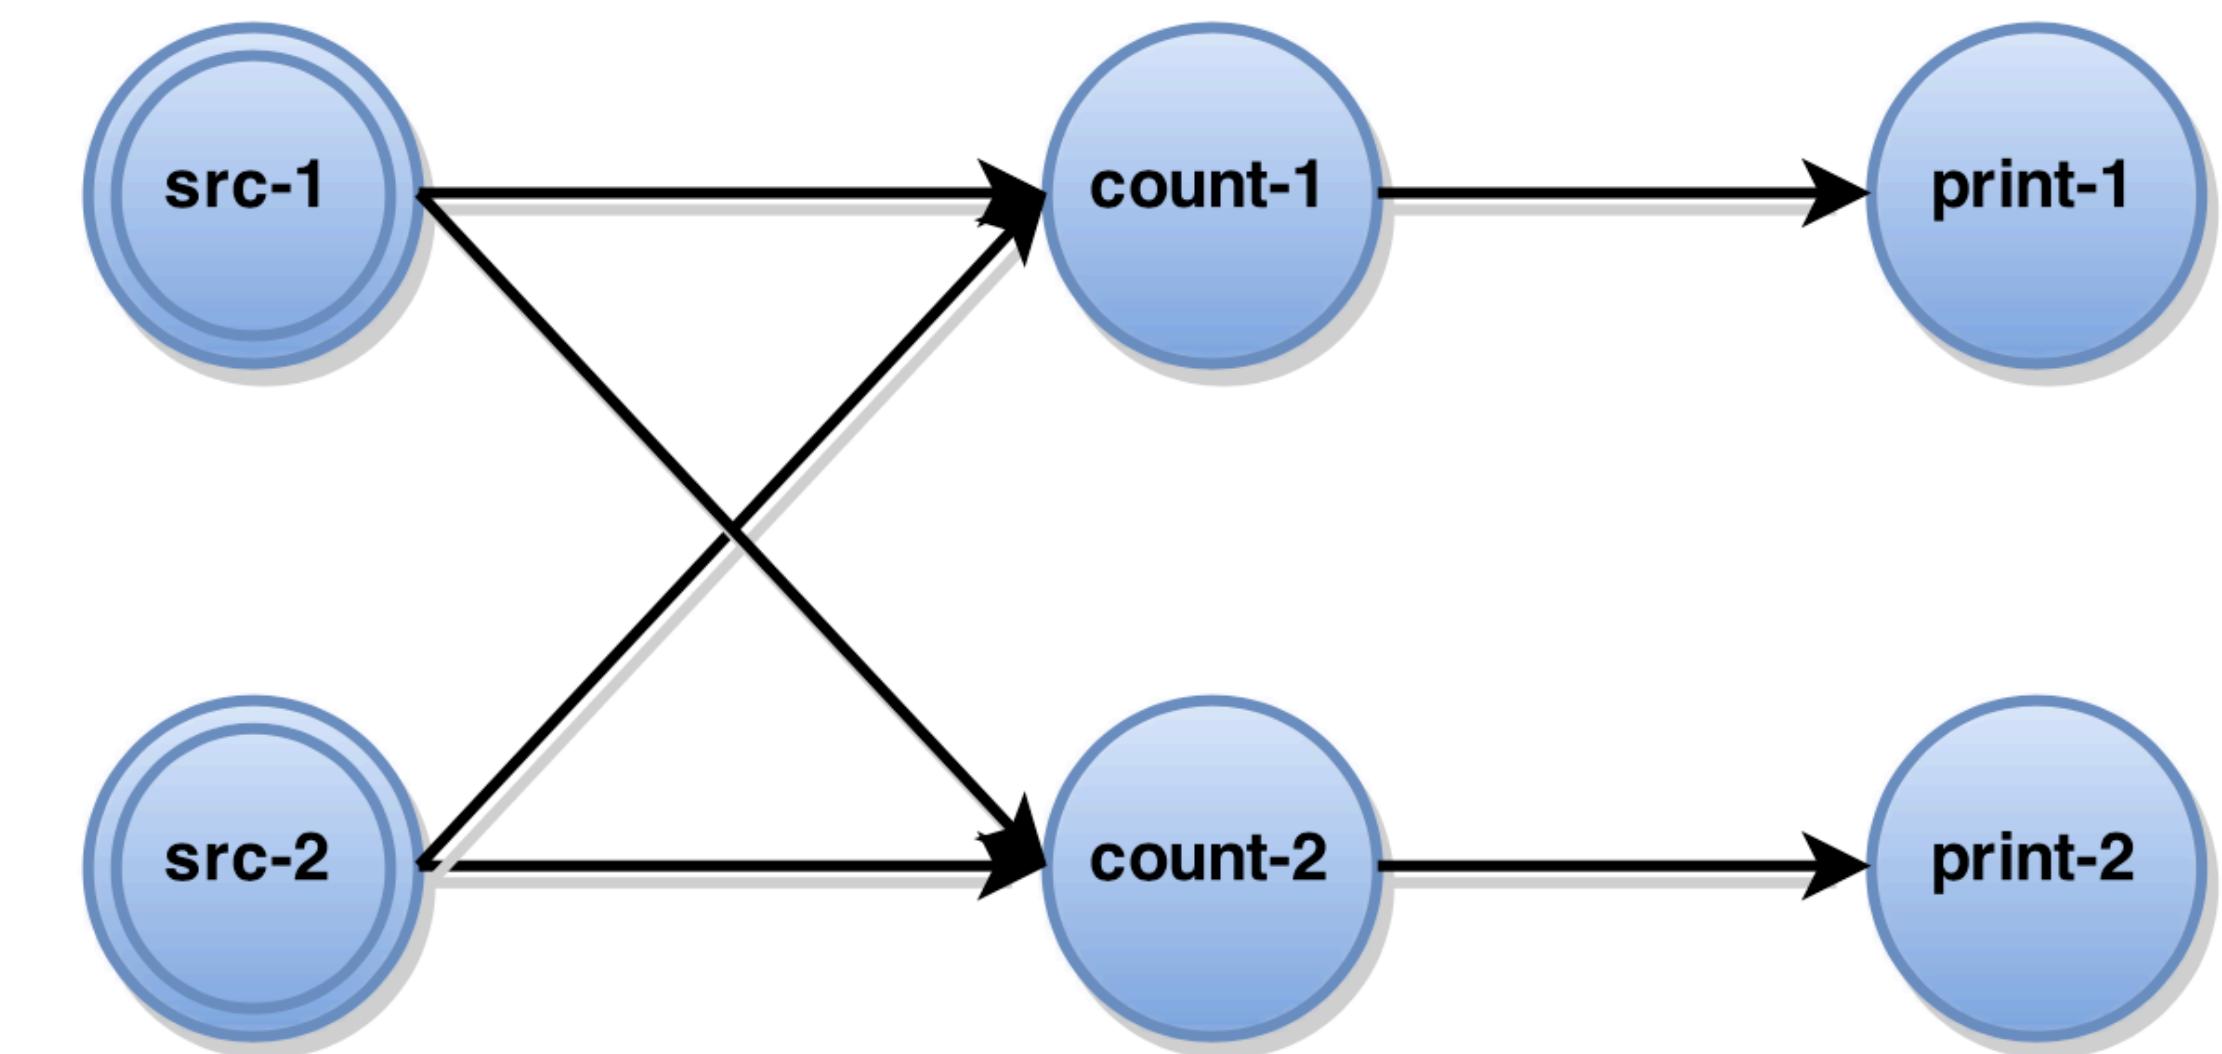
<http://sites.computer.org/debull/A15dec/p28.pdfs>

Flink Dataflow Graph

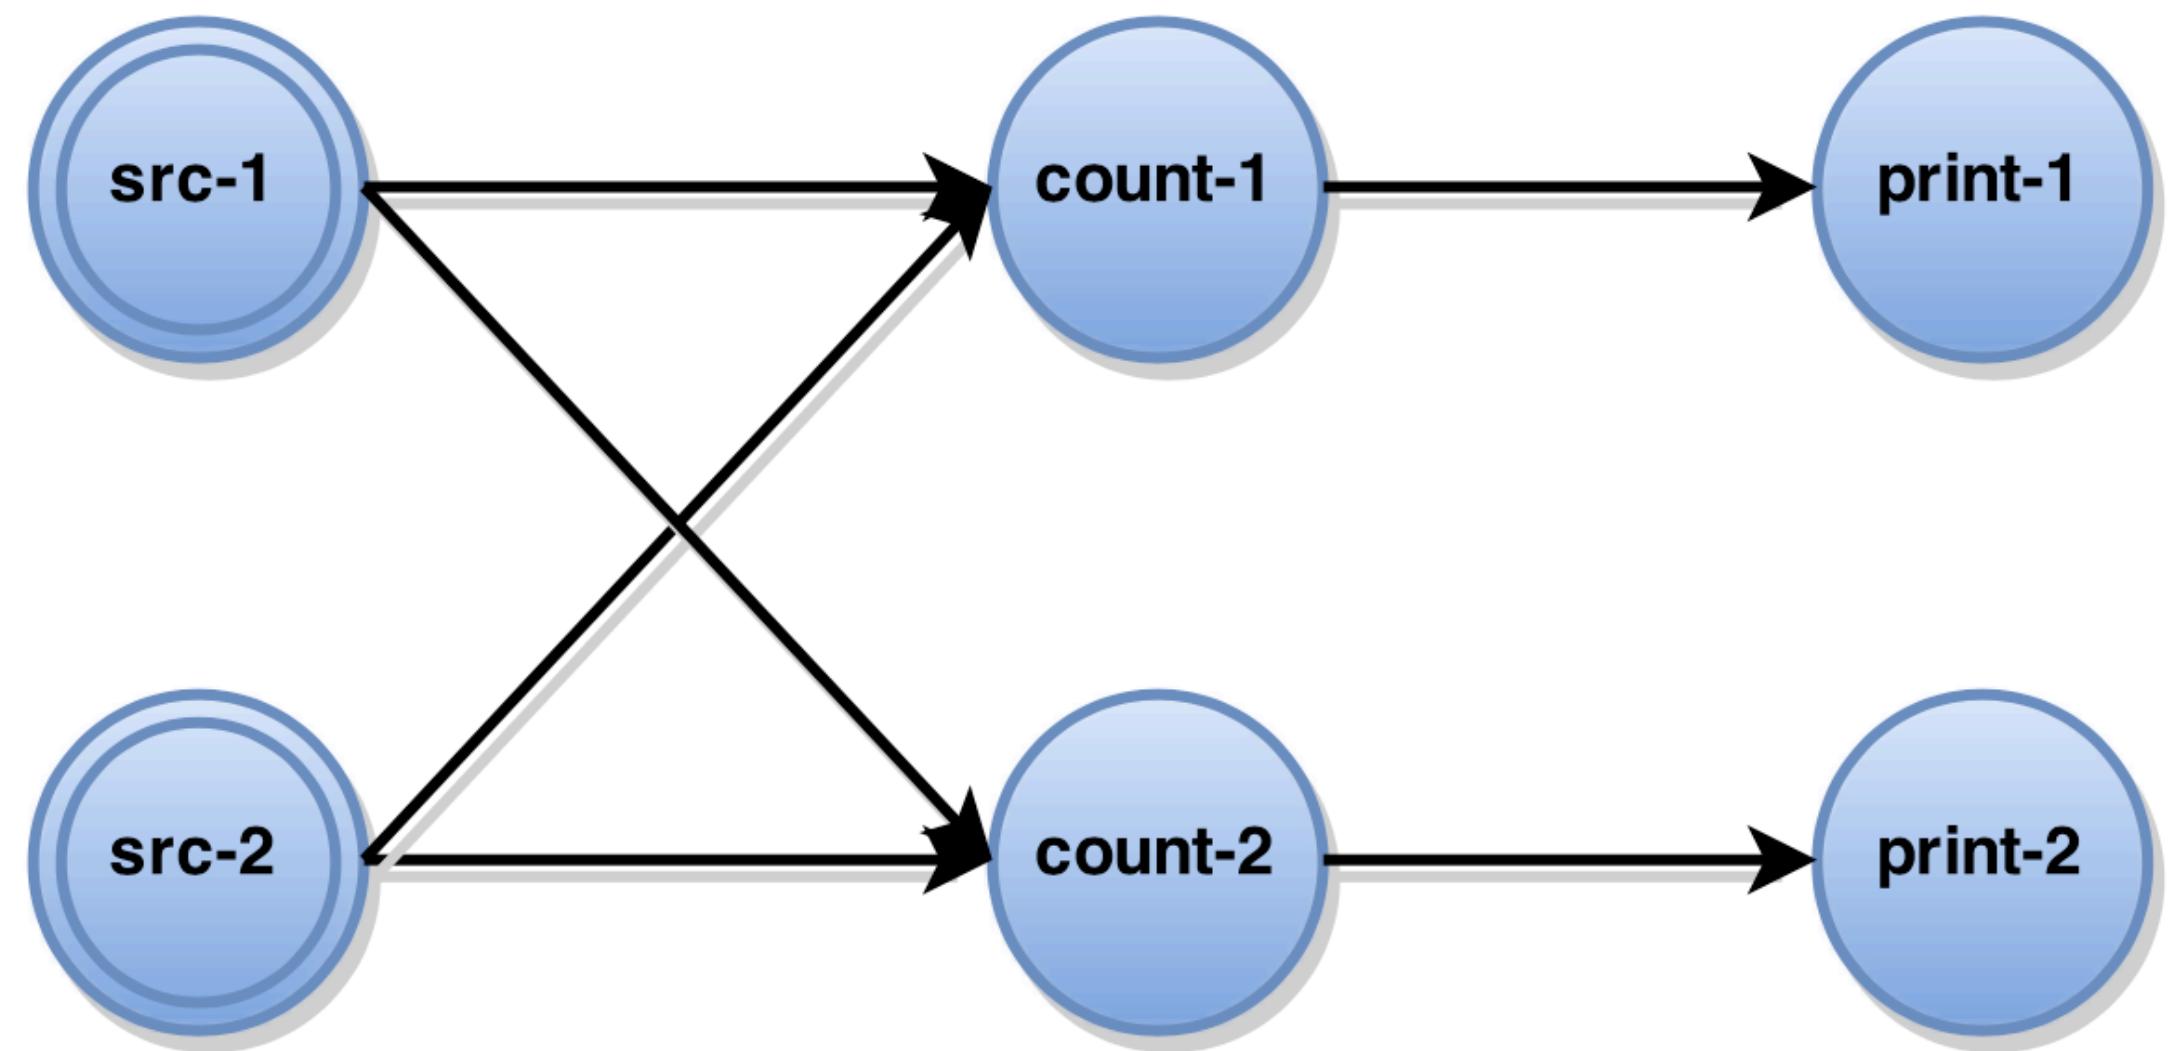


Word Count Example

```
1 val env : StreamExecutionEnvironment = ...
2 env.setParallelism(2)
3
4 val wordStream = env.readTextFile(path)
5 val countStream = wordStream.groupBy(_).count
6 countStream.print
```



Problem Definition



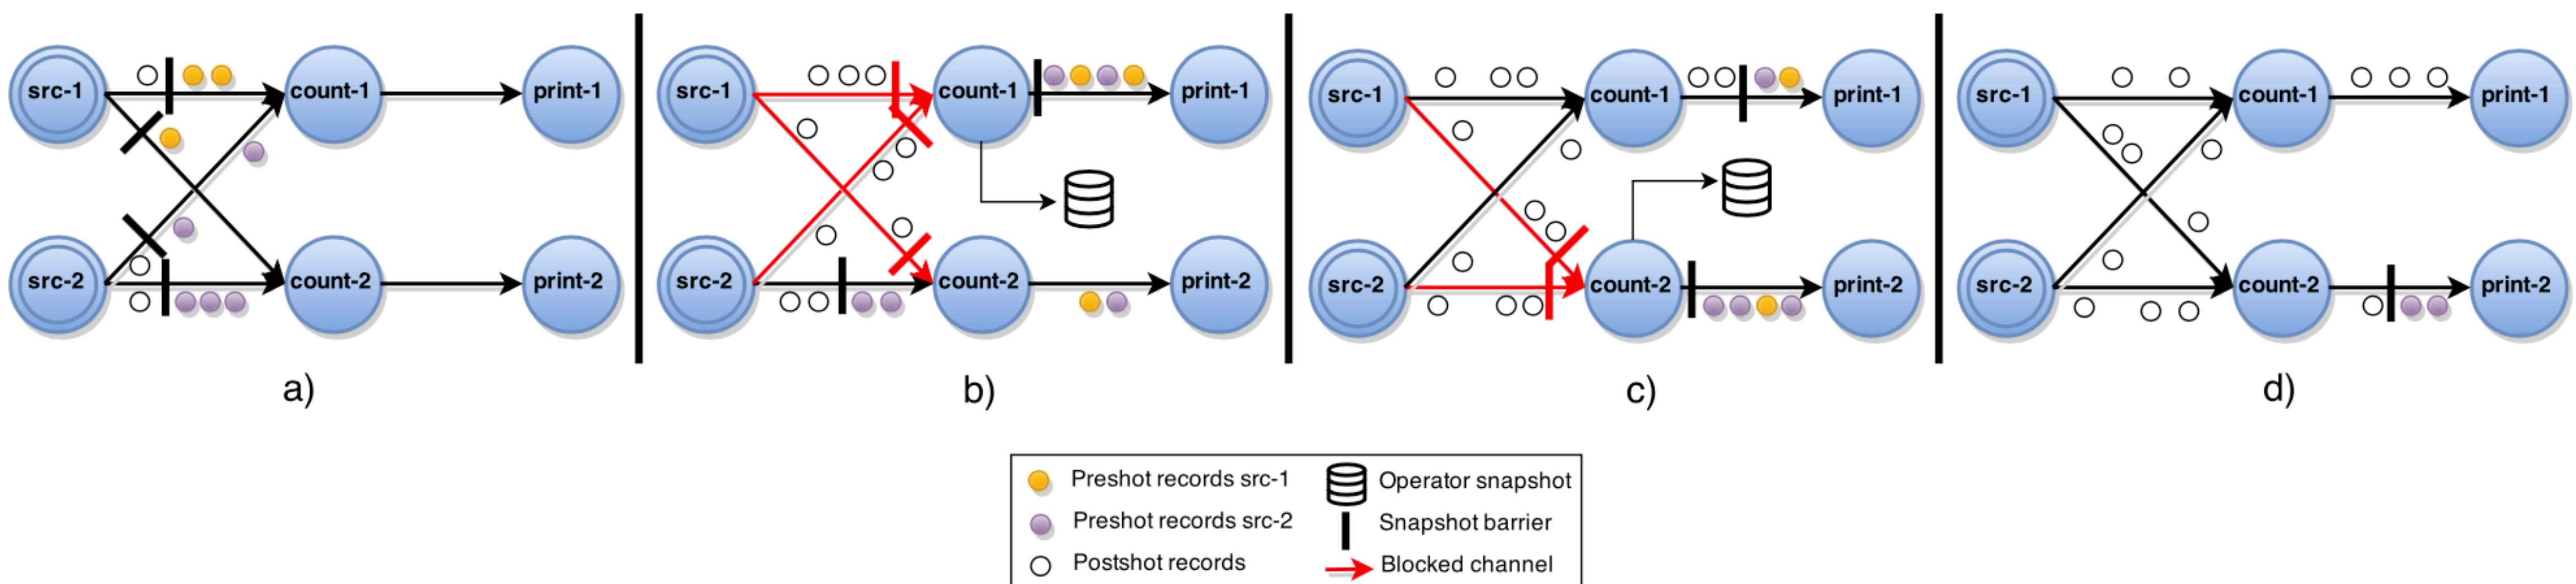
$$G=(T, E)$$

Vertex — Task
Edge — channel

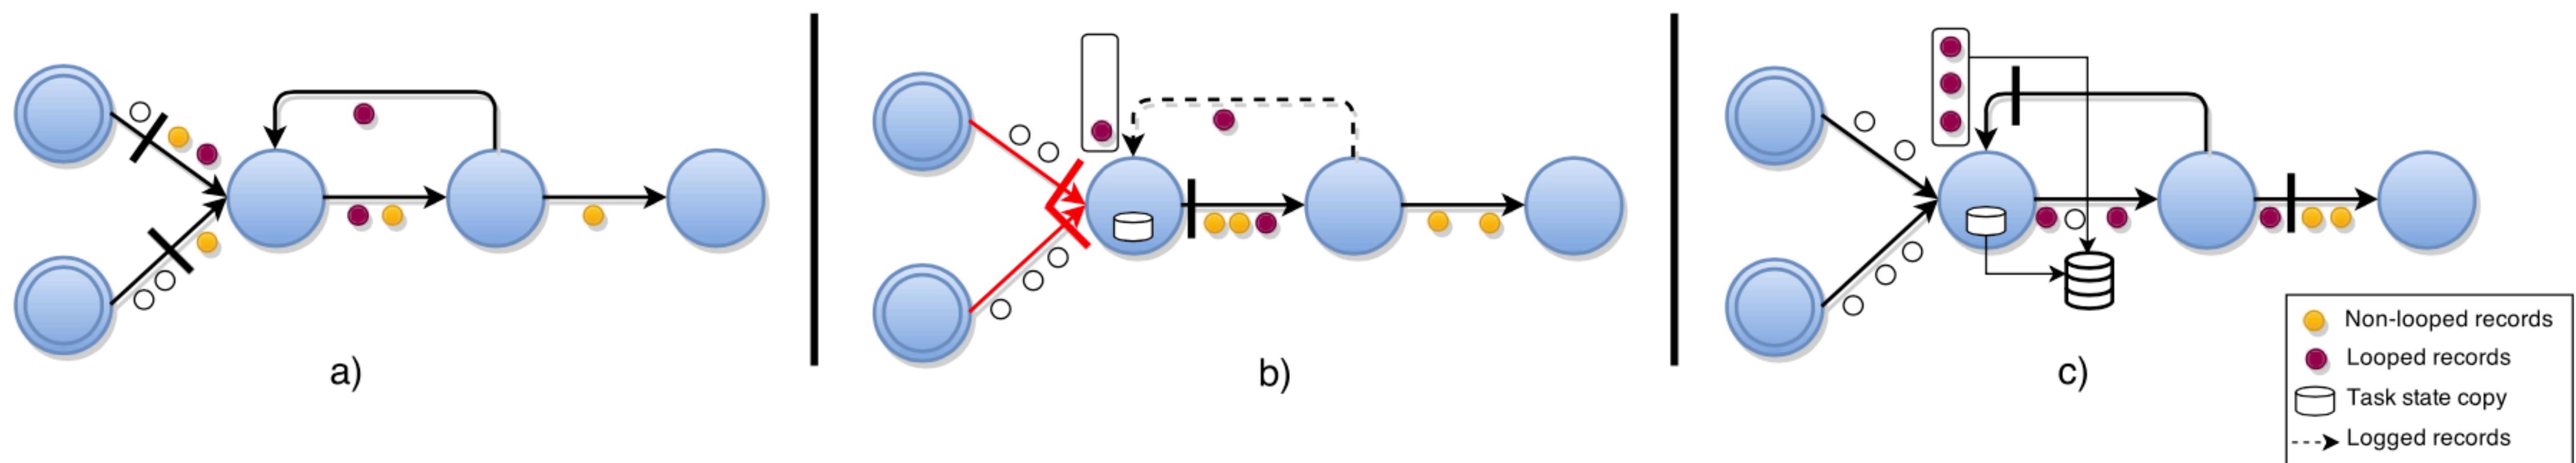
Asynchronous Barrier Snapshotting

- ABS for Acyclic Dataflows
- ABS for Cyclic Dataflows

ABS for Acyclic Dataflows



ABS for Cyclic Dataflows



Correctness Summary

- Exactly-Once
 - At-least-once Transmission + Deduplication
 - Distributed Snapshots
 - Asynchronous Barrier Snapshotting

Thanks