

SJTU SAIL (Software Architecture and Infrastructure Lab)



**SHANGHAI JIAO TONG
UNIVERSITY**

The State of Serverless Art

Yalun Lin Hsu

2020-11-15





Background

What is Serverless Computing?

Serverless computing is a programming abstraction that enables users to upload programs, run them at any scale, and pay only for resources used.

Functions-as-a-Service (FaaS)

- AWS Lambda, Google Cloud Functions, OpenWhisk (IBM), Azure Functions, OpenLambda, OpenFaaS, kNative...
- Optimized for simplicity – register functions, enable triggers, and scale transparently

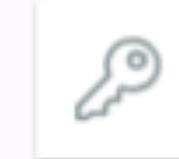
Function code [Info](#)

Code entry type [Edit code inline](#) Runtime Python 3.6

```
1 import boto3
2 import pickle
3 import random
4 import time
5
6 client = boto3.resource('dynamodb')
7 table = client.Table('vsreekanti')
8 candidates = client.Table('candidates')
9 THRESHOLD = 50
10
11 def main(thisid):
12     print("My invocation's id is: " + thisid)
13     thisid = int(thisid)
14     vote_round_count = 0
15     am_leader = False
16
17     while True:
18         time.sleep(.25)
```

Add triggers

Choose a trigger from the list below to add it to your function.



API Gateway

AWS IoT

Alexa Skills Kit

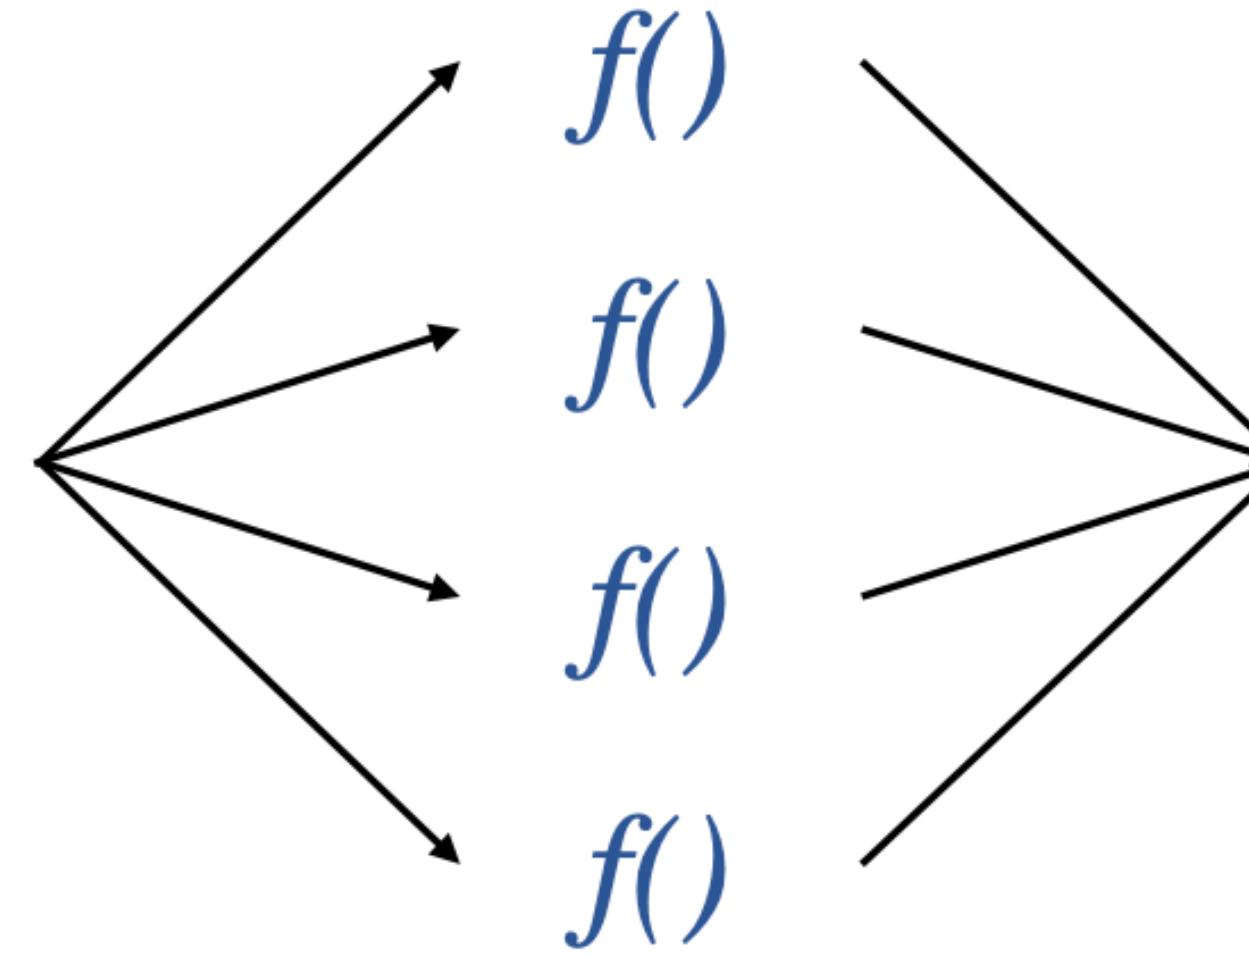
Alexa Smart Home

Application Load Balancer

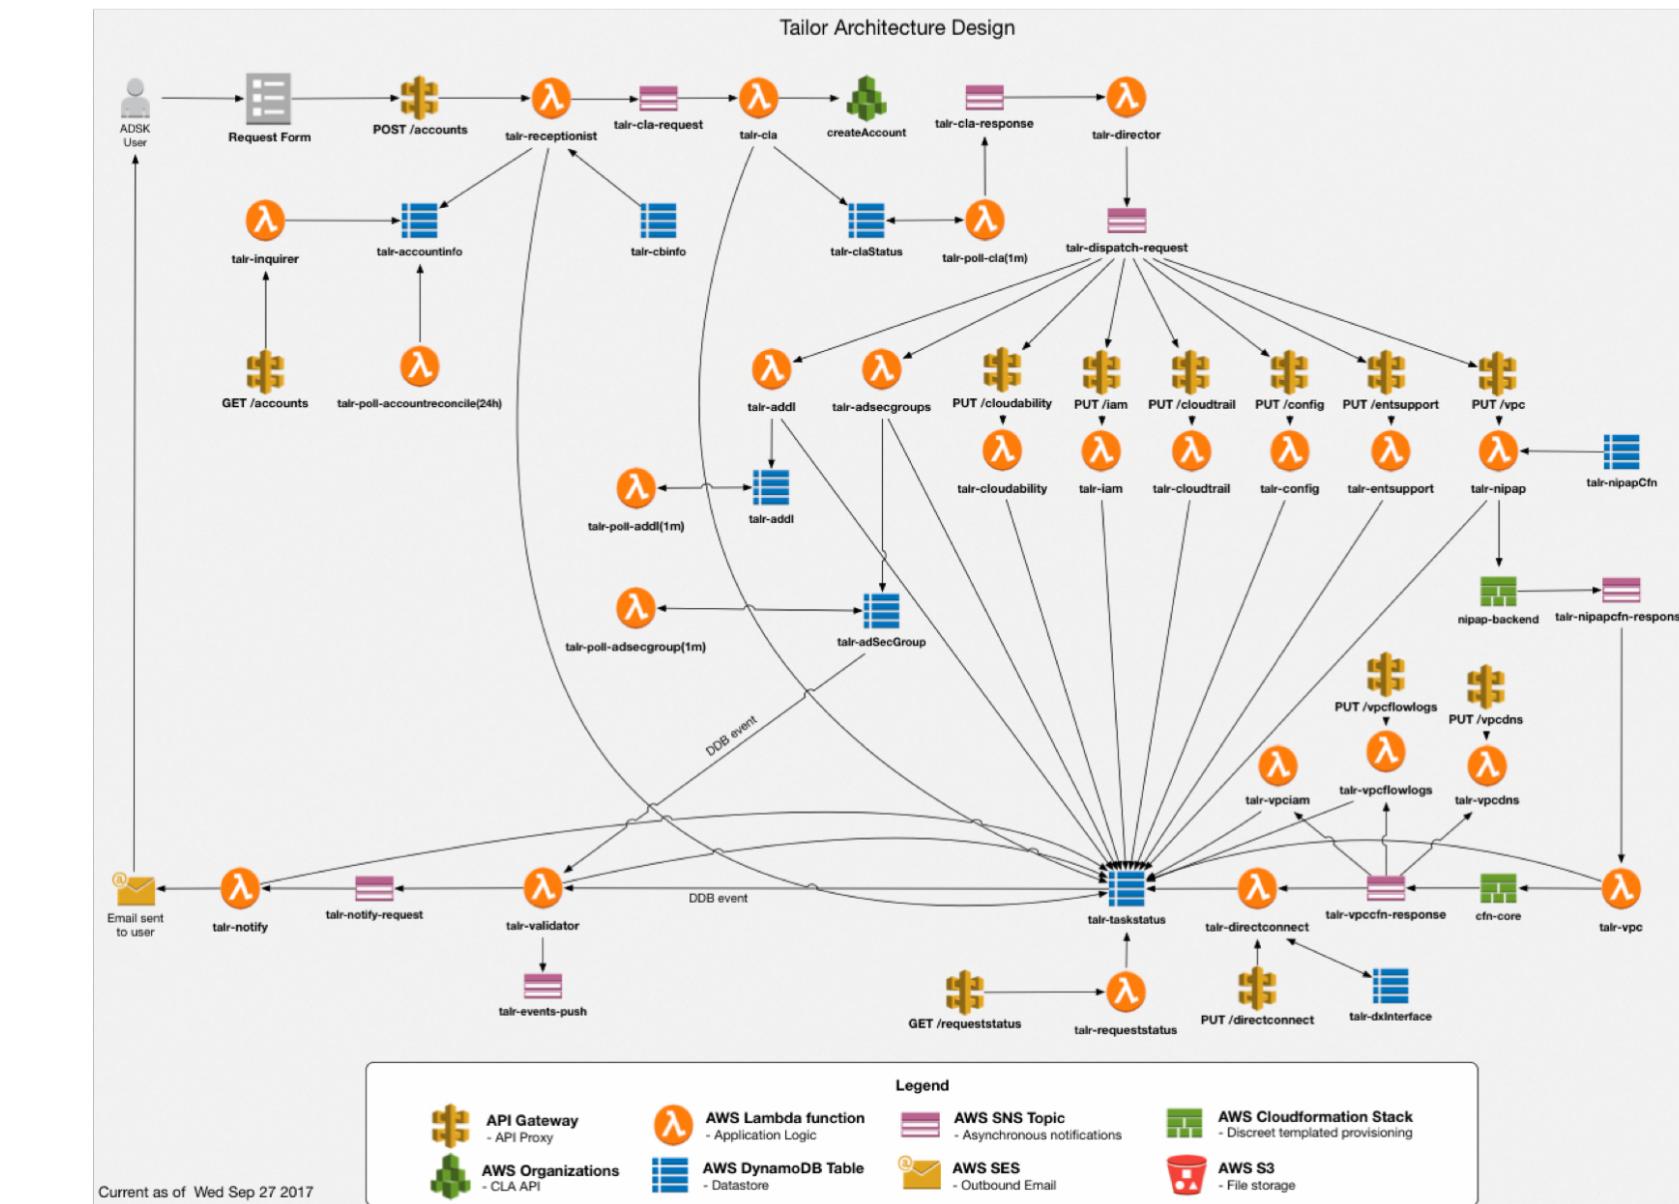
CloudFront

Add triggers from the list on the left

What is Serverless Good at Today?



Embarrassingly parallel tasks

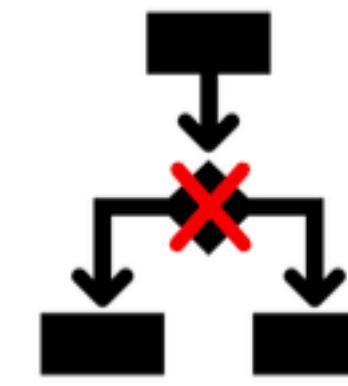


Workflow orchestration

Limitations on FaaS Today



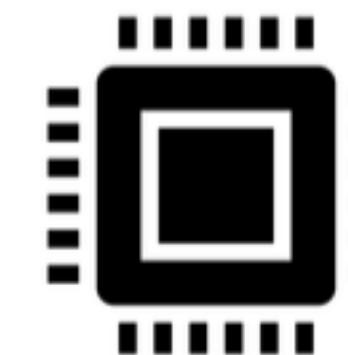
Limited execution lifetimes



No inbound network connections



IO is a bottleneck



No specialized hardware

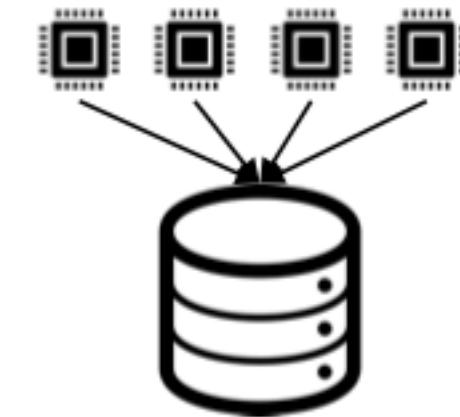
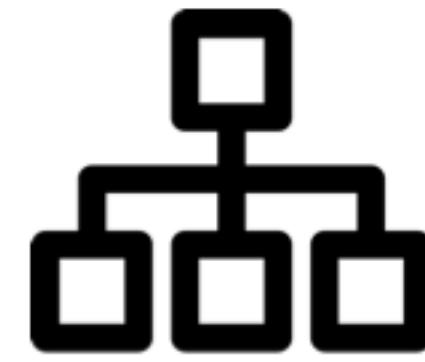
But that's^{NOT} okay: Everything is^{NOT} functional!

- Functional programs don't have side effects or mutable state!
- And it *is called* AWS Lambda

Dysfunction-as-a-Service

- FaaS is not designed for functional programming because *real applications share state*

$$f(g(x))$$

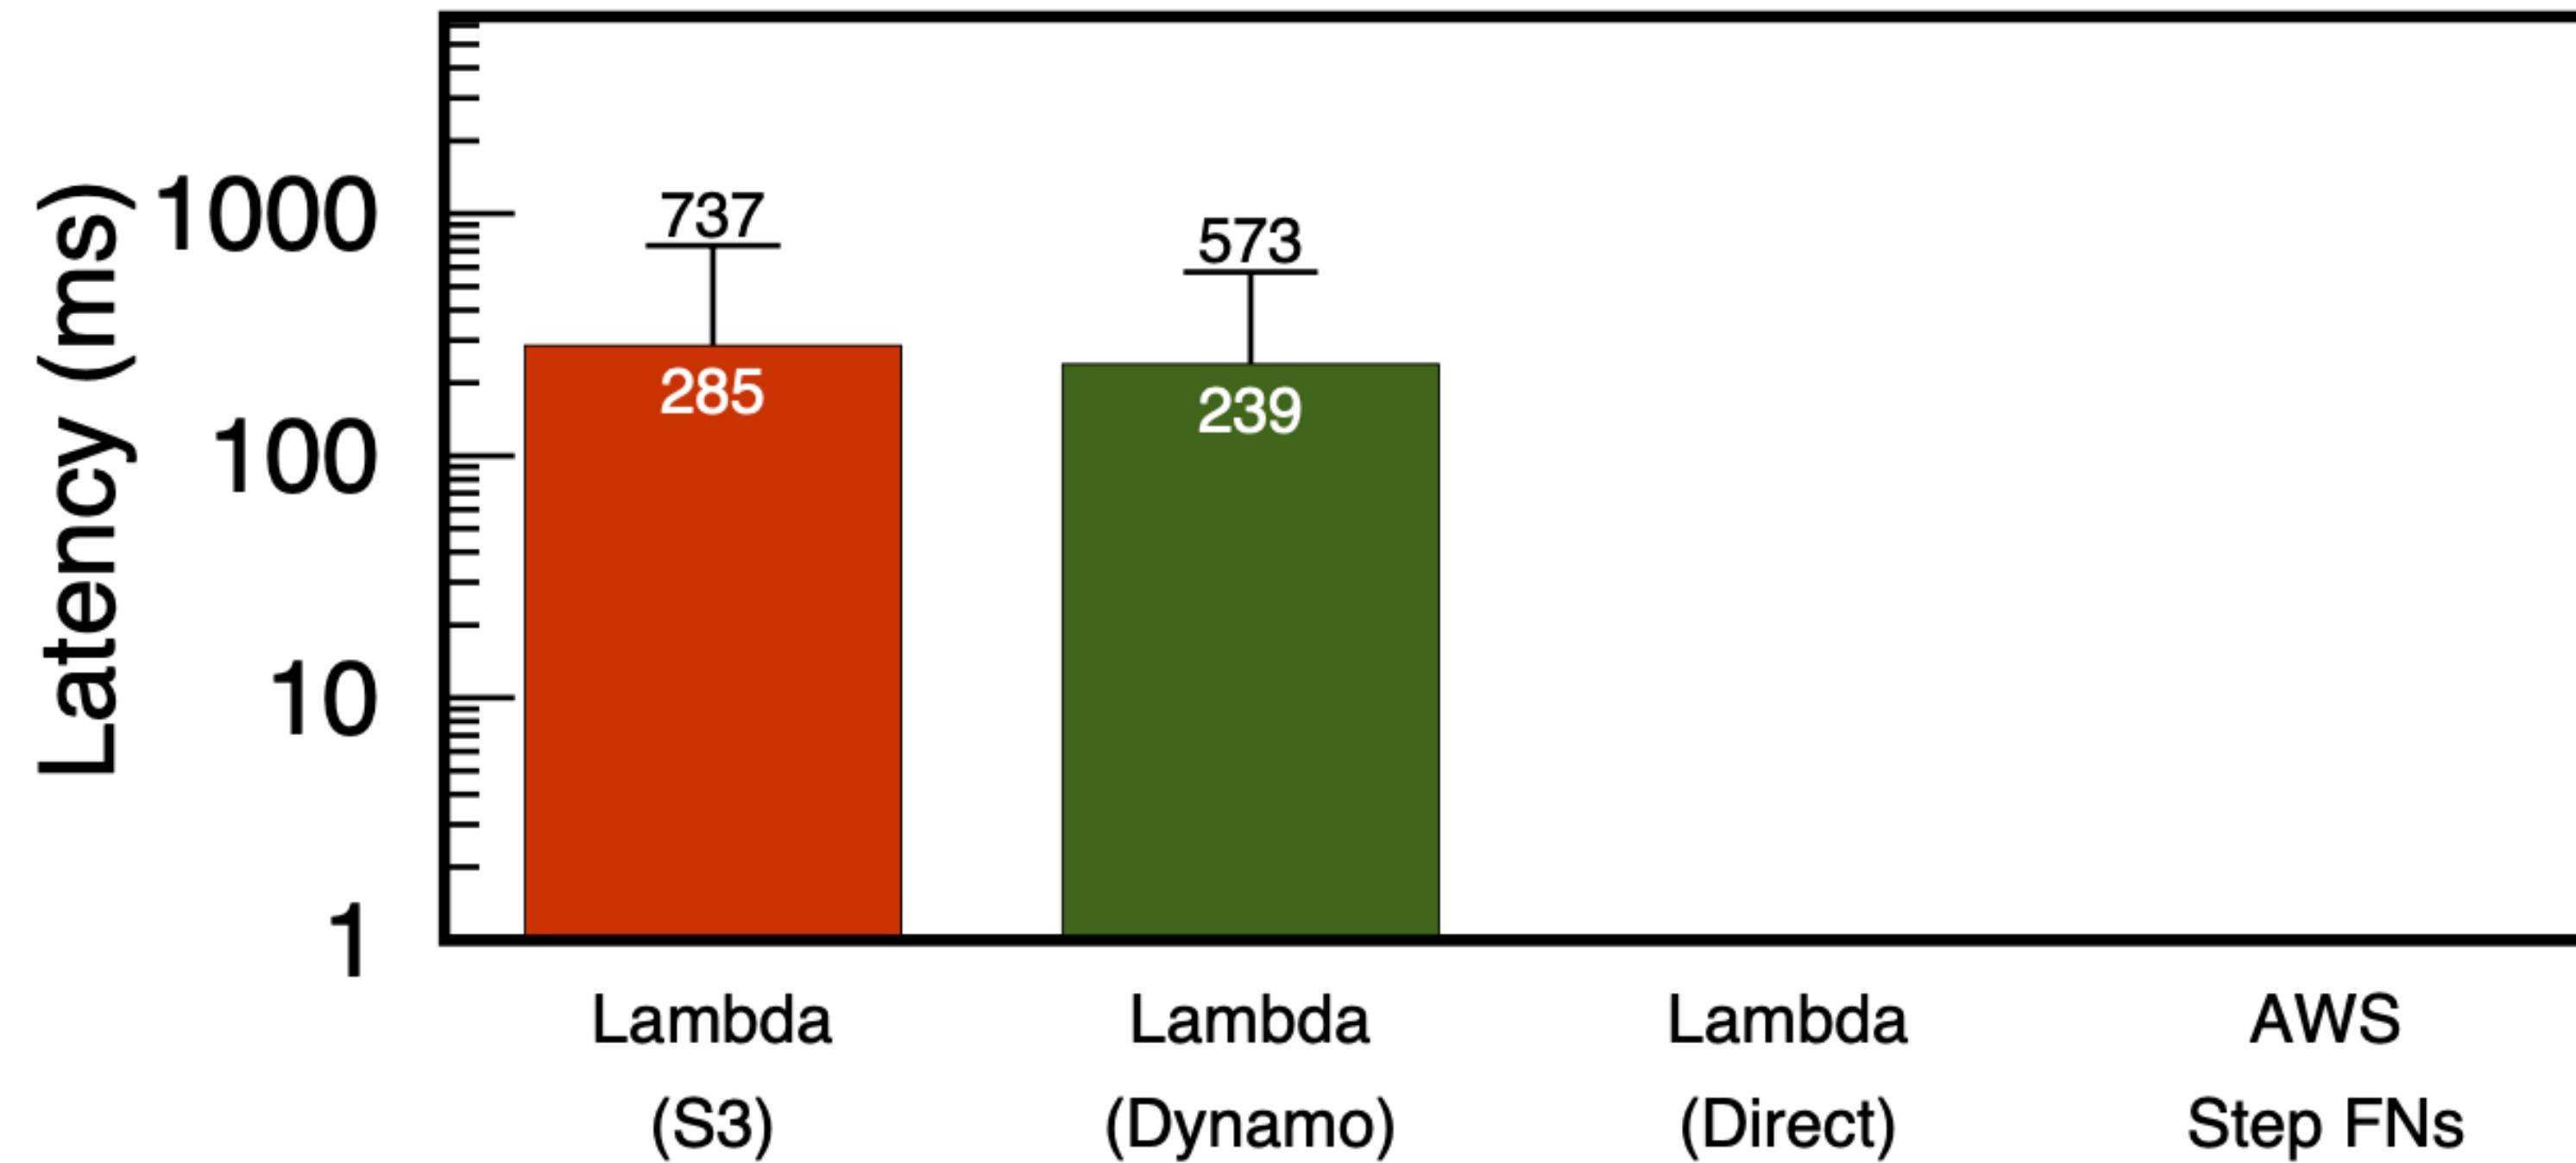


- FaaS is poorly suited for all of these

Quantifying The Pain of FaaS

Even Functional Programming is Slow!

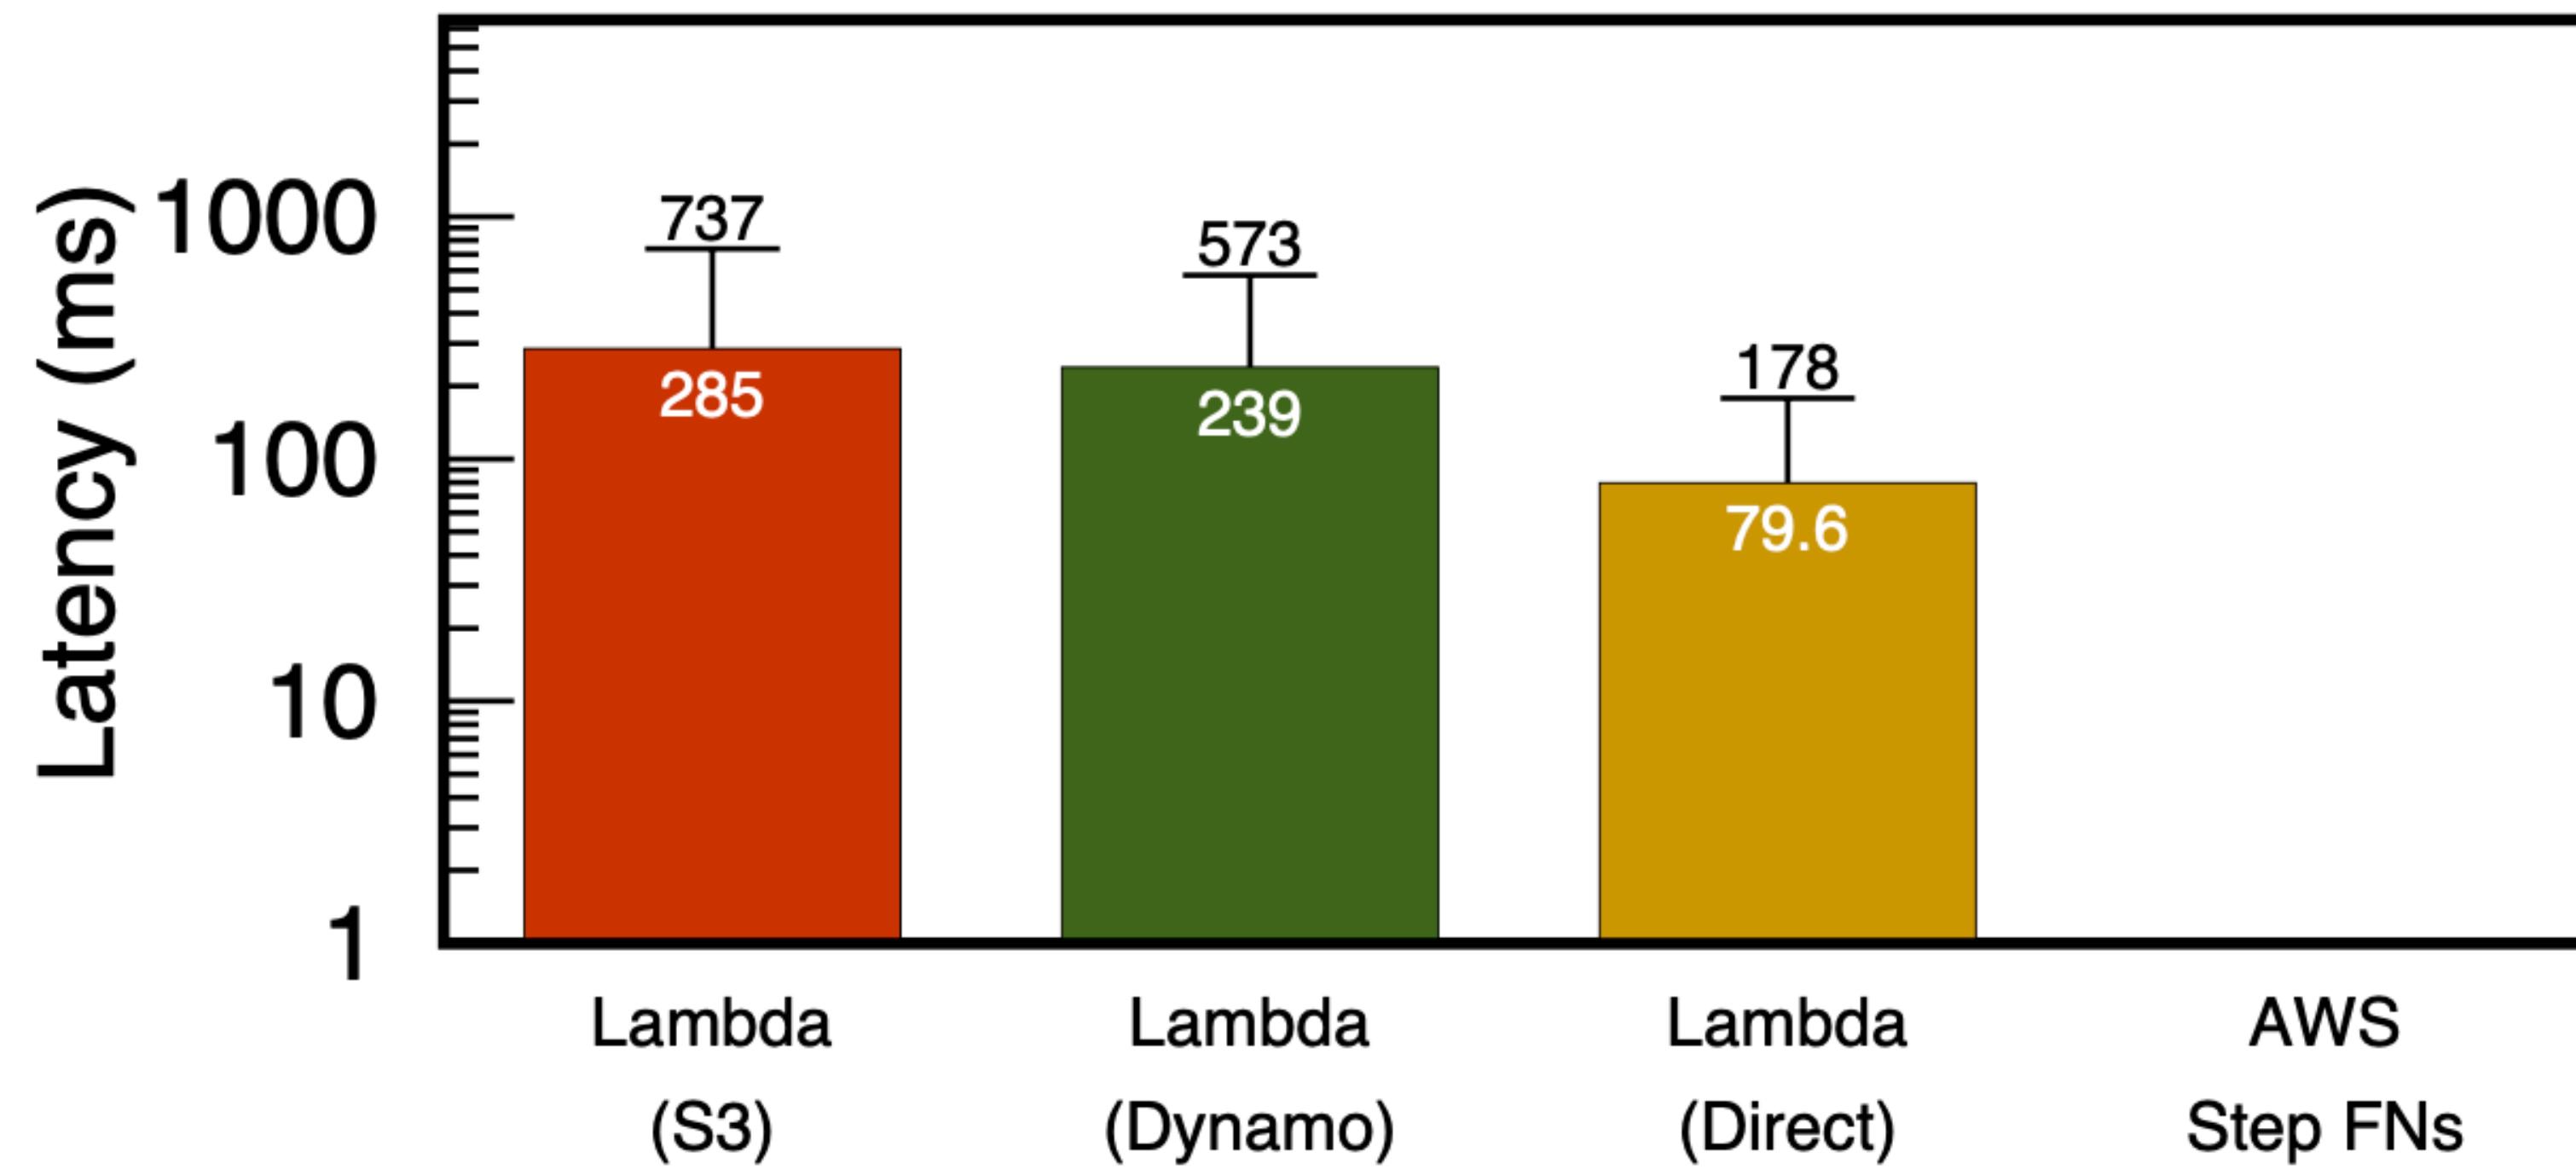
$f(g(x))$



Median and 99th percentile latencies for composing two arithmetic functions on AWS Lambda.

Even Functional Programming is Slow!

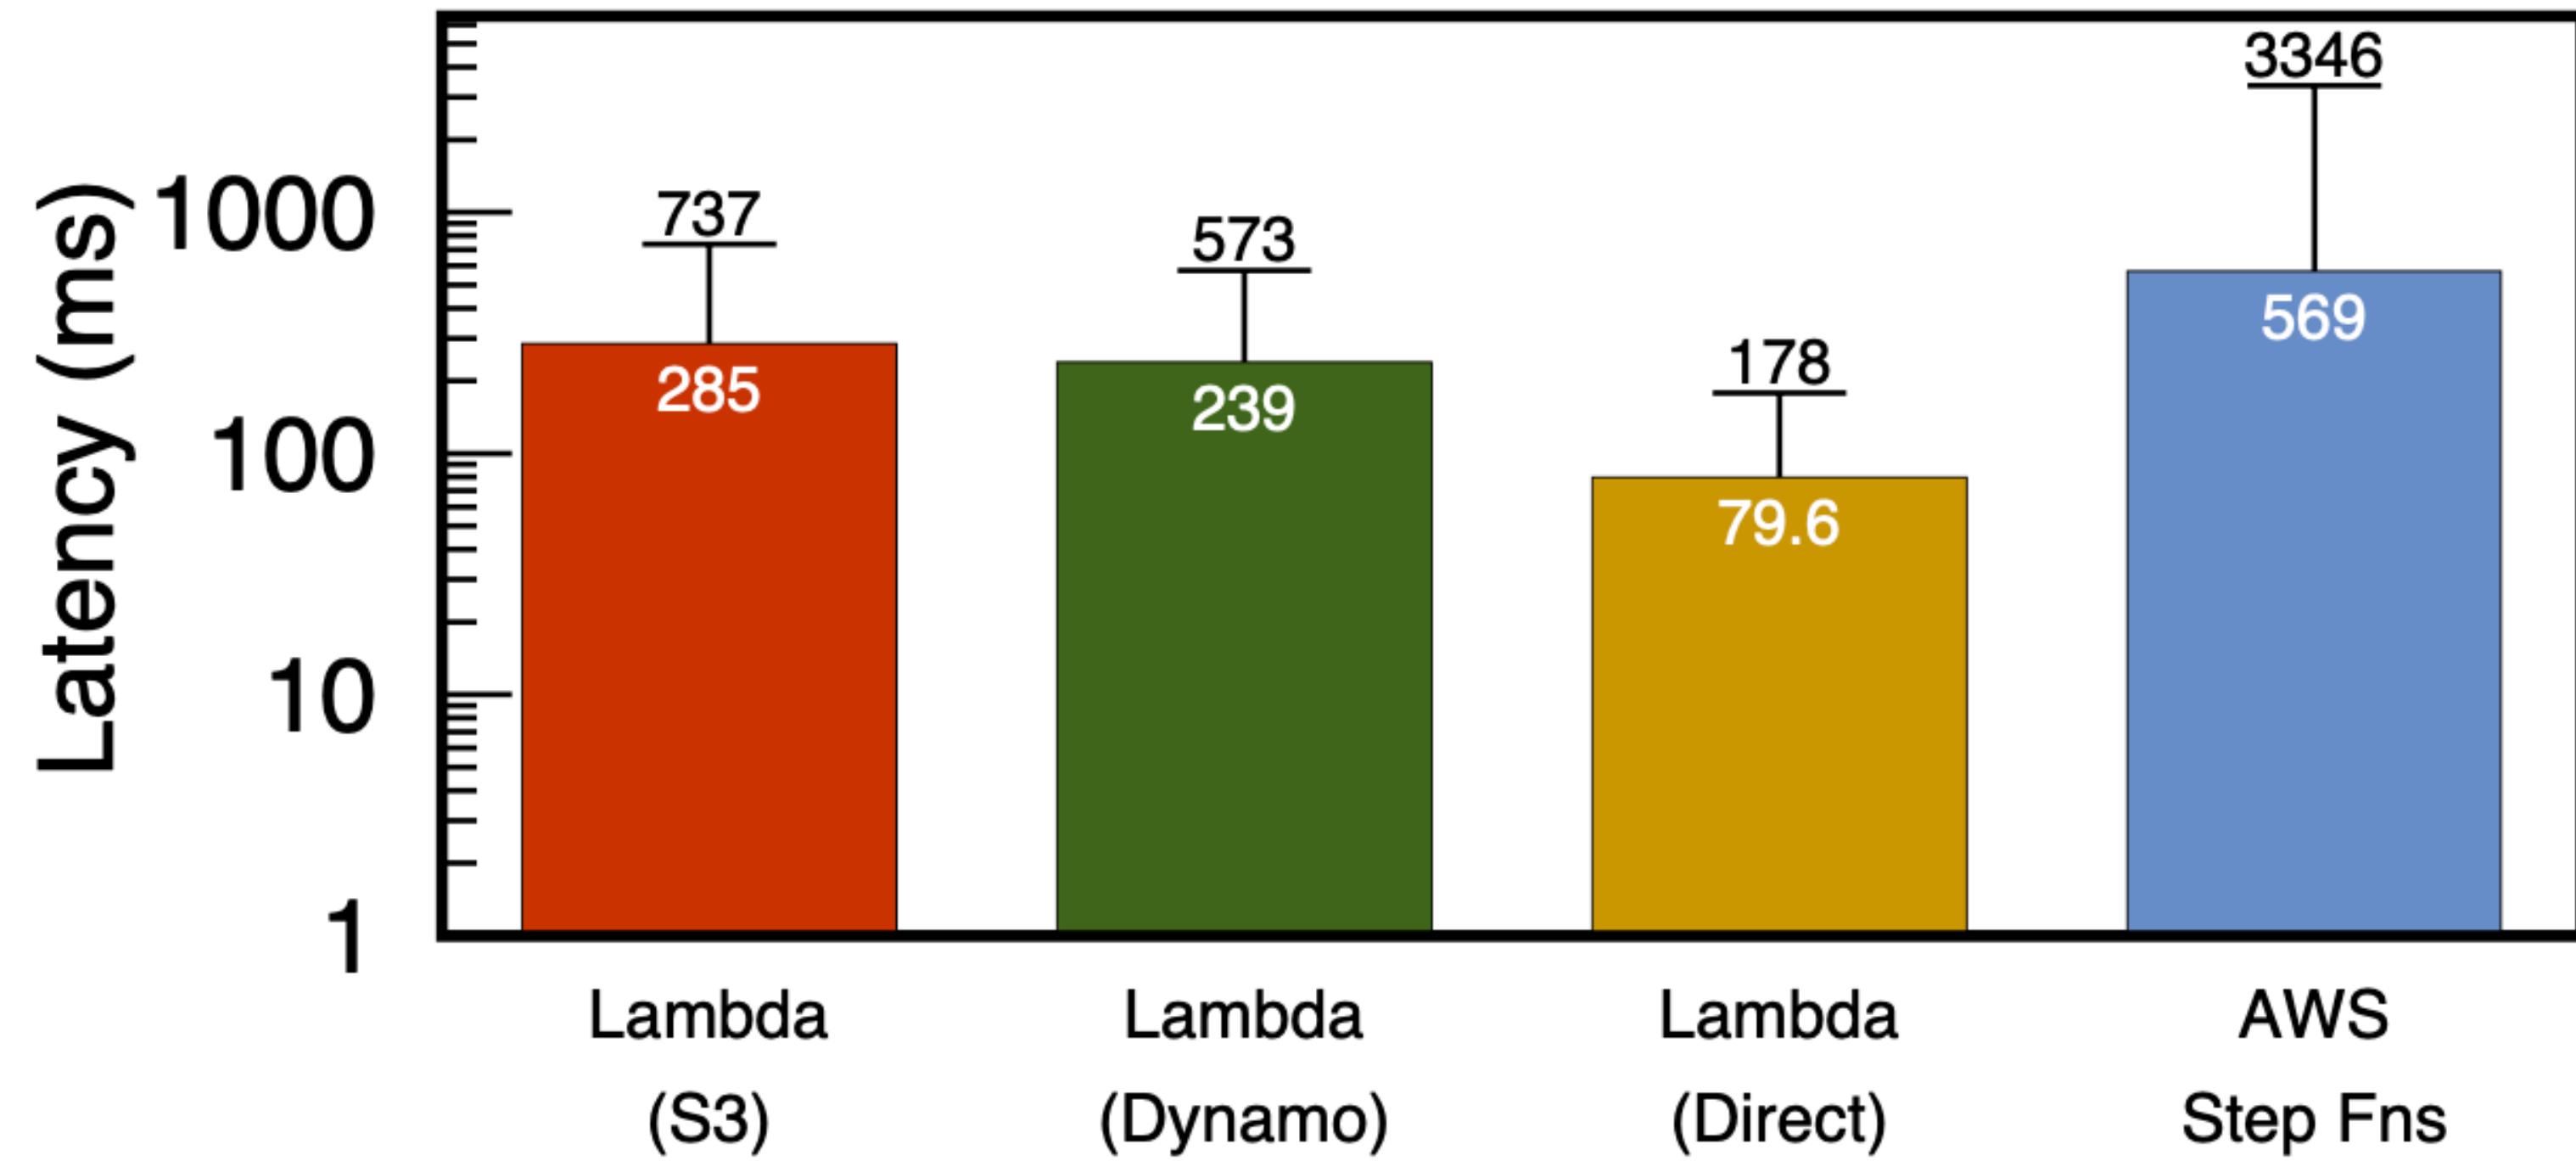
$f(g(x))$



Median and 99th percentile latencies for composing two arithmetic functions on AWS Lambda.

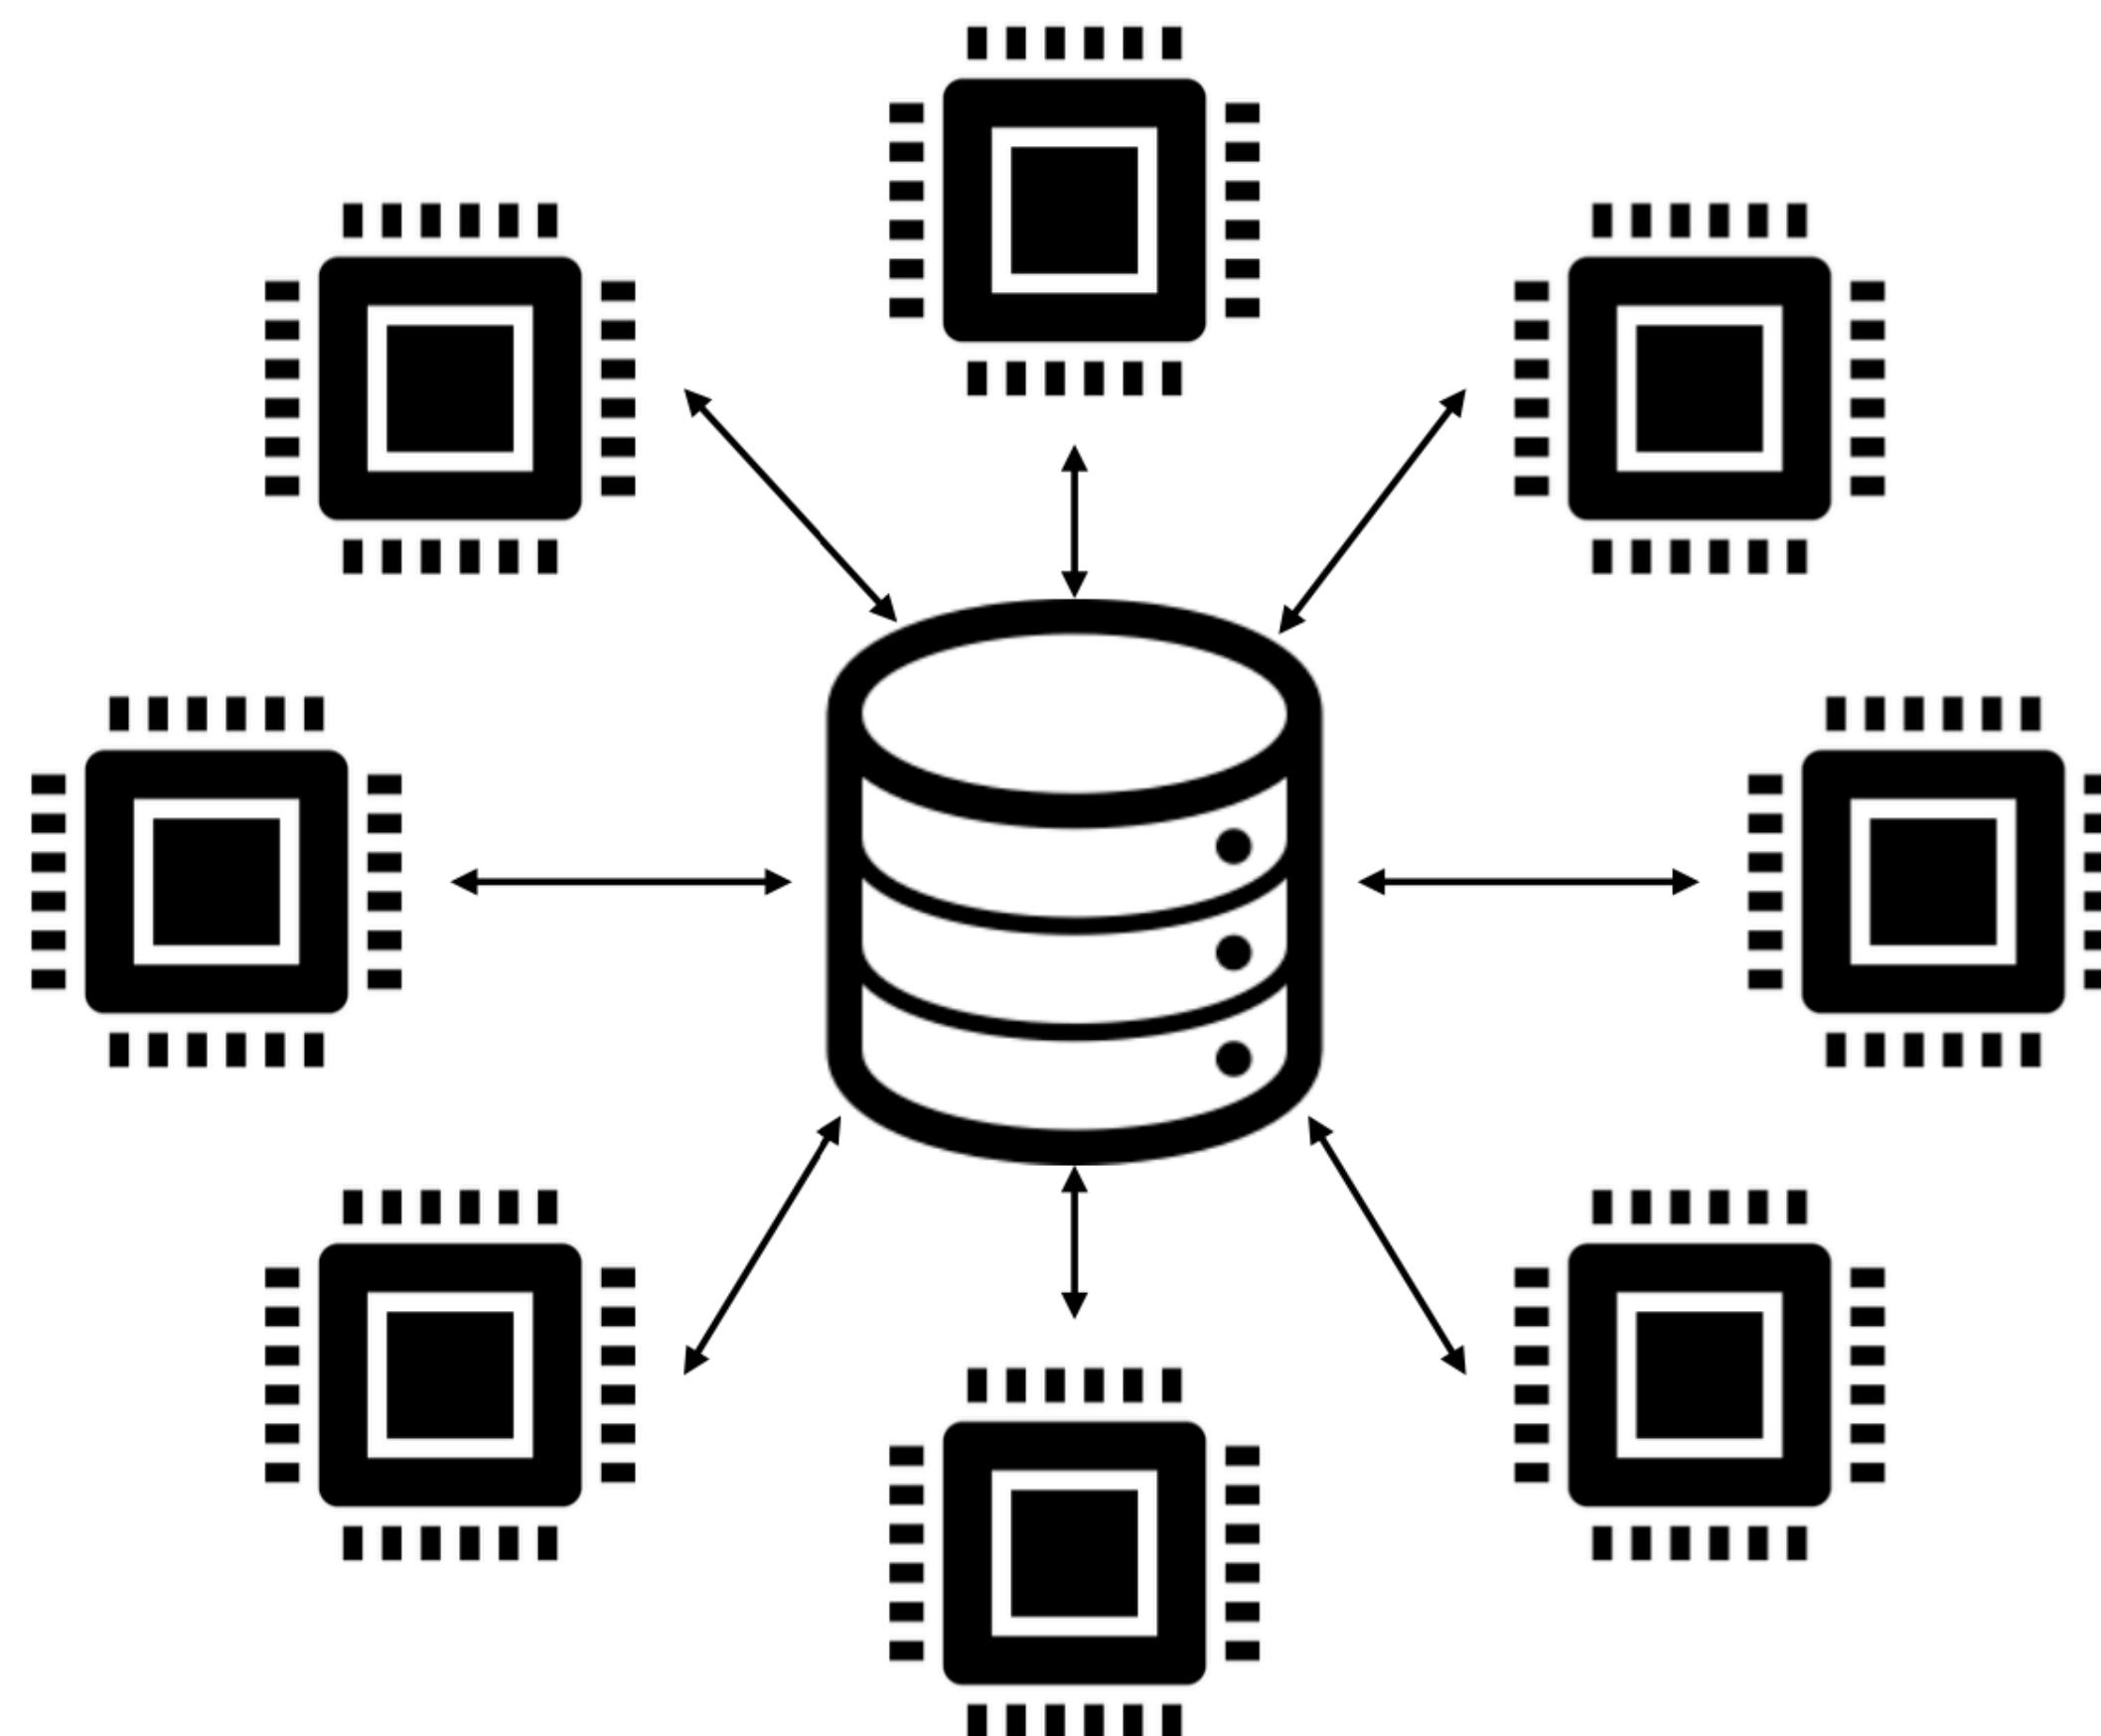
Even Functional Programming is Slow!

$f(g(x))$



Median and 99th percentile latencies for composing two arithmetic functions on AWS Lambda.

Shared Mutable State

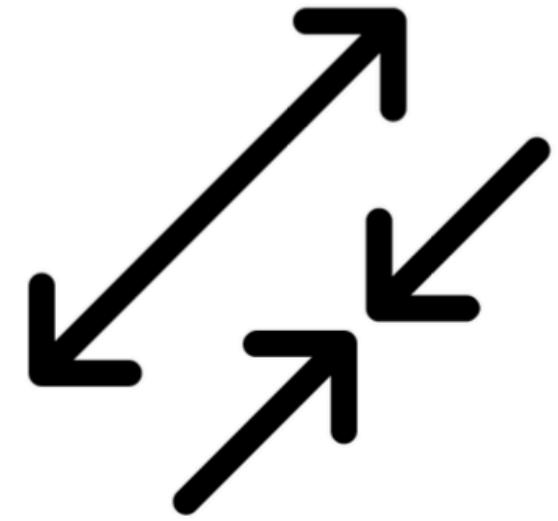


Shared Mutable State





Serverless Storage



Autoscaling

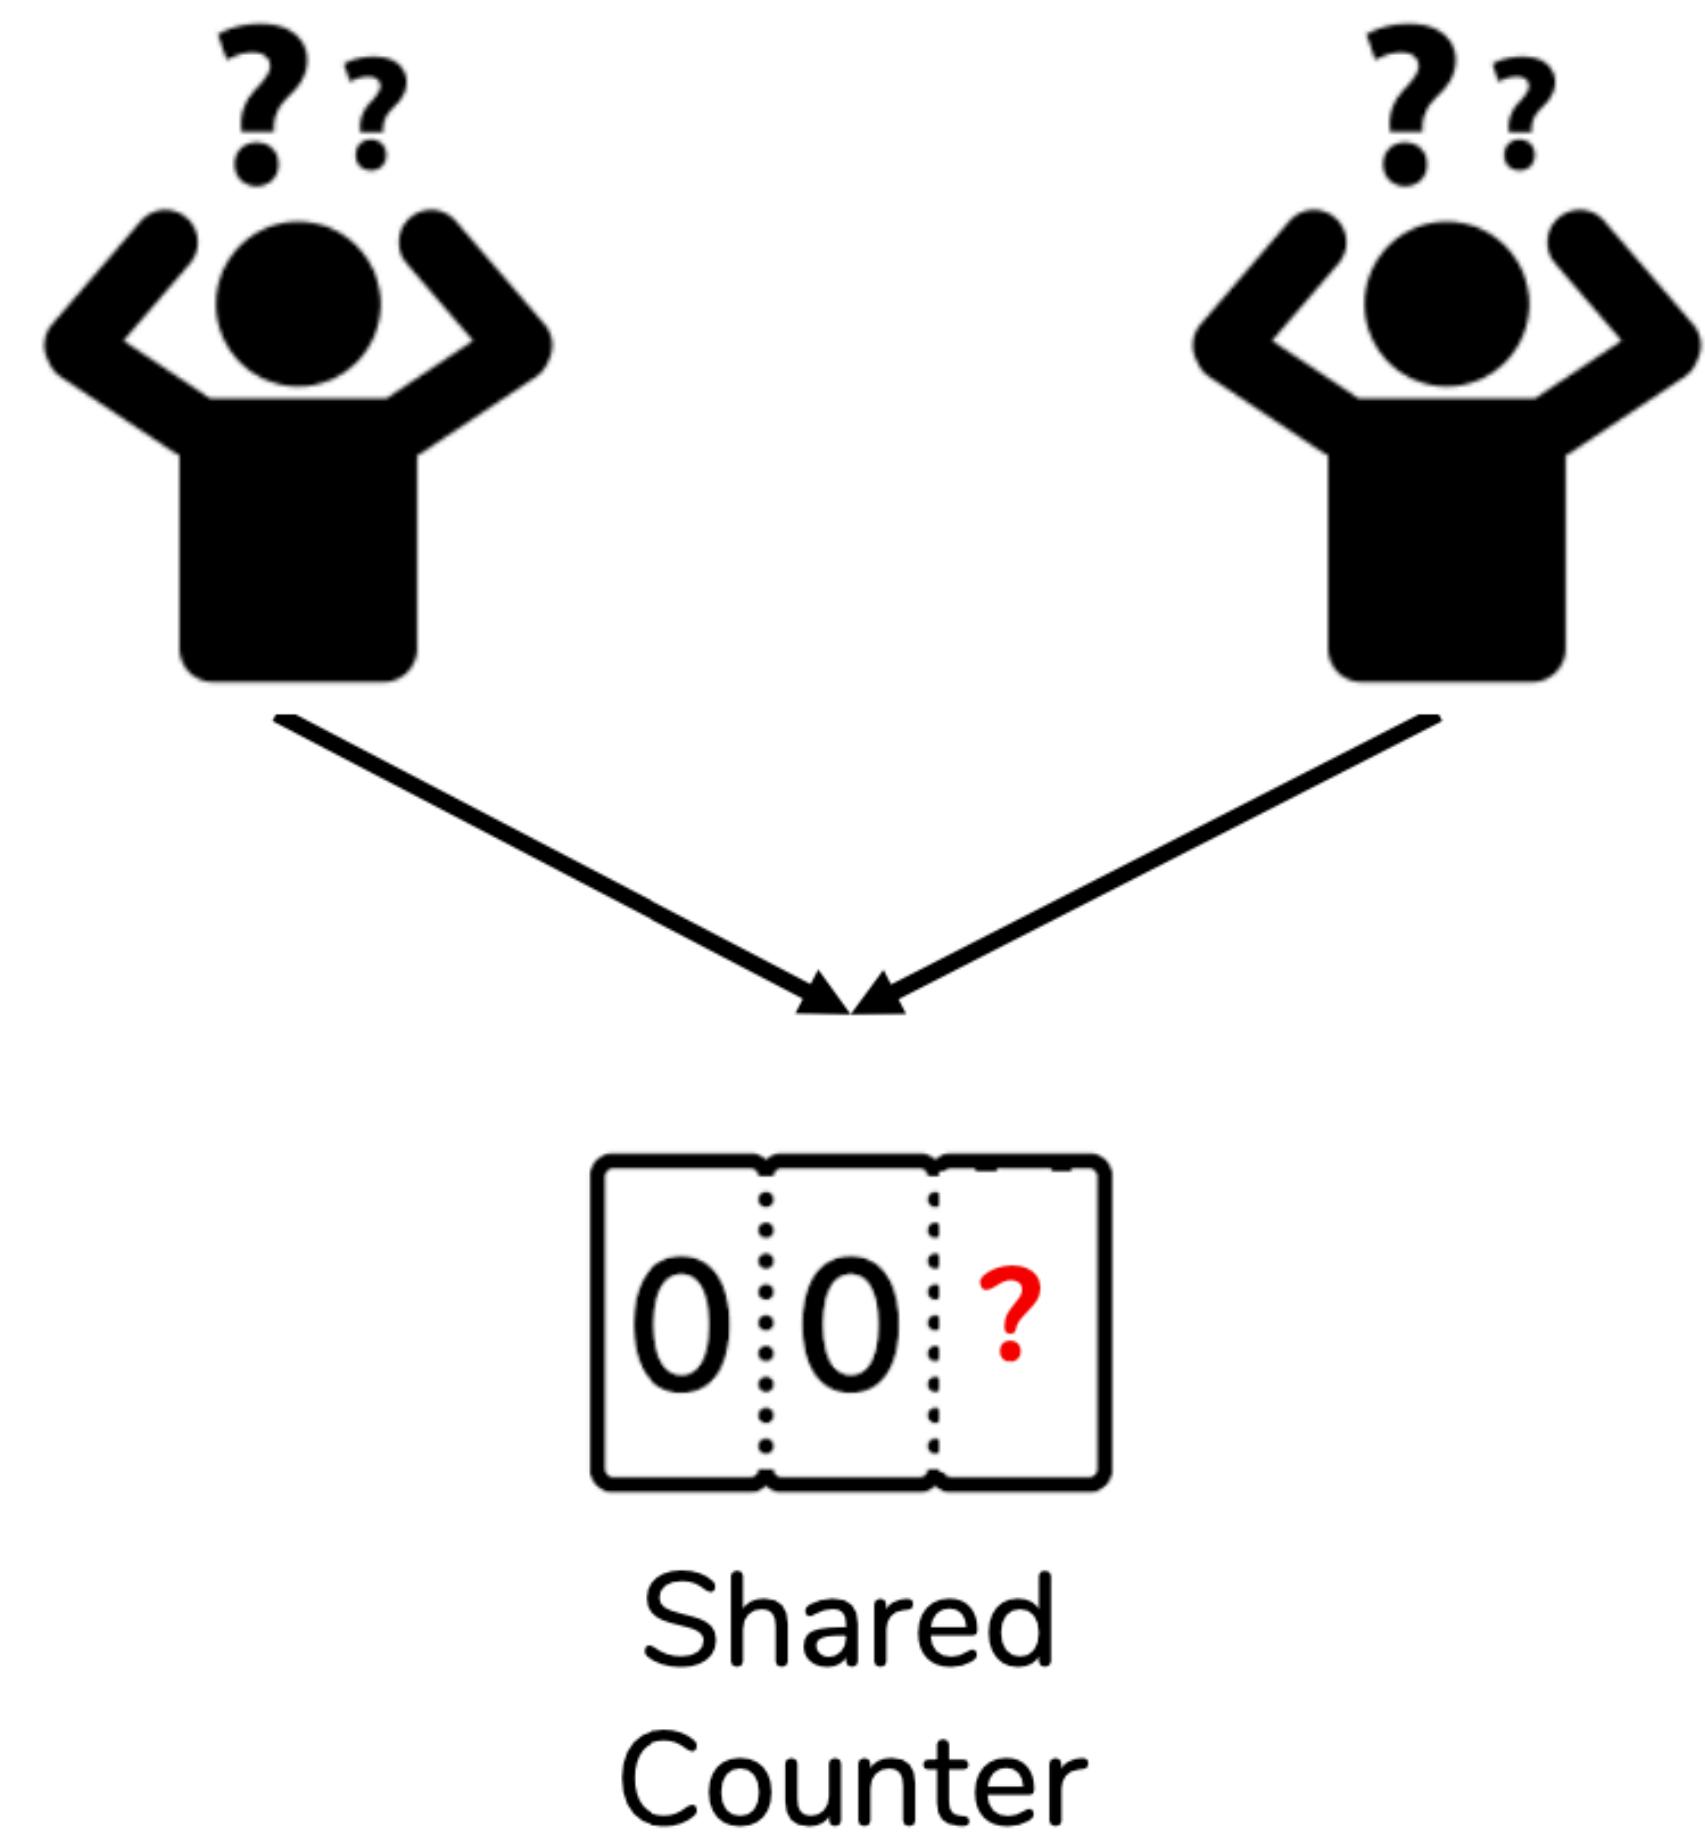


TRADEOFF!



Low Latency

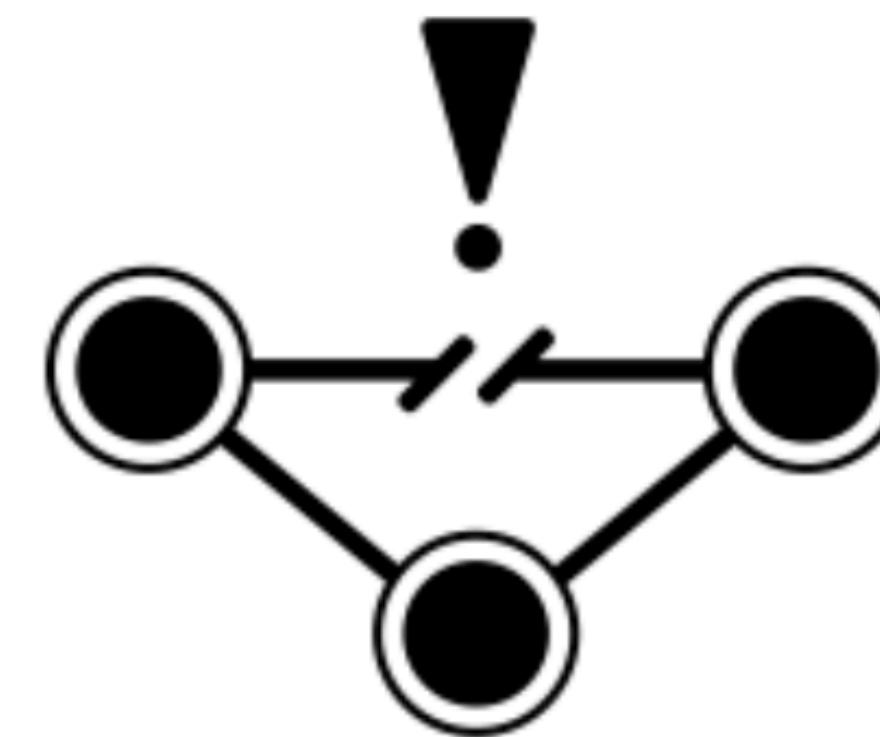
(In)Consistency Guarantees



No Inbound Network Connections

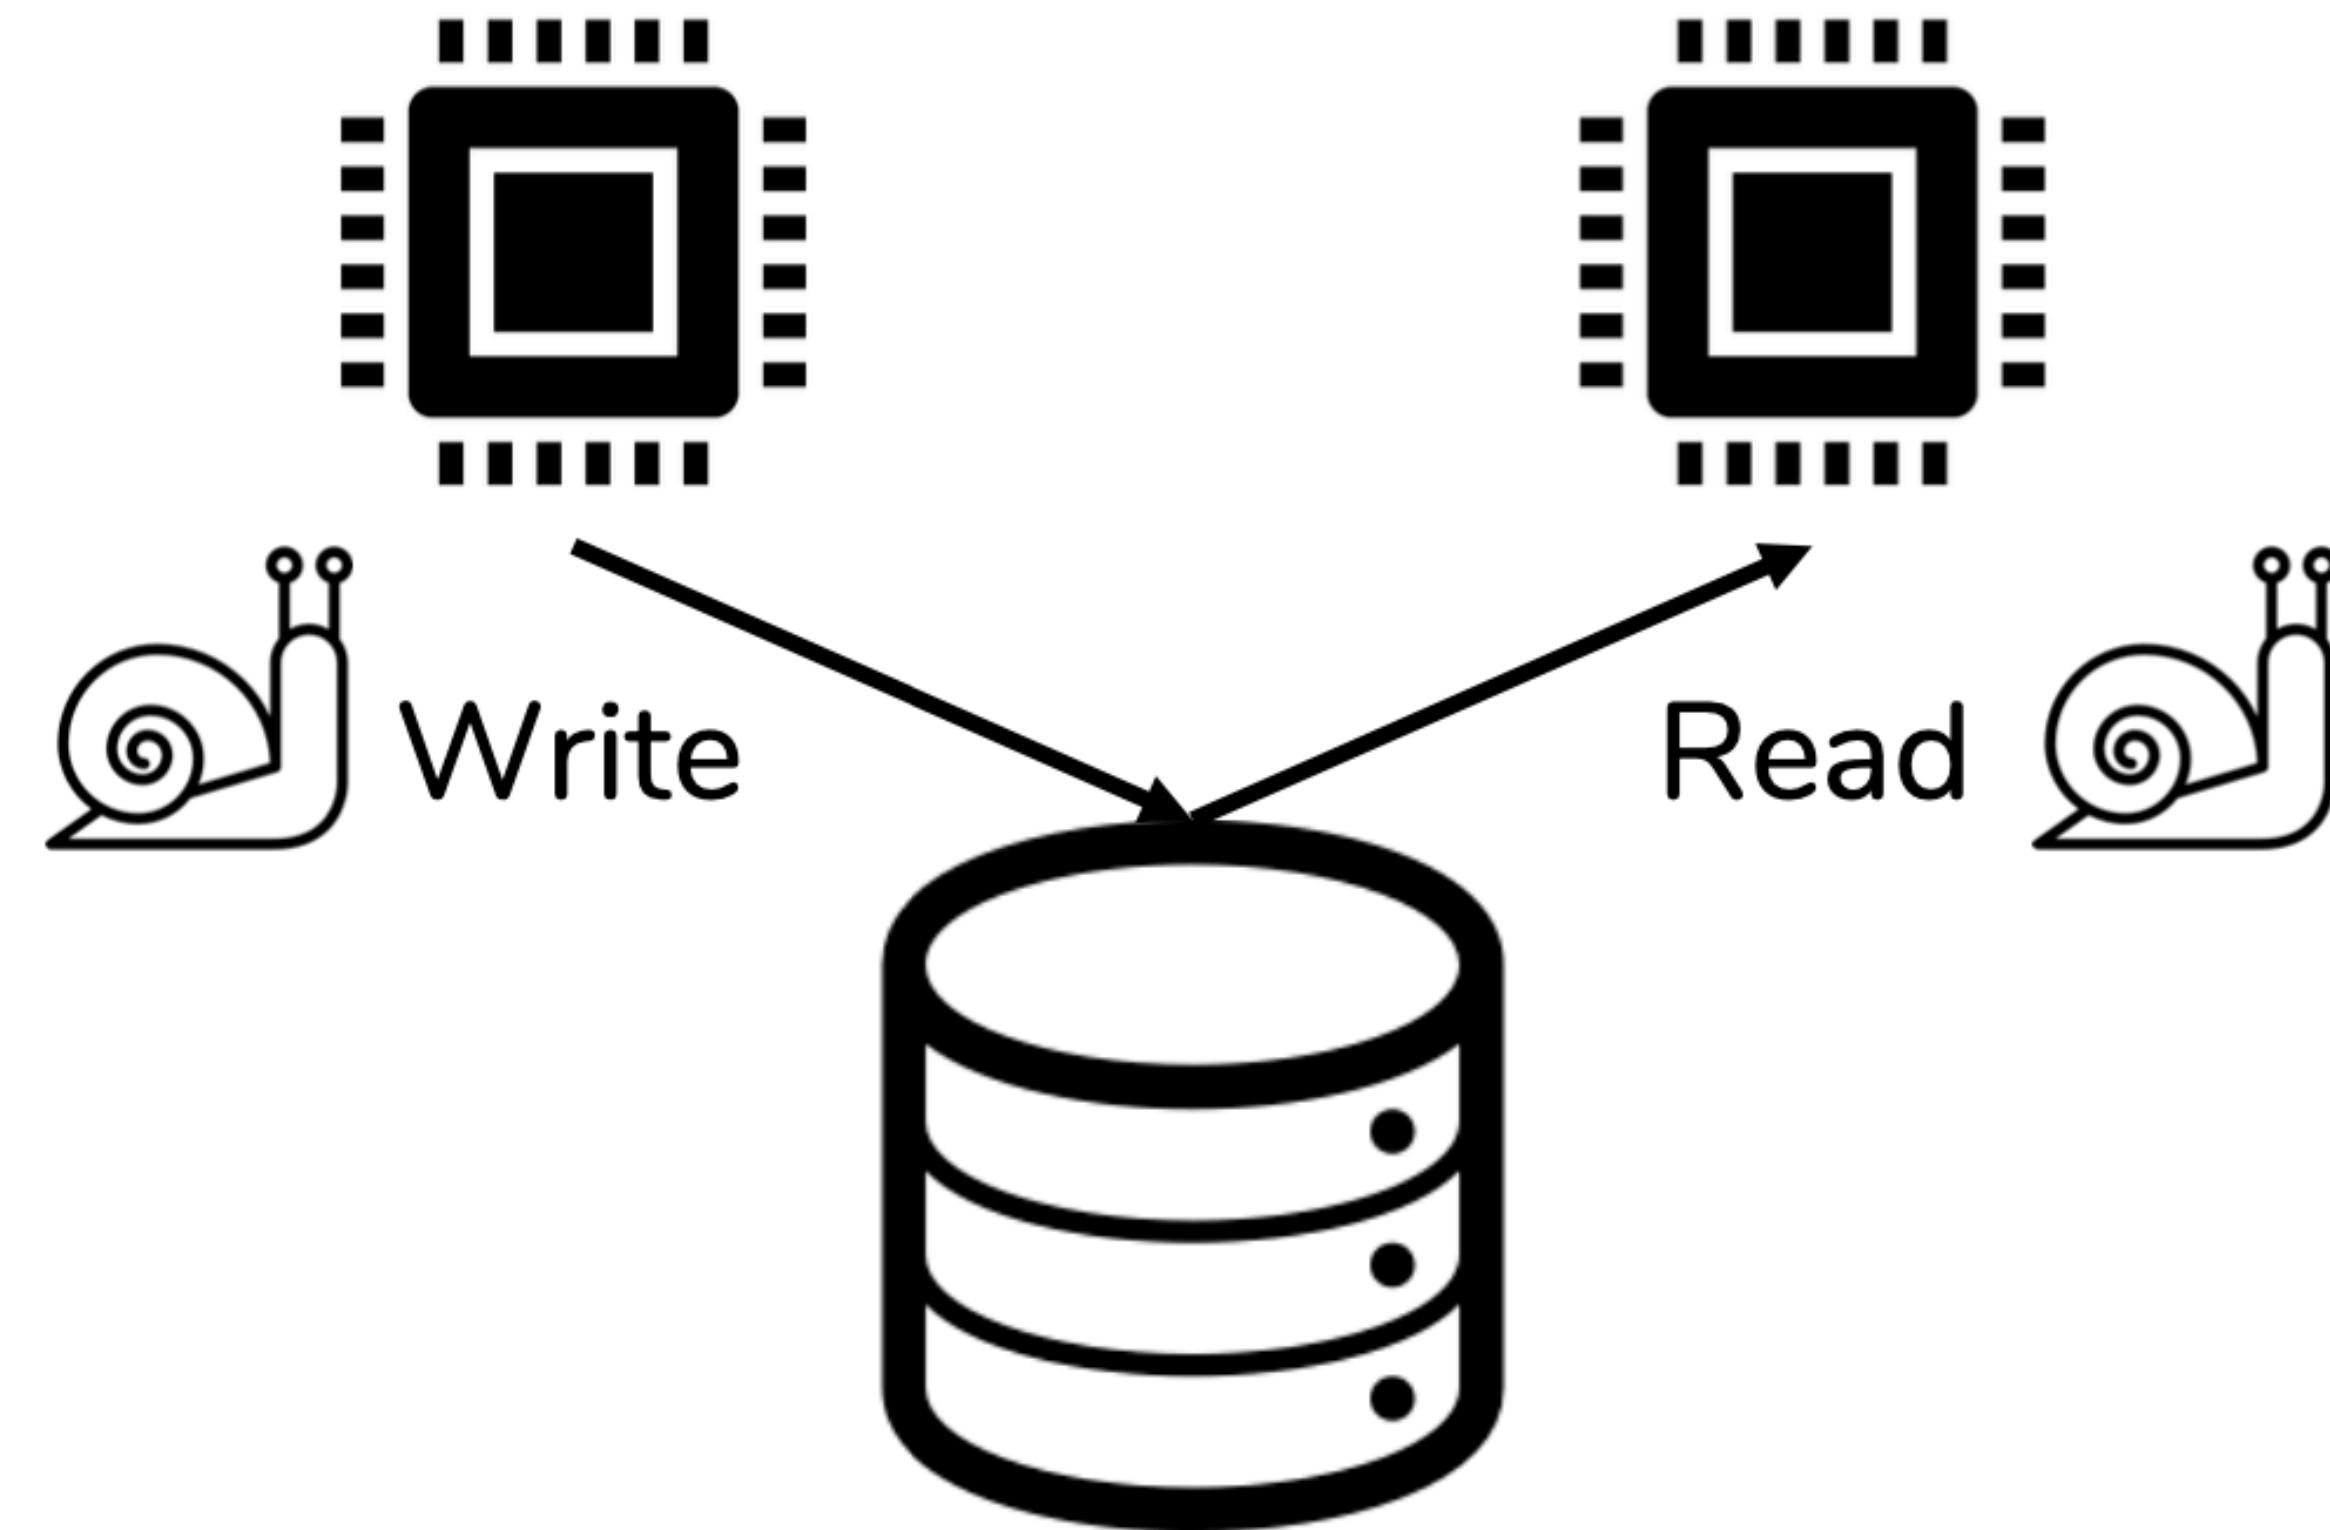


Enables Process
Migration



Easy Fault
Tolerance

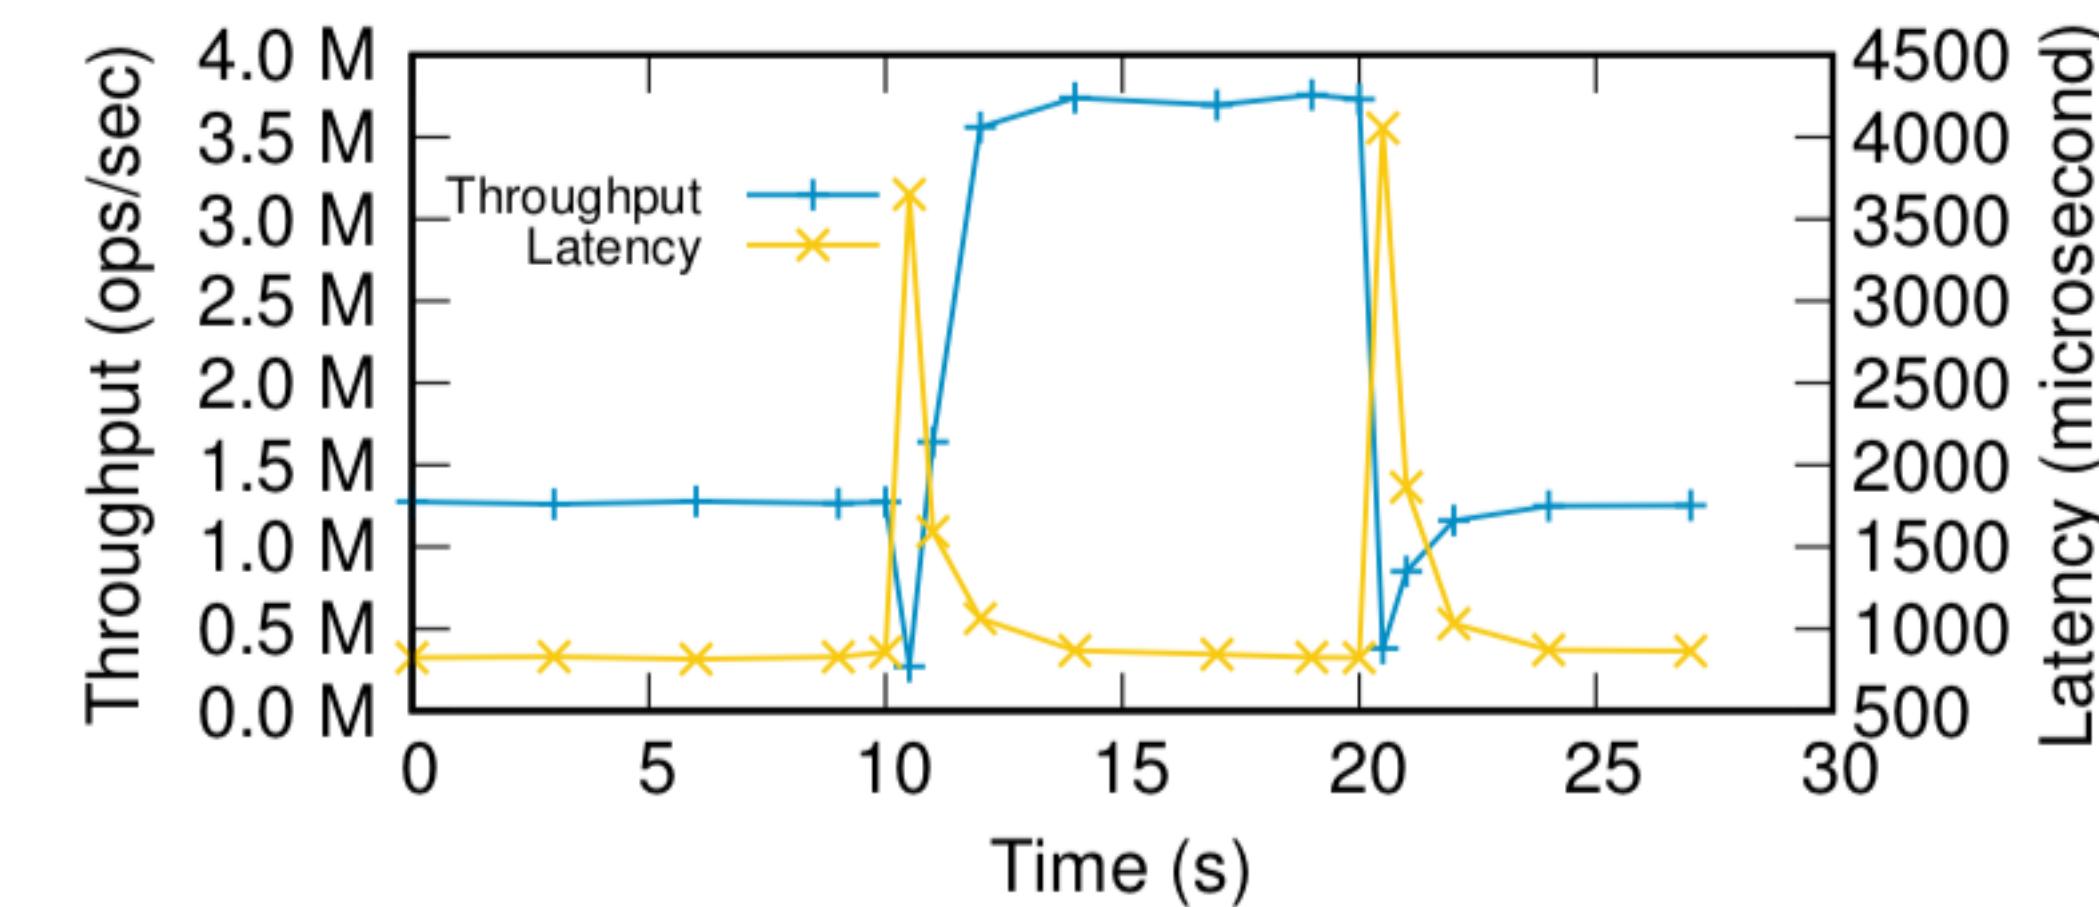
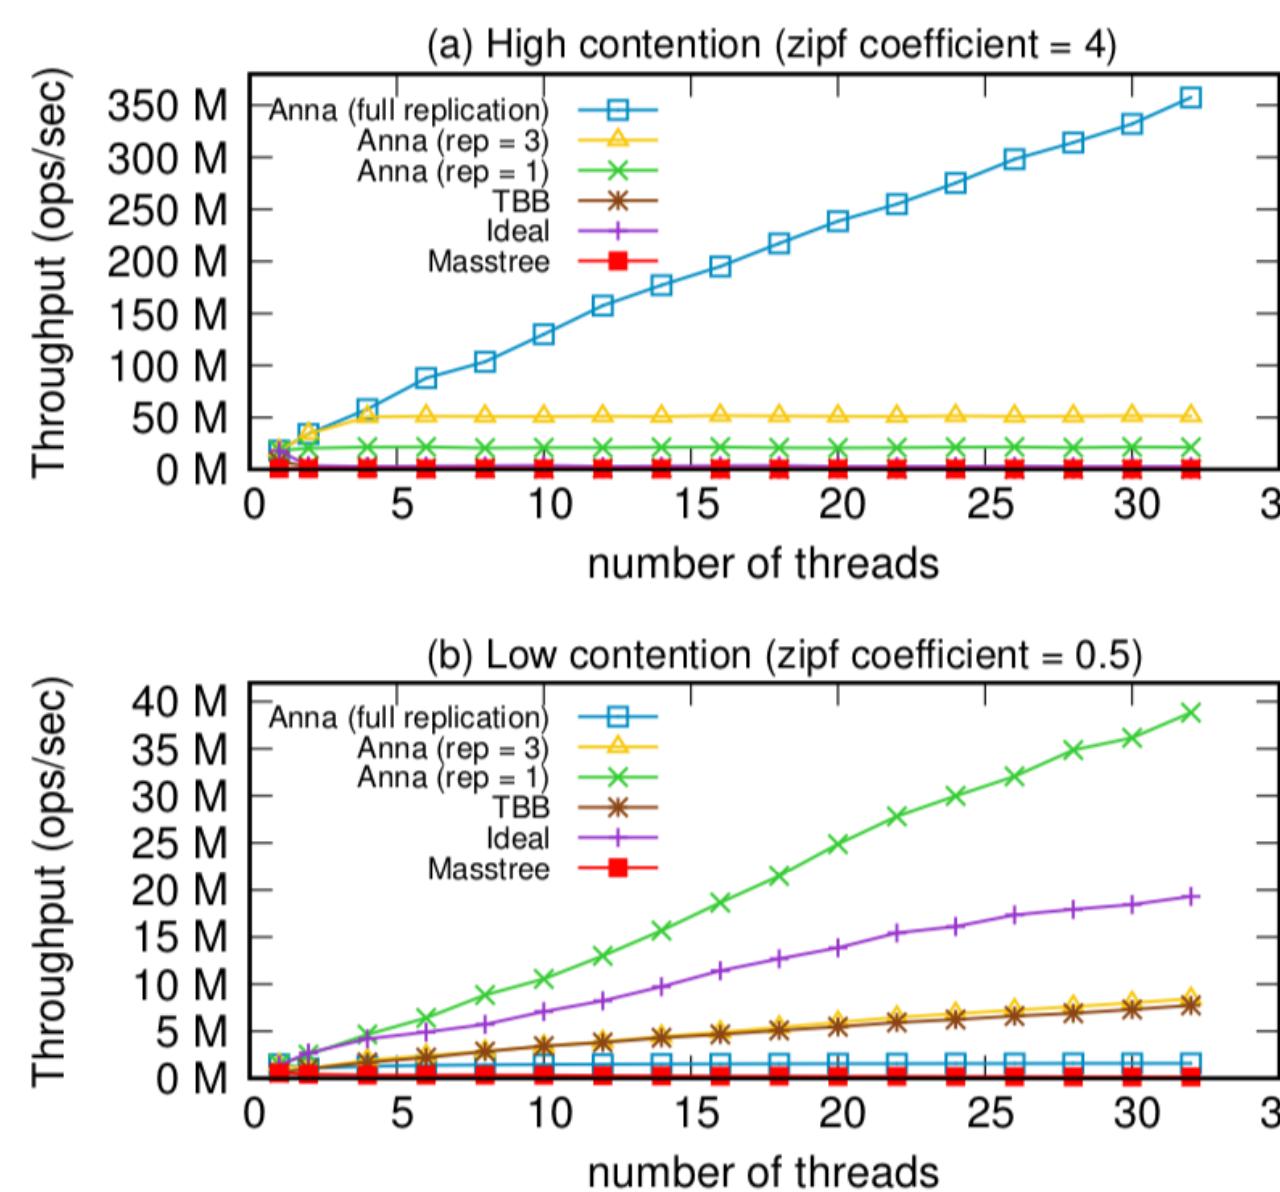
Indirect Communication



Cloudburst: Stateful FaaS

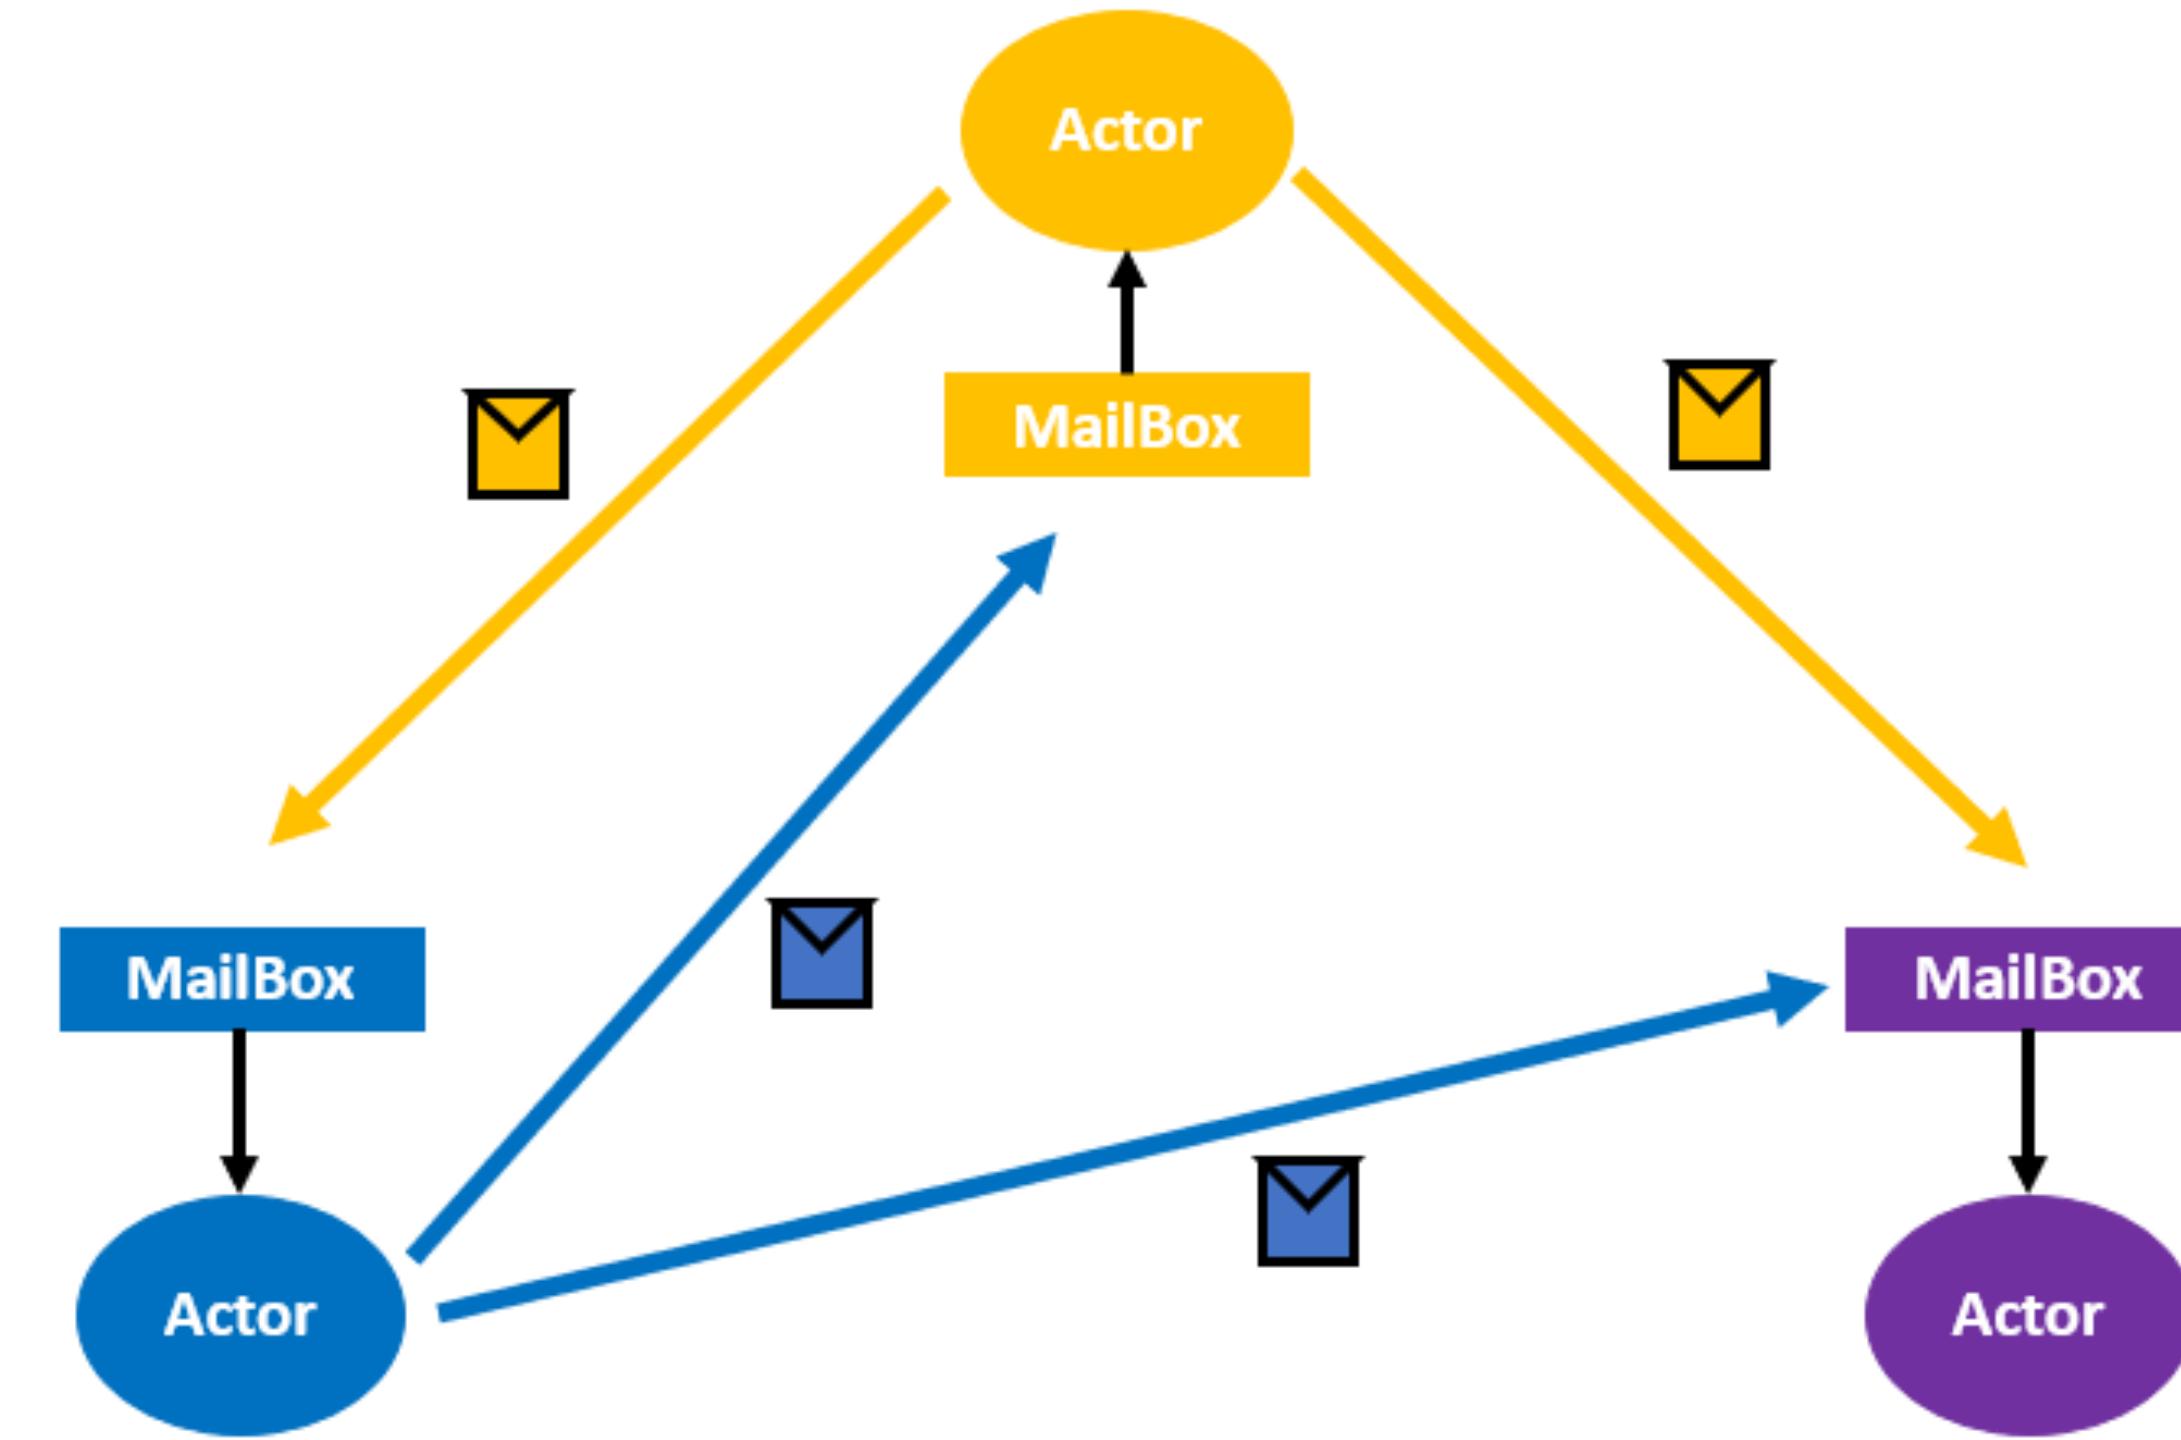
Background: Anna Overview

- High performance across orders of magnitude in scale
 - ✓ 10x faster than Redis/Cassandra in a geo-distributed deployment
- Autoscaling & cost-efficient
 - ✓ 500x faster than Amazon DynamoDB for the same cost



Coordination-free consistency. No atomics, no locks, no waiting ever!

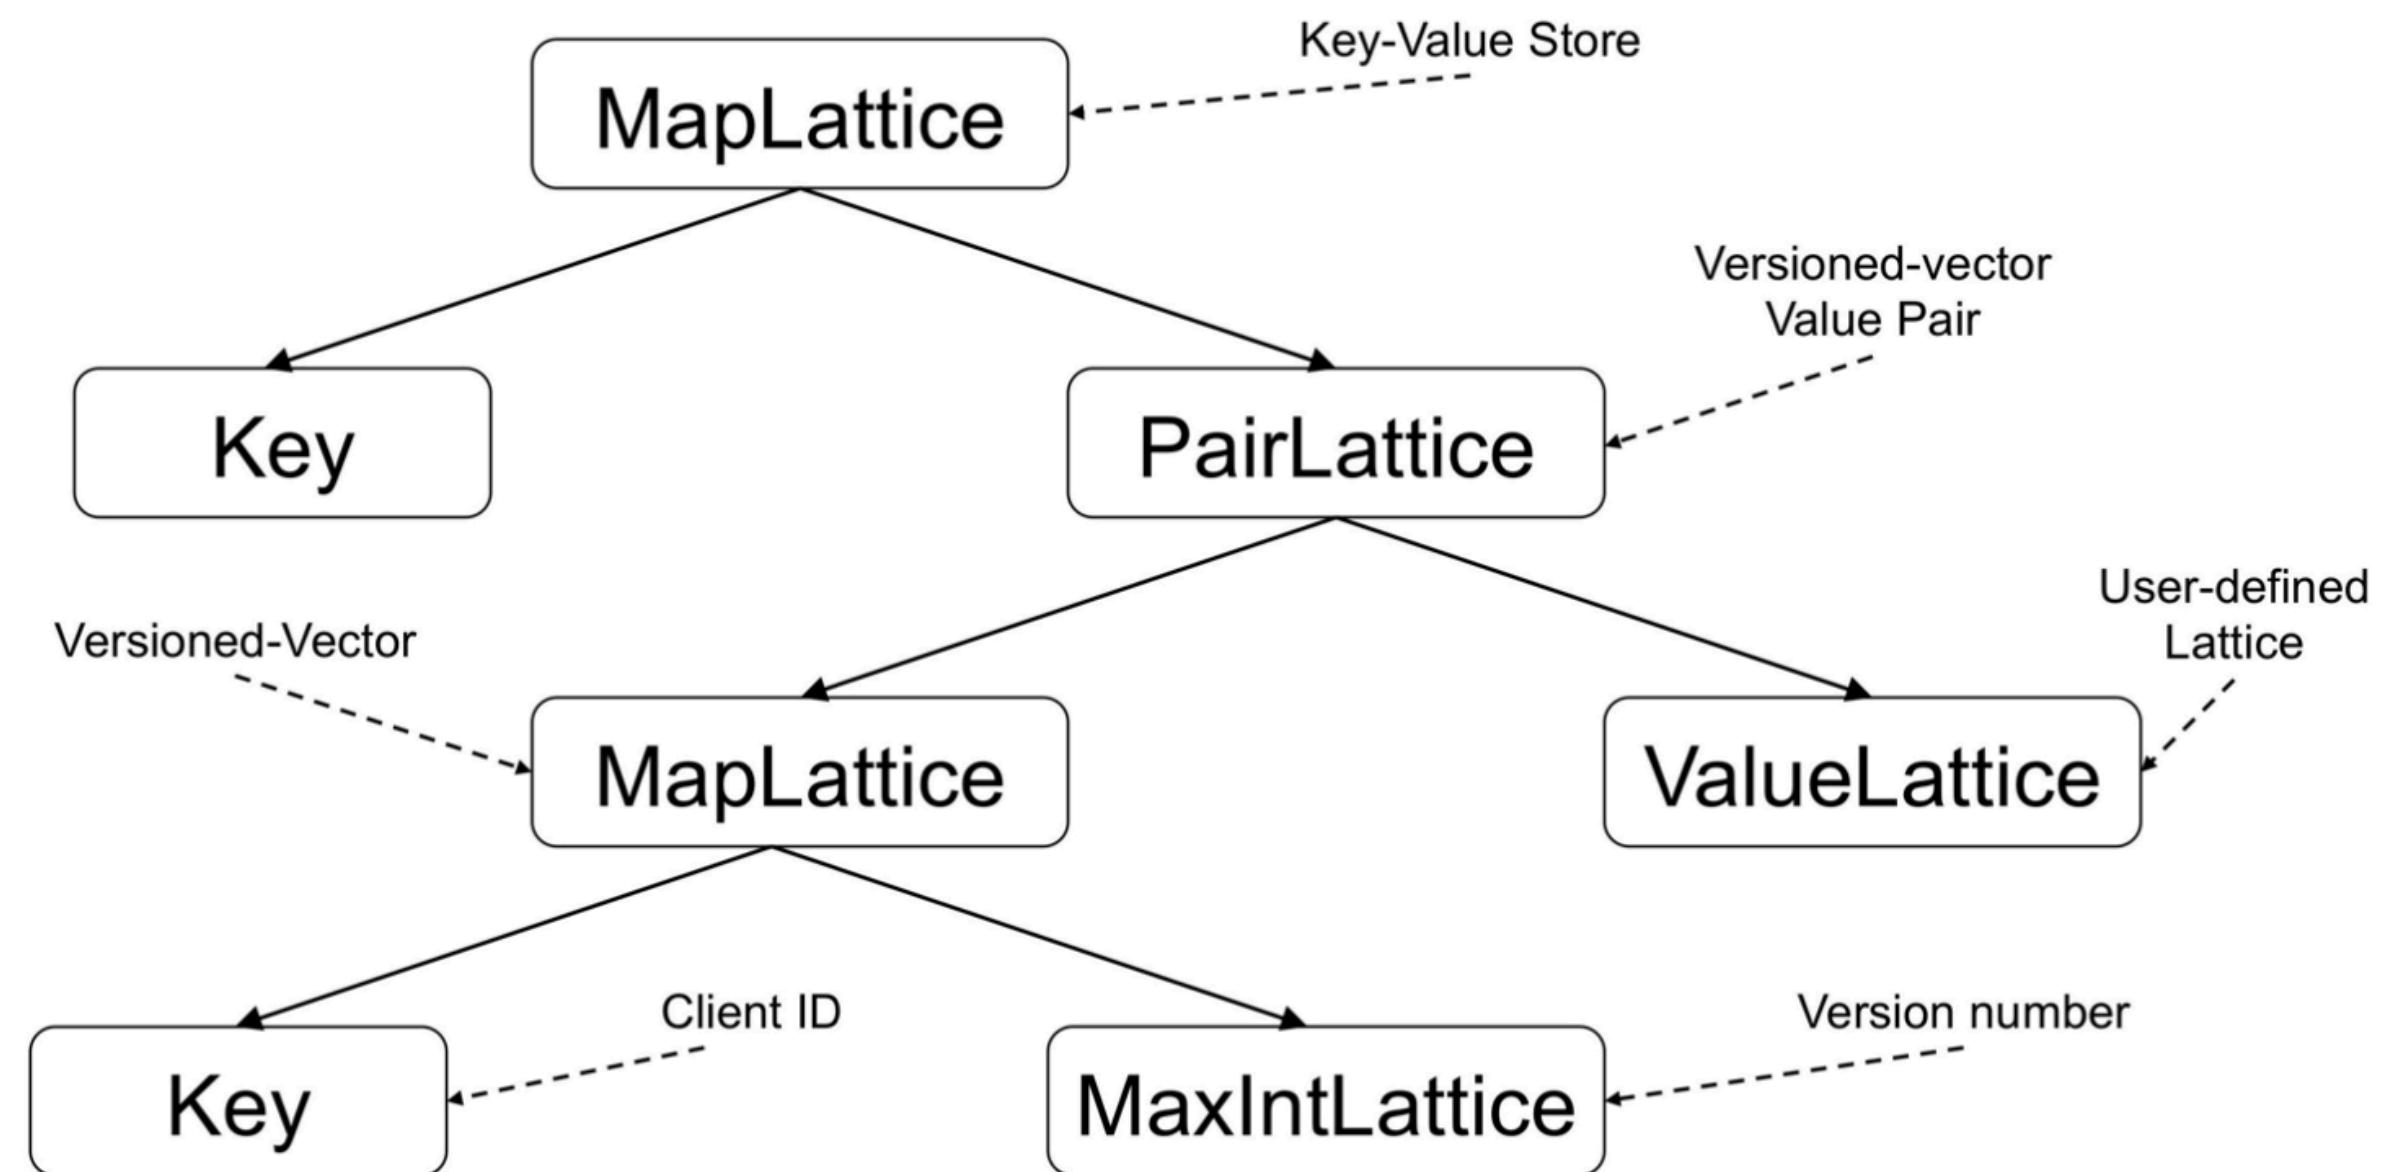
Background: Actor Model



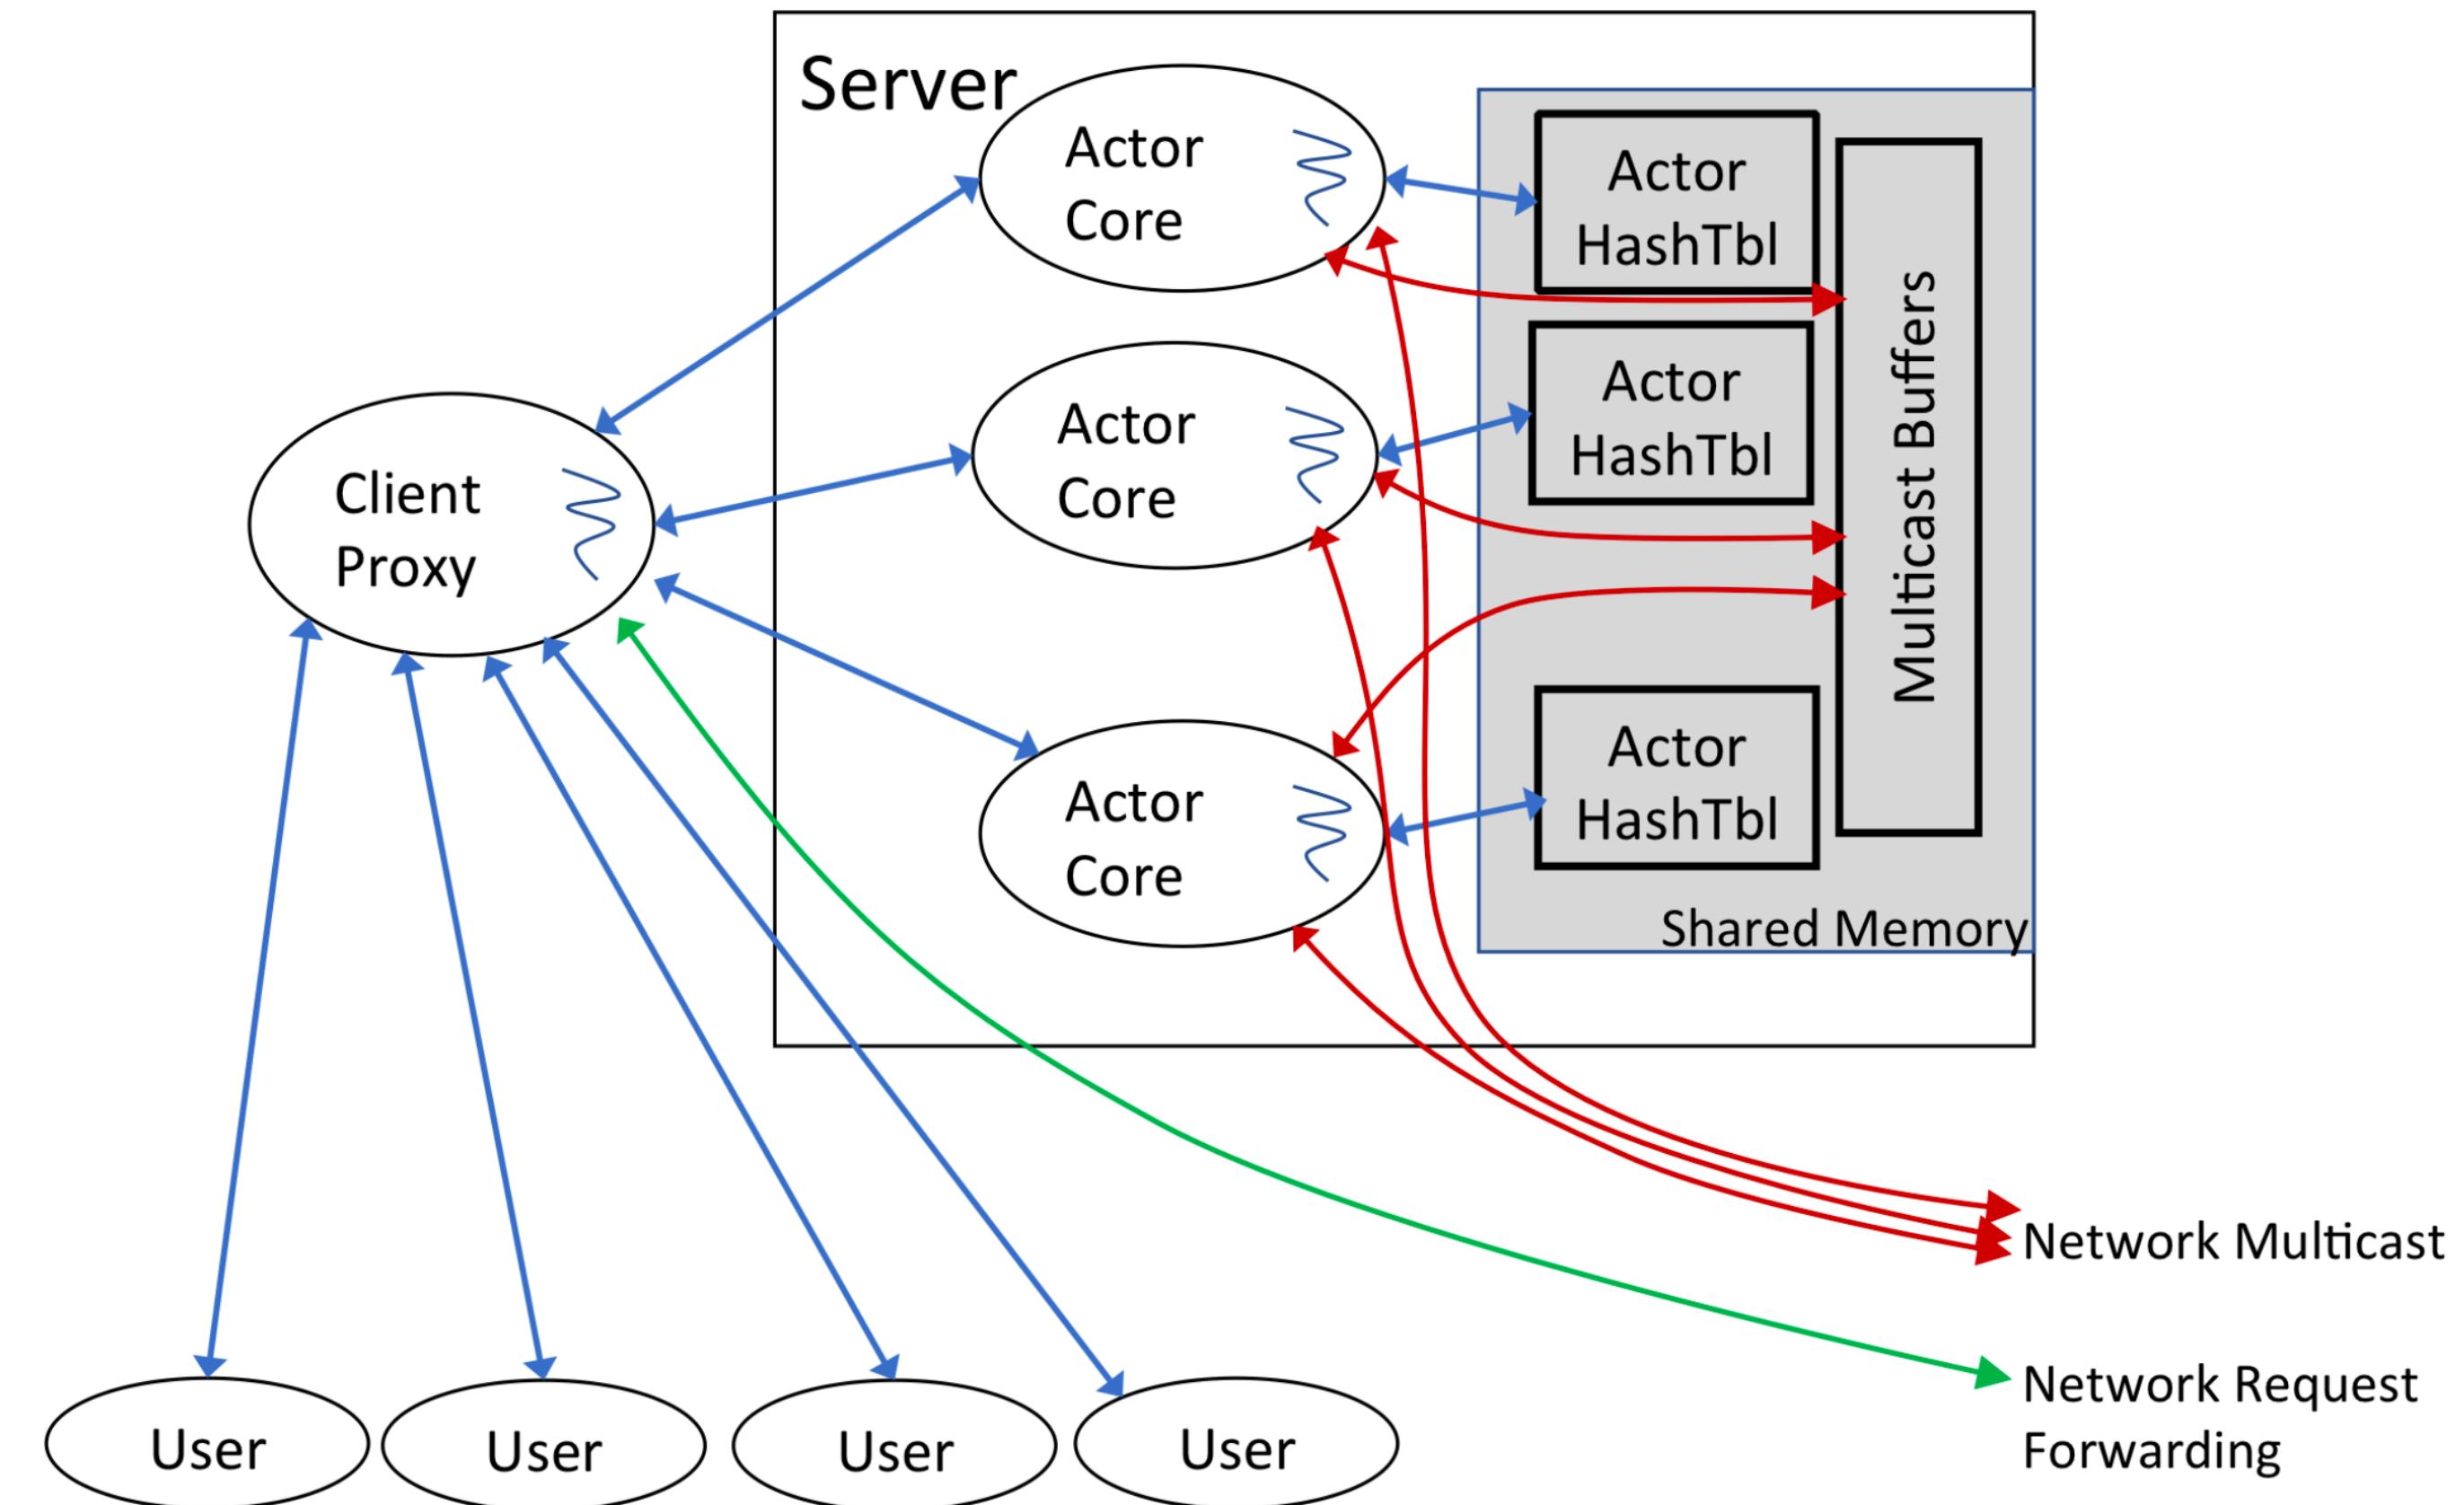
Background: Lattice

- ✓ Associativity: $(a+b)+c = a+(b+c)$
- ✓ Commutativity: $a+b = b+a$
- ✓ Idempotence: $a+a = a$

Background: Causal Lattice

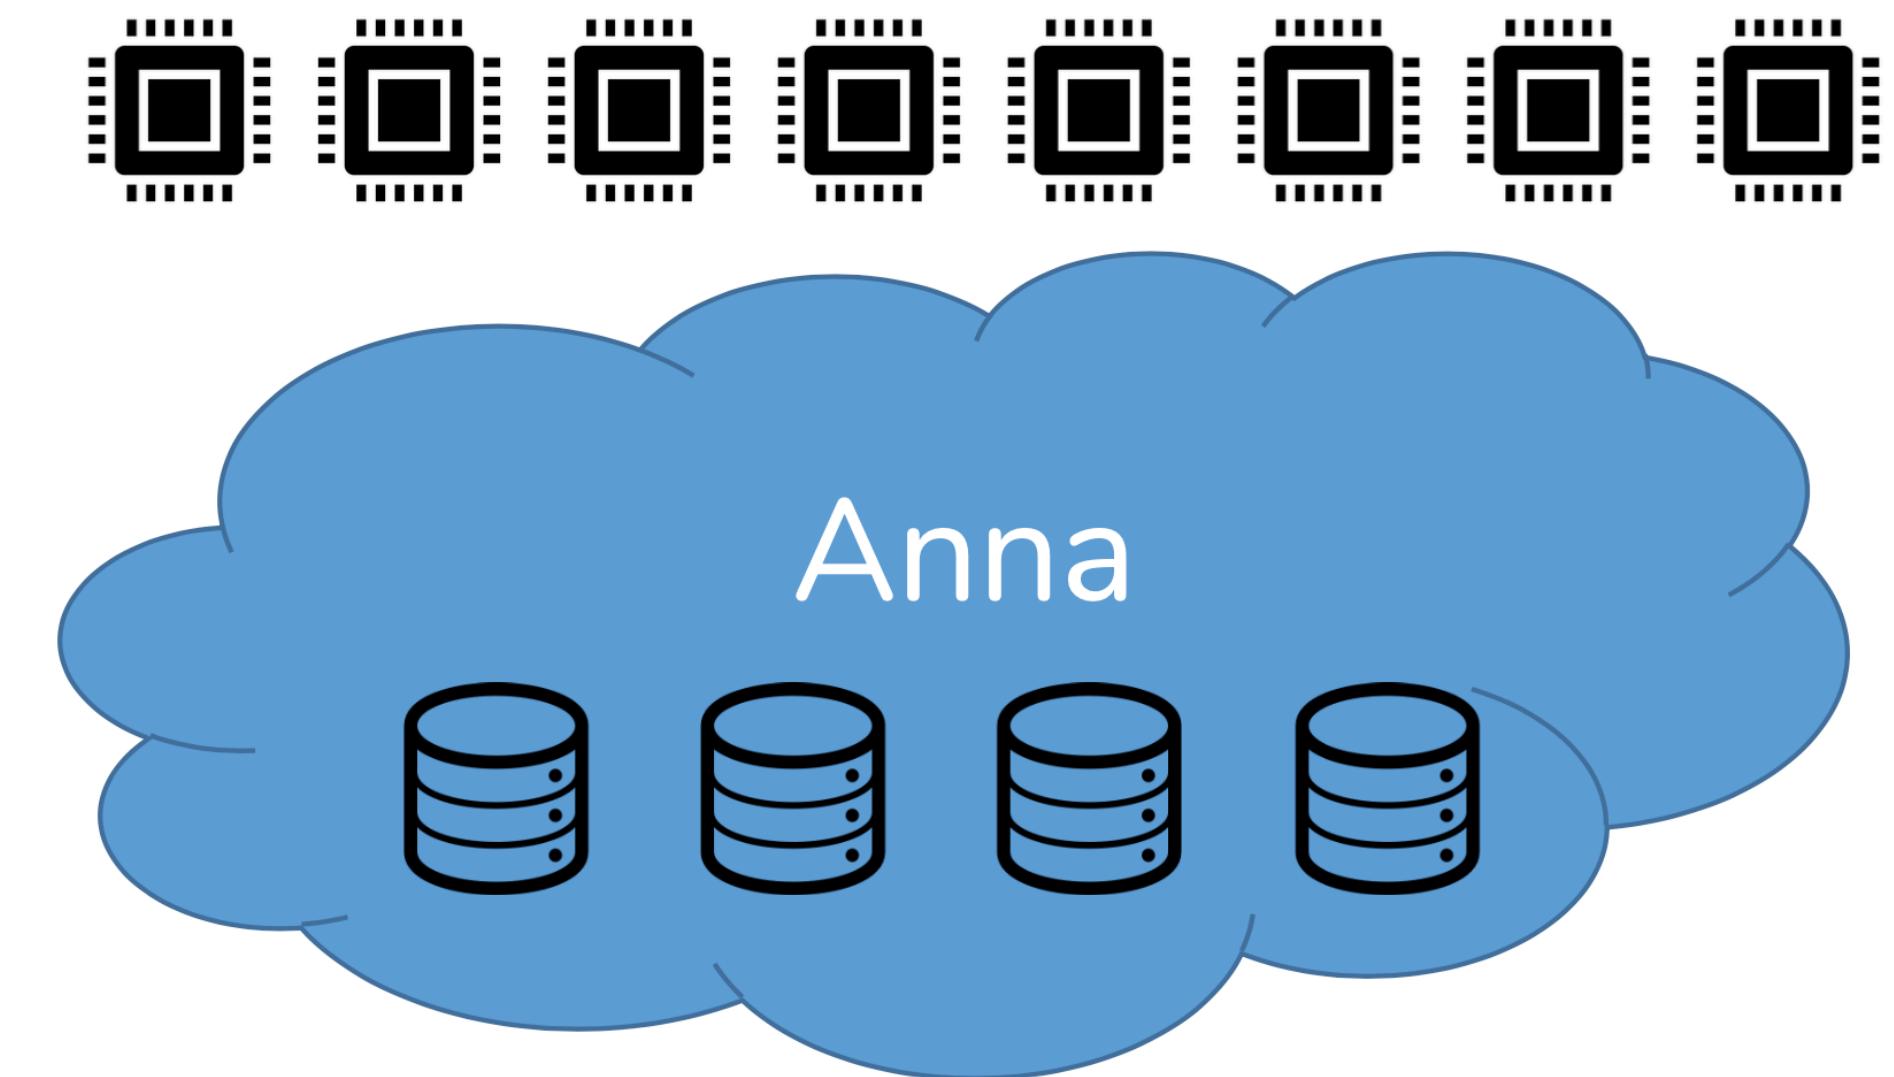


Background: Anna Architecture

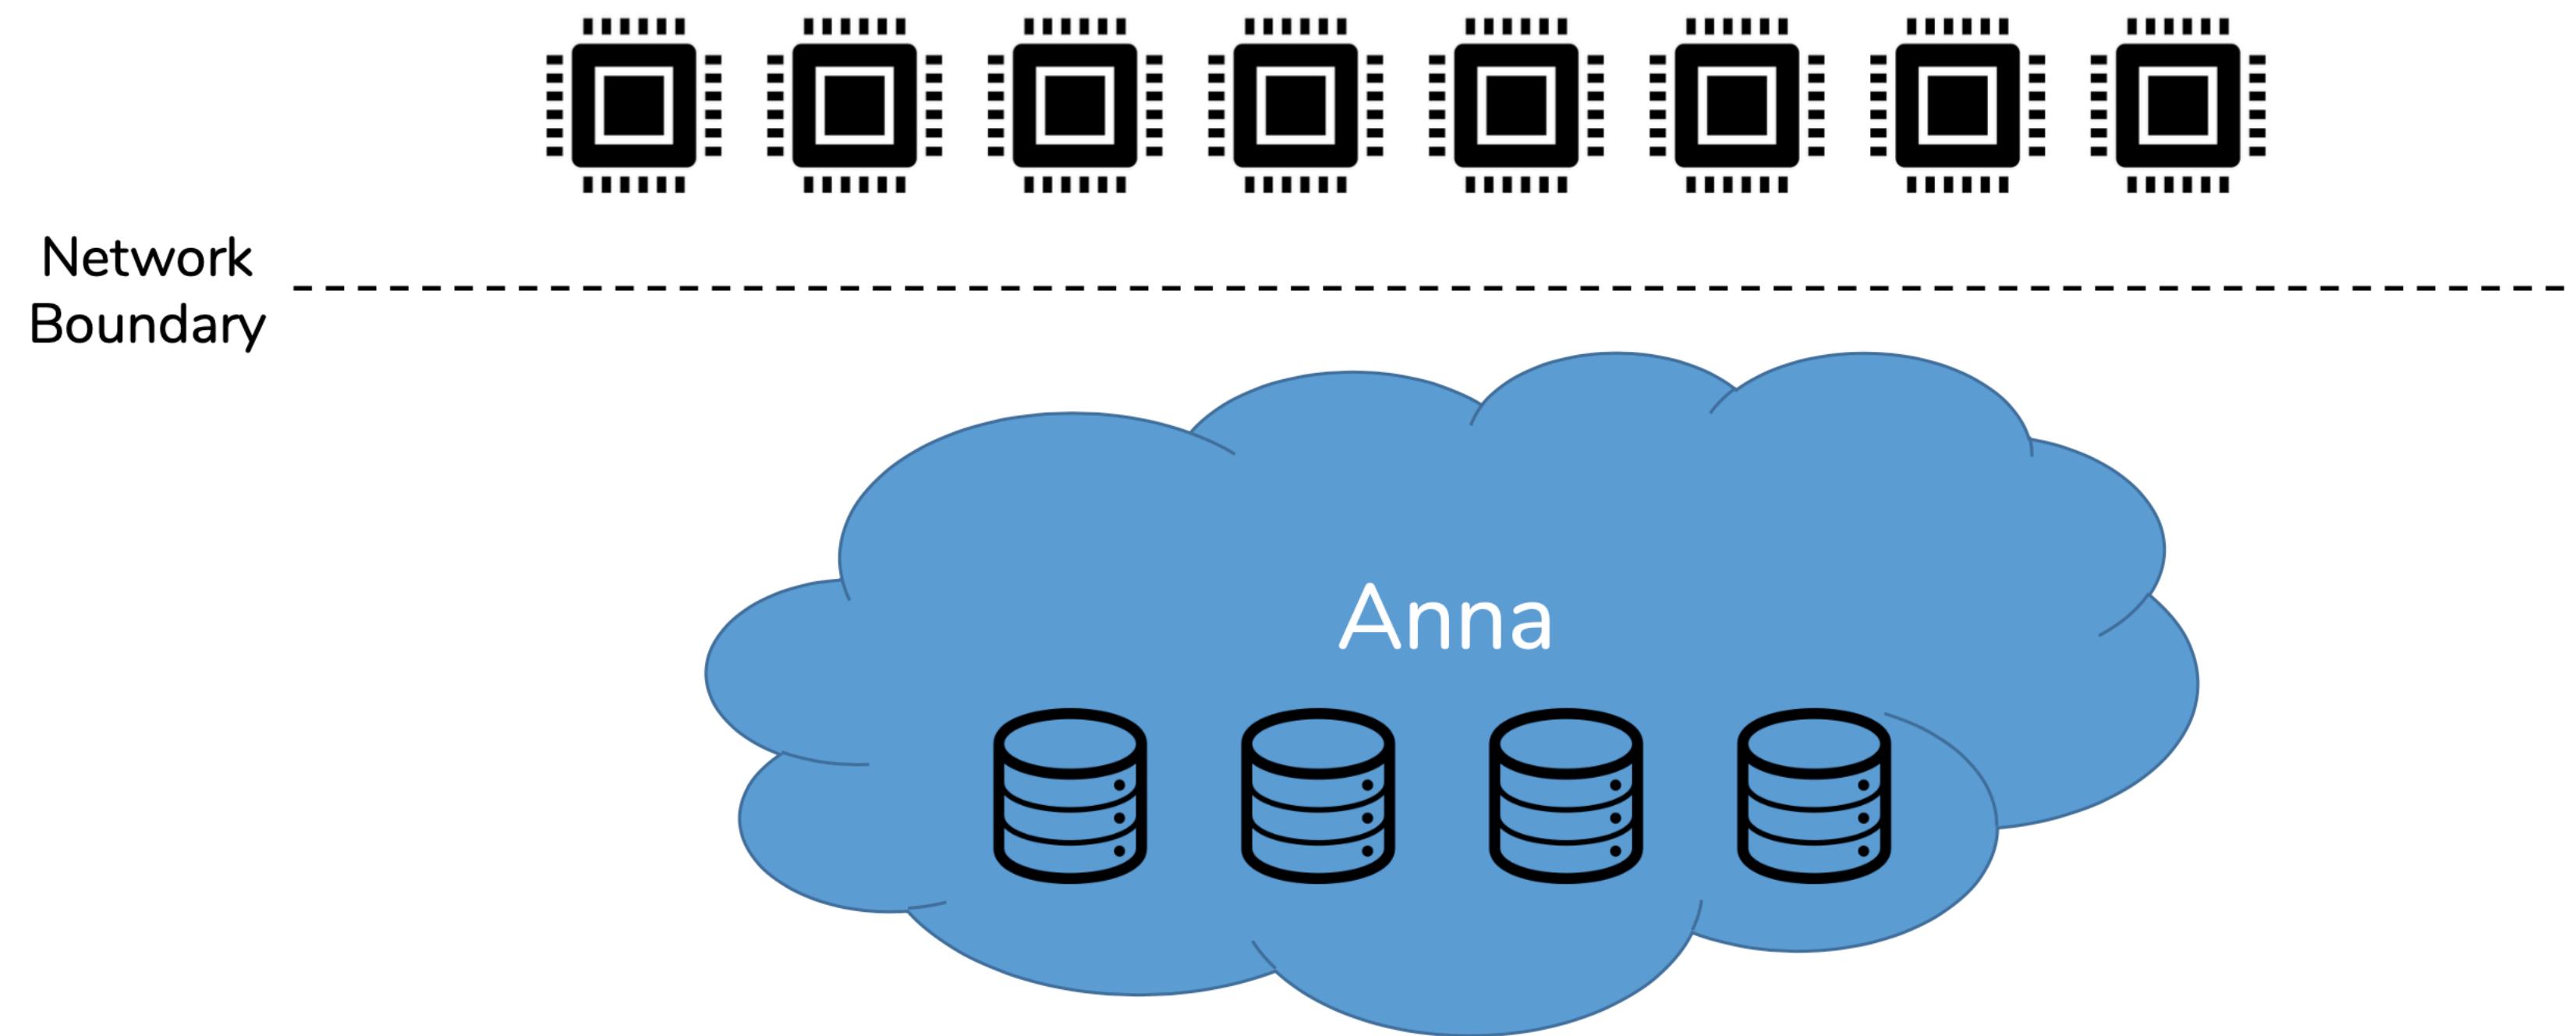


Cludburst: FaaS over Anna

- Maintain disaggregation of compute & state
- Make serverless a viable option for stateful applications
- Use Anna for both storage and communication



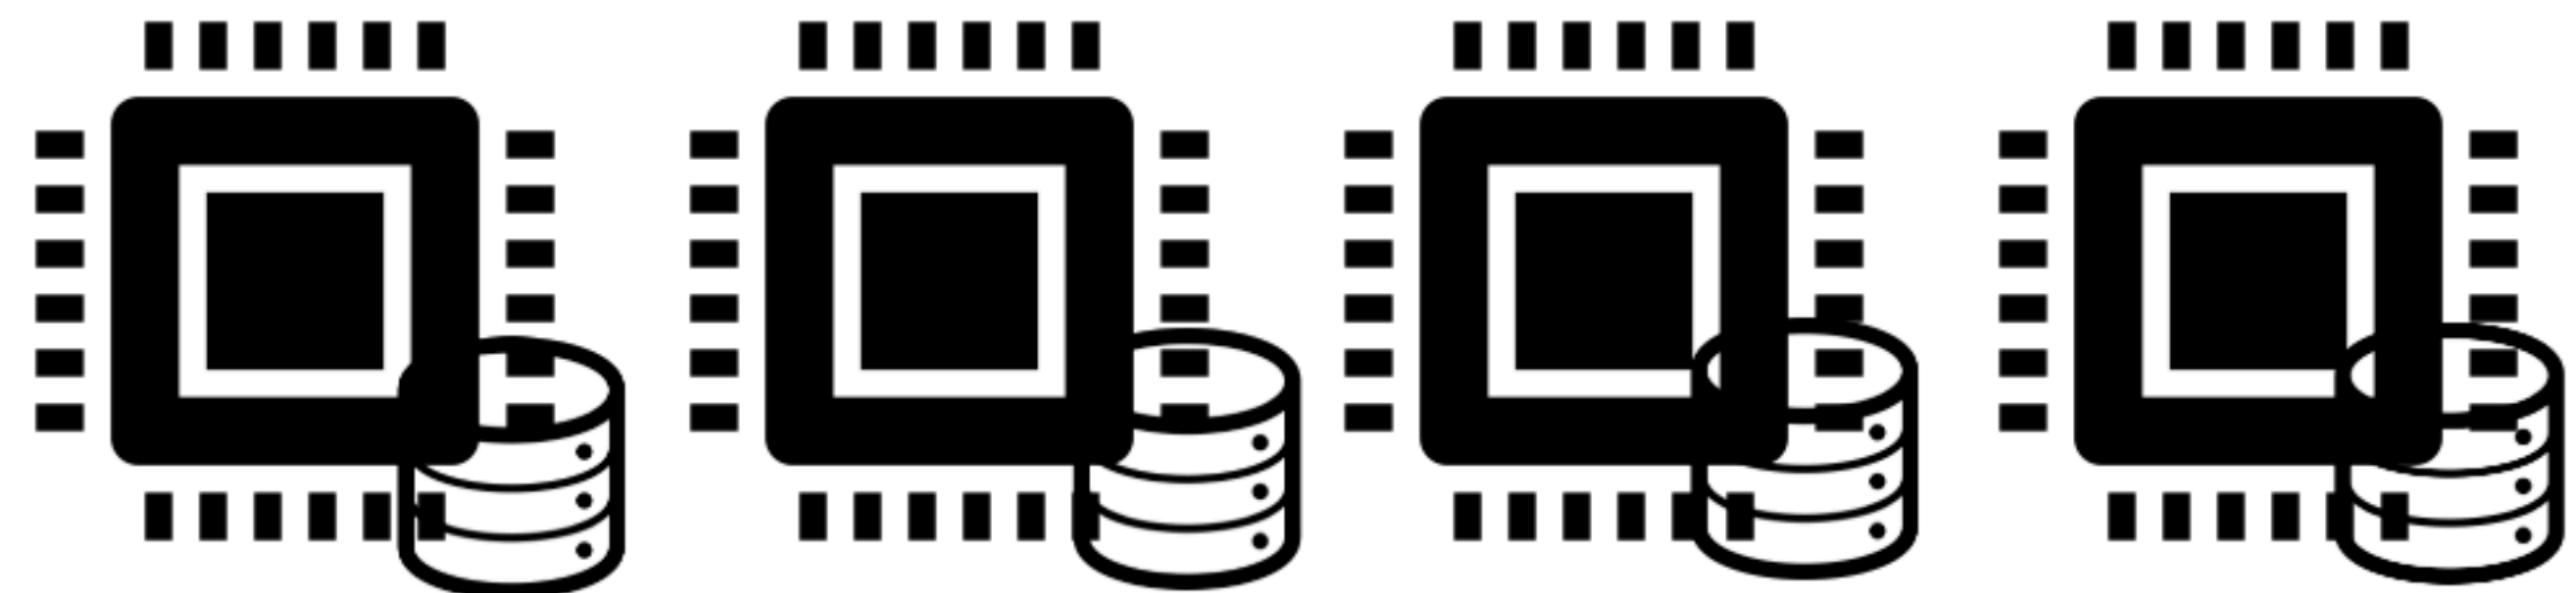
Cludburst: FaaS over Anna



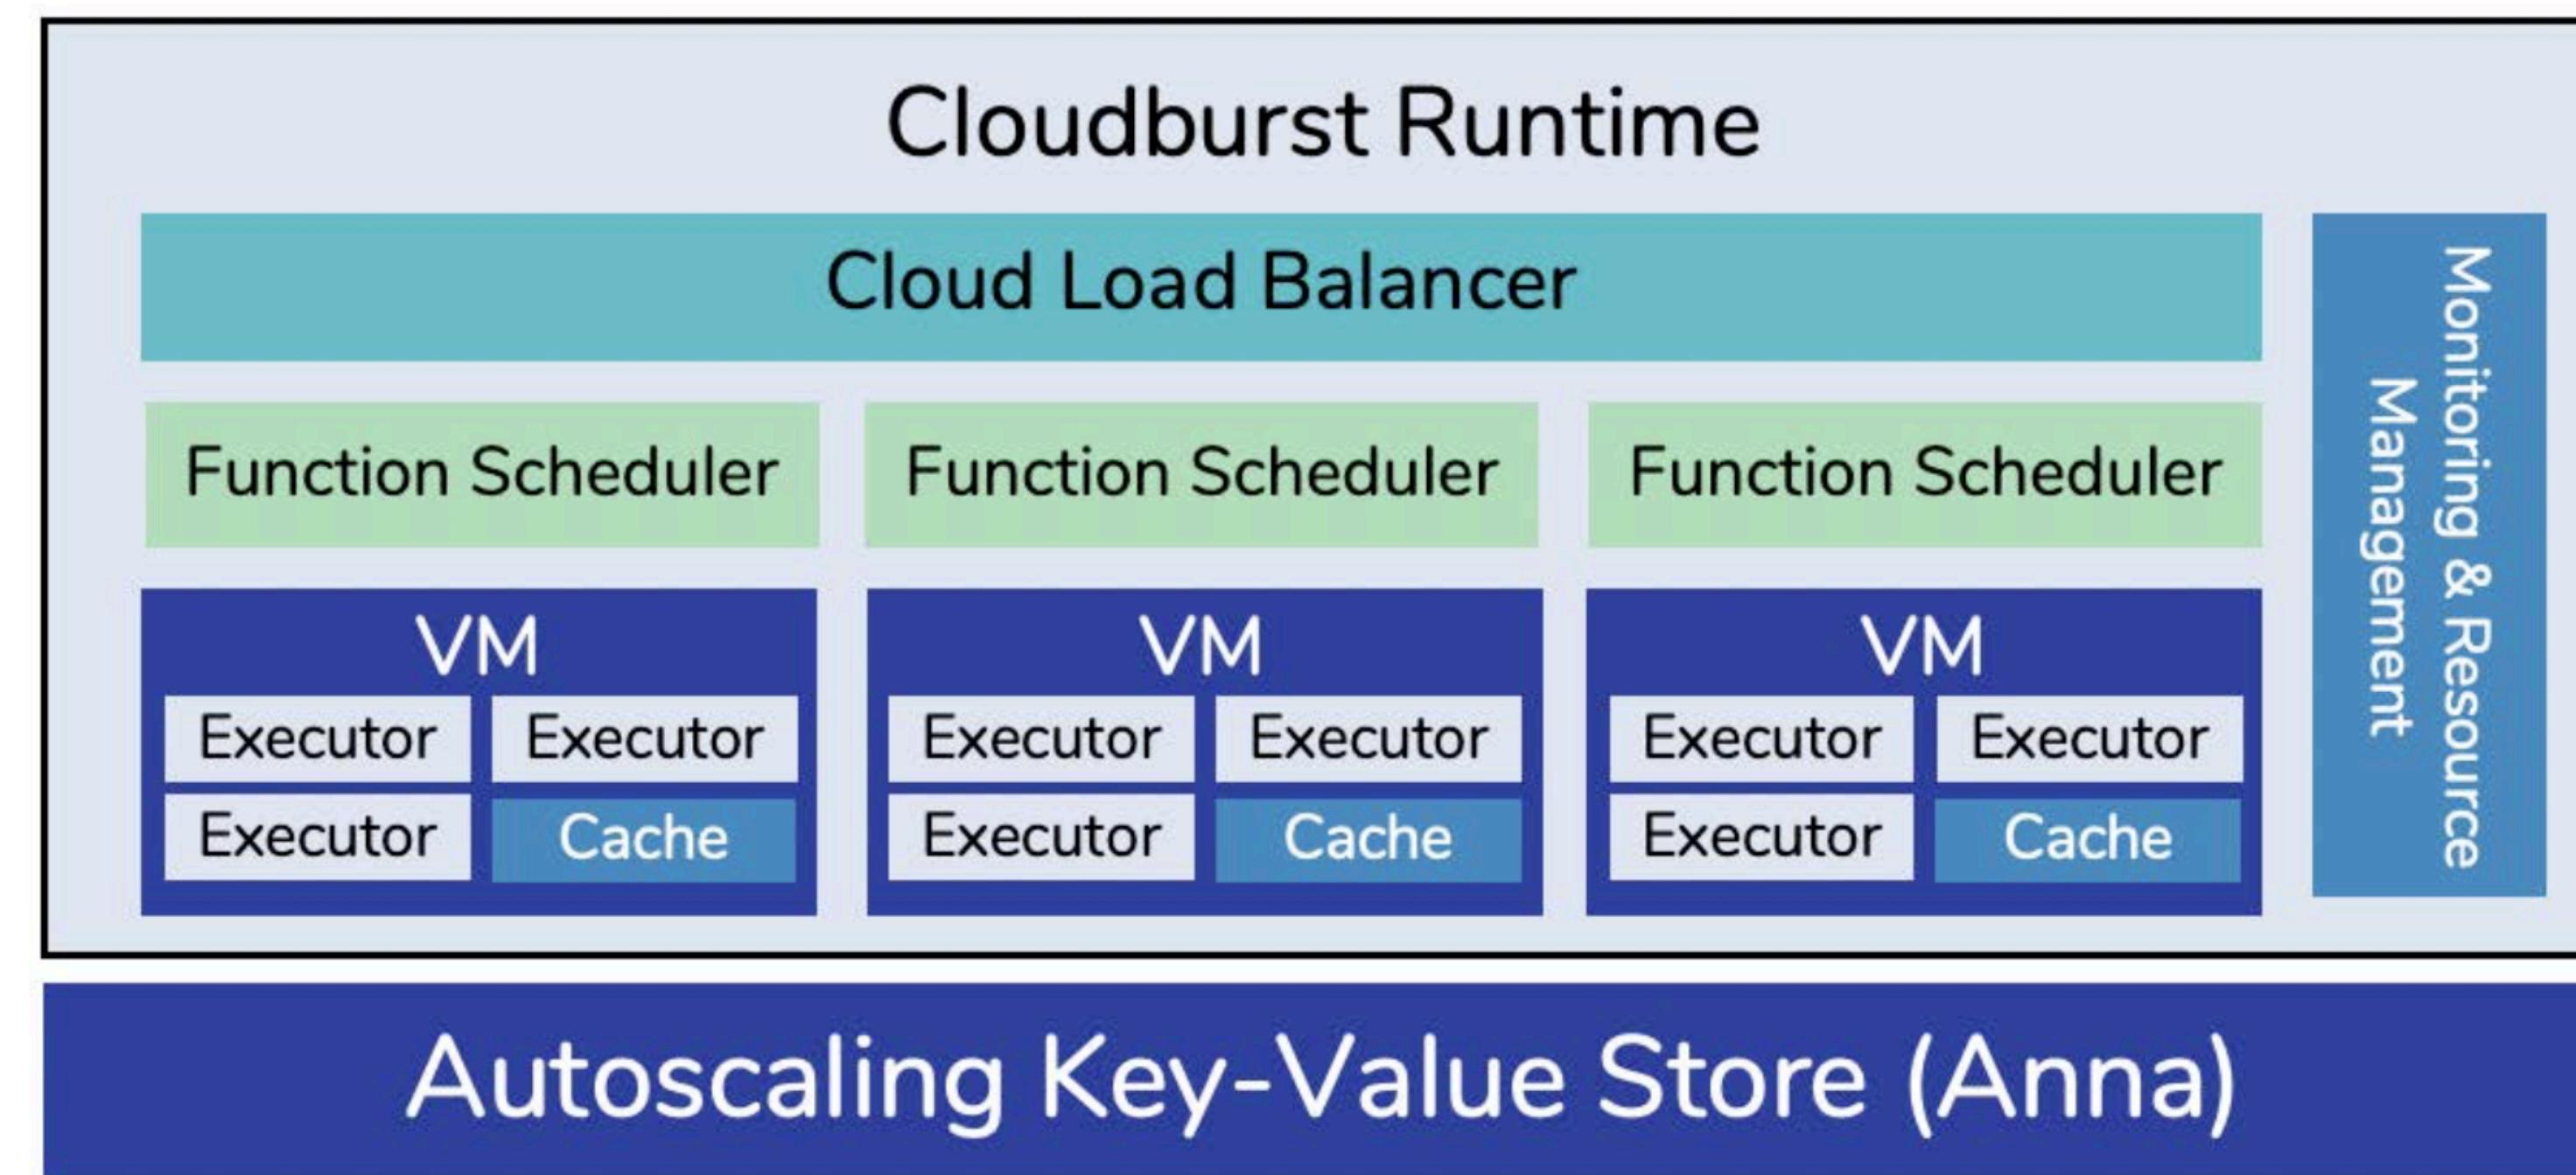
SOLUTION: Logical disaggregation with physical colocation

Key Idea: Caching

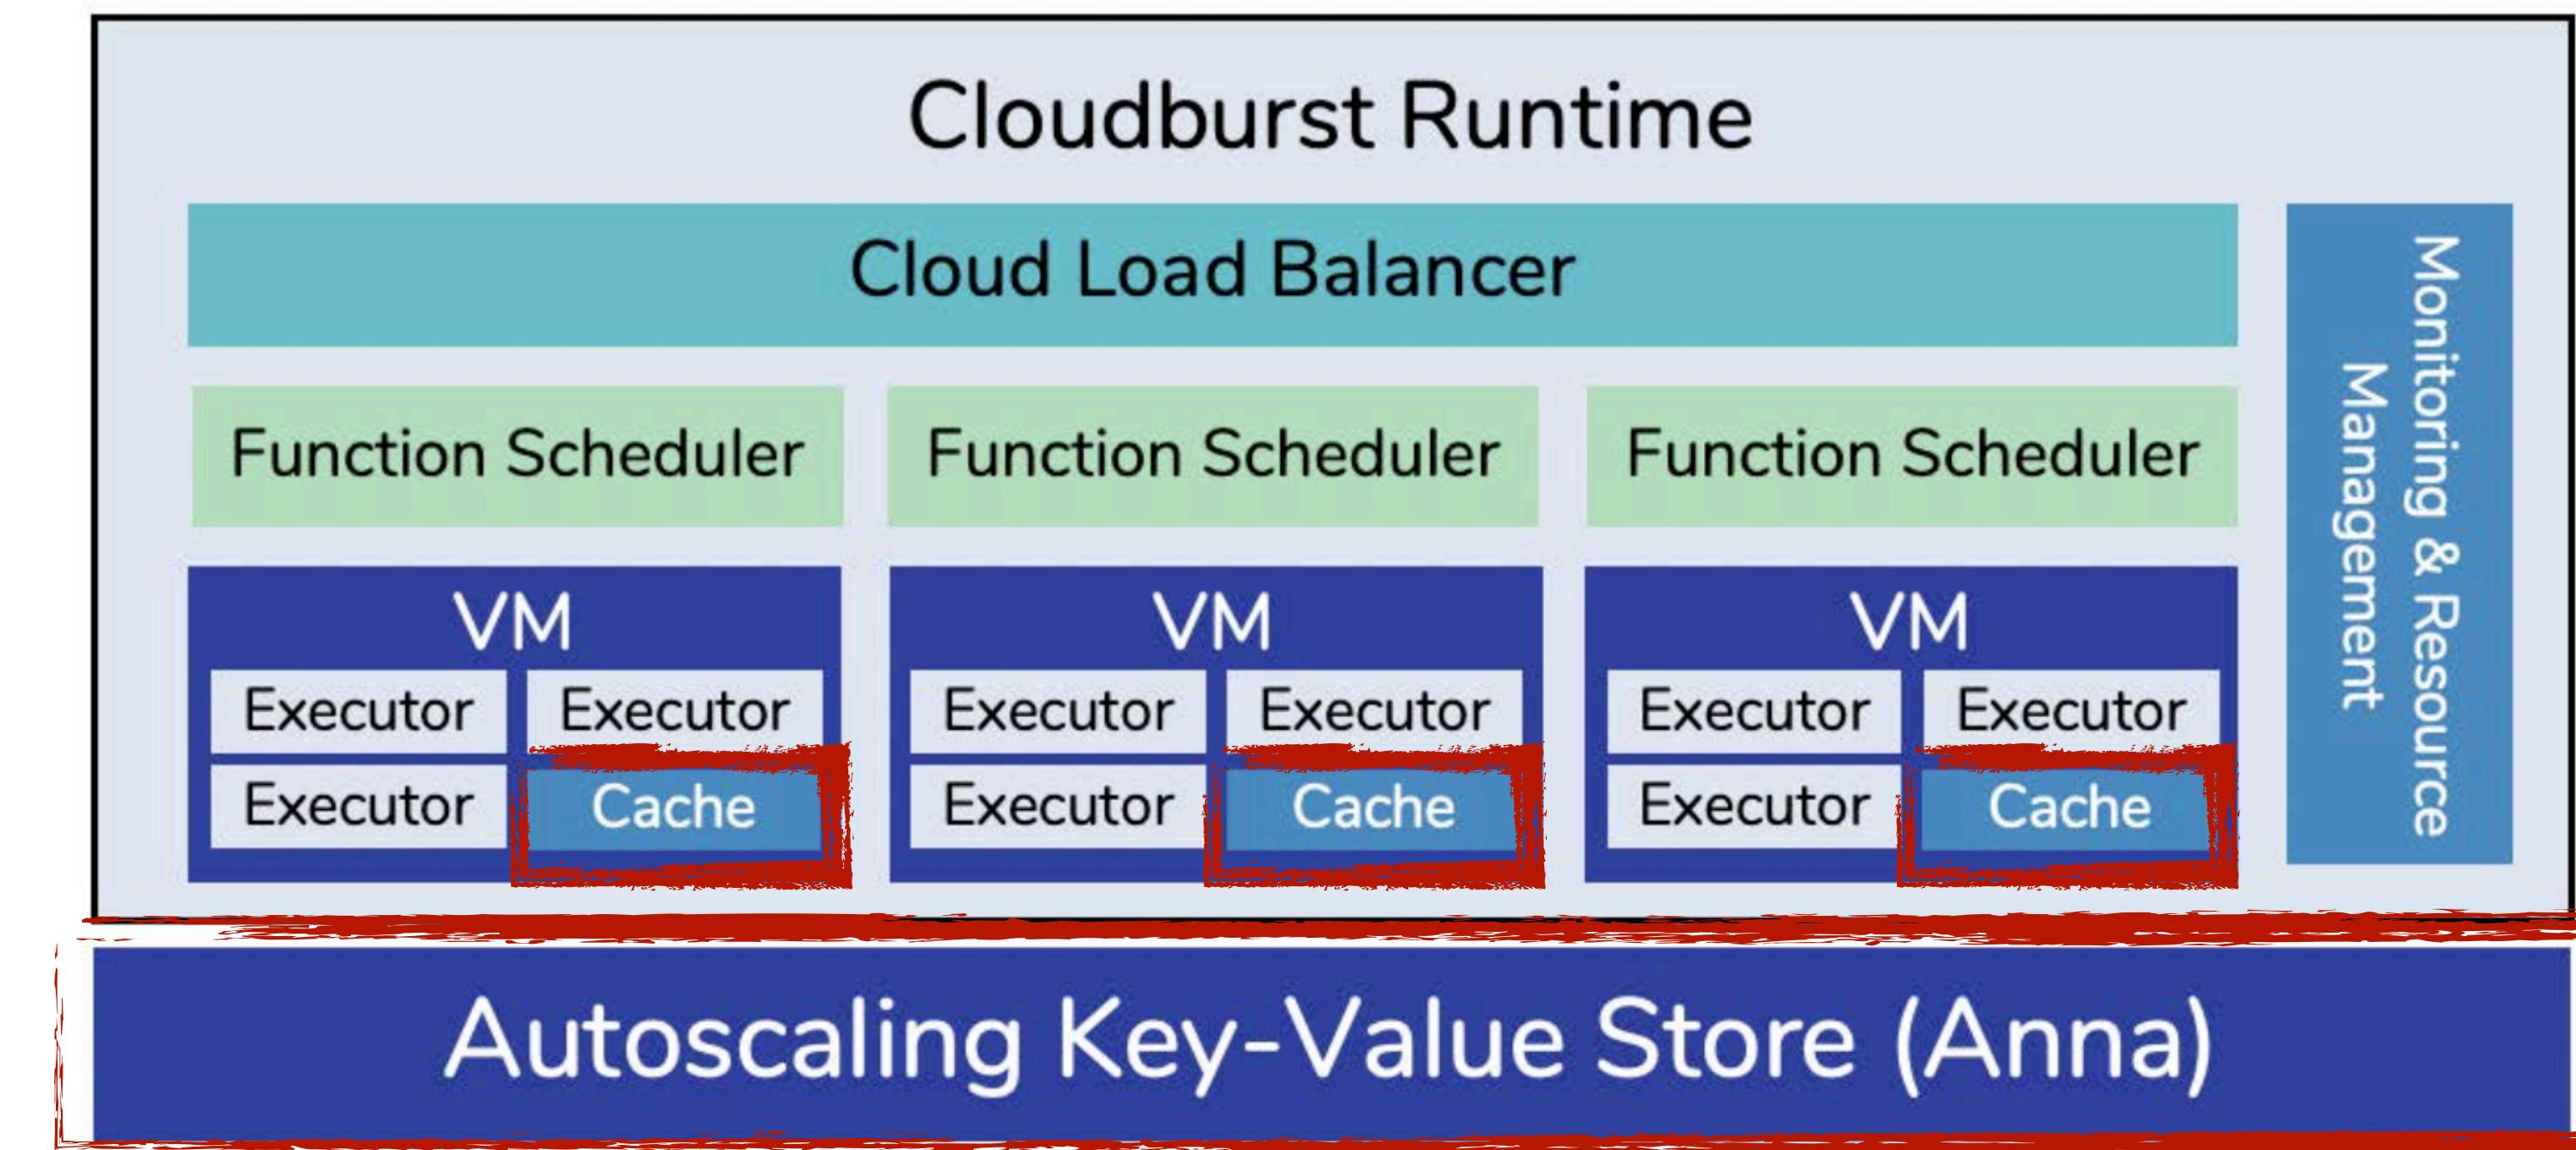
- Enable low-latency data access by caching data close to code execution
- Communication (and composition) is achieved via a fast-path on top of KVS puts and gets



Anna Architecture

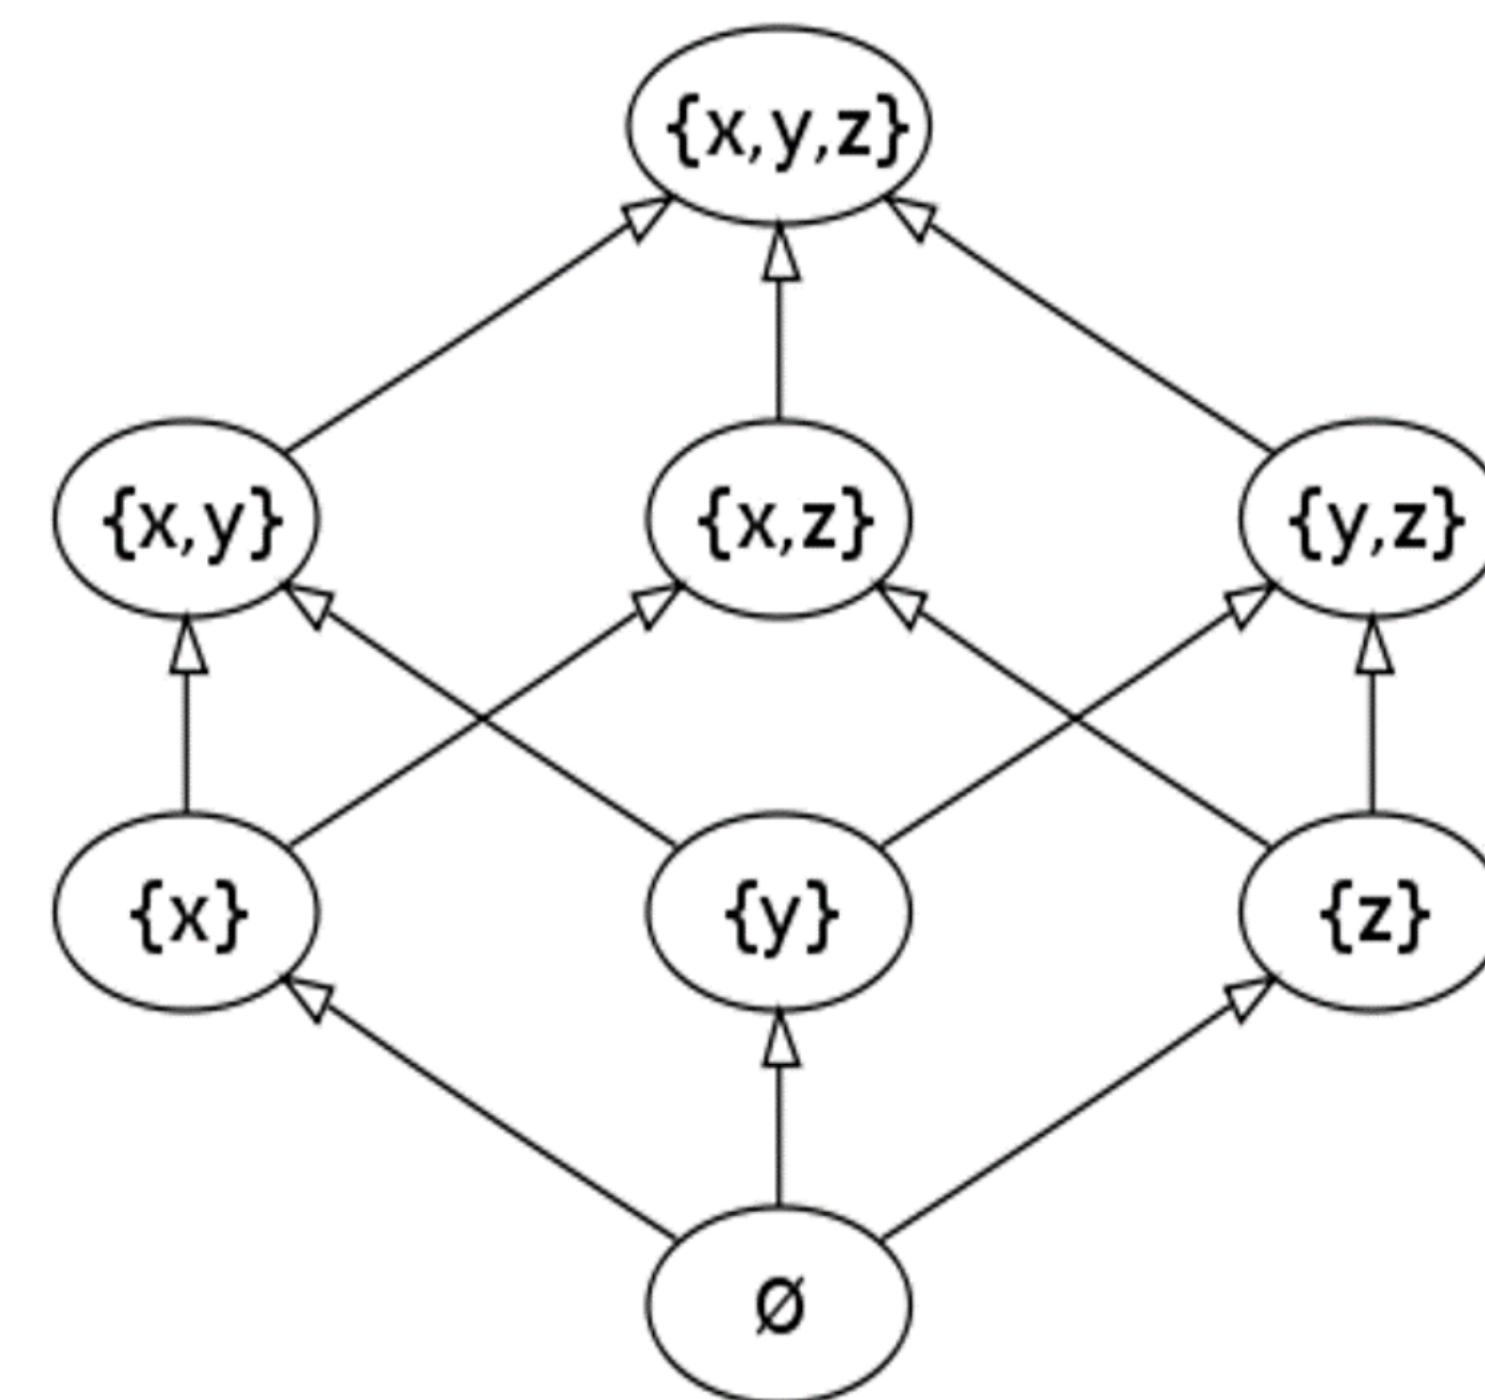


Caching is a Problem



Lattice

- Data structure that accepts incoming update in a way that is associative, commutative, and idempotent (ACI)
- Achieves eventual replica convergence



Cache Consistency in DAG

- The Scheduler creates a schedule for a DAG
- Scheduler persist DAG topologies in KVS
- DAG is a session – **CACHE CONSISTENCY NEEDED**
 - Run the entire DAG in a single thread
 - ✗ Autoscaling
 - ✗ Flexible
 - We need a distributed session consistency.

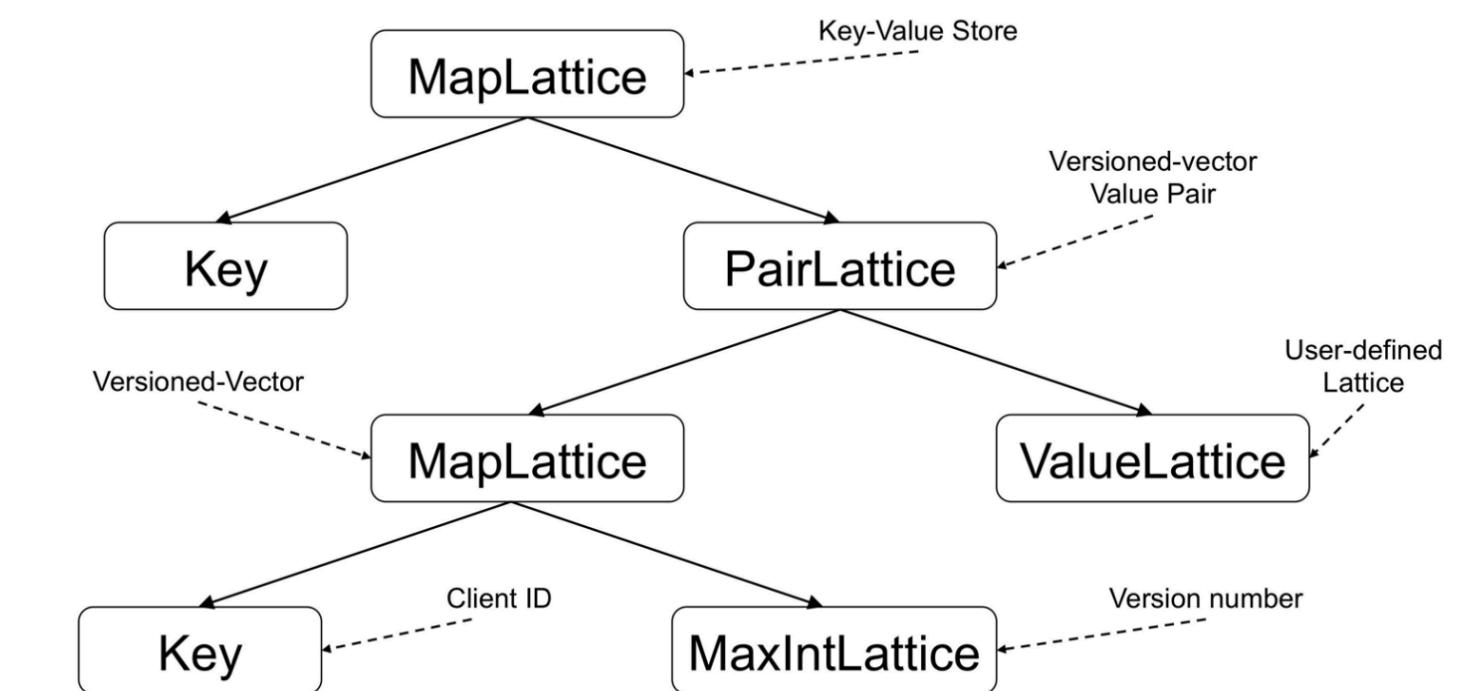
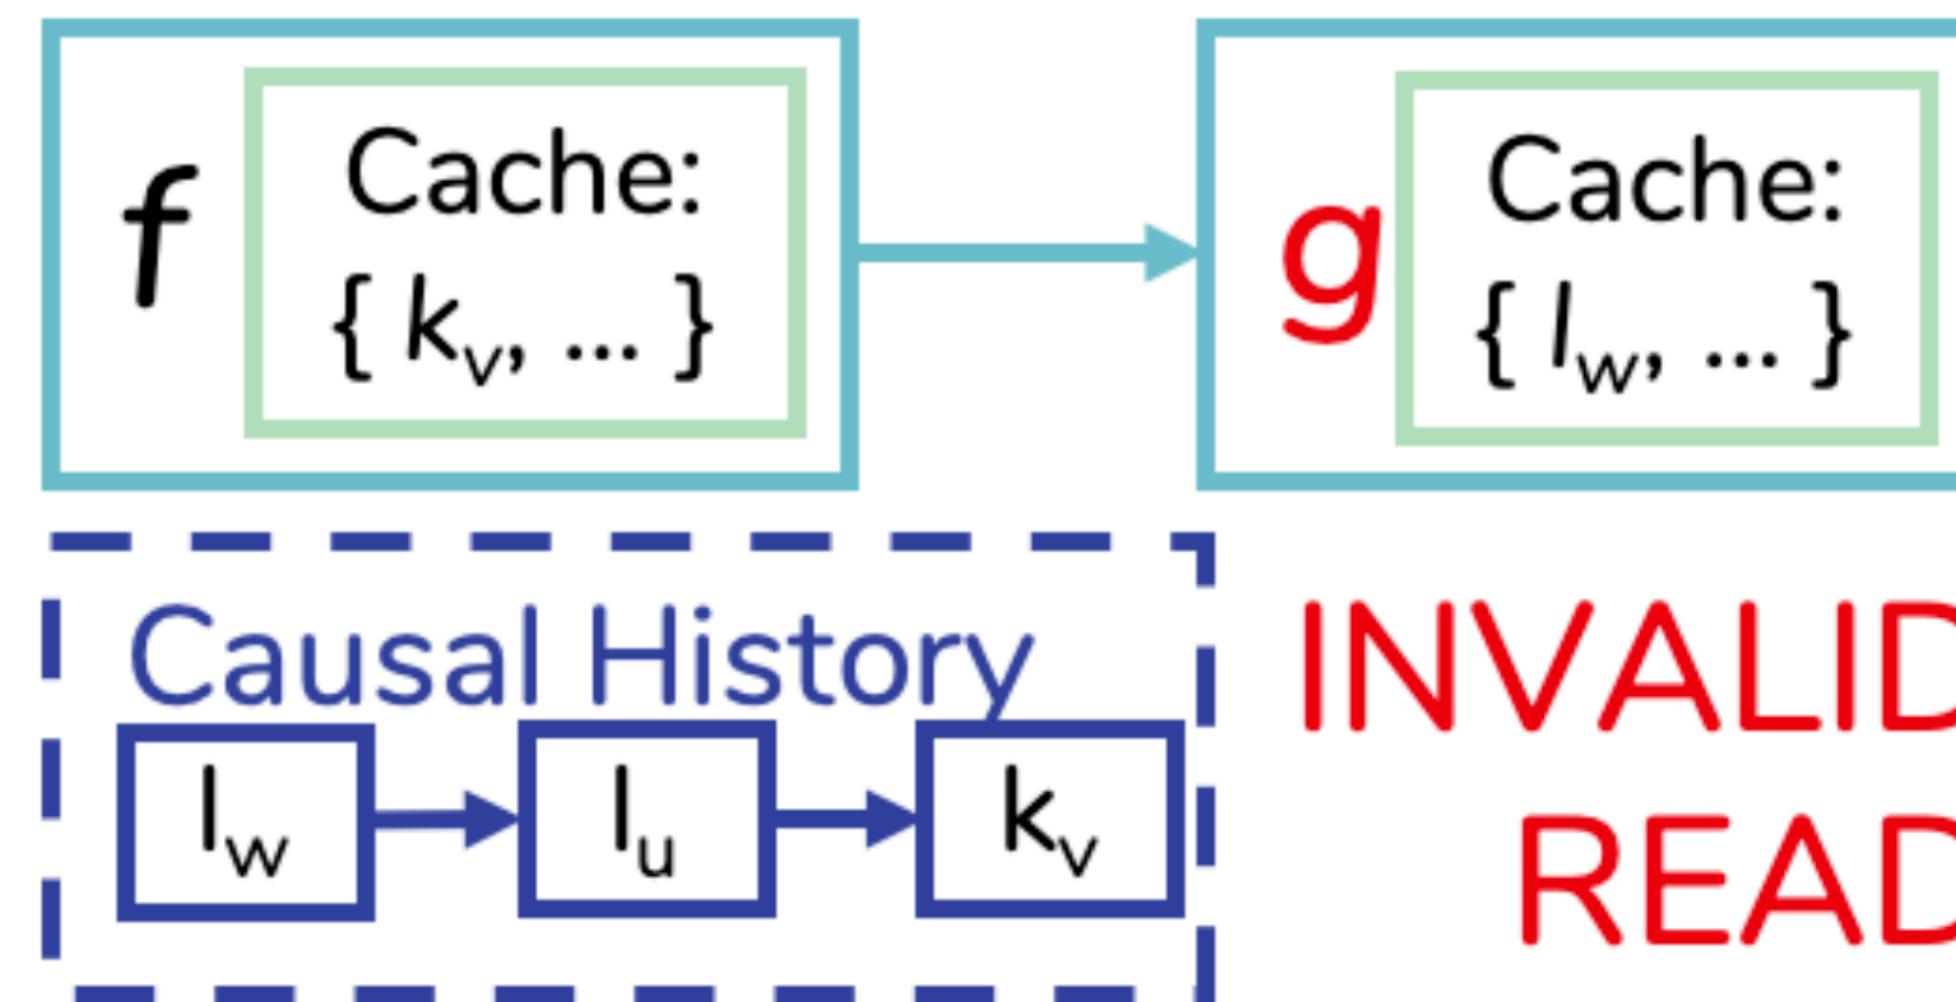
Distributed Session Repeatable Read

Algorithm 1 Repeatable Read

Input: k, R

```
1: //  $k$  is the requested key;  $R$  is the set of keys previously read  
    by the DAG  
2: if  $k \in R$  then  
3:    $cache\_version = cache.get\_metadata(k)$   
4:   if  $cache\_version == NULL \vee cache\_version \neq R[k].version$  then  
5:     return  $cache.fetch\_from\_upstream(k)$   
6:   else  
7:     return  $cache.get(k)$   
8: else  
9:   return  $cache.get\_or\_fetch(k)$ 
```

Distributed Session Causal Consistency



Distributed Session Causal Consistency

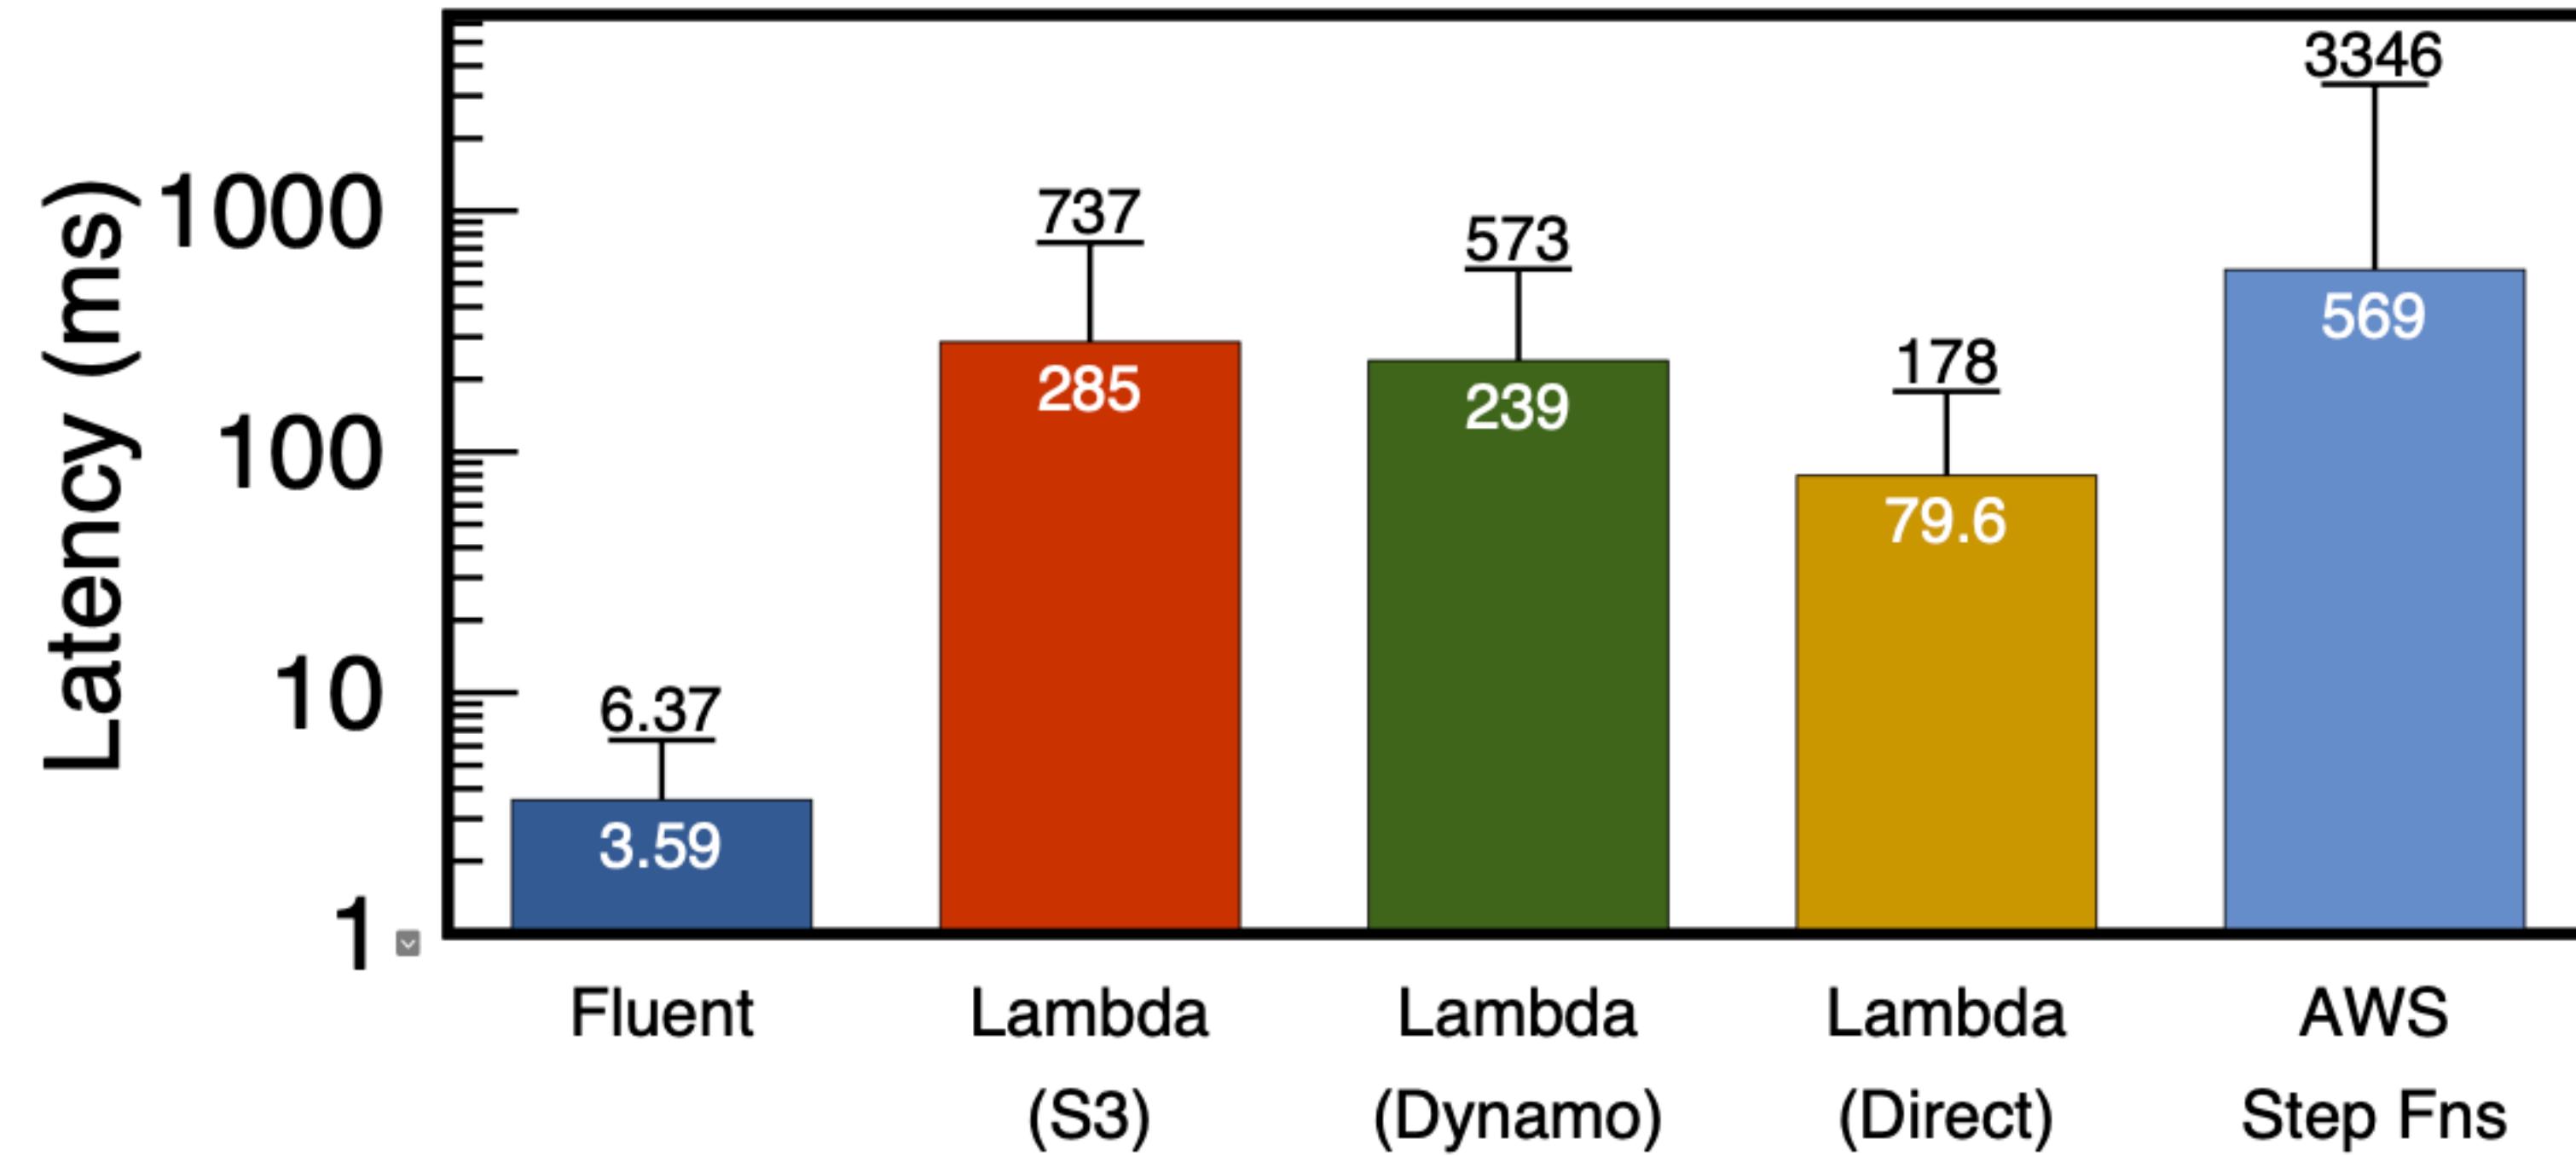
Algorithm 2 Causal Consistency

Input: $k, R, dependencies$

```
1: //  $k$  is the requested key;  $R$  is the set of keys previously read  
   by the DAG;  $dependencies$  is the set of causal dependencies  
   of keys in  $R$   
2: if  $k \in R$  then  
3:    $cache\_version = cache.get\_metadata(k)$   
4:   //  $valid$  returns true if  $k \geq cache\_version$   
5:   if  $valid(cache\_version, R[k])$  then  
6:     return  $cache.get(k)$   
7:   else  
8:     return  $cache.fetch\_from\_upstream(k)$   
9: if  $k \in dependencies$  then  
10:    $cache\_version = cache.get\_metadata(k)$   
11:   if  $valid(cache\_version, dependencies[k])$  then  
12:     return  $cache.get(k)$   
13:   else  
14:     return  $cache.fetch\_from\_upstream(k)$ 
```

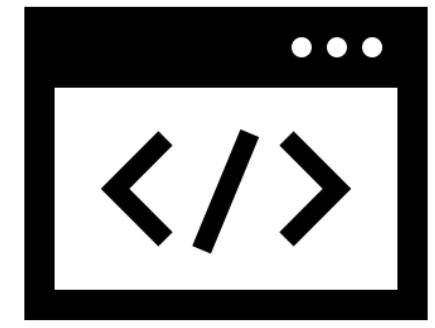
Function Composition, Revisited

$f(g(x))$

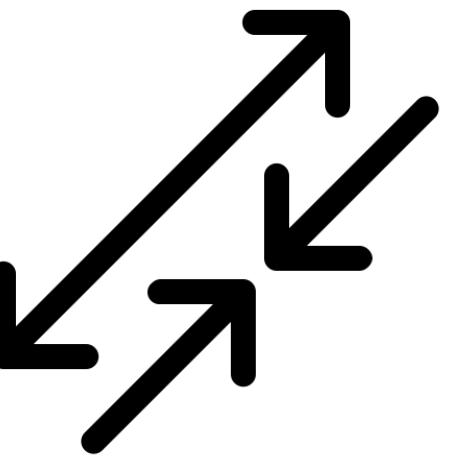


Median and 99th percentile latencies for composing two arithmetic functions on AWS Lambda.

Moving Forward from FaaS



Building
Developer
Tools



Developing
Autoscaling
Policy



Designing
SLOs & SLAs

Thanks